

VayuNetra — Complete Project Reference

Team DaGoats · ET AI Hackathon 2.0 · Problem Statement 5 — AI-Powered Urban Air Quality Intelligence for Smart City Intervention

Revision 19 August 2026. Generated from the code it describes.

How to read this document

This is the whole project in one place: what it does, how every part works, which technology does what, what the system is measurably worth, and where it is weak. It assumes you are technical but have never seen this codebase and do not necessarily know the air-quality domain. Terms are defined the first time they appear.

It is a source of truth, which means it contains the bad news too. Chapter 14 is a catalogue of known limitations with numbers attached. Individual chapters state what is unverified rather than rounding up. If a claim here is weaker than the pitch deck, this document is the one that is right.

Every number came from the code or a generated artifact, not from memory. Chapter 15 tells you how to re-derive any of them yourself — the exact commands. If you only trust one chapter, trust that one, because it lets you check the rest.

Which chapter answers which question

If you want to know	Read
What this is and why it exists	1
Where to open it, and the keyboard shortcuts	2
What a user actually sees and does	3
How a sensor reading becomes something on screen	4
What technology is used and why that one	5
Every table, column and access policy	6
Where the data comes from and how often	7
How the forecast works and what it is worth	8
How pollution is blamed on a source	9
What the "agents" actually are	10
Every API endpoint	11
How the web app is built	12
How it is deployed, tested and operated	13
Where it is weak	14
How to verify any claim here	15
What a judge will ask, and the honest answer	16
How to present it	17
Something is broken	18

A note on the screenshots

Every screenshot was captured by an automated Playwright walkthrough against the running application (`web/scripts/qa/full-walkthrough.mjs`), not staged or edited. Where an officer action appears — approving a recommendation, dispatching it, closing it — that action really happened against the live database and was reset afterwards; the audit trail retains the record.

1 • What VayuNetra is

The problem, stated plainly

An Indian city knows its air is bad. A monitoring station reports $180 \mu\text{g}/\text{m}^3$ of PM2.5 and a dashboard turns red. What the city does *not* know is the part that would let it act: **which square kilometre, caused by what, for how long, and who should be sent where.**

That gap is not hypothetical. India's National Clean Air Programme requires city action plans, and a 2024 CAG performance audit found that a large majority of monitored cities could not demonstrate that their monitoring translated into targeted action. The instrumentation exists. The decision layer between a reading and an inspection does not.

VayuNetra is that decision layer for ten Indian cities.

What it produces

Four things, each per ~1 km cell rather than per city:

1. **Where the air is bad now** — a dense field interpolated from sparse stations onto a hexagonal grid, so every cell has an estimate rather than only the handful containing a monitor.
2. **What is causing it** — an apportionment across traffic, industry, construction dust, biomass burning, waste burning and transported pollution, with a stated method and a confidence.
3. **What it will be** — PM2.5 at +24, +48 and +72 hours, with an uncertainty band that is calibrated rather than decorative, and a probability of crossing each CPCB threshold.
4. **What to do about it** — a ranked enforcement worklist, each item carrying satellite evidence and citations to the specific regulation, ending in a draft legal notice; plus citizen advisories in eight languages across four delivery channels.

The fourth is the one that distinguishes this from a dashboard. A recommendation has a lifecycle — proposed, approved, dispatched, closed — and every transition is recorded, so the question "did anything happen because of this?" has an answer.

The vocabulary you need

You will meet these terms throughout. None is invented here; all are standard in their field.

H3 cell. A hexagonal spatial index from Uber. The system uses **resolution 8**, where a cell is roughly 0.74 km^2 — about a kilometre across. Hexagons are used rather than squares because every neighbour is equidistant, which matters when modelling how pollution moves between adjacent areas. Every measurement, forecast and attribution in the system is pinned to one cell. Delhi is about 3,466 cells.

PM2.5. Particulate matter under 2.5 micrometres. Small enough to reach the alveoli and cross into the bloodstream, which is why it dominates health-burden estimates. Measured in $\mu\text{g}/\text{m}^3$. The system forecasts PM2.5 specifically, and reports other pollutants without forecasting them.

CPCB bands. India's Central Pollution Control Board defines the national AQI. For PM2.5 the 24-hour breakpoints are Good 0–30, Satisfactory 30–60, Moderate 60–90, Poor 90–120, Very Poor 120–250, Severe 250+. These thresholds are not cosmetic — they trigger real regulatory action under GRAP (the Graded Response Action Plan) in the Delhi NCR.

Persistence. The forecasting baseline: "in 24 hours it will be what it is now." It is a much harder baseline than it sounds, because air quality is strongly autocorrelated. A model that cannot beat persistence is not adding information. Skill is reported as $1 - \text{RMSE_model} / \text{RMSE_persistence}$, so +0.10 means ten percent less error than doing nothing.

Prediction interval and coverage. The forecast reports a band intended to contain the true value 80% of the time. *Coverage* is how often it actually does. A band claiming 80% and delivering 62% is not conservative — it is wrong, and it is wrong in the direction that makes a user overconfident.

Conformal prediction. The technique used to make that band honest. Explained from first principles in chapter 8; the short version is that it calibrates the interval against the model's own historical errors rather than trusting the model's self-reported uncertainty.

Source apportionment. Deciding what fraction of the pollution came from which kind of source. In the literature this is done with receptor models on chemically speciated filter samples. This system does it from routine pollutant monitoring plus satellite markers, which is cheaper and continuous but less direct — a trade-off chapter 9 is explicit about.

Attribution confidence and abstention. The system has two apportionment methods and will refuse to use the stronger one when it has not earned it. Which method produced a given number is recorded in the data and visible in the interface.

What it is not

- It is not a monitoring network. Every reading comes from public sources — CPCB, OpenAQ, ERA5, Copernicus, NASA FIRMS, OpenStreetMap. No hardware was deployed.
- It is not deployed with a real pollution control board. No official has used it operationally, and no recommendation it produced has resulted in a real inspection.
- It does not claim causal evidence that any intervention worked. Chapter 14 is precise about what the intervention analysis does and does not establish.
- It does not use a language model to write anything a citizen reads. That is a deliberate design choice, explained in chapter 10, and it is a safety property rather than a missing feature.

The shape of the system in one paragraph

Public data lands via scheduled connectors into a PostgreSQL database, binned onto H3 cells. Two model families run over it: a gradient-boosted quantile forecaster with conformal calibration, and a source-attribution model that blends chemical-signature priors with a SHAP-explained gradient-boosted model. A LangGraph state machine orchestrates a run — attribute, forecast, and if a spike is detected, generate enforcement recommendations — then always produces citizen advisories. A FastAPI service serves the results; a React console presents them to an officer over a deck.gl map, and a public site presents them to a citizen. Telegram, an IVR phone line and a public-display renderer carry advisories to people who will never open either.

2 • Getting in

Front page	https://vayunetra-aqi.vercel.app — what VayuNetra is, with a live ten-city board. Two doors: Check your city's air (public app) and Open console (operations).
Public app	<code>/city/<id></code> · <code>/map</code> · <code>/forecast</code> · <code>/rankings</code> · <code>/about</code> . Opens on Delhi. No login: reads are public under Postgres row-level security; every write goes through the server with the service role.
Console	<code>/console</code> — the operations surface. Seven sections, addressed by <code>?section=</code> .
API	https://vayunetra-c8i8.onrender.com — <code>GET /health</code> reports <code>DEMO_MODE</code> . Every endpoint returns <code>{success, data, error, meta}</code> . Interactive docs at <code>/docs</code> , generated by FastAPI from the code.


Run locally	<code>make dev</code> starts the API on :8000 and the web app on :5173. <code>make prewarm</code> wakes and verifies the live stack (GO / NO-GO). <code>make test</code> runs the 341 backend tests.
Demo mode	<code>DEMO_MODE=true</code> serves bundled fixtures with an identical interface, labelled in <code>/health</code> . Independently, the web app falls back to bundled fixtures whenever the API is unreachable, showing an amber <i>"backend waking up — showing bundled demo snapshot"</i> notice. The free-tier API host sleeps after inactivity; the notice clears on the first successful call.
Keyboard	In the console: 1–7 switch sections · [] cycle cities · P presentation mode (larger type for a projector) · ⌘K / Ctrl-K jump to any city or section · Esc closes the Cell Story.
Deep links	Every console state is a URL: <code>/console?city=mumbai&section=forecast&cell=88608b56cbffff&mode=pm25&layers=sources,plumes,wards,freight,fires</code> . This is how you hand someone the exact view you are describing.

3 • What you actually see


Two surfaces share one backend. The **public site** is for a resident; the **operations console** is for an officer. Everything below is a real screenshot of the running application.

2.1 The landing page

The entry point states what the system does and, unusually, what its numbers are worth — the validation table is on the front page rather than buried.



The landing page




BUILT ON PUBLIC DATA INFRASTRUCTURE

CPCB / CAAQMS Sentinel-5P Sentinel-2 Open-Meteo · ERA5 NASA FIRMS OpenStreetMap GPW v4.11


LIVE RIGHT NOW

Ten cities, this minute.


Indian National AQI (CPCB) — the maximum of each city's pollutant sub-indices. Open any city.



THE DATA, AT A GLANCE




City PM2.5 — now vs forecast +24h



Running scale



The data at a glance — live counts, not illustrations



Live from the production system, as of 19 Aug, 08:46 pm — aggregated from real station measurements and the latest attribution run; refreshed every 10 minutes. Open the console for the cell-level numbers.

HOW IT WORKS

From raw signal to a signed notice, in one pipeline.

CPCB measures. SAFAR forecasts. The missing layer is operational: who is responsible in this square kilometre, and what should be done before tomorrow.



Trace

A gradient-boosted model with SHAP explanations assigns PM2.5 blame to traffic, construction, industry and burning — per square-kilometre H3 cell, cross-checked, bucket by bucket, against published apportionment studies for Delhi and Bengaluru — agreements and disagreements both shown.



Predict

Quantile forecasts at 24, 48 and 72 hours for every cell, with conformal-calibrated intervals and a calibrated probability of crossing Very Poor / Severe. Benchmarked on a strict temporal split against persistence, seasonal-naïve and climatology – the numbers, including the weak spots, are printed in the product.



Act

A ranked enforcement worklist scores 647 registered and satellite-detected sources across 10 cities. One click opens an evidence dossier with a real Sentinel-2 patch, regulatory citations, and a draft notice PDF.



Protect

Health advisories generated from the forecast and targeted at the most vulnerable zones — 11,000+ mapped hospitals, 7,700+ schools, elder-care homes and outdoor-work sites — in eight Indian languages (Hindi, Kannada, Marathi, Tamil, Telugu, Bengali, Gujarati, English) over the web app, Telegram, IVR calls and public displays.

ARCHITECTURE

Two seams. Five agents and a gate. Zero blocking.

Everything decouples through the Supabase schema and the API contract: models write rows, the API reads

How the system works, end to end

VALIDATION

Every number is checked — including the failures.


Each claim below is reproducible from the repository — the evaluation notebook and `python -m ml.eval.benchmark`, whose artifacts the API serves and the console prints. Where a method underperformed, that result ships too.

CLAIM	RESULT	METHOD
Attribution checked against published apportionment	cosine 0.88 / 0.90 / 0.93	vs SAFAR-Delhi 2018, CSTEP-Bengaluru 2022 (verified from the report), Urban-Emissions Mumbai — plus bucket-by-bucket tables vs TERI-ARAI and Guttikunda et al.
Attribution behaves physically	2.30× traffic signal in rush hours	IST rush vs off-peak SHAP; weather controlled
Forecast beats real baselines	Delhi +9 / +13 / +12% · Mumbai +17 / +19 / +21% at 24/48/72 h	multi-season temporal split (2025-26 winter + summer 2026), monthly refit on the trailing 90 d; served forecast = LightGBM blended with persistence; vs persistence and weekly seasonal-naive — the raw model's weaker numbers ship too
It warns before the air turns	51-54% of Very-Poor onsets flagged 1-3 days ahead	alarm on the calibrated probability $P \geq 0.3$ (precision 0.64-0.68); persistence catches 0% by construction; the Severe tail stays weak — stated, not hidden
Uncertainty is calibrated, and where it is not we say so	80% band — 0.783 Delhi · 0.749 Kolkata, published by predicted level	conformal calibration; we report coverage per predicted band rather than one marginal number, because Kolkata drops to 0.55 where it matters most
Enforcement is equitable	no socio-economic inputs, by construction	fairness audit on every live recommendation (n=390 at the July audit)
Model choice was earned	TFT trained on GPU — and rejected	LightGBM won every launch city on held-out skill
The loop is fast	seconds from signal to cited recommendation	measured pipeline latency with live per-node agent traces — the time to PRODUCE the recommendation, not a municipality's response time; approval, dispatch and closure are timestamped per action


Validation on the front page: the numbers, with their baselines

2.2 The public city page


What a resident sees. The design goal is that someone with no technical background can answer "is it safe to go out?" in under five seconds, and then go deeper if they want to.




The city overview




Every measured pollutant, with the prominent one marked




A pollutant's reference detail — what it is and where it comes from




The trend graph and the worst areas right now




The air-quality calendar — a year at a glance




Health guidance, including the cigarette equivalent



Condition-specific guidance, illustrated per condition




Comparison across the ten cities




The same air on three scales: CPCB, US EPA, WHO

2.3 The map


The map is the core of the console. Every ~1 km H3 cell is coloured by its dominant source or by concentration, and clicking one opens its story.




Map coloured by PM2.5 concentration




Map coloured by dominant source — the blame map




Satellite NO₂ column, from Sentinel-5P




Layer control: sources, plumes, wards, freight corridors, fires



Overlays combined



Time-scrub replay of the last 24 hours



A cell's story — who is to blame here, and how confident we are



The same cell's forecast, with its calibrated band

2.4 The officer loop — enforcement

This is the part that closes. A recommendation is ranked, evidenced, cited, drafted into a legal notice, dispatched to a ward team and closed with a finding. Every transition is recorded.

1

Morning brief

2026-08-18

PDF

Send to Telegram

?

City mean **34.3 $\mu\text{g}/\text{m}^3$** (Satisfactory) ▲ 2.6 vs yesterday · worst near **Gharoli** 48 · 24 h outlook **40.2** (Satisfactory)

WHERE THE AIR IS ABOUT TO TURN

Sahibabad Daulat Pur 90% Very Poor · +48h · 140

TOP ACTIONS TODAY

- Construction dust** · Ndmc Charge 3 — 77% of local PM2.5, ~15,000 exposed
- Construction dust** · Malviya Nagar — 66% of local PM2.5, ~15,000 exposed
- Construction dust** · R. K. Puram — 51% of local PM2.5, ~15,000 exposed

+47 more in the worklist below

YESTERDAY'S DISPATCHES

Ndmc Charge 3: collecting post-dispatch measurements (0 d)

Ndmc Charge 3: collecting post-dispatch measurements (0 d)

Advisories: worst tier moderate in 4 zone(s), 2 language(s) · templated from stored rows, no language model

2

Enforcement Worklist

?

All

Construction

Industrial

Waste

Search recommendations...

acting as

officer name

Construction dust · Ndmc Charge 3 priority 16 · rubric 9/10

Construction dust contributes approximately 77.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #14 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

77% contribution · 15,000 exposed +11 similar sites here

The morning brief — what changed overnight and what to do about it

1

Morning brief

2026-08-18

PDF

Send to Telegram

?

City mean **34.3 $\mu\text{g}/\text{m}^3$** (Satisfactory) ▲ 2.6 vs yesterday · worst near **Gharoli** 48 · 24 h outlook **40.2** (Satisfactory)

WHERE THE AIR IS ABOUT TO TURN

Sahibabad Daulat Pur 90% Very Poor · +48h · 140

TOP ACTIONS TODAY

- Construction dust** · Ndmc Charge 3 — 77% of local PM2.5, ~15,000 exposed
- Construction dust** · Malviya Nagar — 66% of local PM2.5, ~15,000 exposed
- Construction dust** · R. K. Puram — 51% of local PM2.5, ~15,000 exposed

+47 more in the worklist below

YESTERDAY'S DISPATCHES

Ndmc Charge 3: collecting post-dispatch measurements (0 d)

Ndmc Charge 3: collecting post-dispatch measurements (0 d)

Advisories: worst tier moderate in 4 zone(s), 2 language(s) · templated from stored rows, no language model

2

Enforcement Worklist

?

All

Construction

Industrial

Waste

Search recommendations...

acting as

Construction dust · Ndmc Charge 3

priority 16 · rubric 9/10

Construction dust contributes approximately 77.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #14 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

77% contribution · 15,000 exposed +11 similar sites here

The ranked enforcement worklist

Where to send an inspector today, with the evidence to back it — from the worst places on the map to a signed notice and a tracked outcome.

- 1 Morning brief
- 2 **Ranked worklist**
- 3 Evidence & notice
- 4 Dispatch by ward
- 5 Track outcomes

acting as

Construction dust · Ndmc Charge 3

priority 16 · rubric 9/10

Construction dust contributes approximately 77.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #14 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

77% contribution · 15,000 exposed [+11 similar sites here](#)


Hide dossier

Notice PDF

Approve

Dismiss

SATELLITE EVIDENCE



Sentinel-2 patch - Satellite-detected construction site #14

construction · 90% detection confidence

REGULATORY CITATIONS (RAG)

Graded Response Action Plan (GRAP)

match 24%

hospitals: additional respiratory care wards activated. - Helpline numbers published; mobile health units deployed.

An evidence dossier: satellite patch, contribution, cited regulation

acting as

Guide walkthrough

Construction dust · Ndmc Charge 3

priority 16 · rubric 9/10

Construction dust contributes approximately 77.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #14 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

77% contribution · 15,000 exposed · +11 similar sites here

Evidence dossier

Notice PDF

Dispatch team

History

Construction dust · Malviya Nagar

priority 12 · rubric 8/10

Construction dust contributes approximately 66.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #39 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

66% contribution · 15,000 exposed · +3 similar sites here

Evidence dossier

Notice PDF

Approve

Dismiss

Construction dust · R. K. Puram

priority 9 · rubric 8/10

Construction dust contributes approximately 51.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #28 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

51% contribution · 15,000 exposed · +7 similar sites here

Evidence dossier

Notice PDF

Approve

Dismiss

Construction dust · near Madanpur Khadar

priority 6 · rubric 8/10

Approving a recommendation

acting as

Guide walkthrough

Construction dust · Ndmc Charge 3

priority 16 · rubric 9/10

Construction dust contributes approximately 77.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #14 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

77% contribution · 15,000 exposed +11 similar sites here

Evidence dossier

Notice PDF

✓ Dispatched · tracking armed

Close case

History

Field finding

violation found ▼

note (optional): what

Record & close

Construction dust · Malviya Nagar

priority 12 · rubric 8/10

Construction dust contributes approximately 66.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #39 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

66% contribution · 15,000 exposed +3 similar sites here

Evidence dossier

Notice PDF

Approve

Dismiss

Construction dust · R. K. Puram

priority 9 · rubric 8/10

Construction dust contributes approximately 51.1% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #28 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

51% contribution · 15,000 exposed +7 similar sites here

Closing the case with a finding

VayuNetra — User Guide & Project Reference · 20/170

Where to send an inspector today, with the evidence to back it — from the worst places on the map to a signed notice and a tracked outcome.

- 1 Morning brief
- 2 Ranked worklist
- 3 Evidence & notice
- 4 Dispatch by ward
- 5 Track outcomes

Construction dust · Ward PUSA priority 3 · rubric 8/10

Construction dust contributes approximately 34.3% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #20 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

34% contribution · 15,000 exposed [+6 similar sites here](#)

[Evidence dossier](#)
[Notice PDF](#)
[Approve](#)
[Dismiss](#)

4 Ward dispatch queues EMPTY ?

Nothing dispatched yet. Approve or dispatch a worklist item above and it lands in its ward's field queue here.

5 Intervention tracking ?

rec #12902 · cell 9fffff measuring

measuring — 0 days since dispatch · collecting post-dispatch measurements

rec #12901 · cell 9fffff measuring

measuring — 0 days since dispatch · collecting post-dispatch measurements

Export as NCAP action-plan evidence (PRANA-ready CSV) ↓ — each dispatched intervention with its measured effect, mapped to the NCAP spending head the city reports against. We feed the official portal, not compete with it.

Dispatch queues by ward

Where to send an inspector today, with the evidence to back it — from the worst places on the map to a signed notice and a tracked outcome.

- 1 Morning brief
- 2 Ranked worklist
- 3 Evidence & notice
- 4 Dispatch by ward
- 5 Track outcomes

Construction dust · Ward PUSA priority 3 · rubric 8/10

Construction dust contributes approximately 34.3% of PM2.5 in this cell (~15,000 residents exposed); Satellite-detected construction site #20 is the registered construction source at this location. Site inspection required: verify dust suppression norms compliance (anti-smog gun, water sprinkling, green net coverage). Regulatory basis: Graded Response Action Plan (GRAP).

34% contribution · 15,000 exposed [+6 similar sites here](#)

[Evidence dossier](#)
[Notice PDF](#)
[Approve](#)
[Dismiss](#)

4 Ward dispatch queues EMPTY ?

Nothing dispatched yet. Approve or dispatch a worklist item above and it lands in its ward's field queue here.

5 Intervention tracking ?

rec #12902 · cell 9fffff measuring

measuring — 0 days since dispatch · collecting post-dispatch measurements


rec #12901 · cell 9fffff measuring

measuring — 0 days since dispatch · collecting post-dispatch measurements

Export as NCAP action-plan evidence (PRANA-ready CSV) ↓ — each dispatched intervention with its measured effect, mapped to the NCAP spending head the city reports against. We feed the official portal, not compete with it.

Outcome tracking after dispatch

2.5 Forecast



The 72-hour outlook with its 80% band and the persistence baseline

What the air will do in the next 72 hours, how much to trust that, who it will affect — and what happened before.

- 1 72-hour outlook
- 2 How good is it, really
- 3 Real orders, in hindsight
- 4 Who is in the forecast
- 5 The past




2 Forecast validation MEASURED multi-season live 90d ?

Strict temporal split: trained on 2025-02-17 → 2025-11-01, tested on 2025-11-01 → 2026-08-15 (39 station cells, 234,491 test hours @24h). The served forecast (LightGBM median blended with persistence, weight from the calibration tail) vs persistence, weekly seasonal-naive and hour-of-day climatology on one shared support mask.

horizon	vs persistence	vs seasonal	winter	>120 hours	80% PI
+24h	+9%	+26%	+7%	+7%	+78%
+48h	+13%	+20%	+11%	+11%	+78%
+72h	+12%	+16%	+11%	+10%	+77%

The benchmark behind the forecast, shown in the product

2.6 Advisories — the citizen loop



Ward-level advisories, tiered by risk and vulnerability

Tell residents what to do, in their language, on the channel they actually use — and let them tell you where it smells like smoke.

- 1 Advisories by ward
- 2 Send it
- 3 Clean-air routes
- 4 Citizen reports

1 Citizen Advisory

Hindi ▼ ?

App Telegram IVR call Big screen

how the same advisory reaches citizens on each channel

zone-00b3

moderate

Delhi zone-00b3: अगले 24 घंटों में हवा मध्यम रहने का अनुमान है. संवेदनशील लोग (बच्चे, बुजुर्ग, दमा रोगी) लंबी बाहरी मेहनत कम करें.

zone-02c1

moderate

Delhi zone-02c1: अगले 24 घंटों में हवा मध्यम रहने का अनुमान है. संवेदनशील लोग (बच्चे, बुजुर्ग, दमा रोगी) लंबी बाहरी मेहनत कम करें.

zone-02c3

moderate

Delhi zone-02c3: अगले 24 घंटों में हवा मध्यम रहने का अनुमान है. संवेदनशील लोग (बच्चे, बुजुर्ग, दमा रोगी) लंबी बाहरी मेहनत कम करें.

zone-0e3d

moderate

Delhi zone-0e3d: अगले 24 घंटों में हवा मध्यम रहने का अनुमान है. संवेदनशील लोग (बच्चे, बुजुर्ग, दमा रोगी) लंबी बाहरी मेहनत कम करें.

2 Send it TELEGRAM · IVR · DISPLAY

?

 Broadcast latest alert (Telegram + IVR)

The same advisory, rendered per channel above; a WhatsApp-ready card for any place is one click away in its cell story.

3 Cleanest air right now LOWEST ~1 KM CELLS

?

48

Good

zone-8dd5

48

Good

zone-8dd7

The same advisory in Hindi

Tell residents what to do, in their language, on the channel they actually use — and let them tell you where it smells like smoke.

- 1 Advisories by ward
- 2 Send it
- 3 Clean-air routes
- 4 Citizen reports

1 Citizen Advisory

English ?

App Telegram IVR call Big screen

how the same advisory reaches citizens on each channel

VayuNetra Bot

⚠ zone-00b3 · moderate

Delhi zone-00b3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

⚠ zone-02c1 · moderate

Delhi zone-02c1: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

⚠ zone-02c3 · moderate

Delhi zone-02c3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.



Live two-way bot: **/start** → pick a city → auto-receive advisories. Scan to subscribe on your own phone.

2 Send it TELEGRAM · IVR · DISPLAY

?

🔔 Broadcast latest alert (Telegram + IVR)

The same advisory, rendered per channel above; a WhatsApp-ready card for any place is one click away in its cell story.

3 Cleanest air right now LOWEST ~1 KM CELLS

?

48

Good

48

Good

Rendered for Telegram

Tell residents what to do, in their language, on the channel they actually use — and let them tell you where it smells like smoke.

1 Advisories by ward
2 Send it
3 Clean-air routes
4 Citizen reports

1 Citizen Advisory
English ?

App
Telegram
IVR call
Big screen

how the same advisory reaches citizens on each channel

zone-00b3 moderate

Delhi zone-00b3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-02c1 moderate

Delhi zone-02c1: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-02c3 moderate

Delhi zone-02c3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-0e3d moderate

Delhi zone-0e3d: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

2 Send it TELEGRAM · IVR · DISPLAY ?

Broadcast latest alert (Telegram + IVR)

The same advisory, rendered per channel above; a WhatsApp-ready card for any place is one click away in its cell story.

3 Cleanest air right now LOWEST ~1 KM CELLS ?

48 Good zone-8dd5

48 Good zone-8dd7

Broadcasting: choose the ward and the language before sending

Tell residents what to do, in their language, on the channel they actually use — and let them tell you where it smells like smoke.

1 Advisories by ward 2 Send it 3 Clean-air routes 4 Citizen reports

1 Citizen Advisory English ?

App Telegram IVR call Big screen

how the same advisory reaches citizens on each channel

zone-00b3 moderate

Delhi zone-00b3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-02c1 moderate

Delhi zone-02c1: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-02c3 moderate

Delhi zone-02c3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-0e3d moderate

Delhi zone-0e3d: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

2 Send it TELEGRAM · IVR · DISPLAY ?

 Broadcast latest alert (Telegram + IVR)

The same advisory, rendered per channel above; a WhatsApp-ready card for any place is one click away in its cell story.

3 Cleanest air right now LOWEST ~1 KM CELLS ?

48 Good

zone-8dd5

48 Good

zone-8dd7

Clean-air routes — the flip side of the blame map

Tell residents what to do, in their language, on the channel they actually use — and let them tell you where it smells like smoke.

- 1 Advisories by ward
- 2 Send it
- 3 Clean-air routes
- 4 Citizen reports

1 Citizen Advisory

English ?

App Telegram IVR call Big screen

how the same advisory reaches citizens on each channel

zone-00b3

moderate

Delhi zone-00b3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-02c1

moderate

Delhi zone-02c1: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-02c3

moderate

Delhi zone-02c3: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

zone-0e3d

moderate

Delhi zone-0e3d: air is forecast moderate in +24h. Sensitive people (children, elderly, asthma) should limit long outdoor exertion.

2 Send it TELEGRAM · IVR · DISPLAY

?

 Broadcast latest alert (Telegram + IVR)

The same advisory, rendered per channel above; a WhatsApp-ready card for any place is one click away in its cell story.

3 Cleanest air right now LOWEST ~1 KM CELLS

?

48

Good

zone-8dd5


48

Good


zone-8dd7

Citizen reports coming back in

2.7 Cities, simulator, impact and pipeline



The ten-city scoreboard, ranked by the same figure the badge derives from



The what-if simulator

Before spending money: run an intervention on the model and read the $\Delta\text{PM}_{2.5}$, the people protected and the health ₹ — with citations.

1 Choose an intervention
2 Run & read the result
3 Best bundle for a budget

1 Choose an intervention WHAT-IF
?

Counterfactual over attribution × forecast, with cited health & carbon impact.

Intervention

Crop-residue / waste burn ban

Horizon +24h +48h +72h

Run simulation

2 Result +24H · 23 CELLS
?

Near-zero effect — honest result: this source contributes ~0% of the city's current $\text{PM}_{2.5}$ mix (see the blame map), so banning it changes little today. Try an intervention matching the dominant source, e.g. a traffic or construction measure.

avg ΔAQI 0	best cell ΔAQI 0
cells affected 23	confidence 40%

PEOPLE PROTECTED 0	HEALTH COST AVOIDED ₹0
DEATHS AVERTED 0.00 premature, over the window	CO ₂ E CO-BENEFIT 1,242 t vs cited emission factor

[sources \(5\) — every figure is cited](#)

3 Best bundle for a budget OPTIMIZER
?

Ranked actions under the inspector-hour budget.

Inspector-hour budget

20h

Rank packages

A simulated intervention, with its cited assumptions

Before spending money: run an intervention on the model and read the $\Delta\text{PM}_{2.5}$, the people protected and the health ₹ — with citations.

- 1 Choose an intervention
- 2 Run & read the result
- 3 Best bundle for a budget

1 Choose an intervention WHAT-IF

Counterfactual over attribution $\times$ forecast, with cited health & carbon impact.

Intervention

Crop-residue / waste burn ban

Horizon

+24h

+48h

+72h

Run simulation

2 Result +24H · 23 CELLS

Near-zero effect — honest result: this source contributes ~0% of the city's current $\text{PM}_{2.5}$ mix (see the blame map), so banning it changes little today. Try an intervention matching the dominant source, e.g. a traffic or construction measure.

avg ΔAQI 0

cells affected 23

best cell ΔAQI 0

confidence 40%

PEOPLE PROTECTED

0

HEALTH COST AVOIDED

₹0

DEATHS AVERTED

0.00

premature, over the window

$\text{CO}_2\text{E CO-BENEFIT}$

1,242 t

vs cited emission factor

[sources \(5\) — every figure is cited](#)

3 Best bundle for a budget OPTIMIZER

Ranked actions under the inspector-hour budget.


20h

Inspector-hour budget



Rank packages

Allocating a fixed inspector-hour budget



Health and economic burden

The funding case: this city's annual health burden, what the NCAP target would avert, where the money should go — and whether the system is fair.

1 The funding case

2 Where funds should go

3 Is it fair?

1 City ROI — the funding case

ATTRIBUTABLE DEATHS / YR

73,395

ANNUAL HEALTH BURDEN

₹3.67 lakh cr

AVERTABLE (~30% NCAP)

18,219

AVERTABLE ₹ / YR

₹91,095 cr



Annual mean $92 \mu\text{g}/\text{m}^3$ vs WHO $5 \mu\text{g}/\text{m}^3$ · pop 2,06,00,000

At $92 \mu\text{g}/\text{m}^3$ annual PM2.5, ~73,395 premature deaths/yr are attributable to pollution above the WHO guideline (₹366,977 cr/yr). A 30% NCAP-style cut would avert ~18,219 deaths/yr (₹91,095 cr/yr) — the funding case.

[sources \(5\) — every figure is cited](#)

2 Where the funds should go

ATtribution-WEIGHTED

34% of PM2.5 is construction dust → prioritise **C&D dust control**

21% of PM2.5 is traffic → prioritise **vehicular emission control**

20% of PM2.5 is industrial → prioritise **industrial emission control**

NCAP cities spent 67% of funds on road dust and <1% on industry (CREA, 2026) — because allocation wasn't attribution-led. This is the per-city fix.

3 Fairness audit

N=390 LIVE RECS


What drives enforcement priority — measured on every live recommendation across all cities (2026-07-19):

Pollution contribution

$r = 0.95$

the dominant driver — the system targets pollution, by design

Fairness: who is exposed, and who is protected



The pipeline view — every agent run, reconstructed from its trace

The six AI agents, run live in front of you — every step traced, every number sourced.

1 Run the agents 2 Read the trace

1 Agent Pipeline

1.4s end-to-end ?

last multi-agent run · detect → decide → issue

No trace yet — press "Run agents live" to watch the pipeline think.


► Run agents live




© OpenStreetMap © CARTO | MapLibre ⓘ

A single run, node by node, with latencies


2.8 Dark theme and mobile




The console in dark theme



The city scoreboard in dark theme




The public city page in dark theme



The public page on a phone



The console on a phone





The compact navigation sheet

4 • Architecture and data flow, end to end

The path from a pollutant reading to an officer's enforcement decision follows a complete pipeline: data enters through city-agnostic connectors, normalizes to a canonical schema, flows through a multi-agent graph that adds intelligence at each step, and surfaces results via a REST API to web and mobile clients. This chapter traces that end-to-end flow and explains the spatial encoding, ML feature construction, and orchestration that power the system.

The Data Flow at a Glance





Ingestion Layer: Connectors

Each data source has a connector that maps raw payloads to the canonical measurement schema. No per-city code: a city is defined by a YAML config file specifying its bbox, CAAQMS station IDs, wards GeoJSON, and supported languages.

Connectors and sources (connectors/*.py):

Connector	Source	Mapping	Cadence	Fallback
cpcb.py	CPCB CAAQMS via data.gov.in	PM2.5, PM10, NO ₂ , SO ₂ , CO, O ₃ per station → H3 cell	~1 hour (snapshot via API)	OpenAQ historical backfill
openaq.py	OpenAQ free API	Same pollutants, hourly series	1 hour	Complements CPCB if gov feed is down
openmeteo.py	Open-Meteo (free, no key)	Wind u/v, BLH, temp, RH, precip, AQ forecast	Hourly forecast + nowcast	Built-in fallback (no rate limits)

Connector	Source	Mapping	Cadence	Fallback
earth_engine.py	Google Earth Engine (Sentinel-5P, MODIS/VIIRS)	NO ₂ column, AOD, active-fire counts, Sentinel-2 tiles	Daily nightly precompute (cron)	Caches in Storage; features pre-aggregated
osm_sources.py	OpenStreetMap	Industrial zones, major roads (traffic corridors)	Static (updated monthly if needed)	Vendor-maintained
mobility.py	OSM road network + GTFS + time-of-day model	Proxy traffic density per corridor and hour	Time-of-day static model, no real-time feeds	Known limitation (ARCHITECTURE.md §24)
population.py	WorldPop + Census OSM	Population density, vulnerability (hospitals, schools)	Static	Used for exposure weighting

Canonical measurement schema (what lands in the `measurements` table):

```
{
  "city_id": "delhi",
  "h3_cell": "883da1a3a1fffff",
  "ts": "2026-08-19T15:00:00+00:00",
  "variable": "pm25",
  "value": 145.3,
  "unit": "µg/m³",
  "source": "caaqms",
  "confidence": 0.95,
  "ingested_at": "2026-08-19T15:15:22+00:00"
}
```

The connectors handle schema migration transparently: data.gov.in occasionally changes field names; the code tries both. Missing data is marked with `confidence < 1.0`, allowing downstream ML to weight accordingly (CPCB.py:64–90).

Spatial Layer: H3 Hexagonal Grid

Every reading is mapped to a **Hexagonal Hierarchical Spatial Index (H3) cell** at resolution 8, the system's universal spatial key. Resolution 8 yields cells of approximately **0.74 km²** (edge length ~0.46 km), matching the brief's "~1 km grid" requirement (core/spatial/h3_utils.py:12).

This choice is deliberate:

- 1. **Uniform:** One math for every city, regardless of ward boundaries or station density.
- 2. **Hierarchical:** Cells nest into coarser resolutions (res 6 ≈ 42 km) so results can be aggregated or presented at ward level without recomputing.
- 3. **Efficient for spatial joins:** Native to Deck.gl's `H3HexagonLayer`, rendering millions of cells interactively.
- 4. **Trivial for forecasting:** Space-time features include the 6 spatial neighbors; lagged values from upwind cells encode transport.

Conversion (`core/spatial/h3_utils.py`):

```
latlng_to_cell(lat=28.61, lng=77.21, res=8) # → "883da1a3a1fffff"
cell_to_latlng("883da1a3a1fffff")          # → (28.611, 77.206)
k_ring(cell, k=1)                           # → 7 neighbors for advection feature
```

Ward mapping is computed once per city at startup:

```
ward_to_cells(ward_geojson, res=8) # → {ward_id: [list of H3 cells]}
```

This lets officers see results at the familiar ward level ("Ward 12: 68% construction dust") while the pipeline computes on uniform cells, avoiding polygon aggregation artifacts.

Feature Engineering: From Raw Measurements to ML Input

The **forecast model** (Agent 2) is the pivot point where raw measurements become a supervised learning problem. The feature table is built in `ml/forecast/features.py`:

Input: long-format measurements (one row per city/cell/timestamp/variable).

Output: wide format with engineered features per (cell, timestamp):

```
h3_cell | ts | pm25 | pm10 | temp | wind_u | ventilation | advected_pm25 | pm25_lag24 | y_24h
-----|----|-----|-----|-----|-----|-----|-----|-----|-----
883da1a3a1fffff | 2026-08-18 15:00 UTC | 145.3 | 210 | 28.5 | 2.1 | 1680 | 142.1 | 151.2 | 152.0
```

Feature categories (`ml/forecast/features.py`):

1. **Pollutants** (PM2.5, PM10, NO₂, SO₂, CO, O₃): the measured mixture signature.
2. **Meteorology**: wind_u/v, BLH, temp, RH, precip — broadcast from city-scale observations, merged by (city_id, ts).
3. **Physics-informed**: - **Ventilation** = wind_speed × BLH: directly quantifies pollution clearance. - **Advected PM2.5**: traces wind backward 6 hours, encodes transport.
4. **Calendar**: hour-of-day, day-of-week, flags for stubble season and Diwali window (PRD §12.2).
5. **Lags**: PM2.5 at t-1h and t-24h per cell (autoregressive signal).
6. **Data quality**: Values outside instrument range (PM2.5 > 1500 µg/m³) are dropped before pivoting.

Target construction:

```
y_24h = df.groupby("h3_cell")["pm25"].shift(-24) # the value at t+24h
```

This ensures strict temporal order (no future leakage) and per-cell isolation (no information bleeding across space).

Feature Table to Forecast: The ML Pipeline

MVP model: LightGBM with quantile loss (alpha=0.1, 0.5, 0.9) trained on historical (city, cell, ts) tuples:

- **Horizons**: 24, 48, 72 hours.
- **Baselines**: Persistence ($y(t) = y(t-24)$) and climatology stored alongside every forecast for honest comparison.
- **Intervals**: Quantile regression outputs prediction intervals (pi_low, pi_high).
- **Serving**: Batch inference via GitHub Actions cron 6-hourly; results written to `forecasts` table.

The forecast is then used by Agent 3 (spike detection) and Agent 4 (advisory generation), and validated by the evaluation harness with a skill score target ≥ 0.25 (ARCHITECTURE.md §13).

Agent Graph: Orchestration and Decision Flow

The multi-agent system is a **LangGraph** state machine (`agents/graph.py`) coordinating five agents. The graph executes nodes via a typed shared state:

```
START
↓
```

```

orchestrator (Agent 0)
  • Reads latest PM2.5 from DB or fixture
  • Identifies spiking cells (PM2.5 > 120 µg/m³)
  • Stamps trace entry with duration_ms
  ↓
attribution (Agent 1)
  • Loads source shares for spiking cells
  • Per cell: {construction: 0.68, traffic: 0.22, confidence: 0.83}
  ↓
forecast (Agent 2)
  • Loads 24/48/72h forecasts
  • Detects if any forecast > 300
  ↓
spike_gate (conditional edge)
  • if (focus_cells OR forecast_spike): → enforcement
  • else: → advisory
  ↓
[if spiking: enforcement]
enforcement (Agent 3)
  • Calls agents.enforcement.run_enforcement()
  • Scores sources: priority = contribution × pop_exposed × actionability
  • RAG-retrieves CPCB/GRAP rules; builds dossier
  ↓
advisory (Agent 4)
  • Generates ward-level advisories
  • Localizes to language (hi/en/kn/mr)
  ↓
END
state["latency_ms"] = duration from first trace entry to last

```

Each node stamps a trace entry with node name, timestamp, and duration. The state accumulates citations from the RAG retrieval. End-to-end latency is computed as the difference between the first and last trace timestamps (agents/graph.py:88–97).

DEMO_MODE (default: `true`; agents/graph.py:32, api/main.py:36):

When enabled, each agent node loads from `demo/fixtures/<name>.json` instead of the live database. The fixtures are real data snapshots: a frozen, deterministic city state for all 10 cities. This means the entire demo runs offline, eliminating any dependency on CAAQMS (which is flaky) or Earth Engine availability. **DEMO_MODE is transparent to the frontend:** the API response shape is identical whether data came from fixtures or live DB.

Database Schema

The data platform is a single Supabase project (Postgres 15 + PostGIS + pgvector). Key tables:

```

measurements (
  city_id, h3_cell, ts, variable, value, unit, source, confidence
  -- indexes: (city_id, variable, ts), (h3_cell, ts)
);

attribution (
  city_id, h3_cell, ts_window,
  source_category, share, confidence, method_version, evidence
  -- Agent 1 output: source-apportionment results
);

forecasts (
  city_id, h3_cell, issued_at, horizon_h, value, pi_low, pi_high,
  persistence_value, model_version, p_over_120, p_over_250

```

```

-- Agent 2 output: 24/48/72h forecasts + baselines
);

enforcement_recs (
  city_id, h3_cell, ts, source_id, priority_score, contribution,
  pop_exposed, rationale, rag_citations, rubric_score, status
  -- Agent 3 output: officer worklist
);

advisories (
  city_id, ward_id, h3_cell, issued_at, horizon_h, risk_tier,
  audience_segment, language, channel, message
  -- Agent 4 output: citizen alerts
);

kb_chunks (
  doc_id, title, source_url, modality, chunk_text, image_ref,
  embedding vector(384), metadata
  -- RAG: NCAP/GRAP/CPCB regulations + (E6) Sentinel-2 patches
  -- index: ivfflat on embedding for cosine similarity
);

action_traces (
  city_id, signal_ts, attribution_ts, forecast_ts, enforcement_ts,
  advisory_ts, total_latency_ms, trace jsonb
  -- North-Star metric: signal → action latency
);

```

All timestamps are in UTC. Indices are chosen for access patterns: (city_id, variable, ts desc) for measurements; (h3_cell, ts desc) for spatial queries.

API: The Contract with Clients

Every response conforms to a standard envelope (api/main.py):

```

{
  "success": true,
  "data": [...],
  "error": null,
  "meta": {"ms": 123}
}

```

Key endpoints (fully implemented):

Method	Endpoint	Cache	What it returns
GET	/aqi/current?city=delhi	45s	Latest PM2.5 + pollutants per H3 cell; CPCB and EPA AQI indices
GET	/attribution?city=delhi&cell=<id>	60s	Source split + confidence + evidence + method_version
GET	/forecast?city=delhi&horizon=24	60s	AQI forecast + intervals + persistence baseline
GET	/enforcement?city=delhi	NOT cached	Officer worklist (live reads)
GET	/enforcement/{rec_id}/dossier	NOT cached	Evidence packet with RAG citations + satellite patch
GET	/advisory?city=delhi&lang=hi	NOT cached	Ward alerts in specified language
POST	/agent/query	NOT cached	Full multi-agent output via orchestrator graph

Caching headers (api/main.py:92–126) are set per endpoint via middleware. Officer-facing endpoints are explicitly uncached (`no-store`) so feedback is live. Public endpoints cache aggressively with `stale-while-revalidate` for instant repeats.

Concrete Request Path: "A User Opens the Console for Delhi and Sees the Map"

Step 1: Load city list

```
GET /cities → api/main.py:243
├ DEMO_MODE=true: fixture("cities")
└ DEMO_MODE=false: db.table("cities").select(...).execute()
Response: [{"city_id": "delhi", ...}, {"city_id": "bengaluru", ...}, ...]
```

Step 2: Fetch current AQI for Delhi

```
GET /aqi/current?city=delhi → api/main.py:256
├ DEMO_MODE=true: fixture_rows("aqi_current", "delhi")
└ DEMO_MODE=false: Query DB for measurements in last 24h
Response: {
  "h3_cell": "883da1a3a1fffff",
  "pm25": 145.3,
  "ts": "2026-08-19T14:45:00+00:00",
  "aqi_in": 218, # CPCB National AQI
  "aqi_us": 201 # US EPA AQI
}
```

Step 3: Frontend renders map using Deck.gl H3HexagonLayer

Step 4: User clicks a cell → fetch attribution

```
GET /attribution?city=delhi&cell=883da1a3a1fffff → api/main.py:587
Response: {
  "shares": {
    "construction_dust": 0.68,
    "traffic": 0.22,
    "transported": 0.10
  },
  "confidence": 0.83,
  "method_version": "hybrid-gbm-shap-v2"
}
```

Step 5: Fetch 72-hour forecast

```
GET /forecast?city=delhi&horizon=24 → api/main.py:638
Response: {
  "value": 152,
  "persistence_value": 145,
  "pi_low": 125,
  "pi_high": 185,
  "p_over_120": 0.91
}
```

Step 6: Officer views enforcement recommendations

```
GET /enforcement?city=delhi → api/main.py:670
Response: [{
  "id": 1001,
  "priority_score": 0.87,
```

```
"contribution": 0.68,
"pop_exposed": 18000,
"rationale": "Active construction site (CV-detected), high dust attribution, 18k residents nearby.",
"rag_citations": [{"title": "GRAP Stage 2: Construction Sites", "source_url": "..."}],
"rubric_score": {"attribution_match": 2, "actionability": 2, "exposure": 2, ...}
}]
```

The console displays a worklist ranked by `priority_score`, linked to satellite patches and regulatory citations.

Latency Measurement: The North-Star Metric

The latency trace is stamped at each agent node (agents/graph.py:79–85):

```
orchestrator → ts: 15:00:00.123
attribution  → ts: 15:00:00.145 (duration: 22ms)
forecast     → ts: 15:00:00.312 (duration: 167ms)
enforcement  → ts: 15:00:01.856 (duration: 1544ms – RAG retrieval + scoring)
advisory     → ts: 15:00:02.234 (duration: 378ms)

total_latency_ms = 2234 - 123 = 2111 ms ≈ 2.1 seconds
```

This is **pipeline latency**: how long it takes to compute a recommendation. It does not include citizen reaction time, officer dispatch time, or inspector travel time. The brief says "demonstrate reduction in response time from signal to intervention"; VayuNetra's interpretation is precise: we measure the system's response (typical ~2 seconds, target < 5 minutes), not the municipality's administrative response. This is honest given that many cities have no documented process today (CAG 2024 audit).

DEMO_MODE: The Resilient Offline Demo

DEMO_MODE is controlled by an environment variable (default: `true`). When enabled:

- All DB queries are replaced with JSON fixture reads from `demo/fixtures/` (api/main.py:167–172).
- Fixtures are real data snapshots: 10 cities' current AQI, attribution, forecasts, enforcement, advisories.
- The API behavior is identical; clients cannot tell the difference.
- The demo is **reproducible**: 100 runs yield the same map, latency, and recommendations.
- It is **resilient**: CPCB down? Earth Engine throttled? Fixtures are local.

The key fixtures that will be populated post-model-training are `attribution.json`, `forecast.json`, `enforcement.json`, and `advisory.json`. Until then, they are empty, and endpoints return `[]` gracefully. Once models are trained, a snapshot command exports live data to fixtures, and every demo replay uses the same data.

Summary

Data flows from external sources (CPCB, OpenAQ, Earth Engine, Open-Meteo) → normalized measurements → spatial H3 grid (res 8, ~0.74 km²) → ML feature table (pollutants, meteorology, physics-informed features, calendar) → multi-agent graph (orchestrator → attribution → forecast → spike gate → enforcement/advisory) → database (Postgres + PostGIS + pgvector) → REST API (standard envelope, cached reads, role-gated writes) → console/PWA. The entire pipeline is traced for latency per node; DEMO_MODE swaps live DB for local JSON fixtures, enabling a reproducible, offline demo. End-to-end request for a Delhi map takes 2–5 seconds of backend compute and returns a list of H3 hexagons with source attribution, forecast intervals, and ranked enforcement recommendations.

5 · The technology stack, and why each piece is there

VayuNetra's architecture is built from five distinct layers, each chosen for specific operational constraints: the need to run on free compute, produce explainable outputs, support real-time decisioning on sparse sensor data, and scale across multiple Indian cities without proprietary dependencies.

Backend Framework and API

The API runs **FastAPI with Uvicorn** (api/main.py:1-100). FastAPI provides automatic OpenAPI documentation, built-in Pydantic validation, and async/await support for real-time WebSocket connections (line 27 imports WebSocket). Every response follows a standard envelope (core/schemas/canonical.py:83-100): `{success: bool, data?: T, error?: ApiError, meta?: Meta}`. This contract is enforced at the Pydantic model level, making invalid responses impossible at parse time.

Middleware (api/main.py:69-83) includes: - **CORSMiddleware** with regex-based localhost allowlist. Originally pinned to ports 5173 and 4173, this was changed to accept any localhost port because pinning caused silent failures where the browser would block responses and the frontend would fall back to bundled fixtures, making the app appear to work when the API was actually down (line 72-75 comments). - **GZipMiddleware** with 1 KB floor. The `/coverage` endpoint alone is 278 KB uncompressed; even on a free-tier host, transfer time dominates latency. Compression saves roughly nine-tenths of payload size.

HTTP caching (api/main.py:92-100) is an allowlist: `/cities` (300s, city configs change only on deploy), `/metrics/*` (300s, published artifacts), `/history/*` (120s, closed days). Deliberately absent: `/enforcement`, `/brief`, `/agent/*` — officers' own actions must always read live. The browser's HTTP cache is uncontrollable from JavaScript (can never be cleared programmatically), so the frontend maintains its own 10-second in-memory cache for read-after-write consistency.

Datastore: Supabase PostgreSQL with PostGIS and pgvector

The canonical data model (core/schemas/canonical.py:18-68) defines 17 measurement variables: air quality (PM2.5, PM10, NO2, SO2, CO, O3), satellite-derived (AOD, NO2 column), source signals (Fire detections, traffic), meteorology (wind components, temperature, humidity, precipitation, boundary layer height), and derived (population density). Every measurement is pinned to an H3 cell (resolution 8, ~1 km resolution) and a timestamp.

Supabase tables (core/supa.py, inferred from the schema) store measurements, attribution, forecasts, enforcement recommendations, advisories, and action traces. PostGIS enables geospatial joins: associating emission sources to wards, buffering industrial sites for population exposure, retrieving admin boundaries. pgvector stores RAG embeddings for regulation retrieval.

Database access uses **SQLAlchemy + psycopg[binary]** (requirements.txt:24-25). The Supabase client (core/supa.py:1-50) wraps the REST API rather than using direct Postgres drivers; this isolation allows DEMO_MODE to serve fixtures without a live database, critical for local development and CI pipelines.

Geospatial Indexing: H3

H3 (requirements.txt:13; web/package.json:20) is Uber's hexagonal hierarchical spatial index. VayuNetra uses **resolution 8** cells (~1 km across), providing: - Uniform area (unlike square grids that distort near poles, though not relevant in India) - Efficient neighbor queries (ring buffers for exposure zones) - Discrete binning that matches the forecast model's training granularity

The Python library powers connectors (geospatial binning of station observations), ML preprocessing, and SQL queries. The frontend uses h3-js to convert H3 cell IDs to lat/lng for map rendering (web/src/BlameMap.tsx:7).

Machine Learning: Forecast and Attribution

VayuNetra employs **LightGBM for quantile regression** (ml/forecast/train.py:26-38). The forecast model predicts PM2.5 at three quantiles ($\alpha=0.1, 0.5, 0.9$) for 24, 48, and 72-hour horizons, producing a prediction interval. LightGBM was chosen over neural networks because: - **Explainability**: SHAP provides feature contributions per prediction, required for officers to trust enforcement recommendations. - **Free compute**: Trains on Google Colab T4 GPUs in minutes. - **Asymmetric loss**: Quantile regression at multiple alphas handles the asymmetric cost of forecast errors (missing a spike is worse than predicting a false spike).

Attribution (ml/attribution/shap_attribution.py:1-54) is a hybrid approach: LightGBM trained on source-marker features (NO₂, CO → traffic; SO₂ → industrial; PM₁₀/PM_{2.5} ratio → construction; fire detections → biomass; advected PM_{2.5} → transported) produces SHAP contributions, normalized into source shares and blended with chemical-signature priors (0.6 ML, 0.4 prior). This blend dampens circularity and maintains minimum samples threshold; below 400 observations, the system falls back to signature priors.

Both models run on **CPU only** (requirements.txt comment lines 1-2, 27-28): scikit-learn, numpy, pandas. There is no CUDA or heavy PyTorch dependency in the runtime requirements.

Agent Orchestration: LangGraph

The multi-agent pipeline (agents/graph.py:1-14) is orchestrated by **LangGraph** (requirements.txt:33):

```
START → orchestrator → attribution → forecast → spike_gate
                                     └─ enforcement (if spiking)
                                     └─ advisory (always)
enforcement → advisory → END
```

Each node stamps a timestamp into state["trace"], enabling end-to-end latency measurement (signal to action). The target is <5 minutes.

LangGraph provides typed state (GraphState: TypedDict, agents/graph.py:47-66), conditional routing via spike_gate (only runs enforcement when focus_cells exist), and composable pure functions that serialize and run in message queues or inline.

In DEMO_MODE (the default), all nodes read fixtures (demo/fixtures/*.json); this allows the frontend, tests, and CI to run without a live backend.

LLM usage: google-generativeai (requirements.txt:35) is listed but **NOT used at runtime**. The only LLM integration is an optional, operator-gated fluency polish script (scripts/llm_polish_advisories.py:1-20) that uses Gemini 2.0 Flash to rephrase advisories, with strict fact validation: zone IDs, time horizons, "N95", and all digits must survive verbatim; any deviation keeps the original template. This script is not wired into any cron—its use is explicit operator choice (docs/AI_METHODODOLOGY.md:23). Stage 1 advisories are deterministic templates (agents/advisory.py:1-6), preventing hallucinated medical advice.

Frontend: React, Vite, MapLibre, deck.gl

The web app (web/package.json) runs on **React 18.3.1 + TypeScript 5.6**, compiled by **Vite 5.4.0**. Vite was chosen for instant HMR and vendor chunk splitting: MapLibre/deck.gl libraries bundle together and stay cached across deploys; the landing page never pulls them.

Charting: **Recharts 3.9.2** renders time series, bar charts, and tooltips. Recharts is lightweight, composable, and works well with small datasets.

Mapping: **MapLibre GL 4.7.1** (open-source fork of Mapbox GL v1) provides raster tiles. **deck.gl 9.0** (web/package.json:15-18) layers vector overlays: - **H3HexagonLayer**: colored hexagons for PM_{2.5} AQI - **GeoJsonLayer**: ward boundaries, emission source points, industrial zones - **LineLayer**: wind-oriented Gaussian plume footprints - **ScatterplotLayer**: station locations

Both integrate via @deck.gl/mapbox, making deck.gl a shader-based layer atop MapLibre.

State management: none, in the library sense. City, section, selected cell, map mode and layer toggles are held in plain `useState` inside `web/src/App.tsx` (17 `useState` calls) and passed down as props; the URL is the shared source of truth for anything worth deep-linking. `zustand` and `@tanstack/react-query` are declared in `web/package.json:19,25` but **imported nowhere in web/src** — verify with `grep -r zustand web/src`, which returns nothing. They are leftovers, not architecture. Because nothing imports them, Vite never bundles them, so they cost nothing at runtime; they should still be removed so the manifest stops describing a system that does not exist.

Data fetching: a hand-rolled envelope-aware client in `web/src/api.ts`. It unwraps the `{success, data, error}` envelope, carries the Supabase anon key, applies a 25-second timeout for reads (raised from 12s because the slowest live endpoint runs ~9s warm on a free-tier host) and a separate 240-second budget for agent runs, and — the part that matters on demo day — falls back to bundled JSON fixtures when the backend is unreachable, emitting an `api-fallback` event so the UI can show a banner rather than rendering blank panels.

Styling: **TailwindCSS 3.4.0** with **PostCSS 8.4.0** for responsive design. The app is a **Progressive Web App** (vite-plugin-pwa:1.3.0), with offline support via service workers and precached assets up to 4 MB.

Testing: **Playwright 1.61.1** runs E2E smoke tests against bundled fixtures (.github/workflows/ci.yml:81-86).

Messaging Channels

Telegram (requirements.txt:43; channels/telegram.py:1-50): python-telegram-bot sends advisories as SMS-like messages to citizen chat groups.

Twilio IVR (requirements.txt:44; channels/ivr.py:1-45): Outbound calls place real phone calls that read advisories in natural voice. Inbound calls (webhook at /ivr/inbound) let citizens dial a number, choose a city on the keypad, and hear the latest advisory. Voice synthesis uses **AWS Polly** (English and Hindi; line 34) and **Google Cloud TTS** (six other Indian languages: Marathi, Kannada, Tamil, Telugu, Bengali, Gujarati; lines 36-42). The switch to Google was deliberate: Polly covers only two of eight languages; Google's Wavenet voices support all eight and accept SSML for pacing (essential for elderly listeners).

Hosting and Deployment

API Backend: Deployed to **Render.com** (free tier, vayunetra-c8i8.onrender.com, keepalive.yml:15). Free-tier containers sleep after 15 minutes; .github/workflows/keepalive.yml pings /health every 10 minutes to keep it warm, avoiding cold starts.

Web App: Hosted on **Vercel** (vercel.json:1-16). Vercel's edge caching and global CDN are free for hobby projects. The rewrites section proxies /api/* calls to \$VITE_API_BASE_URL , enabling localhost:5173 (dev) and vercel.app (prod) to share code.

Deployment: Google Cloud Build (cloudbuild.yaml:1-28) orchestrates CI/CD: 1. Build container image 2. Push to Google Container Registry 3. Deploy to **Google Cloud Run** (region: asia-south1, managed, unauthenticated, port 8080)

This separates CI/CD from runtime hosting: code on GitHub, builds on Cloud Build, deployment to Cloud Run with auto-rollback on failure.

Continuous Integration

GitHub Actions (.github/workflows/ci.yml:1-93) runs on every push: 1. **Python checks:** lint with ruff, pytest with 55% coverage gate, DEMO_MODE=true 2. **API smoke test:** start uvicorn, curl /health, /cities, /agent/query 3. **Web build:** cd web && npm install && npm run build 4. **E2E Playwright:** 7 critical flows against offline fixtures; failures upload report

Every commit must pass to merge to main.

Summary: Technology Choices and Constraints

Layer	Technology	Version	Why This One	Where in Repo
API Framework	FastAPI + Uvicorn	(standard)	Type-safe validation, async, built-in OpenAPI, low latency	api/main.py
Database	Supabase (PostgreSQL + PostGIS + pgvector)	(managed)	Serverless, free tier, built-in geospatial/embedding support	core/supa.py, core/schemas/canonical.py
Geospatial Index	H3 (Uber)	h3@latest, h3-js@4.5	Uniform hexagons, efficient neighbor queries	requirements.txt, web/package.json
Forecast Model	LightGBM Quantile Regression	lightgbm	Explainable via SHAP, free compute (Colab), asymmetric loss at 3 quantiles	ml/forecast/train.py
Attribution Model	GBM + SHAP + Signature Priors	lightgbm, shap	Prevents hallucination via prior blending, SHAP-transparent contributions	ml/attribution/shap_attribution.py

Layer	Technology	Version	Why This One	Where in Repo
Agent Orchestration	LangGraph	langgraph	Typed state graph, composable nodes, conditional routing, DEMO_MODE compatible	agents/graph.py
LLM (Optional)	Gemini 2.0 Flash	google-generativeai	Free API, optional offline polish only with strict fact validation; NOT runtime	scripts/llm_polish_advisories.py
Frontend Framework	React + TypeScript	react@18.3, typescript@5.6	Type safety, component reuse, large ecosystem	web/src/*.tsx
Build Tool	Vite	vite@5.4	Instant HMR, vendor chunk splitting (map libs stay cached)	web/vite.config.ts
Charting	Recharts	recharts@3.9	Lightweight, composable, responsive to state	web/src/*.tsx
Base Map	MapLibre GL	maplibre-gl@4.7	Open source, raster tile layer	web/src/BlameMap.tsx
Vector Overlays	deck.gl	@deck.gl/*@9.0	GPU-accelerated hexagon rendering, large GeoJSON support	web/src/BlameMap.tsx
State Management	React <code>useState</code> + the URL	(no library)	Deep-linkable state beats a store the app never needed	web/src/App.tsx
Data Fetching	hand-rolled fetch client	(no library)	Needed envelope unwrapping + fixture fallback, which a generic cache layer does not give	web/src/api.ts
Styling	TailwindCSS + PostCSS	tailwindcss@3.4, postcss@8.4	Rapid responsive design, small CSS bundle	web/src/*.tsx
PWA	vite-plugin-pwa	vite-plugin-pwa@1.3	Offline support, service worker precaching, installable	web/vite.config.ts
E2E Testing	Playwright	playwright@1.61	Fast, real browser, headless in CI	.github/workflows/ci.yml
Telegram Messaging	python-telegram-bot	(standard)	Official bot API, async, simple message API	channels/telegram.py
IVR (Voice)	Twilio + Polly + Google TTS	twilio	Managed infrastructure, multi-language synthesis, SSML support	channels/ivr.py
API Backend Hosting	Render.com	(free tier)	Auto-deploy from GitHub, kept warm by cron ping	.github/workflows/keepalive.yml
Web App Hosting	Vercel	(free tier)	Global CDN, auto-deploy on push, edge rewrite for API proxy	web/vercel.json
Build & Deploy	Google Cloud Build + Cloud Run	(managed)	asia-south1 region, unauthenticated endpoint, auto-rollback	cloudbuild.yaml
CI/CD	GitHub Actions	(standard)	Native to GitHub, free for public repos	.github/workflows/*.yml

Deliberately NOT Used

LLMs at runtime: The project deliberately avoids using large language models for core decision logic. Advisories (Agent 4) are deterministic templates (agents/advisory.py:6), preventing hallucinated medical guidance. Attribution uses explainable gradient boosting instead of black-box deep learning. The single LLM pathway—optional advisory fluency polish—is fact-gated, off by

default, and disclosed to users. This reflects a principle: a system recommending law enforcement must be auditable, not opaque.

Async Python for the agent pipeline: The API uses `async/await` (FastAPI is `async-first`), but agents run sequentially. This is deliberate: latency is measured end-to-end (signal to action <5 min), and concurrent execution would make timing opaque and harder to debug.

NoSQL databases: All data lives in relational Postgres. Schema enforcement prevents silent bugs: missing fields crash the model, catching errors early.

Commercial LLM APIs at scale: Gemini's free tier caps requests; a deployed system calling LLMs for every advisory would exceed quotas. The fallback-to-template design ensures the system never breaks due to API limits.

6 • The data model: every table, column and policy

VayuNetra's data contract is a Postgres schema (with PostGIS geometry and pgvector embeddings) enforced at the database layer via row-level security. Every agent and connector writes to one or more tables; every API endpoint reads from them. This chapter documents the complete schema, retention policy, and access control.

Core Data Tables

cities

Configuration-driven city onboarding. No per-city business logic lives in code; everything is parameterized from this table.

Column	Type	Purpose
city_id	text, PK	Lowercase identifier (delhi, bengaluru, mumbai)
name	text	Display name
state	text	State code (DL, KA, MH)
bbox	geometry(Polygon, 4326)	PostGIS bounding box (WGS84) for map clipping and aggregation queries
center	geometry(Point, 4326)	City center point (WGS84) for map initial zoom
languages	text[]	Localization array (e.g., {hi,en,kn}) used by Agent 4 for advisory generation
caaqms_station_ids	text[]	Reference to regulatory station IDs in the CAAQMS network
ward_geojson_ref	text	URI/reference to ward boundary GeoJSON for spatial filtering
active	boolean	Soft-delete flag for inactive cities

Written by: Admin onboarding endpoint (`POST /admin/cities`).

Read by: API cities endpoint, all agents for scoping.

RLS: Public SELECT; admin-only INSERT/UPDATE.

measurements

Universal single-source-of-truth for all readings: ground stations, satellites, weather, traffic, mobility, and static spatial data (population density stored as a variable).

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text, NOT NULL	Foreign key to cities; partitions reads by city
h3_cell	text, NOT NULL	H3 grid cell at resolution 8 (~1 km²); enables spatial aggregation and heatmaps

Column	Type	Purpose
station_id	text, NOT NULL	Source-specific identifier (empty string "" if no station). Part of the deduplication key to prevent duplicate readings from overlapping ingests
ts	timestamptz, NOT NULL	Observation timestamp; indexed for time-range queries
variable	text, NOT NULL	Pollutant or meteorological variable: pm25, pm10, no2, so2, co, o3, aod, fire, wind_u, wind_v, blh, temp, rh, precip, traffic, population
value	double precision, NOT NULL	Measurement value in the specified unit
unit	text	Unit string (µg/m³, m/s, %, etc.)
source	text	Data source identifier: caaqms, openaq, s5p (Sentinel-5P), modis, s2 (Sentinel-2), openmeteo, osm_gtfs
confidence	double precision	Measurement quality score (0–1); defaults to 1.0 (implicit machine-read trust)
ingested_at	timestamptz	Insertion timestamp for audit trail

Indexes: - `idx_measurements_city_var_ts` on (city_id, variable, ts) — filters readings by city/pollutant/time - `idx_measurements_cell_ts` on (h3_cell, ts) — maps time-scrub queries - `uq_measurements_reading` unique on (city_id, h3_cell, station_id, variable, ts, source) — enforces one row per reading; enables idempotent upserts

Data Retention: Default 180 days of live raw readings. Older monthly blocks are rolled up into `pm25_daily_rollup` and archived to Supabase Storage, then deleted. Non-PM2.5 pollutants and static variables (population) are never archived. The daily rollup ensures trend and calendar views never lose history.

Written by: Connectors (openaq, caaqms, openmeteo, Sentinel, MODIS, OSM); Agent 2.
Read by: API (trend/map views), all agents, ML pipelines.
RLS: Public SELECT.

attribution

Source-attribution model output (Agent 1). One row per city, grid cell, and time window with the estimated contribution of each pollution source category.

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text, NOT NULL	City scope
h3_cell	text	Grid cell at H3 resolution 8
ts_window	tstzrange	Time window over which attribution was computed
source_category	text, NOT NULL	traffic, construction_dust, industrial, biomass_burning, transported, other
share	double precision, NOT NULL	Contribution as a fraction (0–1); sums to 1.0 per (city, cell, window)
confidence	double precision	Model confidence in the attribution
method_version	text	Version tag of the attribution model for reproducibility
evidence	jsonb	SHAP values and other explainability data

Indexes: `idx_attribution_city_cell` on (city_id, h3_cell) — query cell-level attributions for the map.

Written by: Agent 1; runs hourly per city.
Read by: Officer dashboard (/attribution), metrics, public dashboard.

RLS: Officer+ SELECT scoped to their city; public SELECT (transparency).

forecasts

PM2.5 and AQI forecasts with uncertainty bounds and calibrated exceedance probabilities (Agent 2). One row per city, grid cell, issue time, and forecast horizon.

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text	City scope
h3_cell	text	Grid cell at H3 resolution 8
issued_at	timestampz	When the forecast was generated
horizon_h	int	Forecast lead time: 24, 48, or 72 hours
target_var	text	Variable being forecast (default 'aqi')
value	double precision	Median forecast (AQI or PM2.5 µg/m³)
pi_low , pi_high	double precision	Prediction interval (quantile regression at α=0.1 and α=0.9)
persistence_value	double precision	Persistence baseline (today's reading) for honest comparison
climatology_value	double precision	Long-term mean for the same day-of-year and hour
p_over_120 , p_over_250	double precision	Exceedance probabilities from split-conformal prediction (CPCB Very Poor/Severe thresholds)
calibration_n	integer	Number of held-out residuals used for calibration
model_version	text	Version tag of the forecast model

Indexes: - idx_forecasts_city_cell_issued on (city_id, h3_cell, issued_at) — primary query path -
idx_forecasts_city_horizon on (city_id, horizon_h, issued_at DESC) — filters by horizon

Written by: Agent 2 (LightGBM quantile regression); runs hourly.

Read by: API /forecast , advisories, officer dashboard.

RLS: Public SELECT.

emission_sources

Registry and CV-detected industrial/construction sources, with impact footprints. Used by Agent 3 to target enforcement.

Column	Type	Purpose
id	bigserial, PK	Referenced by enforcement_recs
city_id	text	City scope
geom	geometry(Geometry, 4326)	PostGIS geometry (point, polygon, buffer) for spatial joins
type	text	construction, industry, waste_burn, diesel_corridor
name	text	Human label (e.g., "Badarpur Power Plant")
registry_ref	text	External reference (CPCB license, permit ID)
source_origin	text	'registry' (pre-loaded) or 'cv_detected' (E1 computer vision)

Column	Type	Purpose
detection_confidence	double precision	For CV-detected sources, confidence (0–1)
attributes	jsonb	Additional metadata (operating hours, stack height, etc.)

Indexes: - `idx_emission_sources_city` on (city_id) - `idx_emission_sources_origin` on (city_id, source_origin) — distinguish registry vs. detected

Written by: Admin registry uploads; Agent 3 (E1); service role.

Read by: Officer console, enforcement generator, map layers.

RLS: Public SELECT; admin-only INSERT/UPDATE.

enforcement_recs

Enforcement recommendations generated by Agent 3 and tracked by officers. Transitions: proposed → approved → dispatched → closed.

Column	Type	Purpose
id	bigserial, PK	Referenced by enforcement_status_log and intervention_tracking
city_id	text, NOT NULL	City scope
h3_cell	text	Recommended cell for intervention
ts	timestamptz	Recommendation timestamp
source_id	bigint	Foreign key to emission_sources; NULL if diffuse pollution
priority_score	double precision	Prioritization rank; sorted descending
contribution	double precision	Estimated contribution to cell PM2.5
pop_exposed	bigint	Population in the recommended cell; used for equity weighting
rationale	text	Plain-text explanation for the officer
evidence	jsonb	SHAP/explainability data for debugging
rag_citations	jsonb	References from the KB (regulations, procedures)
rubric_score	jsonb	Structured compliance scoring against CPCB/GRAP rubric
status	text, NOT NULL	proposed, approved, dispatched, dismissed, closed
closed_at	timestamptz	When the officer marked it closed
closure_finding	text	violation_found, compliant, inaccessible, not_applicable
closure_note	text	Officer's free-form note at closure

Indexes: - `idx_enforcement_city_ts` on (city_id, ts) - `idx_enforcement_city_priority` on (city_id, priority_score DESC) - `idx_enforcement_city_status` on (city_id, status)

Audit Trail: Every status change is appended to `enforcement_status_log` with the actor (officer name), timestamp, and note. The log survives deletion of the recommendation.

Written by: Agent 3 (initial INSERT), officers (PATCH status/closure).

Read by: Officer console, public dashboard, intervention tracking.

RLS: Officer+ SELECT scoped to their city; public SELECT.

advisories

Citizen-facing health warnings generated by Agent 4. One row per city, ward, cell, audience segment, and language.

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text	City scope
ward_id	text	Ward or neighborhood ID for geographic targeting
h3_cell	text	Grid cell at H3 resolution 8
issued_at	timestamptz	Advisory generation timestamp
horizon_h	int	Forecast horizon (24, 48, 72 hours)
risk_tier	text	Health risk level (Green, Yellow, Orange, Red, Severe)
audience_segment	text	general, outdoor_worker, elderly, school, respiratory
language	text	hi, en, kn, mr
channel	text	pwa, telegram, ivr, display
message	text	Localized HTML/plain text advisory

Indexes: idx_advisories_city_ward on (city_id, ward_id)

Written by: Agent 4; scheduled job per ward × audience × language.

Read by: API /advisories , Telegram webhook, IVR, public displays.

RLS: Public SELECT.

citizen_reports

Citizen complaint entry point (SLA-tracked): photos and location data linked to enforcement when verified.

Column	Type	Purpose
id	bigint, generated always as identity, PK	Row identifier
city_id	text, NOT NULL	City scope
h3_cell	text	Cell snapped from lat/lng; NULL if outside service area
lat , lng	double precision, NOT NULL	Raw GPS coordinates from mobile app
category	text, NOT NULL	waste_burning, construction_dust, industrial_smoke, vehicle_smoke, other
description	text	User's free-text note
photo_url	text	URI to the uploaded image (Supabase Storage)
status	text, NOT NULL	received, verified, actioned, resolved, rejected
sla_hours	integer, NOT NULL	Service-level agreement window (default 72)
source_id	bigint	Foreign key to emission_sources once verified
created_at , updated_at	timestamptz	Audit timestamps
resolved_at	timestamptz	When the report was closed

Indexes: idx_citizen_reports_city_status on (city_id, status, created_at DESC)

Written by: Mobile app / API (public POST /report); officer updates.

Read by: Officer console, public dashboard (aggregate statistics).

RLS: Public SELECT; writes via service role.

vulnerability

Population vulnerability layer: facilities (hospitals, schools, eldercare) and population density per H3 cell, ingested from OSM and Gridded Population of the World v4.11.

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text, NOT NULL	City scope
h3_cell	text, NOT NULL	Grid cell at H3 resolution 8
zone_id	text, NOT NULL	Human-readable label (e.g., "zone-a18e")
population	bigint, NOT NULL	GPW residents; default 0
hospitals	int, NOT NULL	Count of OSM amenity=hospital or clinic
schools	int, NOT NULL	Count of OSM amenity=school, college, kindergarten
eldercare	int, NOT NULL	Count of OSM social_facility nursing homes
outdoor_sites	int, NOT NULL	Outdoor-work anchors (markets, hubs, construction)
vulnerability_index	double precision, NOT NULL	Composite score (0–1) weighting population and facilities; used by Agent 4 to adjust advisory severity
updated_at	timestamptz, NOT NULL	Last update from connector

Unique constraint: (city_id, h3_cell)

Indexes: idx_vulnerability_city_index on (city_id, vulnerability_index DESC)

Written by: Connector (connectors/vulnerability.py); service role.

Read by: API /static-layers , Agent 4 (advisory adjustment).

RLS: Public SELECT.

coverage_field

Dense 1 km PM2.5 field produced by downscaling models (Stage 2, Agent E2). One row per city and H3 cell with estimate and uncertainty.

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text	City scope
h3_cell	text	Grid cell at H3 resolution 8
pm25	double precision	Downscaled PM2.5 estimate (µg/m³)
pm25_stations	double precision	Sparse interpolation baseline (inverse-distance-weighted from CAAQMS)
uncertainty	double precision	MC-dropout standard deviation (µg/m³)
model_version	text	Version tag of the downscaling model
generated_at	timestamptz	Generation timestamp; recomputed hourly

Indexes: idx_coverage_city_cell on (city_id, h3_cell)

Written by: Batch job (GitHub Actions, ml/coverage); service role.

Read by: API /coverage , frontend map interaction.

RLS: Public SELECT.

kb_chunks

RAG knowledge base: text and multimodal image-patch embeddings for retrieval during enforcement and advisory generation (Agent 3 and Agent 4).

Column	Type	Purpose
id	bigserial, PK	Row identifier
doc_id	text	Source document identifier (e.g., "cpcb_standards_2024")
title	text	Human-readable title (e.g., "CPCB Stack Emissions Limits")
source_url	text	URI to the original document
modality	text, NOT NULL	'text' for regulatory corpus, 'image' for Sentinel-2 patches (E6)
chunk_text	text	Extracted text or OCR'd content; searched for semantic similarity
image_ref	text	Supabase Storage URI when modality='image'
embedding	vector(384)	BGE-Small embedding (384 dimensions) indexed for cosine similarity
metadata	jsonb	Additional context (source_id, facility_type, pollution_class, etc.)

Indexes: idx_kb_chunks_embedding on (embedding) using ivfflat and vector_cosine_ops; idx_kb_chunks_modality on (modality)

Embedding Dimension: Default 384 (BGE-Small via sentence-transformers). Environment variable EMBEDDING_DIM=384 in .env controls this. To switch to Gemini (768 dims): 1. Set EMBEDDING_DIM=768 in .env . 2. Migrate schema: alter table kb_chunks alter column embedding type vector(384) . 3. Re-embed via rag/ingest.py.

Written by: rag/ingest.py (batch ingestion); service role.

Read by: Agent 3, Agent 4 (RAG retrieval); officer+ via API.

RLS: Officer+ SELECT; inserts via service role only.

action_traces

Latency telemetry for the North-Star metric: end-to-end time from measurement to enforcement recommendation or advisory.

Column	Type	Purpose
id	bigserial, PK	Row identifier
city_id	text	City scope
signal_ts	timestamptz	Timestamp of the input measurement
attribution_ts	timestamptz	Agent 1 completion time
forecast_ts	timestamptz	Agent 2 completion time
enforcement_ts	timestamptz	Agent 3 completion time
advisory_ts	timestamptz	Agent 4 completion time
total_latency_ms	int	End-to-end latency in milliseconds
trace	jsonb	Structured log of each agent's runtime (memory, CPU, cache hits, DB rows read)

Indexes: idx_action_traces_city on (city_id, signal_ts)

Written by: Agent orchestration.
Read by: API /latency , performance dashboards.
RLS: Officer+ SELECT scoped to their city; public SELECT.

intervention_tracking

Before/after effect measurement for dispatched enforcement actions. Baseline PM2.5 is frozen at dispatch; effect computed on read once post-window has data.

Column	Type	Purpose
id	bigint, generated always as identity, PK	Row identifier
rec_id	bigint, NOT NULL	Recommendation ID (survives rec deletion)
city_id	text, NOT NULL	City scope
h3_cell	text, NOT NULL	Cell where the intervention occurred
dispatched_at	timestamptz, NOT NULL	When the action was dispatched
baseline_pm25	double precision	PM2.5 mean over baseline_days before dispatch
baseline_days	int, NOT NULL	Lookback window for baseline (default 7)
created_at	timestamptz, NOT NULL	Row creation timestamp

Unique constraint: (rec_id)

Written by: API endpoint (PATCH enforcement status to 'dispatched').
Read by: API /interventions endpoint.
RLS: Public SELECT.

enforcement_status_log

Immutable audit trail of every status transition and closure for an enforcement recommendation. Survives deletion of the recommendation itself.

Column	Type	Purpose
id	bigserial, PK	Row identifier
rec_id	bigint, NOT NULL	Recommendation ID (no cascade delete; survives parent deletion)
city_id	text, NOT NULL	City scope (denormalized for fast filtering)
from_status	text	Previous status (nullable for initial creation)
to_status	text, NOT NULL	New status (proposed, approved, dispatched, dismissed, closed)
actor	text	Officer name as free-text
note	text	Officer's commentary on the transition
finding	text	Closure finding (violation_found, compliant, etc.)
created_at	timestamptz, NOT NULL	Timestamp of the change

Indexes: - idx_status_log_rec on (rec_id, created_at DESC) - idx_status_log_city on (city_id, created_at DESC)

Written by: API enforcement status PATCH endpoint.
Read by: Officer console (recommendation history), audit dashboards.
RLS: Public SELECT.

pm25_daily_rollup

Long-term archive of daily PM2.5 per cell, enabling trend and calendar views to survive raw measurement retention prune.

Column	Type	Purpose
city_id	text	City scope
h3_cell	text	Grid cell at H3 resolution 8
day	date	Calendar day
pm25	double precision	Average PM2.5 for the day (µg/m³)
n	integer	Number of hourly readings; quality indicator

Primary key: (city_id, h3_cell, day)

Written by: scripts/archive_measurements.py (before deleting raw rows).

Read by: API /history/trend (via pm25_daily_trend RPC); calendar views.

RLS: Public SELECT.

advisory_subscribers

Telegram subscription registry for Agent 4 two-way messaging.

Column	Type	Purpose
chat_id	text, PK	Telegram chat ID
city_id	text	Subscribed city
language	text	Preferred language (default 'en')
active	boolean	Subscription status (soft delete)
created_at , updated_at	timestamptz	Audit timestamps

Indexes: idx_advisory_subscribers_city on (city_id, active)

Written by: Telegram webhook (/telegram/webhook); subscription/unsubscription.

Read by: Telegram broadcast job (Agent 4 delivery).

profiles

User profile mapping Supabase Auth users to roles and scoped cities. Auto-created by a trigger on auth.users.

Column	Type	Purpose
user_id	uuid, PK	Foreign key to auth.users (cascade delete)
role	text, NOT NULL	admin, officer, inspector, citizen (default 'citizen')
city_id	text	Scoped city for officer/inspector roles (NULL for citizen/admin)
created_at	timestamptz	Profile creation timestamp

Trigger: on_auth_user_created inserts a row with role='citizen' on new user signup.

Written by: Auth trigger; admin role assignment.

Read by: RLS helpers (current_role_name, current_city).

RLS: Users can SELECT/UPDATE their own profile only (role cannot be self-elevated); admin can manage all.

Stored Procedures (RPC Functions)

Every stored procedure is callable via PostgREST and exposed to the API. They aggregate data server-side to reduce transfer volume and query latency.

pm25_daily_trend(p_city text, p_days integer default 90, p_cell text default null)

Daily PM2.5 averages over a window, falling back to the archived rollout for days outside the retention window. Used by /history/trend .

Logic: 1. Query raw measurements for the last p_days . 2. Group by day; average the value; count rows. 3. For any day with no raw data (older than retention), query pm25_daily_rollup . 4. Union and sort.

Optional p_cell restricts to a single H3 cell (for drill-down); NULL returns city-wide mean.

pm25_hourly_cells(p_city text, p_hours integer default 24)

Hourly PM2.5 per cell over a trailing window. Feeds the map time-scrub feature.

Aggregates raw measurements by cell and hour; returns cells×hours rows for efficient map rendering.

city_pollutants_hourly(p_city text, p_hours integer default 24)

Hourly city-wide means for all index pollutants (pm25, pm10, no2, so2, co, o3, nh3). Feeds the city overview pollutant chips.

city_pollutants_daily(p_city text, p_days integer default 365)

Daily city-wide means for all index pollutants, with fallback to pm25_daily_rollup for PM2.5 older than retention. Used by the AQI calendar.

All four procedures are granted EXECUTE to roles anon , authenticated , service_role .

Row-Level Security (RLS) and Roles

VayuNetra enforces access control via Postgres RLS, applied at the database layer before any row is returned to the application.

Roles:

Role	Purpose	Typical User
citizen	Default for all authenticated users; read-only access to public data	End users
inspector	Field officer; inspects sources and records closures	CPCB/municipal field staff
officer	Air quality manager; approves/dispatches enforcement; reads attribution	Air quality department official
admin	System administrator; onboards cities; manages users	VayuNetra team
anon	Unauthenticated public access; read-only under RLS	Public website visitors
service_role	Server-side trusted pipeline; bypasses RLS entirely	Agents, connectors, batch jobs

Helper Functions: - current_role_name() → text: Returns the role from the caller's profile. - current_city() → text: Returns the city_id from the caller's profile.

Policy Summary:

Table	Policy	Applies To	Condition
cities	public read	SELECT	true
cities	admin insert/update	INSERT, UPDATE	current_role_name() = 'admin'
measurements	public read	SELECT	true
forecasts	public read	SELECT	true
advisories	public read	SELECT	true
coverage_field	public read	SELECT	true
pm25_daily_rollup	public read	SELECT	true

Table	Policy	Applies To	Condition
intervention_tracking	public read	SELECT	true
vulnerability	public read	SELECT	true
attribution	officer read	SELECT	current_role_name() in ('admin','officer','inspector') and (admin or scoped to current_city())
attribution	public read	SELECT	true
enforcement_recalls	officer read	SELECT	current_role_name() in ('admin','officer','inspector') and (admin or scoped to current_city())
enforcement_recalls	public read	SELECT	true
enforcement_status_log	public read	SELECT	true
emission_sources	public read	SELECT	true
emission_sources	admin write	INSERT, UPDATE	current_role_name() = 'admin'
kb_chunks	officer read	SELECT	current_role_name() in ('admin', 'officer', 'inspector')
action_traces	officer read	SELECT	current_role_name() in ('admin','officer','inspector') and (admin or scoped to current_city())
profiles	own profile read/update	SELECT, UPDATE	user_id = auth.uid() and role cannot be self-elevated
profiles	admin manage	SELECT, INSERT, UPDATE, DELETE	current_role_name() = 'admin'
citizen_reports	public read	SELECT	true

Key Design Decisions:

- 1. **Service Role Bypasses RLS:** All writes from agents, connectors, and batch jobs use SUPABASE_SERVICE_ROLE_KEY with role = 'service_role' , which is not subject to any RLS policy. This enables trusted server-side code to update all tables without policy overhead.
- 2. **Anon Key for Public Dashboard:** The deployed frontend (Vercel) authenticates with the public SUPABASE_ANON_KEY (role='anon'). Postgres RLS allows anon SELECT on public-interest data (cities, measurements, forecasts, attribution, enforcement, advisories). No POST/PATCH/DELETE allowed for anon; all writes happen server-side via the API using the service role.
- 3. **City Scoping for Officers:** Officers and inspectors see data only for their assigned city. The RLS policy checks city_id = current_city() in addition to the role check. Admins see all cities.
- 4. **Citizen Role is Limited:** Authenticated citizens have no special row-level access beyond anon — they cannot see officer-only data (enforcement_recalls, attribution, kb_chunks). The role exists for future expansion.

Service Role vs. Anon Key

Service Role (SUPABASE_SERVICE_ROLE_KEY): - Used server-side only (agents, connectors, API backend). - Bypasses all RLS policies. - Has full read/write on all tables. - Never shared with the client or frontend.

Anon Key (SUPABASE_ANON_KEY): - Shared with the frontend (embedded in the client code). - Subject to RLS policies (Postgres enforces row-level access control). - Can SELECT public-interest data (cities, measurements, forecasts, advisories, attribution,

enforcement). - Cannot INSERT/UPDATE/DELETE (all writes forbidden for anon). - Used by the Vercel dashboard when DEMO_MODE=false.

API Usage: The FastAPI backend uses service_client() (service role key) for all writes and admin operations. It does not directly expose the anon key to the frontend; instead, the frontend makes requests to the API (which then uses the service role). This keeps secrets on the server.

PostGIS Usage

Geometries are stored in WGS84 (EPSG:4326) for web-native lat/lng support.

Column	Table	Usage
bbox	cities	ST_MakeEnvelope-derived bounding box for map clipping and spatial aggregation
center	cities	ST_SetSRID(ST_MakePoint(...), 4326); map initial zoom target
geom	emission_sources	Point/polygon/buffer for spatial joins; answers "which cells are downwind" via plume model

No active spatial queries in current sprint; geometries are pre-computed and stored. Future enhancements may use ST_Distance, ST_Intersects, ST_Buffer for dynamic impact modeling.

pgvector and RAG Embeddings

The kb_chunks.embedding column is type vector(384) , indexed with ivfflat for approximate nearest-neighbor search.

Embedding Model: BGE-Small (BAAI/bge-small-en-v1.5) produces 384-dimensional vectors, normalized. Ingested via rag/ingest.py using sentence-transformers.

Dimension Configuration: Environment variable EMBEDDING_DIM=384 (default) controls the vector dimension. To switch to Gemini embeddings (768 dims): 1. Set EMBEDDING_DIM=768 in .env . 2. Migrate schema: alter table kb_chunks alter column embedding type vector(384) . 3. Re-embed via rag/ingest.py --gemini.

Retrieval: rag/retrieve.py executes a cosine-similarity query after embedding the query text:

```
select * from kb_chunks order by embedding <=> query_vector limit k
```

This enables Agent 3 (enforcement) and Agent 4 (advisory) to retrieve relevant regulations, procedures, and vulnerability patterns for SHAP and citation.

Data Retention and Archival

Live Window: 180 days of raw measurements kept in the database (configured via --keep-days in scripts/archive_measurements.py, default 180).

Archival Process: 1. For each city × complete calendar month older than the cutoff (now - 180 days): a. Rollup daily PM2.5 means into pm25_daily_rollup . b. Export raw rows as gzip CSV to Supabase Storage (bucket archive). c. Read back and verify row counts match the query. d. Delete the raw rows from measurements . 2. Non-PM2.5 variables and static variables (population) are never archived (see STATIC_VARIABLES in archive_measurements.py). 3. pm25_daily_rollup is kept forever, so trend and calendar views never lose history.

Restore: If raw readings are needed (e.g., for model retraining), they can be restored from the archive bucket using core.supabase.insert_measurements (the unique key makes this idempotent).

Practical Consequence: At any given time, the dashboard shows: - Raw hourly readings for the last 180 days (or whatever --keep-days is set to). - Daily averages for all prior time (via pm25_daily_rollup). - All forecasts (no retention limit). - All attribution, enforcement, advisories (no retention limit).

The retention window is per-variable within each city. All PM2.5 readings older than 180 days are archived, regardless of the city's data volume.

Entity-Relationship Overview

cities (city_id)

└ measurements (city_id, variable, ts, h3_cell)

└ pm25_daily_rollup (city_id, h3_cell, day)

└ forecasts (city_id, h3_cell, horizon_h)

└ attribution (city_id, h3_cell, source_category)

└ enforcement_recs (city_id, h3_cell, status)

└ enforcement_status_log (rec_id)

└ intervention_tracking (rec_id)

└ emission_sources (city_id, source_origin)

└ advisories (city_id, ward_id, h3_cell)

└ vulnerability (city_id, h3_cell)

└ coverage_field (city_id, h3_cell)

└ citizen_reports (city_id, h3_cell)

└ advisory_subscribers (city_id)

kb_chunks (document corpus, indexed by embedding for RAG)

[no city_id; global to all agents]

action_traces (city_id, signal_ts to final agent latency)

profiles (user_id → auth.users, role, city_id)

Summary: Write/Read/Retention Table

Table	Rows Written By	Read By	Retention
cities	Admin API	All systems	Forever (config)
measurements	Connectors (openaq, caaqms, weather, satellite, traffic), Agent 2 trace	All systems, ML pipelines	180d raw + rollup forever
forecasts	Agent 2 (hourly)	API, officers, advisories, public	Forever (no prune)
attribution	Agent 1 (hourly)	API, officers, public	Forever (no prune)
enforcement_recs	Agent 3 (hourly); officers (status updates)	API, officers, public	Forever; closed recs retained
enforcement_status_log	API (on status change)	API, officers, audit	Forever (survives rec deletion)
advisories	Agent 4 (per ward, per horizon, per audience)	API, public, Telegram broadcast	Forever (archive after delivery)
emission_sources	Admin API; Agent 3 (CV detections)	API, officers, enforcement, public	Forever (config + detections)
citizen_reports	Mobile/web API (public submit)	API, officers, public stats	Forever (SLA tracking)
vulnerability	Connectors (OSM, GPW) on city onboarding and refresh	API, Agent 4, static layers	Updated on refresh; row lifetime = infinity
coverage_field	ML batch job (hourly)	API, frontend	Latest snapshot; ~1 refresh/hour kept
kb_chunks	rag/ingest.py (corpus load); E6 multimodal	Agent 3 RAG, Agent 4 RAG, officers	Forever (config; re-embed on model switch)

VayuNetra — User Guide & Project Reference - 65/170

Table	Rows Written By	Read By	Retention
<code>action_traces</code>	Agent orchestration	API latency view, performance dashboards	Forever (optional archive after N days)
<code>intervention_tracking</code>	API (on dispatch status change)	API interventions view	Forever (effect measurement)
<code>pm25_daily_rollup</code>	Archive script (from measurements prune)	API trend/calendar	Forever (post-archival PM2.5 history)
<code>advisory_subscribers</code>	Telegram webhook (subscribe/unsubscribe)	Telegram broadcast job	Soft-delete (active flag)
<code>profiles</code>	Auth trigger (on user signup); admin API	RLS policies, profile endpoints	Soft-delete (cascade on auth.users)

7 • Where the data comes from: connectors and scheduled jobs

VayuNetra consumes data from 13 distinct sources, each shaped by different upstream APIs, spatial grids, and operational constraints. Every connector produces canonical rows for the central `measurements` table or related domain tables (`emission_sources`, `vulnerability`), normalized to H3 resolution 8 cells (roughly 1–2 km across). The ingestion pipeline runs on two GitHub Actions schedules: hourly for real-time ground and weather data, and daily for satellite features, emission registries, and model outputs.

Architecture: canonical measurement format

Every connector writes to one of three tables with a consistent schema. The `measurements` table (connectors/openaq.py, connectors/cpcb.py, etc.) holds:

- `city_id`, `h3_cell`, `station_id`: geospatial keys
- `ts` (ISO-8601 UTC): measurement timestamp
- `variable` (pm25, pm10, no2, so2, co, o3, temp, rh, precip, wind_u, wind_v, population, traffic, etc.)
- `value`, `unit`: the actual reading and its SI/common unit
- `source`: provider name (e.g., "openaq", "caaqms", "s5p", "gpw411")
- `confidence`: a float in [0, 1] reflecting the measurement's fidelity (official government stations = 1.0, community sensors = 0.6, satellite = 1.0, OSM-derived = 0.72)

Emission sources go to `emission_sources` with `type` (construction, industry, waste_burn, diesel_corridor), `name`, `geom` (GeoJSON Point), and `attributes` (YAML dict holding h3_cell, osm_tags, confidence). Population vulnerability lives in `vulnerability` indexed by H3 cell with a computed `vulnerability_index` and facility counts.

Ground and weather: hourly sources (live)

These five connectors run every hour at minute 5 UTC (which is **10:35 IST** due to offset). Each command runs once per tracked city (currently 10: Delhi, Bengaluru, Mumbai, Hyderabad, Chennai, Kolkata, Pune, Ahmedabad, Jaipur, Lucknow).

CPCB CAAQMS via data.gov.in

Upstream: Central Pollution Control Board (CPCB), Ministry of Environment Forest & Climate Change. "Real time Air Quality Index" resource ID `3b01bcb8-0b14-4abf-b6f2-c1bfd384ba69` at <https://api.data.gov.in>.

What it provides: Current AQI and per-pollutant concentration at each certified CAAQMS station—roughly 40 per city, 500+ nationwide. Returns a single snapshot per station (not history; the last-update timestamp may be 1–6 hours old in practice).

Fetched: Every hour, Delhi only. The open endpoint is geographically unreachable from some ISPs but reachable from GitHub Actions runners, providing a redundancy path when OpenAQ lags.

Written: - Table: `measurements` - Variables: `pm25`, `pm10`, `no2`, `so2`, `co`, `o3` (as $\mu\text{g}/\text{m}^3$, or mg/m^3 for CO) - Source: "caaqms" - Confidence: 1.0 (government official) - Unit conversions: none; CPCB reports $\mu\text{g}/\text{m}^3$ for particulate matter and trace gases, mg/m^3 for CO—matched directly.

Failure handling: The endpoint is frequently flaky (502, timeouts). The connector retries up to 6 times with exponential backoff (sleep 2s, 4s, 6s, ... 10s max). If all retries fail, the run logs a warning but does not crash the workflow.

Required secret: `DATA_GOV_IN_API_KEY` (free registration at <https://data.gov.in>). The workflow fails cleanly if missing.

Code: `connectors/cpcb.py:1–173`

OpenAQ v3 REST API

Upstream: OpenAQ Foundation, aggregator of ~50,000 sensors globally via provider integration (partners include governmental CPCB/CAAQMS feeds, PurpleAir networks, AirGradient, StateAir, AirNow). <https://api.openaq.org/v3>.

What it provides: Hourly historical timeseries for PM2.5, PM10, NO2, SO2, CO, O3 at tens of thousands of stations. Each query fetches up to 12 monthly pages of hourly measurements per sensor. Includes sensor metadata (location, last-update time, parameter availability).

Fetches: Every hour, all 10 cities. For each city, the connector finds sensors within 25 km of the city centre, allocates a budget of 4 sensors per pollutant (so ~24 sensors total per city to stay within rate limits on shared GitHub Actions IPs), prioritizes active stations by recency, and pulls the last 1 day of data per sensor.

Written: - Table: `measurements` - Variables: `pm25`, `pm10`, `no2`, `so2`, `co`, `o3` - Source: "openaq" - Confidence: 1.0 (curated government + official networks) - Units: $\mu\text{g}/\text{m}^3$ for particulates and trace gases, as reported by OpenAQ.

Failure handling: API returns HTTP 429 (rate limited) frequently. The connector respects the Retry-After header and backs off exponentially (up to 2^{attempt} seconds, max 60s), retrying up to 5 times before raising an exception. The workflow catches the error and continues to the next city.

Required secret: `OPENAQ_API_KEY` (free sign-up at <https://openaq.org>). Free tier allows ~60 requests/min.

Code: `connectors/openaq.py:1–229`. Budget allocation comment explaining the per-pollutant cap is at lines 162–172.

Open-Meteo weather forecast + archive

Upstream: Open-Meteo, a free (no-key) weather API backed by NOAA/ECMWF model data. <https://api.open-meteo.com/v1/forecast> and <https://archive-api.open-meteo.com/v1/archive>.

What it provides: - Forecast: 2 days ahead + 3 days of recent history, hourly resolution. - Archive: ERA5 historical re-analysis (1940-present), on-demand for any date range.

Variables fetched: air temperature (2 m), relative humidity (2 m), precipitation, boundary-layer height, wind speed and direction (10 m). No API key required; rate limit is generous (~100k requests/day per IP).

Fetches: Every hour, all 10 cities. Each call fetches 3 days of past observations and 2 days of forecast for the city's centre point. Wind speed/direction are converted to u/v components (eastward/northward m/s).

Written: - Table: `measurements` - Variables: `temp` (degC), `rh` (%), `precip` (mm), `blh` (m), `wind_u` (m/s), `wind_v` (m/s) - Source: "openmeteo" - Confidence: 1.0 - Units: SI (Celsius, %, mm, m, m/s). Wind components are derived: $u = -\text{speed} * \sin(\text{bearing})$, $v = -\text{speed} * \cos(\text{bearing})$, where bearing is the FROM direction (meteorological convention).

Failure handling: None special beyond HTTP error logging; missing fields are skipped row-by-row without crashing.

Required secret: None.

Code: `connectors/openmeteo.py:1–156`. Wind conversion at lines 41–46.

Community sensors via OpenAQ

Upstream: Non-governmental sensor networks already aggregated by OpenAQ (PurpleAir, AirGradient, StateAir, AirNow). The connector filters for any station NOT operated by CPCB/CAAQMS, so it captures low-cost sensor clusters.

What it provides: Same pollutant timeseries as OpenAQ but explicitly marked as community-grade (reduced confidence for attribution and enforcement).

Fetches: Not invoked by the hourly cron; run manually or as a backfill via `scripts/fetch_history.py`.

Written: - Table: `measurements` - Variables: `pm25`, `pm10`, `no2`, `so2`, `co`, `o3` - Source: "community" - Confidence: 0.6 (explicitly lower to signal lower instrument precision and calibration) - Units: `ug/m³`.

Failure handling: Same as OpenAQ (rate limiting, retry with backoff).

Required secret: `OPENAQ_API_KEY` (same as OpenAQ).

Status: Configured but **NOT LIVE** in the hourly workflow. Available for manual backfill or future scheduled ingest.

Code: `connectors/community_sensors.py:1–109`.

Satellite and spatial data: daily sources (live)

Earth Engine: Sentinel-5P NO2 + MODIS/VIIRS active fire

Upstream: Google Earth Engine. Sentinel-5P L3 NO2 (Copernicus/ESA, daily revisit) and FIRMS (MODIS/VIIRS active fire detections, NASA).

What it provides: - **Sentinel-5P NO2:** Tropospheric NO2 column density (`mol/m²`), 5.5×10 km pixel grid. - **FIRMS:** Count of active-fire pixels per cell (indicator of biomass burning, proxy for open-waste and construction dust).

Earth Engine collection IDs: - S5P NO2: `COPERNICUS/S5P/OFFL/L3_NO2`, band `tropospheric_NO2_column_number_density` - FIRMS: `FIRMS` dataset, band `T21` (temperature 21, VIIRS)

Fetches: Daily at 01:30 UTC (= **07:00 IST**), all 10 cities. Each run pulls a 30-day rolling window and samples the image at the centre point of each H3 cell that has ground measurements in that city (cells are discovered by querying existing `measurements` rows for the city).

Written: - Table: `measurements` - Variables: `no2_sat` (`mol/m²`), `fire` (count) - Source: `s5p`, `modis` - Confidence: 1.0 - Units: `mol/m²` for NO2, count (unitless) for fire. - Timestamp: Set to end-of-window date at midnight UTC (ISO-8601 +00:00).

Failure handling: The connector requires Earth Engine credentials (service account JSON + project ID). If `GEE_SERVICE_ACCOUNT`, `GEE_PROJECT`, or `GEE_KEY_JSON_B64` are missing, the daily workflow logs a warning but skips cleanly (does not fail the run), leaving `no2_sat` and `fire` at their last ingested value.

Required secrets: `GEE_SERVICE_ACCOUNT` (email), `GEE_PROJECT` (Google Cloud project ID), `GEE_KEY_JSON_B64` (base64-encoded service account key JSON).

Status: **LIVE in production but was missing from the daily schedule for 3 months** (Feb–May 2026). The connector existed and had been run once by hand, but cron was never added, so `no2_sat` and `fire` read 0.0 across all attribution rows. Fixed in this cycle (ingest.yml lines 100–126).

Code: `connectors/earth_engine.py:1–142`. Collection IDs at lines 30–33.

OpenStreetMap emission-source registry via Overpass

Upstream: OpenStreetMap via the free Overpass API (<https://overpass-api.de/api/interpreter> and mirror <https://overpass.kumi.systems/api/interpreter>). Tags queried: `landuse=construction`, `landuse=industrial`, `landuse=landfill`, `power=plant`.

What it provides: Real-time geometry and metadata for emission sources (construction zones, industrial areas, landfills, power plants) within a city's bounding box. Nightly refresh allows auto-add/remove as OSM edits accumulate.

Fetches: Daily at 01:30 UTC (= **07:00 IST**), all 10 cities. The query fetches all named elements (or landfills, which may be unnamed). Results are ranked by substantiality (area-mapped elements > point-of-interest, keywords like "industrial", "refinery", "plant" as tiebreakers) and capped per type (8 construction, 8 industry, 4 landfill) to keep the registry curated and Overpass queries cheap.

Written: - Table: `emission_sources` - Fields: `name`, `type` (construction, industry, waste_burn), `source_origin` ("osm"), `registry_ref` (e.g., "osm:way/12345"), `detection_confidence` (0.9), `geom` (GeoJSON Point at element centre), `attributes` (`h3_cell`, `osm_tags` dict, `pop_exposed_estimate` heuristic) - Deduplicated by (name, type) within a city; sources that vanish from OSM are soft-deleted (flagged `attributes.stale_in_osm=true`) if they have officer-acted enforcement recommendations, otherwise hard-deleted.

Failure handling: Overpass is rate-limited and flaky. The connector tries up to 4 times, cycling through two mirrors, with 30s backoff between attempts (30s, 60s, 90s, 120s). If all fail, it raises a `RuntimeError`; the daily workflow catches it and continues (OSM sources stay at their last-known state).

Required secret: None (Overpass is free and open).

Status: **LIVE in production.** Runs daily (ingest.yml line 96–99). The daily cron then immediately regenerates all enforcement recommendations (line 157–169) to ensure the worklist always reflects current OSM state.

Code: connectors/osm_sources.py:1–279. Scoring logic at lines 61–73; cap configuration at line 46.

Population distribution (GPW v4.11 via Earth Engine)

Upstream: CIESIN/SEDAC GPW v4.11 Population Count (2020 census-calibrated ~1 km raster, NASA). Accessed via Earth Engine.

What it provides: Gridded population per 1 km cell, summed over each H3 res-8 cell polygon. Used to populate the `population` variable in `measurements` and to power the vulnerability index (which blends population percentile with facility density).

Fetches: Not run on a schedule; invoked ad-hoc by `scripts/run_stage1_writes.py` during initial data setup. Once ingested, population is treated as static (2020 snapshot, ts=2020-01-01T00:00:00+00:00).

Written: - Table: `measurements` - Variable: `population` - Source: `gpw411` - Unit: people - Confidence: 1.0 - Only cells that have attribution, forecast, or emission-source rows are sampled (to avoid sparse/meaningless data).

Failure handling: Requires Earth Engine credentials (same as `earth_engine.py`). If missing, fails.

Required secrets: `GEE_SERVICE_ACCOUNT` , `GEE_PROJECT` , `GEE_KEY_JSON_B64` .

Status: Configured, ingested once at setup, **NOT RUN ON A SCHEDULE**. Population does not change daily, so re-ingesting is unnecessary.

Code: connectors/population.py:1–112.

Population vulnerability index (OSM + GPW)

Upstream: OpenStreetMap (hospitals, clinics, schools, colleges, kindergartens, elder-care facilities, outdoor-work anchors like markets and bus stations) + GPW population already in `measurements` .

What it provides: Per-zone (H3 cell) vulnerability score blending population density percentile (50% weight) and facility-count density (50% weight). Formula: $\text{index} = 0.5 * \text{pop_percentile} + 0.5 * \min(1, \text{facility_score} / \text{FACILITY_CAP})$, where $\text{facility_score} = 2 * \text{hospitals} + 1 * \text{schools} + 2 * \text{eldercare} + 1 * \text{outdoor_sites}$ and $\text{FACILITY_CAP} = 14.0$.

Fetches: Not scheduled; invoked by setup and by `scripts/fetch_wards.py` for manual vulnerability analysis.

Written: - Table: `vulnerability` - Fields: `city_id` , `h3_cell` , `zone_id` (readable label like "zone-8d5a"), `vulnerability_index` (0–1), facility counts (hospitals, schools, eldercare, outdoor_sites), population.

Failure handling: Same as OSM (Overpass flakiness, backoff and retry).

Required secret: None.

Status: Configured, **NOT RUN ON A SCHEDULE**. Vulnerability is relatively static and mainly used for dashboard and ward-level summaries.

Code: connectors/vulnerability.py:1–216. Heuristic weights at lines 40–43.

Static layers (Stage-1 seed data)

Upstream: Hardcoded seed data for development/demo (emission sources, roads, vulnerability zones). Deterministic fixtures enabling the system to run end-to-end without live OSM/Earth Engine during early development.

What it provides: JSON blobs with city-specific emission sources (construction, industry, waste_burn, diesel_corridor), vulnerability ward summaries, and road segments with traffic weights.

Fetches: Not scheduled; invoked by `scripts/run_stage1_writes.py` or manually via the `--push` flag to seed the database.

Written: `emission_sources` (hardcoded per city in connectors/static_layers.py:26–54) and indirectly consumed by mobility (via `live_scale`) and advisories refresh.

Status: LIVE as a fallback. If Overpass or Earth Engine fail, the static layers and live traffic proxy keep the UI responsive. The advisories script (line 81 in ingest.yml) loads them; the enforcement pipeline uses OSM sources when available.

Code: connectors/static_layers.py:1–182.

Traffic and mobility (live with optional real-time upgrade)

Traffic proxy: OSM + time-of-day patterns

Upstream: Hardcoded road weights (from static_layers.py) + periodic heuristics (weekday morning rush = +35%, evening = +45%, night = -65%, weekend = -45%).

What it provides: Synthetic traffic congestion index per road segment, updated hourly. Can be scaled by live TomTom data when available (otherwise passes through at 1.0x).

Fetches: Not scheduled; data is generated on-demand by scripts/run_stage1_writes.py or when the mobility module is imported (e.g., in attribution code).

Written: - Table: measurements - Variable: traffic - Source: osm_gtfs (or osm_gtfs×tomtom_live if TomTom key is set) - Unit: index (0–200+, relative to baseline) - Confidence: 0.72 (heuristic, not observed)

Status: Configured, generated on-demand, **NOT scheduled** but integrated into attribution logic.

Code: connectors/mobility.py:1–110. Time-of-day multipliers at lines 19–30; live_scale integration at lines 39–53.

TomTom Traffic Flow API (optional real-time upgrade)

Upstream: TomTom Traffic Flow API (free tier: 2,500 req/day). Returns current/freeFlowSpeed per road segment.

What it provides: Live congestion ratio (0 = free-flow, 1 = standstill) derived from currentSpeed / freeFlowSpeed.

Fetches: On-demand by the mobility module when the TomTom key is set. Queries 5 sample points (city centre + quadrant midpoints) to get representative congestion.

Written: Not directly; the congestion ratio is multiplied into the traffic index when present (scale factor = $0.6 + 0.8 * \text{congestion_ratio}$).

Failure handling: Any TomTom API failure (missing key, no coverage, timeout, 4xx/5xx) is caught silently; the system falls back to the OSM proxy without breaking attribution.

Required secret: TOMTOM_API_KEY (optional; defaults to fallback if absent).

Status: Optional/configured but **NOT active** (no TOMTOM_API_KEY in production secrets). The proxy alone powers traffic in emission attribution.

Code: connectors/traffic_live.py:1–106. Live_congestion function at lines 34–59.

Construction permits (prepared, awaiting open data)

Upstream: Municipal permit registries (DPCC, MCD, BBMP, MCGM, etc.). Currently no Indian city publishes permits as open data; this is a placeholder for when they do.

What it provides: If a CSV with columns (permit_id, site_name, lat, lon, valid_from, valid_to, area_sqm, dust_plan) is dropped into data/permits/{city}.csv, the connector parses it and upserts rows into emission_sources.

Fetches: Never; no permit CSV files exist in the repository.

Written: - Table: emission_sources - Fields: city_id, name, type ("construction"), source_origin ("permit_registry"), registry_ref (permit_id), detection_confidence (1.0, ground truth), geom, attributes (valid_from, valid_to, area_sqm, dust_plan boolean) - Deduplication on registry_ref.

Status: Configured, **NOT LIVE** (no permit data available). Ready to ingest the moment a municipality publishes.

Code: connectors/permits.py:1–86.

Scheduled ingestion: GitHub Actions workflows

The `.github/workflows/ingest.yml` defines two recurring jobs triggered by cron schedule. GitHub Actions runs these at minute 0 of the specified hour UTC.

Job 1: Hourly ground + weather (every hour at minute 5 UTC)

Cron: `5 * * * *`

UTC: Every hour at XX:05 (e.g., 00:05, 01:05, ... 23:05)

IST: Every hour at XX:35 (e.g., 05:35, 06:35, ... 04:35 next day)

Duration: ~2–5 minutes per city (throttled OpenAQ + retries)

Steps: 1. Iterate over all tracked cities. 2. Run `python -m connectors.openaq --city {c} --days 1 --per-var 4 --push` with 20s sleep between cities (rate-limit courtesy). 3. Run `python -m connectors.cpcb --city delhi --push` (Delhi only; attempts direct CAAQMS path; if unreachable from the ISP, it will fail gracefully; this is redundancy). 4. Run `python -m connectors.openmeteo --city {c} --push` for all cities. 5. Ping Supabase once (`SELECT * FROM cities LIMIT 1`) to prevent project auto-pause.

Failure modes: - If OpenAQ rate-limits, the job backs off and retries; a timeout or network error is logged as a GitHub Actions warning (visible in the run summary, not fatal to the workflow). - If Supabase URL/keys are missing, the job fails cleanly. - If Open-Meteo is unreachable, the weather data for that hour is skipped; forecasting continues with older weather.

Secrets required: `SUPABASE_URL` , `SUPABASE_SERVICE_ROLE_KEY` , `DATA_GOV_IN_API_KEY` , `OPENAQ_API_KEY` .

Job 2: Daily forecast + enforcement + satellite (daily at 01:30 UTC)

Cron: `30 1 * * *`

UTC: 01:30 (every day)

IST: 07:00 (every day)

Duration: ~10–20 minutes (parallelized across cities by Python loops, not GitHub parallel jobs)

Steps: 1. **RAG ingest** (line 85): `python -m rag.ingest` — prepares language-model context corpus. 2. **OSM emission-source refresh** (lines 96–99): Refresh `emission_sources` with live OSM data for all cities. Unconditional; clears OSM-based enforcement recommendations to re-generate them fresh (lines 94–99). 3. **Satellite ingest** (lines 100–126): If Earth Engine secrets are present, sample Sentinel-5P NO2 + MODIS fire over the last 30 days at all existing cells. If secrets missing, log warning and skip cleanly. 4. **Forecast + attribution recompute** (lines 127–137): For each city, retrain the forecast model and recompute attribution with the latest measurements. 5. **Multi-agent enforcement pipeline** (lines 138–153): Run the agent graph to generate enforcement recommendations. 6. **Enforcement regeneration** (lines 157–169): Unconditionally regenerate all enforcement rows (even if the agent spike gate skips, the worklist must not be empty). 7. **Citizen advisories refresh** (line 177): `python scripts/refresh_advisories.py` — generates user-facing guidance text. 8. **Morning brief → Telegram** (line 187): `python scripts/send_morning_brief.py` — sends officer summary (skips cleanly if bot token missing). 9. **Retention: archive + prune** (line 196): `python scripts/archive_measurements.py --keep-days 180 --apply` — rolls up raw measurements older than 180 days into daily rollups, archives to S3, then deletes. Keeps only daily PM2.5 rollup forever.

Failure modes: - If Earth Engine secrets are absent, satellite ingest is skipped; `no2_sat` and `fire` become stale (days old). - If Gemini API key missing, advisory/enforcement generation fails; the agent runs but advisories do not refresh (line 177, line 145). - If Telegram token missing, the morning brief is silently skipped (line 184, `|| true`). - Archive failure does not halt the cron (line 196, `|| true`).


Secrets required: `SUPABASE_URL` , `SUPABASE_SERVICE_ROLE_KEY` , `SUPABASE_ANON_KEY` , `GEE_SERVICE_ACCOUNT` , `GEE_PROJECT` , `GEE_KEY_JSON_B64` , `GEMINI_API_KEY` , `TELEGRAM_BOT_TOKEN` , `PUBLIC_API_BASE_URL` , `PUBLIC_WEB_URL` , `SUPABASE_DB_URL` .

Live vs. configured-but-not-running connectors

Connector	LIVE in cron?	Invocation	Notes
cpcb.py	Yes, hourly (Delhi only)	Ingest.yml line 48	Redundancy when OpenAQ lags; falls back gracefully if data.gov.in unreachable.
openaq.py	Yes, hourly (all cities)	Ingest.yml lines 40–44	Primary ground-truth source for air quality. 4 sensors/pollutant.

Connector	LIVE in cron?	Invocation	Notes
openmeteo.py	Yes, hourly (all cities)	Ingest.yml lines 50–51	Forecast drivers (temperature, humidity, wind, boundary layer).
community_sensors.py	No	Manual or scripts/fetch_history.py	Code ready; not scheduled. Would add 0.6-confidence low-cost sensors.
earth_engine.py	Yes, daily (all cities, gated on secrets)	Ingest.yml lines 100–126	Satellite NO2 (Sentinel-5P) + fire (FIRMS). Was missing from schedule for 3 months (bug fixed).
osm_sources.py	Yes, daily (all cities)	Ingest.yml lines 96–99	Emission-source registry refresh. Feeds enforcement pipeline.
population.py	No (one-time setup)	scripts/run_stage1_writes.py	Ingests once; population is static (2020). Queried by counterfactual & enforcement.
vulnerability.py	No (one-time or manual)	scripts/fetch_wards.py (manual)	Same; mostly static. Ingested once, then referenced.
static_layers.py	No (seed/fallback)	scripts/run_stage1_writes.py	Hardcoded demo data. Acts as fallback if OSM/EE fail.
mobility.py	No (generated on-demand)	ml/attribution or scripts	Traffic proxy generated at runtime, not pre-ingested.
traffic_live.py	No (optional enhancement)	connectors/mobility.py::live_scale()	TomTom key unset; falls back to OSM heuristic.
permits.py	No (awaiting open data)	Manual if permit CSV provided	Ready for CSV at data/permits/{city}.csv but no Indian city publishes yet.

Data flow summary



Every measurement row is timestamped in UTC and indexed by (city_id, h3_cell, variable, ts, source) for rapid aggregation by the attribution and forecasting pipelines. Duplicates (same sensor, same hour) are silently ignored on insert. Population and vulnerability are queried from `measurements` and a separate `vulnerability` table, respectively, by the downstream agents.

Error resilience and retry logic

All connectors implement bounded retry loops with exponential backoff: - **OpenAQ**: 5 retries, 429-aware (respects Retry-After header), max 60s backoff. - **CPCB**: 6 retries, 10s max backoff (endpoint recovers quickly). - **Overpass (OSM/vulnerability)**: 4 retries across two mirrors, 30–120s backoff. - **Earth Engine**: No retry logic (once-per-day batch); missing secrets cause graceful skip. - **Open-Meteo**: No retry (timeout = fail; re-run next hour).

The GitHub Actions workflow wraps each connector with a shell `|| echo "::warning::"` catch, so failures are logged as warnings (visible in the Actions summary) but do not halt the job. The `|| true` tail on steps like Telegram and archive ensures the daily run completes even if an optional downstream step fails.

Time windows and staleness

- **Ground/weather**: Updated every hour, typically 10–30 min old by the time it reaches the database (API latency + ingest time). Attribution runs on the latest hour's data.
- **Satellite**: Updated daily, samples a 30-day rolling mean; NO₂ may be 1–2 days old (Sentinel-5P revisit period). Fire counts are live within the day.
- **Emission sources**: Updated nightly; sources appear within 24 hours of OSM edit.
- **Population, vulnerability**: Static (2020, 2020).
- **Forecasts, attribution**: Recomputed daily at 01:30 UTC (07:00 IST) using the latest measurements.

A stale or missing upstream source is flagged in the workflow summary and does not break the system. Forecast/attribution use the most-recent available data, so a gap of a few hours is absorbed gracefully.

8 · Forecasting: the model, the calibration, and what it is worth

The forecasting engine produces PM_{2.5} predictions per H3 cell at horizons of 24, 48, and 72 hours ahead. It is a trained model (LightGBM), not a physics simulation, and it is deliberately designed to sit next to a hard baseline—persistence—so that when it fails, that failure is visible and measurable. This chapter explains what the model does, how it is built, how it knows its own uncertainty, and what the benchmark numbers say about whether any of it matters.

The target and horizons

The target is the PM_{2.5} concentration ($\mu\text{g}/\text{m}^3$) observed within a single H3 cell at a specific future hour. Each training sample is anchored to a location and a time (city, cell, timestamp) and asks: given what we know now and earlier, what will PM_{2.5} be at `now + h` hours? The model trains separately for each city and for each horizon `h ∈ {24, 48, 72}` (ml/forecast/features.py:92–108). Data is held and represented chronologically to prevent leakage; all training happens strictly before each test origin (ml/eval/benchmark.py:73–78).

Early observations in the production dataset are sparse or missing; the model requires at least 60 measured rows per horizon to train (ml/forecast/train.py:48). Rows are dropped if the target is NaN or if the most recent observation (persistence feature) is NaN, because the model cannot work without both (ml/forecast/features.py:92–108). This gives the model a built-in data quality floor.

The feature set

Features fall into five categories: pollutant lags, meteorology, calendar events, dispersion, and hour/day-of-week. Every feature comes from a named source; none are constructed ad-hoc.

Pollutant observations (current and lagged)

The model sees the current and recent concentrations of six pollutants: PM2.5, PM10, NO₂, SO₂, CO, and O₃ (ml/forecast/features.py:16). These are drawn from station observations in the same or nearby cells. Two lags are computed: the value 1 hour ago (`pm25_lag1`) and 24 hours ago (`pm25_lag24`), per cell (ml/forecast/features.py:87–88). Lags are the core persistence signal: what happened before is the strongest predictor of what happens next in air quality.

Raw values outside the instrument's plausible range are dropped before any modelling: PM2.5 > 1500, PM10 > 3000, NO₂ > 1000, SO₂ > 1000, CO > 100, O₃ > 1000 (ml/forecast/features.py:20). Negatives are also dropped. This guard catches sensor glitches and errant entries without modifying real data (ml/forecast/features.py:62–64). The same guard is applied to production data (ml/eval/benchmark.py:48).

Meteorology (broadcast)

Meteorological fields come from Open-Meteo hourly reanalysis: temperature, relative humidity, precipitation, eastward wind (u), northward wind (v), and boundary layer height (BLH) (ml/forecast/features.py:17). These are broadcast to all cells in a city for each timestamp—they are regional, not cell-specific. The data lands on misaligned sub-hour offsets (OpenAQ at :30, Open-Meteo at :00), so both are floored to the hour to align them (ml/forecast/features.py:56–57).

Two derived meteorological features are computed from the wind components:

- **Wind speed:** $\sqrt{u^2 + v^2}$ in m/s, the magnitude of horizontal flow (ml/forecast/features.py:79–80).
- **Ventilation:** `wind_speed` × BLH, the product of transport and mixing depth (ml/forecast/features.py:81). Low ventilation (calm wind + shallow boundary layer) traps pollution; high ventilation disperses it. This is the physically grounded feature that captures the stability of the atmosphere.

Meteorology typically accounts for ~1–2% of RMSE improvement in the benchmark (ml/eval/benchmark.py:429–441; seen in Delhi at 24h as ~2%).

Calendar and seasonal event flags

The model captures three seasonal/event signals that drive pollution in India:

- **Stubble season flag:** 1 if the month is October or November, 0 otherwise. Paddy residue burning in northern Punjab drifts into northern cities, especially Delhi (ml/forecast/seasonal.py:29).
- **Winter inversion flag:** 1 if the month is November, December, January, or February, 0 otherwise (ml/forecast/seasonal.py:30). Shallow boundary layer and stagnant air during winter trap emissions near the surface. This is the dominant seasonal signal in Delhi.
- **Diwali window flag:** 1 if the timestamp is within ±3 days of the main Diwali night (Lakshmi Puja), 0 otherwise (ml/forecast/seasonal.py:21). Diwali dates are hardcoded per year: Nov 12 (2023), Nov 1 (2024), Oct 21 (2025), Nov 8 (2026), Oct 29 (2027) (ml/forecast/seasonal.py:12). The window is 3 days; the flag activates during high firecracker pollution episodes.

These flags are added via `add_calendar_features()` (ml/forecast/features.py:85), which applies the calendar logic vectorized across all timestamps.

Dispersion and upwind advection

Pollution transported by the wind field is captured via an advected PM2.5 feature. For each cell and timestamp, the model traces backward 6 hours along the wind vector to find where the air parcel came from, identifies the nearest data cell to that upwind origin, and reads off its PM2.5 (ml/forecast/features.py:23–49). This is pure kinematics: `upwind_origin()` (ml/dispersion/advection.py:21–28) computes

```
origin_lat = lat - v_north * dt / M_PER_DEG_LAT
origin_lng = lng - u_east * dt / (M_PER_DEG_LAT * cos(lat))
```

where `dt = hours × 3600` seconds, and u/v are the wind components (east and north, in m/s). The feature captures pollution that is *headed toward* this cell from upwind, not just what is overhead now. It is a physics-informed feature: if the model learns to weight it, it is learning to account for transport (ml/forecast/features.py:4).

If wind data is missing or the cell grid cannot be parsed (e.g., in unit tests with non-H3 IDs), the feature is skipped silently (ml/forecast/features.py:35).

Hour and day-of-week

The hour-of-day (0–23) and day-of-week (0–6, Monday = 0) are added as categorical signals (ml/forecast/features.py:83–84). These capture traffic cycles and diurnal boundary-layer effects: rush-hour pollution, daytime mixing, and nighttime stability follow predictable patterns.

The model class and hyperparameters

The model is **LightGBM (Light Gradient Boosting Machine)** with quantile regression loss (ml/forecast/train.py:33–37). LightGBM is a decision-tree ensemble that is fast to train, interpretable per feature, and handles missing values naturally. Quantile loss allows the model to predict not just a central value but also confidence bounds.

The exact hyperparameters for the median ($\tau = 0.5$) and quantiles ($\tau = 0.1$ and 0.9) are:

Parameter	Value	Rationale
objective	"quantile"	Enable quantile regression loss. Three separate models are trained: one for each τ .
alpha	0.1, 0.5, 0.9	The quantile levels: lower bound (10th percentile), median, upper bound (90th percentile).
n_estimators	200	Number of boosting rounds. Stops early if validation loss plateaus.
learning_rate	0.05	Shrinkage per iteration; conservative to avoid overfitting.
num_leaves	31	Max leaves per tree; default for LightGBM. Balances expressiveness and generalization.
min_child_samples	5	Minimum samples per leaf; prevents overfitting on tiny subgroups.
verbosity	-1	No training-time logging.

These are the same hyperparameters used in production (ml/forecast/train.py:27); they are not tuned per city or horizon, so they are a fixed, reproducible baseline.

Training protocol: per city, per horizon

The training pipeline is:

- Load measurements** from Supabase for a given city (ml/forecast/train.py:344; ml/eval/benchmark.py:40–62). These include CPCB station data and optionally OpenAQ data, floored to the hour.
- Build feature table** via `build_feature_table()` (ml/forecast/features.py:52–89): pivot pollutants to wide form (per cell), pivot meteorology to wide form (per city $\times$ timestamp), merge them, add derived met, add calendar flags, add advection, add lags.
- Create supervised samples** via `make_supervised()` (ml/forecast/features.py:92–108): shift the target forward by `horizon_h` rows within each cell, drop rows with NaN targets, keep only complete cases.
- Train and validate** using walk-forward expanding-window backtest (ml/forecast/train.py:41–76): divide the ordered data into `n_folds` (default 3) time folds, train on folds 0 to i , test on fold $i+1$, repeat. This avoids the overfitting trap of a single train/test split and averages skill across multiple regimes.
- Retrain cadence:** In production, the model retrains every 24 hours on the trailing 90 days of data (ml/eval/benchmark.py:565). The 90-day window was chosen over an expanding window because skill on winter +24h is higher with the shorter window, where it matters most (docs/BENCHMARKS.md line 49–51).

For the benchmark—a more rigorous evaluation—the model retrains monthly (rolling-origin protocol) with the same 90-day training window, so each month's forecast is trained strictly on data before that month began (ml/eval/benchmark.py:197–271).

Quantile regression and why three quantiles

Standard regression predicts a point: $\hat{y} = E[y | x]$. Quantile regression predicts the quantile: $\hat{y}_\tau = F^{-1}(\tau | x)$, where F is the conditional distribution. Three quantiles are fitted:

- **$\tau = 0.1$ (pi_low):** the 10th percentile, the lower bound of an 80% prediction interval.
- **$\tau = 0.5$ (value):** the median, the central forecast. This is what is served as the point estimate and what appears on the console.
- **$\tau = 0.9$ (pi_high):** the 90th percentile, the upper bound of an 80% prediction interval.

Why three, not one? Quantile regression captures model uncertainty *without* assuming a distribution (e.g., Gaussian). An air quality forecast that says " $48 \pm 20 \mu\text{g}/\text{m}^3$ " (Gaussian confidence interval) is weaker than one that says "the 10th percentile is 35, the median is 55, and the 90th percentile is 90." The latter is *calibrated*: if you look at many forecasts with predicted interval [35, 90], the true value should fall in that interval 80% of the time. But raw quantile models often under-cover (ml/forecast/train.py:151): testing on Delhi data, naive LightGBM $\tau = 0.1/0.9$ intervals cover only 48–63% of observations, not 80% (ml/forecast/train.py:151).

This is where Conformalized Quantile Regression comes in.

Conformalized Quantile Regression (CQR): recalibrating the bands

CQR (Romano et al., 2019) is a post-hoc calibration procedure that takes a raw quantile model's predictions and widens the interval to achieve nominal coverage *in finite samples*. The method is distribution-free and has finite-sample guarantees.

The conformity score

For each training sample i in the calibration set, compute the **conformity score**—how wrong the quantile model was:

$$E_i = \max(\hat{y}_{0.1}(i) - y_i, y_i - \hat{y}_{0.9}(i))$$

The score is the maximum of: - How far below the true value the lower bound fell: $\hat{y}_{0.1}(i) - y_i$ (negative if the bound was too low). - How far above the true value the upper bound went: $y_i - \hat{y}_{0.9}(i)$ (negative if the bound was too high).

In other words, E_i is the *minimum distance you'd have to expand the interval to enclose the true value*. If the interval already encloses y , then $E_i \leq 0$; if not, $E_i > 0$.

The procedure

To calibrate:

1. **Fit on 75% of training data** (in time order, not shuffled): fit the median, lower, and upper quantile models on the first 75% of samples (ml/forecast/train.py:203; ml/eval/benchmark.py:227).
2. **Score on the remaining 25%:** predict on the held-out calibration tail, compute E_i for each sample, sort the scores (ml/forecast/train.py:159–162).
3. **Compute the adjustment Q:** Find the quantile of the sorted scores at level $\lceil (n+1)(1-\alpha) \rceil / n$, where α is the target miscoverage and n is the number of calibration samples (ml/forecast/train.py:134–142):

$$Q = \text{quantile}(E_{\text{sorted}}, \lceil (n+1) * \text{coverage} \rceil / n)$$

For $\alpha = 0.2$ (80% coverage target) and $n = 282$ (Delhi +24h), this is $\lceil 283 \times 0.8 \rceil / 282 = \lceil 226.4 \rceil / 282 = 227 / 282 \approx 0.804$ (ml/forecast/train.py:134–142). The level is slightly *above* $1 - \alpha$ to guarantee coverage in the next sample. This is the finite-sample correction.

4. **Serve the adjusted band:** the prediction interval is $[\hat{y}_{0.1} - Q, \hat{y}_{0.9} + Q]$. Every new forecast uses this same Q (ml/forecast/train.py:303–304).

What CQR guarantees and what it doesn't

Guarantee: Split conformal prediction gives **marginal coverage**—across *all* test samples, the fraction inside the band is $\geq (1 - \alpha)$ in expectation. On Delhi's rolling-origin benchmark with 10 origins and 208k test samples, the 80% band achieves 0.783 empirical coverage, very close to the nominal 0.80 (docs/benchmarks/delhi.md, +24h).

No guarantee: CQR does *not* promise **conditional coverage**—coverage stratified by regime, pollution level, time of day, or any other feature. This matters.

The calibration failure that is honest about its limits

The Kolkata benchmark exposes this limit clearly (docs/benchmarks/kolkata.md):

horizon	Q1 (clean)	Q2	Q3	Q4	Q5 (polluted)	overall
+24h	0.803	0.778	0.761	0.668	0.733	0.749
+48h	0.799	0.793	0.725	0.62	0.687	0.725
+72h	0.812	0.785	0.649	0.547	0.699	0.699

Marginally, the band covers 0.749–0.699, close to nominal. But grouped by *predicted* PM2.5 level (quintiles), the worst predicted quintile—Q4 (the mid-to-upper range, 56–76 $\mu\text{g}/\text{m}^3$ at +24h)—covers only 0.668 at +24h and drops to 0.547 at +72h. This is the operational range where decisions are made: is it Moderate (56–75) or Poor (75–250)? The band is too narrow there.

Why does this happen? The quantile models under-disperse in the mid-to-upper range. Seven conformity scores and Mondrian approaches (per-bin splitting) were tested over four folds to close the gap (ml/forecast/train.py:120–124). The worst-case quintile moved from 0.615 to 0.646 at best, paid for by degrading coverage in clean-air quintiles. So 3 points on an 18-point shortfall—not the lever.

The design choice: keep the simple two-sided conformity score and **report the failure** (ml/forecast/train.py:126–130). The benchmark now emits `pi80_coverage_by_predicted_quintile` so the 0.67 sits next to the 0.75 rather than being averaged into it. Making the problem measurable was the honest fix available.

The persistence blend: why and how

Persistence is the baseline forecast: $\hat{y}_{\text{persistence}}(t+h) = y(t)$. It is the current PM2.5, assumed to hold 24/48/72 hours out. It is a hard baseline—on city-hours where the model's median beats it, the value added is real; where it loses, the loss is visible.

Raw LightGBM often improves on persistence modestly: +1.7% RMSE on Delhi +24h (docs/benchmarks/delhi.md, raw skill). But when blended with persistence, the same model reaches +9.1%. Why?

The blend is a convex combination:

$$\hat{y}_{\text{served}} = w \cdot \hat{y}_{\text{model}} + (1 - w) \cdot \hat{y}_{\text{persistence}}$$

where $w \in [0, 1]$ is chosen to minimize RMSE on a 21-point grid over the calibration tail (ml/forecast/train.py:169–187). The calibration tail is the held-out 25% of training data used to fit the quantile models, in chronological order. This ensures no test leakage: the blend weight for origin M is fit only on data strictly before month M.

Why does the blend help? Persistence is unbiased: $E[\hat{y}_{\text{persistence}}] = E[y]$. The model's median may be biased (too high in clean seasons, too low in severe episodes) or noisy (high variance on sparse cells). The blend trades off model skill (potentially biased but precise) against persistence (unbiased but crude). The weight w is *per-origin*: on Delhi's rolling-origin benchmark, w ranges from 0.25 to 1.0 across the 10 monthly refit origins (docs/benchmarks/delhi.md, +24h: [0.4, 0.5, 0.55, 0.9, 0.25, 0.85, 0.9, 0.85, 1.0, 0.65]).

Risk: w can land at 1.0 (full model) or near 0.0 (mostly persistence) depending on the calibration tail's regime. If the tail is calm, w tends toward 1 because the model wins then. If the tail is severe, w drops because persistence is hard to beat during spikes. This is not a bug; it is the blend adapting to what the training data says. But it means the blend *can* become the pure model in calmer periods, losing the robustness that persistence provides.

Exceedance probabilities and CPCB thresholds

The model serves two exceedance probabilities:

- **P(>120):** probability that PM2.5 exceeds 120 µg/m³. CPCB "Very Poor" band.
- **P(>250):** probability that PM2.5 exceeds 250 µg/m³. CPCB "Severe" band.

These are **calibrated** using the residuals from the same 25% calibration tail (ml/forecast/train.py:190–224; ml/eval/benchmark.py:344–357):

$$P(y > \text{threshold}) = 1 - \hat{F}(\text{threshold} - \hat{y}_{\text{median}}) = (\text{count of } r < \text{threshold} - \hat{y}_{\text{median}}) / n_{\text{calibration}}$$

where *r* are the sorted residuals from the blended forecast (ml/forecast/train.py:216–223). In other words, the empirical CDF of residuals is used as the predictive distribution. This is calibration by construction: if you collect many forecasts with predicted $P(>120) = 0.30$, the true exceedance rate in that sample should be ≈ 0.30 .

CPCB bands: India's Central Pollution Control Board (CPCB) defines air quality ranges: - Satisfactory: 0–50 µg/m³ - Moderate: 51–100 µg/m³ - Poor: 101–250 µg/m³, broken at 120 for Very Poor (120–250) and below (101–120) - Severe: >250 µg/m³

The two thresholds capture the boundaries where intervention decisions change: crossing 120 means "activate warnings"; crossing 250 means "emergency measures."

What the forecasts are actually worth

The benchmark is built on the strictest rules (ml/eval/benchmark.py, lines 6–10): strict temporal split (no future leakage), production-faithful retraining (monthly on 90 days), hard baselines (persistence, weekly seasonal-naive, climatology), one shared support mask (every forecaster on the same rows), regimes kept separate (full test, winter, non-winter, high-pollution hours).

Delhi: strong overall, weaker on the worst tail

Setup: 39 station cells, 449k hourly rows, test from Nov 1, 2025 (the 2025-26 winter plus spring/summer 2026). Rolling monthly refit, 90-day training window.

horizon	served skill vs persistence	winter only	very-poor hours (>120)	80% PI coverage	coverage @ worst quintile
+24h	+9.1%	+6.7%	+6.5%	0.783	0.704
+48h	+12.9%	+11.4%	+10.6%	0.781	0.731
+72h	+12.1%	+10.6%	+10.0%	0.775	0.705

Reading: The served forecast (blended median) beats persistence at every horizon (+9/+13/+12%), and it holds up in winter (+7/+11/+11%). On high-pollution hours (Very Poor, >120 µg/m³), it still wins (+7/+11/+10%) but margins are thinner. The 80% interval covers 78–78% of observations marginally, but on the predicted mid-to-upper range (Q4–Q5), coverage drops to 0.70–0.82. On the Severe tail (>250), the model roughly matches persistence (–0.5% / +0.7% / –2.0%); on the worst hours, there is no added value from the median forecast.

Onset recall: The metric that matters for intervention is onset recall—detecting clean→bad transitions. On Very Poor events, with the median alarm (forecast >120), only 24% of clean→very-poor transitions 1 day ahead are caught. But with the **calibrated probability** at $\tau = 0.3$, recall jumps to 54% with $F1 = 0.765$, better than persistence's static 0% (docs/benchmarks/delhi.md, +24h).

Mumbai: best-in-class, but coastal clean-season masking

Setup: 27 station cells, 276k test rows, same protocol. Mumbai's coastal regime is where the model earns most of its skill.

horizon	served skill	winter	very-poor hours	80% PI coverage
+24h	+16.5%	+15.2%	–2.6%	0.816

horizon	served skill	winter	very-poor hours	80% PI coverage
+48h	+19.1%	+17.1%	−0.0%	0.817
+72h	+21.3%	+20.1%	−0.1%	0.794

Reading: The model *loses* on high-pollution hours (Very Poor, >120): −2.6% at +24h, −5.0% on Severe. Why? Mumbai is coastal and has fewer severe episodes; 1.6k of 117k test samples are >120 vs 62.8k for Delhi. Persistence is hard on rare events, and the model hasn't learned to beat it there. Overall skill is strong because *most* hours are clean, and the model does better than persistence on clean air. Meteorology contributes only 0.8% at +24h and near-zero thereafter (docs/benchmarks/mumbai.md, +99–100).

The exceedance probability model is actually *mis-calibrated* on Mumbai's rare high-pollution regime: Brier skill is *negative* on $P(>120)$ and $P(>250)$ (−11.8% and −11.2% at +48h). The model learned the wrong distribution on the tail because the tail is so small.

Kolkata: thin network, skill on average but not on episodes

Setup: 10 station cells (the sparsest network of the three), 114k test rows.

horizon	skill	winter	very-poor hours	80% PI coverage	coverage @ Q5 (worst quintile)
+24h	+14.1%	+12.4%	+3.2%	0.749	0.733
+48h	+9.6%	+13.0%	−0.7%	0.725	0.687
+72h	+9.2%	+8.3%	−6.9%	0.699	0.699

Reading: Kolkata beats persistence overall (+14/+10/+9%) and in winter, but loses on the exact hours that matter for intervention (Very Poor, +48/+72h). The thin station network (10 sites) means local features are noisy, and generalization to unseen cells is risky. The 80% interval fails on the worst-predicted regime: at +72h, Q5 coverage is only 0.699 when it should be 0.80. Onset recall at $P \geq 0.3$ is near-zero on severe episodes.

Cities where the model does not beat persistence

On Mumbai's Very Poor hours and Kolkata's high-pollution tail at +48/+72h, the model underperforms. This is reported cleanly in the benchmark (docs/benchmarks/mumbai.md, lines 25–27, 29–33) and kept in the public record (docs/BENCHMARKS.md, line 6: "Negative numbers are kept, not hidden.").

The pattern: rare events and sparse networks. When the training data has few examples of the severe regime, the model learns the median distribution but misses the tails. Persistence—the current value—is a surprisingly hard baseline for extreme air quality, especially 2–3 days out.

The one design finding: 90-day window beats expanding

With an expanding training window (all data up to each test month), Delhi's winter +24h skill falls to +3.7% vs +6.7% on the 90-day rolling window (docs/benchmarks/delhi_expanding.md vs delhi.md). The short window's recency bias buys skill exactly where it matters most: the current winter, the current regime. This motivated production to retrain on 90 days, not forever (docs/BENCHMARKS.md, line 49–51). Production data is GC'd after 90 days anyway (core/supa.py), so the choice is free.

Summary: honest metrics, visible limits

The forecast is worth using for long-range early warning on clean→bad transitions. It beats persistence on the median, in winter, and on high-pollution hours overall (Delhi, Kolkata), though not on rare extreme events (Mumbai's Severe tail). The 80% interval achieves nominal marginal coverage but fails on the predicted mid-to-upper range where decisions happen. The exceedance probabilities ($\tau = 0.3$) are the better operating point for alarming: they trade recall for precision and flag 50%+ of onsets 1 day ahead with $F1 > 0.74$.

What it is not: a replacement for dispersion models on transport episodes, or a tool for predicting the Severe tail (>250) with confidence. Where it excels: hyperlocal, updated daily, calibrated, with confidence intervals on every cell, and honest about where it fails.

9 • Source attribution, dispersion and the dense field

The second ML core answers a different question: given that a cell shows high PM2.5, which sources are responsible? The attribution system apportions the measured pollution mass among sources by combining chemical signature priors with a learned model, then validates its output against published urban apportionment studies.

The problem and the six source buckets

A sensor reading is a blend. When Delhi's Anand Vihar station reports 140 µg/m³ PM2.5, the mass comes from traffic exhausts, construction dust, industrial stacks, biomass burning, pollution advected in from upwind, and unmeasured other sources in unknown proportions. Knowing which source dominates is essential for enforcement. A construction site emitting 20 km downwind of the monitor does not justify a traffic ban; a power plant stack dominating the signal does.

The attribution system estimates a *receptor-side* apportionment: which signals move the reading at that cell. This is not an emission inventory (which counts what is released at the source) — a tall industrial stack 50 km upwind may contribute most of the local PM via transport, and the receptor model will see it as `transported` while an inventory credits it to industry. The system operates under this constraint openly and validates against what published receptor-side studies report.

VayuNetra allocates PM2.5 among six categories (`ml/attribution/signatures.py:11-14`):

Category	Meaning	Key markers
<code>traffic</code>	Vehicle exhaust + road-wear	NO ₂ , CO, satellite NO ₂ column
<code>industrial</code>	Stack emissions (factories, power, refinery)	SO ₂
<code>construction_dust</code>	Soil and demolition dust	PM10 / PM2.5 ratio > 1.8
<code>biomass_burning</code>	Open waste and stubble fires	FIRMS thermal anomalies
<code>transported</code>	Regional advection (pollution from upwind)	Advected PM2.5, satellite fields
<code>other</code>	Unmeasured local or residual	Baseline floor

Method 1: Chemical-signature priors (signature-v1)

The MVP attribution method assigns source shares based on transparent physics and chemistry. Each pollutant that a sensor measures is a marker for specific sources (`ml/attribution/signatures.py:45-81`). The method is deterministic and fully auditable — readers can see exactly why each cell blamed traffic versus dust.

What the signatures are:

For each cell, the system computes a normalized score per source, where normalization clips each pollutant to a reference scale (typically the 90th percentile of that city's measured values, so high readings sit near 1.0). The scores are then renormalized to sum to 1.0 as shares.

Source	Score formula	Rationale
<code>traffic</code>	$0.5 \times \text{norm}(\text{NO}_2) + 0.3 \times \text{norm}(\text{CO}) + 0.2 \times \text{norm}(\text{sat-NO}_2)$	Ground-level NO _x and CO are traffic hallmarks; satellite NO ₂ column (Sentinel-5P) corroborates the ground signal (fusion)
<code>industrial</code>	$\text{norm}(\text{SO}_2)$	SO ₂ is a stack marker; little else emits it in cities
<code>construction_dust</code>	$\max(0, \text{ratio} - 1.8) \times 0.6$, where $\text{ratio} = \text{PM}_{10} / \text{PM}_{2.5}$	Coarse dust skews the PM ₁₀ :PM _{2.5} ratio; receptor modeling standard

Source	Score formula	Rationale
biomass_burning	<code>norm(fire)</code> , fire = FIRMS thermal anomalies	Satellite-detected open fires (375 m hotspots, VIIRS)
transported	0.15 (fixed baseline)	Regional background; refined by advection later
other	0.10 (fixed baseline)	Unmeasured sources; ensures modesty

Where the marker reference scales come from:

The system calibrates references per city from live data. `calibrate_references()` computes the 90th percentile of each pollutant over a city's recent measurement history, then clips to a plausible range. This ensures that the "strong marker" threshold reflects current conditions (monsoon vs. winter burning season), not a fixed seasonal snapshot. If a city has fewer than 10 datapoints for a pollutant, it falls back to the fixed defaults in `_REF` (`signatures.py:19-20`): NO₂ 80 µg/m³, CO 2 mg/m³, SO₂ 30 µg/m³, PM_{2.5} 150 µg/m³, PM₁₀ 300 µg/m³, fire 50, satellite NO₂ 1.5×10⁻⁴ mol/m².

These reference thresholds have no published external source — they are empirically chosen to make a cell at the city's high-pollution percentile a strong marker (≈1.0 normalized) and a typical cell ≈0.5. The design follows receptor-modeling practice (e.g., PMF, CMB studies) where marker strength scales with extremeness.

Satellite NO₂ fusion:

The inclusion of Sentinel-5P satellite NO₂ column is a VayuNetra addition to classical CMB (Chemical Mass Balance). Ground NO₂ alone can be ambiguous (industrial stacks also emit NO_x); adding the satellite column (which sees the integrated tropospheric column at ~10 km footprint) resolves the ambiguity — high ground NO₂ + high satellite NO₂ is traffic, high ground + low satellite is a local stack. The tests confirm this: adding satellite NO₂ raises the traffic share for a cell that would otherwise score lower (`test_satellite_no2_corroborates_traffic`).

Confidence estimation:

For each cell, the system computes a confidence score (0.30–0.95 range) as the maximum share across all categories, clamped between the bounds. The intuition is that if one source dominates with 60% of the mass, we are more confident in the apportionment than if all sources tie at 16%. This simple metric is transparent to users and serves as an honest caveat.

Method 2: Hybrid GBM + SHAP (`hybrid-gbm-shap-v2`)

When a city has sufficient historical data, a supervised gradient-boosted model learns to predict PM_{2.5} from the same marker and condition features, and model-agnostic SHAP explainability attributes each prediction to input features. The hybrid method blends this learned attribution with the signature priors, and falls back automatically if the model fails out-of-sample quality gates.

The model and training (`shap_attribution.py:98-150`):

A LightGBM regressor is trained on per-(cell, hour) rows with marker features (NO₂, CO, SO₂, PM₁₀/PM_{2.5} ratio, fire, advected PM_{2.5}, satellite NO₂), condition features (temperature, RH, precipitation, wind components, boundary-layer height, wind speed, ventilation, hour-of-day, day-of-week, calendar flags for stubble season / Diwali / winter), and calendar features (`is_winter`, season flags, etc.). The target is PM_{2.5} at that row, measured at the cell.

Crucially, the model never sees PM_{2.5} itself as a feature — that would leak the target. Instead, it learns from the source markers and meteorological conditions. The condition features help the fit (better weather => lower pollution), but they do not contribute to the apportionment mass — only marker features do.

Training uses an 80/20 time-ordered split. The model is evaluated on held-out recent data to estimate honest out-of-sample R². If R² is below the minimum threshold (0.15), the model is rejected as having no real predictive skill; all cells revert to signature priors.

What SHAP values are and why they work for apportionment:

SHAP (SHapley Additive exPlanations) is a game-theoretic framework that assigns each input feature a signed contribution to the model's prediction for a single sample. For a cell with PM_{2.5} prediction of 120 µg/m³, SHAP might return: NO₂ → +30 (traffic

signal), $\text{PM}_{10}/\text{PM}_{2.5} \rightarrow +35$ (dust), $\text{SO}_2 \rightarrow +10$ (industry), ventilation $\rightarrow -5$ (weather suppression), baseline $\rightarrow 50$. The sum is exactly 120.

SHAP contributions are used for apportionment because they are:

1. **Attribution by association, not causation** — SHAP reflects "in this model, cells with high NO_2 tend to have high $\text{PM}_{2.5}$, and NO_2 's typical gradient in the data suggests it contributes $+30 \mu\text{g}/\text{m}^3$ to the prediction." This is honest: the model does not prove NO_2 *causes* the PM, but identifies it as a co-varying signal the model relies on.
2. **Locally accurate** — the SHAP contributions sum exactly to the model's prediction, so the apportioned shares are consistent with the model's beliefs.
3. **Comparable to signatures** — both methods (signature + SHAP) map chemical markers to sources, so they can be blended coherently.

The implementation (`_shap_shares`) groups positive SHAP mass by source marker (`SOURCE_MARKERS` mapping: NO_2 / CO / sat- $\text{NO}_2 \rightarrow$ traffic, $\text{SO}_2 \rightarrow$ industrial, $\text{PM}_{10}/\text{PM}_{2.5} \rightarrow$ construction_dust, fire $\rightarrow$ biomass_burning, advected $\text{PM}_{2.5} \rightarrow$ transported). Negative SHAP contributions (e.g., high wind speed reducing local pollution) are excluded from the mass — only sources get blamed. An `other` floor (5%, `OTHER_FLOOR`) ensures shares never claim 100% certainty.

Quality gates (`shap_attribution.py`:48-49):

Two gates must pass for the hybrid method to be trusted:

1. **MIN_SAMPLES = 400**: The city must have at least 400 hourly measurements over the attribution window (typically trailing 72 hours per cell, `RECENT_HOURS`). Below this, the training set is too thin and the learned weights are unreliable.
2. **MIN_HOLDOUT_R² = 0.15**: The model's R^2 on held-out data must exceed 0.15. This is deliberately permissive (a strong forecast model typically scores 0.4–0.7), but even 0.15 means the model beats a naive "pollution is constant" baseline. A model with $R^2 < 0.15$ has no out-of-sample skill and is not allowed to assign blame — reverting to signature priors is the honest choice.

If either gate fails, `apportion_cells()` raises a `ValueError` with the reason, which the runner catches and records in every attribution row's `evidence.fallback_reason` field (line 145, `attribute.py`).

Blending with signature priors (`_blend` , line 218):

The final hybrid share is: `hybrid_share = 0.6 × shap_share + 0.4 × signature_prior_share` (`BLEND_WEIGHT = 0.6`). This weighting keeps the method from over-trusting the learned model. A signature prior derived from published chemistry is preserved in the final output; the model merely tilts the allocation.

Handling marker-less cells (`_apply_hybrid` , lines 72-79):

Delhi's public feed currently has NO_2 at only 4 of its 24 cells. Cells without any marker sensors (no NO_2 , CO, or SO_2 measurement) cannot be apportioned by the hybrid model — SHAP cannot assign sources to features it never sees observed. For these cells, the system shrinks toward the city's hybrid-mean apportionment, weighted 50% city mean + 50% the cell's own signature prior. The method is recorded as `signature-citymean-v1` and the `evidence` field carries `shrunk_toward: "city_hybrid_mean"` .

Confidence calibration (`_calibrated_confidence` , line 224):

For hybrid cells, confidence is a weighted blend of:

- **Agreement between methods** (35% weight): L1 distance between SHAP and signature shares (max 1.0 if identical, 0 if completely opposite).
- **Model fit** (20% weight): the holdout R^2 .
- **Sample depth** (10% weight): how many hours of trailing data the cell has (capped at the full window).
- **Baseline** (35% weight): floor of 0.30, ceiling 0.95.

The formula (`line 231`) yields confidence in the range [0.30, 0.95]. Good agreement between independent methods + a well-fit model + deep history = confidence near 0.95.

Fallback behavior and method identification

When the hybrid model's gates fail (insufficient samples or $R^2 < 0.15$), the runner catches the exception (`attribute.py:134`) and records the reason in every row's `evidence.fallback_reason` field. Users reading the API (e.g., `/attribution` endpoint, `main.py`) see:

- **method_version** (one of `"signature-v1"`, `"hybrid-gbm-shap-v2"`, `"signature-citymean-v1"`) — which method produced the shares.
- **confidence** — how much to trust this apportionment (0.30–0.95).
- **evidence** — a dict containing:
 - For signature: `no2`, `co`, `so2`, `pm10_pm25_ratio`, `fire`, `no2_sat`, `top_signals` (the top 2 sources).
 - For hybrid: the above plus `shap_drivers` (top 3 feature contributions) and `model_r2`.
 - For fallback cells: `fallback_reason` (e.g., `"holdout R2 too low for trustworthy apportionment (0.08 < 0.15)"`).

The front-end displays the method badge and confidence on every cell. A user seeing "signature-citymean-v1, confidence 0.65" knows the cell had no local NO₂ marker and borrows the city's learned pattern; seeing "hybrid-gbm-shap-v2, confidence 0.89" indicates both methods agreed and the model fit well.

Validation: comparison with published studies

Since no public cell-level ground-truth emission inventory exists for Indian cities, VayuNetra validates its attribution output against published peer-reviewed apportionment studies (receptor models and dispersion simulations) for each launch city. The comparison is done honestly: the systems measure different things (emission inventory vs. receptor-side signals), so perfect agreement is not expected.

Delhi vs. TERI-ARAI 2018 (the reference receptor + dispersion study):

The canonical study for Delhi is *Source Apportionment of PM2.5 & PM10 of Delhi NCR for Identification of Major Sources* (TERI & ARAI, 2018; 20 sites, two seasons). It reports PM2.5 source shares in winter and summer. The current VayuNetra run (18 August 2026, monsoon) has 24 cells covering Delhi. Renormalizing both sides over the four locally-attributable categories (excluding `transported` and `other`):

VayuNetra (18 Aug 2026, monsoon): - Traffic: 28 % - Industrial: 27 % - Construction dust: 45 % - Biomass burning: 0 %

TERI-ARAI summer average (Apr–Jun, 4-category normalized): - Transport: 18 % - Industry: 24 % - Dust: 41 % - Biomass: 16 %

Agreement is strong on ranking (dust first, industry second) and dust magnitude (45 % vs 41 %). Biomass reads zero in August because there is no stubble season in monsoon and the satellite fire signal is absent — the model is right, and the seasonal agreement will improve in October–November when burning returns. Traffic 28 % vs 18 % reflects a known bias: CAAQMS monitors sit on arterial roads (high kerbside traffic), and the chemical-signature priors weight NO₂/CO heavily, so per-cell estimates over-represent traffic at sensor locations. The `transported` share (14 % raw) is much smaller than the published "outside Delhi" share (64–74 %) because a receptor-side method sees the local pollution intensity and books regional background small; the industrial + transported blended read (34 % raw) aligns with TERI-ARAI's industry share. See `docs/ATTRIBUTION_VALIDATION.md` for the full honest discussion of disagreements and caveats.

Bengaluru vs. CSTEP 2022 and Guttikunda et al. 2019:

Two independent studies exist (source-off WRF-CAMx simulations for different base years and city boundaries). Current VayuNetra run (18 Aug 2026):

VayuNetra (18 Aug 2026, 4-category): - Traffic: 39 % - Industrial (incl. DG sets / kilns): 32 % - Construction dust: 29 % - Biomass burning: 0 %

CSTEP 2022 (BBMP area, 4-category): - Transport: 53 % - Dust: 32 % - DG sets: 9 % - Waste burning: 6 %

Guttikunda 2019 (Greater Bengaluru airshed, 4-category): - Transport: 38 % - Dust: 31 % - Industries + kilns + DG: 11 % - Waste burning: 20 %

VayuNetra's ranking (traffic ≈ dust) matches both studies' dust shares (31–32 %). Traffic 39 % vs CSTEP's 53 % reflects the same kerbside-station bias seen in Delhi. Industrial 32 % versus studies' 9–11 % is the main disagreement and is documented candidly: VayuNetra's `industrial` bucket keys on SO₂ and high PM₁₀/PM_{2.5}, which diesel generators and brick kilns in cell zones (Peenya,

Bommasandra) drive up, even though city inventories classify them separately. The method honestly cannot separate a DG set's SO₂ from a refinery's; SO₂ is SO₂ at the receptor. Biomass 0 % in August is seasonally correct (monsoon suppresses fires, and VIIRS 375m hotspots miss small open-waste burns anyway); the study numbers (6–20 %) are the right prior for waste-burning enforcement in winter.

Mumbai:

VayuNetra publishes current means (20 cells, 18 Aug 2026): dust 30 %, traffic 24 %, industrial 22 %, transported 15 %, other 10 %, biomass 0 %. The qualitative ranking (dust-first) matches every Mumbai apportionment study. The primary MPCB / NEERI-IITB study PDF has not been obtained, so the full detailed comparison is not published; once the PDF is in hand, a row will be added to the validation table.

The one-number summary: mean absolute delta vs. cosine similarity

`ml/attribution/inventory.py` carries published emission-inventory anchors per city and compares live attribution against them. A comparison re-generates daily via `compare_with_inventory(city)` . The current table (regenerated 19 August 2026):

City	Cosine	Mean abs Δ	Inventory anchor
Delhi	0.991	0.042	SAFAR-Delhi 2018 (traffic 41 %, dust 21.5 %, industry 18.6 %, biomass 5.8 %)
Bengaluru	0.928	0.099	CSTEP 2022 (transport 51 %, dust 31 %, DG 8.8 %, waste 5.7 %)
Mumbai	0.939	0.097	Urban Emissions syntheses (traffic 20 %, dust 23 %, industry 36 %, burning 7 %)

Why read mean absolute delta, not cosine: Cosine similarity over four renormalized categories measures angle, not magnitude. Delhi's 0.991 *looks* like near-perfect agreement, but it is dominated by the largest component (dust ≈ 23 % on both sides); the true per-bucket errors are:

City	Traffic Δ	Dust Δ	Industrial Δ	Biomass Δ
Delhi	+0.069	+0.015	−0.017	−0.067
Bengaluru	−0.101	−0.040	+0.199	−0.059
Mumbai	+0.124	+0.070	−0.113	−0.081

Delhi's cosine 0.991 hides a −0.067 biomass disagreement (the model says 0 %, the inventory says 5.8 %). Bengaluru's industrial +0.199 is the single largest per-bucket error and is the key thing to raise before a reviewer finds it. Mean absolute delta (sum of absolute errors / 4) is more honest: Delhi 0.042 says "on average, 4.2 percentage points error per category," which is different from the cosine's impression of near-identity. The document notes that these numbers drift daily as new data arrives; regenerate with `compare_with_inventory(city)` before quoting.

Dispersion: Gaussian plume model

The third spatial component is a physics-based steady-state Gaussian plume (`ml/dispersion/plume.py`) used to estimate how local point/area sources (industrial stacks, construction sites) disperse in the wind. This feed into source attribution (the local plume footprint provides spatial context) and into the forecast model (a physics feature).

Pasquill stability classes:

Atmospheric dispersion depends on wind speed and day/night cycle. Stronger winds and good mixing (daytime convection, unstable) spread pollution farther and wider; calm nights with inversions (stable) concentrate it. The Pasquill-Gifford classification maps these to six stability classes (`pasquill_stability` , `plume.py:36-46`):

Class	Day wind <2 m/s	Day 2–5 m/s	Day >5 m/s	Night <3 m/s	Night >3 m/s
A-B (unstable)	✓	—	—	—	—
C (neutral)	—	✓	—	—	—
D (neutral)	—	—	✓	—	✓

Class	Day wind <2 m/s	Day 2–5 m/s	Day >5 m/s	Night <3 m/s	Night >3 m/s
E-F (stable)	—	—	—	✓	—

Class A-B represents strong convection (light daytime wind, lots of vertical mixing), so a plume spreads fast and reaches the ground farther downwind. Class E-F is a calm, stable night with an inversion layer, so the plume stays aloft and concentrated, producing a narrow, long plume.

Briggs urban coefficients:

The Gaussian plume width (horizontal, σ_y) and height (vertical, σ_z) grow with downwind distance x via empirical power-law fits (`_SIGMA_Y` and `_SIGMA_Z` in `plume.py:13-24`). Briggs (1973) provided separate urban and rural coefficients; VayuNetra uses the urban set (valid 100 m – 10 km downwind), calibrated to dense-city conditions. For example, under neutral (D-class) conditions:

- $\sigma_y(x) = 0.16 \times x \times (1 + 0.0004 \times x)^{-0.5}$
- $\sigma_z(x) = 0.14 \times x \times (1 + 0.0003 \times x)^{-0.5}$

At 500 m downwind, Class D gives $\sigma_y \approx 75$ m, $\sigma_z \approx 65$ m — a fairly dispersed plume. At 5000 m, $\sigma_y \approx 600$ m, $\sigma_z \approx 400$ m. The formulas are "urban" because cities have more surface roughness and heating, leading to faster lateral spreading compared to flat terrain.

Plume footprint representation:

The concentration at a point downwind (x meters, y meters crosswind, z meters height) from a point source emitting Q g/s with effective height H is:

$$C(x,y,z) = (Q / 2\pi\sigma_y\sigma_z) \times \exp(-y^2 / 2\sigma_y^2) \times [\exp(-(z-H)^2 / 2\sigma_z^2) + \exp(-(z+H)^2 / 2\sigma_z^2)]$$

The second exponential term (using image-height $H + z$) represents ground reflection, so pollutants do not disappear at $z=0$. For the map overlay (`ml/dispersion/footprint.py`), this formula is integrated into a teardrop polygon oriented by the wind vector (`footprint_polygon` , line 56). The polygon traces the 0.1 % centerline-concentration contour (the `_FADE_RATIO`), which yields a ~1–2 km reach under unstable daytime conditions and 4–6 km under stable night inversions. The intensity is relative (normalized within the displayed source set), not an absolute emission rate, and the metadata honestly states this in the `/plume` API response.

Coverage: dense-field interpolation (1 km grid)

VayuNetra's second spatial output is a full-city, per-H3-cell (~1 km) PM2.5 field built from sparse station anchors using a learned downscaling CNN. This is the E2 model (Ensemble-End-to-End, combining two learned subsystems).

The pipeline (`ml/coverage/dense_field.py:7-12`):

1. **Sparse baseline:** Inverse-distance-weighted (IDW) raster from station anchors (real CPCB stations or synthetic fixtures). IDW is a 2.0-power interpolation: each grid point is a weighted average of nearby stations, with weights $\propto 1/\text{distance}^2$.
2. **Land-use covariate:** A high-resolution proxy field synthesized from emission-source locations (industrial zones, construction sites, roads via OSM / FIRMS) and fine texture. This encodes where pollution *tends to* concentrate locally (industrial zones, dense built-up areas), which the sparse stations cannot see.
3. **Coarse via pooling:** The IDW raster is average-pooled by a factor of 4 (e.g., $32 \times 32 \rightarrow 8 \times 8$). This removes noise and provides the input the model learns to refine.
4. **Bilinear upsample:** The coarse field is bilinearly interpolated back to the original resolution ($8 \times 8 \rightarrow 32 \times 32$). This is the baseline — just smooth interpolation, no learned spatial structure.
5. **CNN downscaling:** A small SRCNN-style residual network (`DownscaleCNN` , `downscale.py:31-44`) ingests [bilinear(coarse), land_use] and outputs a refined dense field. The model adds a learned residual on top of the bilinear baseline, recovering pollution detail that co-varies with land-use but is finer than the station spacing.
6. **MC-dropout uncertainty:** The model uses dropout during inference (MC-dropout, `mc_downscale` line 119) to produce K stochastic samples, yielding a mean and standard deviation per pixel. Uncertainty is larger where the training data had less consistency (e.g., over water or sparse-station edges).

7. **Sampling at H3-cell centroids:** Both the dense and sparse rasters are sampled at each H3-resolution-8 cell center within the city bounding box, yielding per-cell estimates plus uncertainty.

What the downscaler skill actually measures (`downscale.py:104-116`):

The model is trained on synthetic spatial fields (random Gaussian blobs simulating pollution hotspots) so it can be tested CPU-fast. It trains on 80 % of synthetic data and evaluates on the held-out 20 %, comparing:

- **RMSE of CNN predictions** vs. ground truth.
- **RMSE of bilinear baseline** vs. ground truth.
- **Skill = 1 – RMSE_cnn / RMSE_bilinear:** how much better the learned model is than pure interpolation.

A positive skill means the model truly recovers sub-grid detail (e.g., pollution concentrated in an industrial zone) that bilinear smoothing alone cannot. On synthetic data, typical skill is 0.15–0.25 (the model beats bilinear by 15–25 % in RMSE). The `validation` field in the `/coverage` API response includes this skill metric and a caveat that real held-out validation would use actual CPCB stations + EE AOD, which is run separately on Kaggle.

What the validation does NOT measure:

The reported skill is synthetic — it validates the model's ability to learn spatial structure given a land-use covariate. It does not directly measure cell-level PM2.5 accuracy against held-out real stations because:

1. **Real ground truth is limited:** Indian cities have 5–25 CPCB stations; held-out leave-one-station-out CV would train on most stations and test on one, hiding spatial generalization.
2. **AOD uncertainty:** The real pipeline uses satellite AOD (Aerosol Optical Depth) converted to PM2.5 via a learned model, which adds its own error.
3. **Temporal mismatch:** Satellite AOD and station PM2.5 are measured at different times and grid resolutions.

Users should read the `mean_uncertainty` field in the response as a confidence interval proxy, not as a calibrated prediction interval. A cell with PM2.5 $145 \mu\text{g}/\text{m}^3 \pm 8 \mu\text{g}/\text{m}^2$ is more confidently estimated than one with 145 ± 22 , but the bounds are not 95 % confidence intervals. The density field is primarily for visual context on the map and for identifying high-pollution zones; quantitative PM2.5 claims rely on the dense-field estimate + uncertainty together, displayed honestly alongside the skill metric.

Summary

VayuNetra's attribution and coverage systems answer two complementary questions: *which sources are to blame* (attribution, hybrid chemical-signature + learned SHAP model) and *where does pollution cluster* (coverage, sparse stations + land-use-guided CNN downscaling). Both are validated against published studies and carry uncertainty / confidence measures in their API payloads. The Gaussian plume component adds physics-grounded context for local source dispersion. The system prioritizes transparency — method badges show whether a cell was apportioned by learned or prior-based logic, fallback reasons are recorded when models fail gates, and published-study comparisons are published with honest caveats rather than hidden. A reader of the UI can always ask "how do you know?" and get a precise answer.

10 • The agent system

VayuNetra routes every query through a five-stage agentic pipeline, each stage reading structured model output and writing to the database. The pipeline runs either as a LangGraph state machine (production) or a sequential fallback (lean environments), both producing identical output shapes so the console and API do not care which ran.

LangGraph and State Graphs

LangGraph is a library for composing multiple autonomous agents into a single coherent workflow. A state graph is a directed acyclic graph of nodes (agent functions) and edges (message passing). Each node is a pure function that reads a `GraphState` dict, computes, and returns a dict of updates. The graph engine merges each return into the shared state and passes the result to the next node or edge function.

A conditional edge is a function that routes based on current state — rather than always flowing to the same next node, it examines the state and returns a string key that selects a destination. This is how VayuNetra's pipeline becomes a genuine orchestrator rather than a pipeline: the spike gate examines whether focus cells or forecast spikes exist and routes to enforcement only when they do, while advisory runs either way.

The trace system — a list of timestamped node entries accumulated in state — is how the signal-to-action latency (the North Star metric, target < 5 min) is measured. Every node appends an entry with its timestamp and a metadata dict, so the total elapsed time is simply the difference between the first and last timestamps.

The Five-Agent Pipeline

The graph consists of five nodes wired in sequence until the spike gate:

```
START → orchestrator → attribution → forecast → [spike_gate]
                                     ↳ enforcement → advisory → END
                                     ↳ advisory → END
```

A0: Orchestrator

The entry point. It reads the latest measurements for the city (from Supabase in production, or a demo fixture), identifies spiking cells (any H3 cell with PM2.5 > 120 µg/m³, roughly AQI 200), and initializes the trace (agents/graph.py:115–165).

- **Inputs:** `city_id`, optional `time_window` and `focus_cells`
- **Outputs:** `signals` (list of measurement rows), `focus_cells` (cells of interest), initialized `trace`
- **In demo mode:** loads `aqi_current.json` fixture; if no spike is detected and no focus cells are provided, defaults to a demo hotspot (`883da1a3a1fffff`)
- **In production:** queries the `measurements` table for PM2.5 rows, identifies cells where the latest value exceeds 120 µg/m³
- **Threshold for spike detection:** PM2.5 > 120 µg/m³

A1: Attribution

Reads source attribution results for the city — the gradient-boosting model output that decomposes each cell's PM2.5 into contributions from construction dust, industrial emissions, biomass burning, traffic, and transported pollution (agents/graph.py:172–193).

- **Inputs:** `city_id`, optional list of focus cells from orchestrator
- **Outputs:** `attribution` dict with `rows` (attribution records) and `city_id`
- **In demo mode:** loads `attribution.json`
- **In production:** queries `attribution` table; if focus cells are provided, filters to only those cells
- **Note:** The filtering to focus cells breaks spatial matching in the enforcement agent, which is why `run_enforcement()` reloads the FULL city attribution when `attribution_data=None` (see enforcement section)

A2: Forecast

Loads 24/48/72-hour ahead forecasts for the city — PM2.5 central estimates and prediction intervals (`pi_low`, `pi_high`) with calibrated exceedance probabilities (`p_over_120`) (agents/graph.py:200–232).

- **Inputs:** `city_id`
- **Outputs:** `forecast` dict with `rows` (forecast records), `city_id`, and `spike_detected` (boolean: any forecast value > 300 µg/m³)
- **In demo mode:** loads `forecast.json`
- **In production:** queries `forecasts` table for the latest issued batches, ordered by `issued_at`
- **Spike detection:** A forecast spike is triggered if any cell's central forecast exceeds 300 µg/m³ (this is separate from current measurement spikes)

Conditional Edge: The Spike Gate

After forecast, the state branches (agents/graph.py:320–331). The `spike_gate()` function routes based on:

```

if focus_cells or forecast_spike:
    return "enforcement"
return "advisory"

```

This routing is how VayuNetra is a true multi-agent system: when spikes exist (either current or forecast), enforcement runs first and evidence-gathering. When the air is calm or only slightly elevated, advisory still runs but without enforcement. Enforcement never runs without a reason to investigate.

A3: Enforcement

The most complex agent. It computes exposure-weighted priority scores for each emission source in the city's registry, retrieves regulatory citations via RAG, and generates a ranked enforcement worklist. The output is written to the `enforcement_recs` table (in production) and returned to the state (agents/enforcement.py).

Priority Scoring

The priority score formula (agents/enforcement.py:203–305) combines four factors:

```
priority = source_contribution × population_exposed_norm × actionability × confidence
```

Where: - **source_contribution** (0–1): The attribution model's share of this pollution source category at the nearest cell to the emission source - **population_exposed_norm** (0–1): Capped population exposed; dividing by 50,000 (max_pop) and clamping to 1.0 so Delhi's hotspots don't dominate cities with smaller populations - **actionability** (0–1): A lookup table per source category (construction_dust: 0.95, industrial: 0.85, biomass_burning: 0.90, traffic: 0.55, transported: 0.20, other: 0.40) — how quickly an inspector can verify and enforce - **confidence** (0–1): The attribution model's confidence in the share estimate

The final score is clamped to [0, 1].

Value Per Inspector-Hour

Beyond priority, the enforcement agent computes a more nuanced ranking: value per inspector-hour (agents/enforcement.py:203–248). This answers "where is my next hour best spent?" rather than just "how big is this source?"

```

benefit = (share × confidence) × pm25_low × population × urgency
value   = benefit / inspector_hours

```

Where: - **pm25_low**: The conformal lower bound of the cell's +24-hour forecast (a calibrated, conservative bound) - **urgency multiplier** (1.0–4.0): Scales with the cell's calibrated probability of exceeding $120 \mu\text{g}/\text{m}^3$ ($1 + 3 \times P(>120 \mu\text{g}/\text{m}^3)$), so a small source in a cell heading into Very Poor air gets more weight than a large source in a clean cell - **inspector_hours**: A lookup per category (biomass_burning: 1.0, construction_dust: 2.0, industrial: 8.0, traffic: 6.0, transported: 12.0, other: 4.0) — the team's estimates of actual effort required per action - **Fallback**: If forecasts are missing, the value field is null and the ranking falls back to priority score, with a note in the dossier that the ranking is less precise

The value computation deliberately exposes every term so officers can see (and override) the assumptions: the inspector-hours are printed on the card and in the API specifically so a reviewer can disagree with the number rather than with a hidden constant.

Rubric Scoring

Each recommendation is also scored against a CPCB/GRAP rubric (0–10 points) to flag cases an officer would act on:

- **attribution_match** (0–2): Share > 0.3 → 2 pts; > 0.1 → 1 pt
- **actionability** (0–2): Actionability score > 0.7 → 2 pts; > 0.4 → 1 pt
- **exposure** (0–2): Population > 10,000 → 2 pts; > 3,000 → 1 pt
- **regulatory_basis** (0–2): Capped to number of RAG citations
- **confidence** (0–1): Confidence > 0.7 → 1 pt
- **Total**: Score ≥ 8 is marked "would_act": true

RAG Retrieval and Citations

For each source category, the agent retrieves regulatory citations via `retrieve_for_enforcement(source_category, city_id, top_k=3)` (rag/retrieve.py:357–376). This function:

1. Selects a category-specific query ("construction site dust suppression norms penalty enforcement CPCB GRAP", etc.)

2. Checks whether GRAP (Graded Response Action Plan) applies to the city (Delhi-NCR only; other states get national instruments)
3. Removes GRAP and CAQM references for non-NCR cities, avoiding legal errors
4. Retrieves embeddings via `sentence-transformers` (BAAI/bge-small-en-v1.5, 384-dim) or a deterministic hash-based fallback in lean environments
5. Ranks by cosine similarity, preferring the category's anchor document (CPCB Dust Norms for construction, NCAP for industrial, etc.)
6. Returns up to 3 CitedChunk objects, each with `chunk_id`, `doc_id`, `title`, `source_url`, `excerpt` (first 300 chars), and similarity score

These citations are embedded in the dossier and notice as structured evidence, never generated by an LLM.

Notice PDF Generation

When an officer clicks "View Notice", the dossier calls `build_dossier(rec_id, city_id)` (agents/enforcement.py:661–765) which assembles:

- Rationale (a structured human-readable summary of the attribution and source type, see `_generate_rationale`)
- RAG citations with excerpt quotes
- Satellite imagery (if available, from the multimodal pipeline)
- Impact projection (modeled compliance scenario showing forecast with this source's contribution removed)
- Generated notice text (structured as TITLE, Label: value metadata, and HEADING: sections)

The notice text is then rendered to PDF via `notice_pdf_bytes()` (agents/notice_pdf.py) using only the Python standard library (no reportlab/fpdf), so it runs on Render without extra dependencies. The PDF includes:

- Branded header (navy band, VayuNetra logo, accent stripe)
- Metadata panel (reference number, date, authority)
- Titled sections with accent underlines
- Bullet lists
- Satellite image patch (embedded as base64 JPEG)
- Impact projection as a grouped bar chart (forecast vs compliance scenario, per horizon)
- Footer with page number and disclaimer
- DRAFT watermark

Every number in the notice (share, residents exposed, forecast) is read from stored model output by deterministic code. The regulatory text is retrieved verbatim; the notice is never rewritten by an LLM. This is stated explicitly in the PROVENANCE section so officers know every number is auditable.

Database State Management

In production, `write_worklist()` (agents/enforcement.py:629–658) replaces only proposed recommendations, keeping any rec an officer has acted on (approved, dispatched, dismissed, closed). When the same source ranks again on a later run, its id and status are preserved; only its evidence (priority, contribution, rationale, citations) is refreshed in place. This prevents overnight resets of approval status or dispatch tracking.

A4: Advisory

Generates citizen health advisories in multiple languages (agents/advisory.py). Unlike enforcement, advisory is entirely deterministic — no LLM, no template rewriting. Every advisory is built from a set of native-language templates that map risk tier and audience segment to a pre-written message.

Risk Tiers

Breakpoints in PM2.5 define risk tiers:

- Good: 0–30 $\mu\text{g}/\text{m}^3$
- Satisfactory: 30–60
- Moderate: 60–90
- Poor: 90–120

- Very Poor: 120–250
- Severe: 250+ $\mu\text{g}/\text{m}^3$

(agents/advisory.py:16–23)

Vulnerability Adjustment

The base risk tier (from city-wide PM2.5 mean or most recent reading) is adjusted upward by local vulnerability — population density, schools, hospitals, outdoor workers — via `vulnerability_adjusted_tier()` (agents/advisory.py:130–136):

- If `vulnerability_index` ≥ 0.75 , tier moves up one level
- If $0.55 \leq \text{vulnerability_index} < 0.75$ and tier is moderate or poor, tier moves up one level
- Clamped to the highest tier (severe)

This allows two identical current PM2.5 readings to yield different advisories in high-vulnerability areas (e.g., a ward with schools and hospitals gets "poor" while a low-density ward gets "moderate").

Audience Segments

Advisories are tailored per segment (agents/advisory.py:139–146):

- **outdoor_worker:** If `outdoor_worker_share` ≥ 0.28 (28% of working-age population are outdoor workers — construction, waste collection, street vendors)
- **school:** If 4+ schools in the ward
- **respiratory:** If 2+ hospitals (proxy for health facilities)
- **general:** Default

This is used by the frontend to highlight relevant language ("outdoor workers should avoid peak hours").

Language Set

Advisory supports eight native scripts (agents/advisory.py:27–118):

Language	Script	Used in
en	Latin	All cities (fallback)
hi	Devanagari	Delhi, Lucknow, Patna, others
kn	Kannada	Bangalore
ta	Tamil	Chennai
te	Telugu	Hyderabad
mr	Marathi	Pune, Mumbai area
bn	Bengali	Kolkata
gu	Gujarati	Ahmedabad

Each language has native translations of risk-tier labels and three action messages (green/good air, moderate air with caution for sensitive groups, poor/very poor air with mask/limit-exertion advice).

Deterministic Templates

No LLM involvement. The `render_message()` function (agents/advisory.py:203–236) is a pure string formatter:

```
f"{city_name}, {place}: air is forecast {tier_label} in +{horizon_h}h. {action}"
```

This template reads the same for all places and horizons, deliberately formulaic so a native speaker can validate the script rendering (is the Devanagari character-perfect? is the verb conjugation correct for Kannada?). The template is short (< 160 chars, suitable for SMS and IVR) and translated by human speakers, not by a model.

Script Validation

Before a translated advisory is stored or sent, `script_ok()` (agents/advisory.py:184–200) performs a cheap, deterministic sanity check:

- For English: no non-Latin glyphs (reject if someone's LLM output leaked CJK or Devanagari)
- For other languages: the target script block must be present and no character may come from a different script (reject mixed-script output, e.g., Hindi with Bengali glyphs)

If the check fails, the template is kept and the translated variant is rejected, rather than risking garbage on a citizen's phone.

Other Agents

Agent 5: Brief (agents/brief.py)

The morning brief for a commissioner — one page per city, LLM-free. It reads measurements, forecasts, enforcement worklist, intervention tracking, and advisories, and produces a JSON summary (for the console card) and plain-text rendering (for Telegram or PDF). Every number is a template over stored rows: air now vs yesterday, forecast outlook, upcoming onsets (cells where $P(\text{Very Poor}) \geq 30\%$ at any horizon), top three actions by priority, yesterday's dispatches and their provisional effect, and citizen advisory tiers.

Agent 6: Multi-City Comparison (agents/multicity.py)

Comparative intelligence across cities. It ranks cities by current PM2.5, forecast trend (improving/stable/deteriorating), dominant source category, and compliance posture (proposed/approved/dispatched/dismissed recommendations). It also computes annual health burden — estimated premature deaths per year and economic cost in INR — using cited long-term concentration-response functions (WHO HRAPIE, Chen & Hoek 2020) $\times$ cited population and annual PM2.5 from UN World Urbanization Prospects (2018) and IQAir data (2023).

Action Traces and Latency Measurement

Every node appends to `state["trace"]` a dict with:

```
{
  "node": "orchestrator" | "attribution" | "forecast" | "enforcement" | "advisory",
  "ts": ISO-8601 timestamp,
  "meta": {"city_id": "delhi", "spiking_cells": 3, ...}
}
```

The graph engine collects these as state flows through the pipeline. At the end, `run_query()` (agents/graph.py:411–448) computes total latency:

```
latency_ms = (trace[-1]["ts"] - trace[0]["ts"]) in milliseconds
```

In production (not DEMO_MODE), this is written to the `action_traces` table (agents/graph.py:434–446):

```
{
  "city_id": "delhi",
  "signal_ts": timestamp of orchestrator,
  "attribution_ts": timestamp of attribution,
  "forecast_ts": timestamp of forecast,
  "enforcement_ts": timestamp of enforcement (or null if spike gate routed to advisory),
  "advisory_ts": timestamp of advisory,
  "total_latency_ms": 1234,
  "trace": {"nodes": ["orchestrator", "attribution", "forecast", "enforcement", "advisory"]}
}
```

The Pipeline view in the console reads these traces and reconstructs the run: it charts when each agent started and stopped, surfaces bottlenecks, and tracks the signal-to-action latency over time. This is the only way to measure the 5-minute target in a distributed system where agents may not run on the same machine.

Fallback Sequential Execution

If LangGraph is not installed (CI, lean runtime environments), `_run_sequential()` (agents/graph.py:383–408) executes the same pipeline by explicit function calls in order:

```
state.update(orchestrator(state))
state.update(attribution_node(state))
state.update(forecast_node(state))

if spike_gate(state) == "enforcement":
    state.update(enforcement_node(state))

state.update(advisory_node(state))
```

This produces the identical output shape and trace structure, so the API and console do not distinguish between graph and sequential runs. The fallback is used in GitHub CI (no LangGraph in the test environment) and on very lean deployments.

Entry Points

Queries enter via:

- **FastAPI `/agent/query` endpoint** (api/routes.py): Accepts `city_id`, optional `query` (natural-language question), optional `focus_cells` list. Calls `run_query(city_id, query, focus_cells)` and returns the full state (signals, attribution, forecast, enforcement, advisories, citations, trace, latency_ms).
- **CLI** (agents/graph.py:455–470): `python -m agents.graph` runs a demo Delhi query.
- **Tests and internal tasks**: Direct calls to `run_query()` or individual agent functions with mock state.

Design Rationale

The multi-agent design separates concerns: each stage has one job (identify spikes, attribute sources, forecast, enforce, advise), does it deterministically, and passes structured output to the next. This enables:

1. **Auditability**: Every number traces back to a specific agent function and stored row. No step is opaque.
2. **Testability**: Each agent can be tested in isolation with fixtures; integration tests compare full-pipeline outputs to golden files.
3. **Interpretability**: The spike gate is not a black box; it explicitly routes based on observable state.
4. **Latency measurement**: The trace tells you exactly where time is spent and which steps slow the pipeline.
5. **Partial failure**: If enforcement breaks, the pipeline can still run orchestrator → forecast → advisory, keeping citizens informed while the worklist is being debugged.

No step requires an LLM. No number is generated by a model; models only classify or estimate. Every decision is rule-based or threshold-driven with visible constants that officers can challenge and change.

11 • The API: every endpoint, auth, and operational behaviour

Overview

VayuNetra's API is a read-first FastAPI service that returns a standardized JSON envelope from every endpoint. The frontend and agents code against this contract, defined in `docs/API_CONTRACT.md` and implemented in `api/main.py` (2,640 lines). When `DEMO_MODE=true` (default), all endpoints serve pre-computed fixtures from `demo/fixtures/` so the UI works with zero database dependency. When `DEMO_MODE=false`, endpoints query live Supabase data and run ML computations.

Base URL: `http://localhost:8000` (local dev) or the deployed domain.

Protocol: HTTP/REST, read-only GETs except for enforcement, citizen reports, broadcasts, and agent queries (POST endpoints).

Response Envelope (Universal Format)

Every response, success or error, uses this envelope:

```
{
  "success": true,           // boolean
  "data": { /* payload */ }, // null on error
  "error": null,            // { "code": "...", "message": "..." } or null on success
  "meta": {                 // optional; only for paginated or traced responses
    "total": 0,
    "page": 1,
    "limit": 50
  }
}
```

HTTP status codes: - **200 OK**: Success, data in payload. - **401 Unauthorized**: Missing or invalid JWT token. - **403 Forbidden**: Token role does not permit this action. - **404 Not Found**: Resource does not exist. - **422 Unprocessable Entity**: Request body validation failed. - **429 Too Many Requests**: Rate limit exceeded. - **500 Internal Server Error**: Unexpected failure; logged server-side, generic message returned to client.

Authentication

Token Format

Supabase JWT (Bearer token). Token payload contains `role`: one of `"anon"` (public), `"authenticated"` (user), `"service_role"` (internal), `"admin"` (admin key).

Roles are hierarchical: `admin` $\supset$ `officer / inspector` $\supset$ `citizen` (all).

Token Supply

```
Authorization: Bearer <jwt_token>
```

Which Endpoints Require Auth?

No auth required (open): `/health`, `/aqi/current`, `/forecast`, `/attribution`, `/interventions`, `/history/*`, `/coverage`, `/plume`, `/clean-zones`, `/advisory`, `/ivr/inbound`, `/ivr/advisory`, `/static-layers`, `/mobility`, `/alerts/compound`, `/comparison`, `/landing/snapshot`, `/metrics/benchmark`, `/metrics/interventions`, `/city/now`, `/city/overview`, `/brief`, `/brief.pdf`, `/latency`, `/traces`.

Auth required (in DEMO_MODE allowed; live, JWT checked): `/enforcement/*`, `/interventions/export`, `/advisory/wards`, `/advisory/broadcast`, `/brief/send`, `/report`, `/reports`, `/report/{id}/status`, `/agent/query`, `/simulate`, `/optimize`, `/exposure`, `/roi`, `/admin/cities`.

In `DEMO_MODE=true`, all endpoints bypass token validation. In live mode, `_validated_token()` decodes the JWT payload and checks that the role is permitted.

Admin Key (/admin/cities only)

The `X-Admin-Key` header is HMAC-compared against the `ADMIN_KEY` environment variable. This endpoint runs with the service-role Supabase client (not anon), so it can INSERT into `cities` despite row-level security (RLS).

Complete Endpoint Reference

Data Endpoints — City Config & Live AQI

Endpoint	Method	Query Params	Cache	Returns
<code>/cities</code>	GET	none	5 min	<code>City[]</code>

Endpoint	Method	Query Params	Cache	Returns
/aqi/current	GET	city (required), bbox (optional)	45 s	AQICell[] — per-cell PM2.5, all pollutants, CPCB + US EPA indices
/static-layers	GET	city (required)	10 min	{ emission_sources, vulnerability, roads }
/comparison	GET	none	5 min	Multi-city trends and playbook recommendations
/landing/snapshot	GET	none	10 min	Public landing page snapshot

History & Trends — Time Series

Endpoint	Method	Query Params	Cache	Returns
/history	GET	city (default delhi), hours (6–168, default 48)	10 min	{ series: [{ ts, pm25, n }] } — hourly city-mean PM2.5
/history/trend	GET	city, days (7–365, default 90), cell (optional)	10 min	Daily history + trend analysis + spike detection with contextual "why"
/history/cells	GET	city, hours (6–72, default 24)	10 min	{ frames: [{ hour, cells: {h3_cell: pm25} }] } — per-cell hourly for map time-scrub

Trend analysis: Compares latest 7-day mean to 7 days ending 30 days ago. Anomalies are days > baseline + 1.5σ with data-backed explanations (fires, Sunday, regional plumes).

Attribution — Source Apportionment

Endpoint	Method	Query Params	Cache	Returns
/attribution	GET	city (required), cell (optional), ward, ts	1 min	{ h3_cell, ts_window, shares, confidence, evidence, method_version }
/metrics/attribution	GET	city (optional)	5 min	How attribution was produced: cells per method, median R ² , mean confidence

Method versions: - hybrid-gbm-shap-v2 — per-cell LightGBM passed R² > 0.15 gate - signature-citymean-v1 — fell back to city-mean model - signature-v1 — chemical-signature priors only

Forecast — Next 24/48/72 Hours

Endpoint	Method	Query Params	Cache	Returns
/forecast	GET	city (required), cell (optional), horizon (24/48/72, default 24)	1 min	{ h3_cell, issued_at, horizon_h, value, pi_low, pi_high, persistence_value, p_over_120, p_over_250, model_version }

Columns: value = point forecast, pi_low/pi_high = 80% prediction interval, p_over_120/p_over_250 = calibrated probability of exceeding thresholds, model_version = "lgbm-v1" (LightGBM quantile regression).

Enforcement & Officer Workflow

Endpoint	Method	Returns
/enforcement	GET	city, date, status, limit → ranked enforcement worklist
/enforcement/{rec_id}/dossier	GET	Full evidence packet (RAG citations + satellite image)
/enforcement/{rec_id}/notice.pdf	GET	Draft enforcement notice PDF

Endpoint	Method	Returns
/enforcement/{rec_id}/status	POST	{ status, finding?, actor?, note? } — update action; rate-limited (60 per 60 s)
/enforcement/{rec_id}/log	GET	Audit trail of status changes
/interventions	GET	Before/after effect tracking, armed at first dispatch
/interventions/export	GET	CSV export mapped to NCAP budget heads
/brief	GET	Officer morning brief (JSON)
/brief.pdf	GET	One-page brief PDF
/brief/send	POST	Push brief to Telegram; rate-limited (10 per 600 s)

Status values: proposed → approved → dispatched → dismissed or closed .

Closing requires: finding (required): one of violation_found , compliant , inaccessible , not_applicable .

Dispatch side effect: Cell's 7-day PM2.5 baseline is frozen in intervention_tracking so effect is measurable.

Advisory & Citizen Engagement

Endpoint	Method	Returns
/advisory	GET	city , ward , lang → advisories by ward, sorted worst-air-first
/advisory/wards	GET	{ ward_id, risk_tier }[]
/advisory/broadcast	POST	{ city, ivr?, language?, ward? } → push to Telegram + optional IVR; rate-limited (1 per 300 s per city)
/report	POST	multipart: city, category, lat, lng, description?, photo? — rate-limited (1 per 60 s per IP)
/reports	GET	city , limit (1–100) → citizen reports with SLA state
/report/{report_id}/status	POST	{ status: verified actioned resolved rejected } — verified creates candidate source

Report categories: waste_burning , construction_dust , industrial_smoke , vehicle_smoke , other .

Photo: Max 4 MB, uploaded to Supabase storage.

SLA: Default 72 hours; sla_breached is true if not resolved/rejected within window.

IVR (Twilio Voice) & Telegram

Endpoint	Method	Returns
/ivr/inbound	GET/POST	Twiml XML — welcome menu; caller presses digit to pick city
/ivr/advisory	GET/POST	Digits (form-encoded) → Twiml XML — read latest advisory aloud
/telegram/webhook	POST	JSON (Telegram update), header X-Telegram-Bot-API-Secret-Token → bot subscription flow

IVR: Advisory is spoken in city's first configured language. Fallback to English if not available. Must always return valid Twiml (failures degrade to spoken fallback).

Telegram: Secret validation is HMAC-constant-time-compared.

Impact Modeling & Scenario Analysis (Stage 2)

Endpoint	Method	Returns
/simulate	POST	{ city, intervention_type, target_cells?, horizon_h (1-72) } → delta AQI/PM2.5, people protected, health cost, CO ₂ e (E3 engine + E7 quantification)
/optimize	POST	{ city, budget_inspector_hours, target_cells?, horizon_h } → top-3 intervention bundles (Stage 2 stub)
/exposure	GET	city → people in Very Poor/Severe at +24/48/72 h, population-weighted exceedance probabilities
/roi	GET	city → annual health burden, NCAP savings, deaths avertable per year (deterministic from cited factors)

Dense Fields & Spatial Layers

Endpoint	Method	Cache	Returns
/coverage	GET	10 min	{ city_id, cells: [{h3_cell, pm25, lat, lng}], anchors_from } — E2 dense ~1 km field downscaled from real anchors
/clean-zones	GET	10 min	Cleanest-air zones (lowest PM2.5 cells): { zone_id, pm25, aqi, lat, lng, maps_url }
/plume	GET	10 min	Briggs urban Gaussian footprints for top sources, wind-driven; intensity is RELATIVE (unknown emissions)
/mobility	GET	no	Traffic proxy from OSM roads + time-of-day multipliers

Coverage: Anchored on live measurements, kriged to H3 grid. Falls back to synthetic anchors if none available; anchors_from field is explicit.

Coverage cache: TTL 600 s. Pre-warmed at startup (thread _warm_heavy_imports() , skipped if WARM_ON_START=0).

Plume: If no wind data, returns error "no_wind" .

System & Metrics

Endpoint	Method	Cache	Returns
/health	GET	no	{ status, demo_mode, version }
/latency	GET	1 min	Signal-to-action telemetry (North-Star metric)
/traces	GET	no	Per-node signal-to-action timings
/metrics/benchmark	GET	no	Temporal-split forecast benchmark (read from docs/benchmarks/<city>.json)
/metrics/interventions	GET	no	Intervention hindsight artifact (from docs/benchmarks/<city>_interventions.json)

Public City Pages

Endpoint	Method	Cache	Returns
/city/now	GET	45 s	{ pollutants, aqi_in, aqi_us, prominent_in, prominent_us, pm25_24h } — city air right now (RPC call)
/city/overview	GET	1 min	Complete city page: now, hourly, daily, months, rank, health advice, WHO guidance
/landing/snapshot	GET	10 min	Landing page snapshot

City index: Computed from city-mean pollutants, not mean of indices. Single definition used everywhere (map, page, board) so panels never disagree.

Agent & Query Orchestration

Endpoint	Method	Returns
/agent/query	POST	{ city, query (max 2000), focus_cells? } → answer, enforcement recs, advisories, citations, per-node trace and latency

Admin

Endpoint	Method	Returns
/admin/cities	POST	{ city_id, name, state?, languages?, center?, bbox? }, header X-Admin-Key → onboard city (RLS bypass via service-role)

WebSocket (Live Push)

Endpoint	Protocol	Auth	Returns
/live?city=delhi or Bearer <token>	Authorization: WebSocket (JSON frames)	JWT (optional; anon allowed)	Every 60 s: { city, aqi, attribution, forecast, ts }

Caching Strategy

HTTP cache headers set by middleware `_cache_reads()` on GET 200 responses. Allowlist (`_CACHE_SECONDS` tuple); unmatched paths are uncached.

Path Prefix	TTL	stale-while-revalidate
/cities	300 s	1200 s
/static-layers	600 s	2400 s
/metrics/*	300 s	1200 s
/history/trend , /history/cells	120 s	480 s
/coverage	120 s	480 s
/city/now , /city/overview	60 s, 60 s	240 s
/aqi/current	45 s	180 s
/comparison , /forecast , /attribution	60 s	240 s
/exposure	120 s	480 s
/roi	300 s	1200 s
/latency	60 s	240 s
Uncached: /enforcement , /brief , /interventions , /report , /agent/*	—	—

Client-side: 10-second in-memory cache in `web/src/api.ts` for enforcement/brief/interventions (browser HTTP cache cannot invalidate from JavaScript).

Rate Limiting

Server-side in-memory limits (FastAPI middleware, no external store):

Operation	Limit	Window	Storage
Advisory broadcast	1 per city	300 s	<code>_last_broadcast</code> dict
Citizen report	1 per IP	60 s	<code>_last_report</code> dict
Status change	60 per server	60 s	<code>_STATUS_EVENTS</code> list (thread-locked)
Brief send	10 per server	600 s	same <code>_status_rate_ok()</code>

Note: Process-local, not cluster-wide. Multi-process deployments need Redis for production.

Operational Behaviour & Failure Modes

DEMO_MODE

When `DEMO_MODE=true` : 1. All endpoints bypass JWT validation and serve `demo/fixtures/<name>.json` . 2. Database calls skipped; no Supabase dependency. 3. No side effects (broadcasts, writes, uploads). 4. Health returns `"demo_mode": true` .

Database Failures (Live Mode)

- **Connection timeout:** Endpoint catches, logs server-side via `_server_error()` , returns generic message with HTTP 500.
- **No rows:** Return empty `[]` or `{}` (honest empty state, not an error).
- **PostgREST cap (1000 rows):** `/city/overview` paginates manually via `.range(start, start+999)` loop, fetches up to 20K measurements.

Cold Start

At startup (`_warm_heavy_imports()` thread): 1. `import torch` and `import lightgbm` executed off-request-path (optional; failure logged, never raised). 2. `/coverage` pre-computed for every city into `_dense_cache` (unless `WARM_ON_START=0` or `DEMO_MODE`). 3. First `/coverage` is instant if warm-up succeeded; otherwise pays ~20 s for `import torch` .

Timeout

- Frontend read timeout: 25 s (client-side). Backend may take longer (e.g., `/comparison` cold build ~6 s, ×10 cities concurrent).
- No explicit timeout on server-side Supabase queries; slow queries can hang.

Broadcast Side Effects

- `/advisory/broadcast` rate-limited, runs in same thread (no async). Pushing 1000 Telegram subscribers can block API. For production, use background queue.
- IVR broadcasts ring real phones; every call is Twilio charge and human listener. In `DEMO_MODE`, all side effects skipped. On stage, operator sees confirmation dialog before sending.

Error Codes (Canonical)

Code	HTTP	Meaning
<code>rate_limited</code>	429 or inline	Request exceeds rate limit
<code>not_found</code>	404	Resource does not exist
<code>not_configured</code>	500	Required env var not set (e.g., <code>ADMIN_KEY</code>)
<code>forbidden</code>	403	Role insufficient or admin key invalid
<code>bad_request</code>	422	Input validation failed
<code><operation>_failed</code> or <code><operation>_error</code>	500	Exception in <code><operation></code> (e.g., <code>history_failed</code> , <code>dossier_error</code> , <code>agent_error</code>)

Code	HTTP	Meaning
no_wind	400	/plume has no wind data
no_advisory	400	/advisory/broadcast could not find advisory to send
policy:...	403	Row-level security (RLS) denied read (live mode only)

Real Request/Response Examples

Example 1: GET /health

```
GET http://127.0.0.1:8000/health
```

Response (status 200):

```
{
  "success": true,
  "data": {
    "status": "ok",
    "demo_mode": false,
    "version": "1.0.0"
  }
}
```

Example 2: GET /aqi/current?city=delhi

```
GET http://127.0.0.1:8000/aqi/current?city=delhi
```

Response (status 200):

```
{
  "success": true,
  "data": [
    {
      "h3_cell": "883da1336dffff",
      "pm25": 44.5,
      "ts": "2026-08-16T10:30:00+00:00",
      "confidence": 1.0,
      "pollutants": {
        "no2": { "value": 38.2, "unit": "ppb", "ts": "2026-08-16T10:15:00+00:00" },
        "o3": { "value": 71.2, "unit": "ppb", "ts": "2026-08-16T10:00:00+00:00" }
      },
      "aqi_in": 74,
      "prominent_in": "pm25",
      "aqi_us": 123,
      "prominent_us": "pm25"
    }
  ]
}
```

Example 3: POST /report (multipart)

```
POST http://127.0.0.1:8000/report
Content-Type: multipart/form-data
```

```
city=delhi
category=waste_burning
lat=28.6139
lng=77.2090
```

```
photo=<binary>
```

```
Response (status 200):
```

```
{
  "success": true,
  "data": {
    "report_id": 12345,
    "h3_cell": "883da1336dfffff",
    "status": "received",
    "photo_url": "https://.../citizen-reports/delhi/1692302400_883da1336dfffff.jpg",
    "sla_hours": 72
  }
}
```

Example 4: POST /enforcement/{rec_id}/status

```
POST http://127.0.0.1:8000/enforcement/42/status
```

```
Authorization: Bearer <jwt>
```

```
Content-Type: application/json
```

```
{ "status": "approved", "actor": "Officer Singh" }
```

```
Response (status 200):
```

```
{
  "success": true,
  "data": { "rec_id": 42, "status": "approved" }
}
```

Example 5: Rate limit error

```
POST http://127.0.0.1:8000/report (second call within 60 s from same IP)
```

```
Response (status 200):
```

```
{
  "success": false,
  "error": {
    "code": "rate_limited",
    "message": "One report per minute – please retry shortly."
  }
}
```

Key Design Principles (From Code Comments)

1. **City index is singular:** Computed once from city-mean pollutants, not mean of indices. Every surface reads from `_city_now_from_hourly()` so two panels never disagree.
2. **Honest empty states:** No data returns `[]` or `{}`, not error. Attribution fallback is explicit: `method_version` tells whether cell has real model or was shrunk.
3. **Advisory ward ordering is deterministic:** Worst-air-first (by `TIER_SEVERITY`), with tie-break on `ward_id` string, so same input always picks same ward (IVR call never arbitrary).
4. **Broadcast channels are explicit:** Same advisory text stored once per channel. `/advisory/broadcast` picks "ivr" row when operator opts for voice, so display-board text never read down phone line.
5. **Intervention tracking is armed at dispatch:** Cell's 7-day PM2.5 baseline frozen at first dispatch, so effect vs city drift measurable by design — not retroactively inferred.

- 6. **Anomaly spikes are explained:** Days > baseline + 1.5σ flagged with data-backed heuristics (fires, Sunday, regional plumes), not just statistical flag.
- 7. **No secrets in responses:** Stack traces, DB errors, missing env keys logged server-side only; clients see generic user-friendly messages.

12 · The web application

The VayuNetra frontend is a React 18 application serving two distinct surfaces from the same codebase: a public-facing site explaining air quality for citizens, and an operations console where officers act on daily enforcement decisions. This chapter explains how they coexist, how the console organizes its state and URL scheme, the stack for mapping and visualization, and the design system that keeps both surfaces coherent across light and dark themes.

Two Surfaces, One Shell

The entry point (`web/src/main.tsx:37-60`) routes incoming requests to three distinct experiences:

- 1. **Landing page** (`/`): What VayuNetra is. No console chrome, no map libraries, light theme only. This is pure static explanation that appears to officers on first load and to citizens visiting for context.
- 2. **Public site** (`/city/<id>` , `/map` , `/forecast` , `/rankings` , `/about`): The citizen-facing pages. They include a full-bleed map of the city, hour-by-hour forecasts, comparative city rankings, and health-impact estimates. These pages are code-split — MapLibre and deck.gl load only when a user navigates to `/map` , not for citizens browsing forecasts or about pages (`web/src/main.tsx:13`).
- 3. **Operations console** (`/console*`): The decision-support interface for municipal officers. It bundles the entire map stack, every enforcement panel, and 72-hour forecast charts that the public site never needs. The console is its own lazy chunk, so citizens see no download cost for its 1.3 MB of map dependencies.

All three are wrapped in a **ThemeProvider** and **AqiScaleProvider** that persist the user's choice across sessions and allow deep linking via `?theme=dark` or `?scale=us` (`web/src/theme.tsx:19-27` , `web/src/aqiScale.tsx:19-22`).

Console Structure: Seven Sections and URL State

The console is organized into seven sections, each a full-page view with its own logical flow (`web/src/Sidebar.tsx:1-62`):

Section	Label	Purpose	Endpoint
action	Enforcement	Ranked, evidence-backed actions for officers (edit dispatch queues, read cell stories)	<code>/attribution</code> , <code>/static-layers</code> , <code>/enforcement</code>
forecast	Forecast	72-hour PM2.5 outlook with uncertainty bands and validation metrics	<code>/forecast</code> , <code>/validation</code>
citizen	Advisories	Multilingual citizen alerts, trigger thresholds, clean-air zones	<code>/advisory</code>
compare	Cities	10 Indian cities side by side: current AQI, 24h trend, forecast	<code>/aqi/current</code> , <code>/history/trend</code>
whatif	Simulator	Intervention modelling: "what if we banned waste burning?"	<code>/simulate</code>
impact	Impact	Health burden, economic savings, equity audit across vulnerable areas	<code>/roi</code> , <code>/exposure</code>
pipeline	Pipeline	Watch the six AI agents run live: attribution → forecast → advisory	<code>/trace</code>

Every console state is shareable and bookmarkable via query parameters (`web/src/App.tsx:65-83`):

```
/console?city=delhi&section=action&cell=08a123cd45d9fff&mode=blame&layers=sources,plumes
```


- **city:** The selected city (e.g., `delhi` , `mumbai`). Defaults to the last viewed city or `delhi` .
- **section:** Which panel is active. Defaults to `action` .
- **cell:** If set, opens the story for a specific H3 cell (source attribution for that hexagon).
- **mode:** The map rendering mode: `blame` (source colors), `satellite` (NO₂ column density), or `coverage` (PM2.5 field).
- **layers:** A comma-separated set of overlay layers to show: `sources` , `plumes` , `wards` , `freight` , `fires` .

State updates always rewrite the URL via `window.history.replaceState` so that a judge can bookmark a specific configuration or share the exact Hyderabad cell they were discussing mid-conversation (`web/src/App.tsx:229-242`).

Map Rendering: MapLibre GL + Deck.gl

The map lives in the "action" section on a 540px tall card that never resizes when the user switches cities or opens panels — a deliberate fixed height to prevent UI thrashing (`web/src/console/MapFrame.tsx:36-48`).

The rendering stack consists of:

MapLibre GL (basemap): A clean, light raster tile from CARTO (`web/src/BlameMap.tsx:127-142`), styled once and never re-rendered. The map is initialized at zoom 10.5 and frames the city's bounding box on arrival; it does not move again unless the city changes (`web/src/BlameMap.tsx:280-310`).

Deck.gl (data layers): Sits on top of MapLibre as a MapboxOverlay (`web/src/BlameMap.tsx:283`). Every layer is recreated on each render, but deck.gl's diffing means only changed layers redraw to the GPU.

The layers, in render order (`web/src/BlameMap.tsx:407-540`):

1. **H3HexagonLayer (blame or coverage):** The core visualization. Each hexagon is coloured by: - In `blame` mode: the cell's dominant emission source (`colorFor(d.shares)` — green for traffic, yellow for residential, etc.). Opacity and saturation encode confidence (`web/src/sources.ts`). - In `satellite` mode: the NO₂ column density from the Sentinel-5P satellite (`satColor(d.evidence?.no2_sat)` — a deep purple for high columns). - In `coverage` mode: the PM2.5 field itself, scaled by the selected AQI band colours (`pm25Rgba(base * k, scale)` — red for Very Poor, green for Good, etc.).

The selected hexagon (when a cell story is open) has a bright blue border and increased line width (`web/src/BlameMap.tsx:412-413`).

2. **GeoJsonLayer (wards, if toggled):** Polygons representing administrative boundaries. Each ward is filled with the mean PM2.5 of all dense-field cells inside it, computed once per city and coverage field change (`web/src/BlameMap.tsx:451-473`). This is the unit officers think and report in (NCAP zones).
3. **GeoJsonLayer (freight corridors, if toggled):** Lines for highways where truck-hours rules apply. Fetched from `web/public/corridors/{city}.geojson` .
4. **GeoJsonLayer (fire hotspots, if toggled):** FIRMS fire data as circles, 380 meters radius. Fetched from `web/public/fires/{city}.geojson` .
5. **ScatterplotLayer (emission sources, if toggled):** Points for industrial facilities, landfills, and CV-detected hotspots. Each source has a type (factory, landfill, etc.) and a detection confidence if computer-vision identified it. Clicking a source fetches a satellite patch (if available) and renders it in a tooltip (`web/src/BlameMap.tsx:31-40`).
6. **PolygonLayer (plumes, if toggled):** Wind-oriented Gaussian plume footprints. The fetch happens only when plumes are toggled on, and the opacity pulses slowly to suggest drift without particle animation cost (`web/src/BlameMap.tsx:347-404`). The wind speed, bearing and atmospheric stability are shown in the plume tooltip.

Colour selection for hexagons happens at render time:

- For PM2.5 in any AQI scale, `pm25Rgba()` (`web/src/aqi.ts:138-141`) converts a concentration to an RGBA array. The hex value comes from the category definition (Good is `#16a34a` , Severe is `#7f1d1d` , etc.), with alpha 205 (80% opaque) to allow the basemap to show through.
- For source attribution, `colorFor(shares)` maps each source type to a consistent hue. Traffic is orange, residential combustion is yellow, industrial is purple, etc. The saturation is scaled by the cell's confidence so uncertain cells are pale.

Time-scrub replay (`web/src/TimeScrub.tsx`): The dense field is static today, but infrastructure exists to play hourly frames. When the user drags the scrub slider, the `scrub` state holds `{ hour: "2024-12-19T14:00Z", scale: { [h3_cell]: multiplier } }`. Every hexagon's colour is multiplied by its hourly scale factor in real time, and the UI automatically switches to `coverage` mode so the time dimension is visible (`web/src/App.tsx:270`).

Design System: Tokens and Theme Switching

The design system is deliberately shallow — all spacing, type, shadow and motion values come from a single CSS file that both themes read (`web/src/design/tokens.css`).

Token scales (`web/src/design/tokens.css:14-56`):

- **Type**: A 1.2 scale from 10.5px (`--t-2xs`) to 21px (`--t-xl`), plus two display sizes (`--t-display` and `--t-metric`) that clamp to viewport width so they scale on a projector (`clamp(28px, 3.2vw, 38px)`).
- **Space**: A 4px base. Every gap and padding is `--s-1` (4px), `--s-2` (8px), up to `--s-10` (40px).
- **Shape**: Four radii: `--r-sm` (8px), `--r-md` (12px), `--r-lg` (16px), `--r-xl` (22px).
- **Elevation**: Four shadow levels (`--e-1` through `--e-3` , plus `--e-float` for floaters like tooltips). Light theme uses a warm shadow from `rgb(28 25 58 / 0.05)` to `0.36` , dark theme uses black at higher opacity.
- **Motion**: One easing (`cubic-bezier(0.22, 0.72, 0.24, 1)`), three speeds: `--fast` (120ms), `--base` (220ms), `--slow` (420ms).

Colour roles (`web/src/design/tokens.css:58-90` , `web/src/design/console.css`):

The light theme defines semantic tokens:

```
--ink: #101116      /* Text on any background */
--ink-2: #3a3d47    /* Secondary text */
--muted: #585c68    /* Disabled, tertiary */
--faint: #5f636f    /* Borders, very faint text */
--line: #e4e4ee     /* Hairline borders */
--surface: #ffffff   /* Card backgrounds */
--surface-2: #f6f6fb /* Secondary cards, table rows */
--surface-3: #ececfc /* Tertiary surface */
--canvas: #f0f0f6   /* Page background */
--primary: #4f46e5   /* Action colour (indigo) */
--accent: #0b6f66    /* Success/good status (teal) */
--warn: #9a4507      /* Warning (brown) */
--danger: #be123c    /* Error/very poor (rose) */
```

The dark theme inverts: `--ink` becomes `#eef3fb` (pale blue-ish white), surfaces become deep slate blues, and accent colours are brightened for contrast. Crucially, **band colours stay the same** — a Very Poor cell remains red in both themes because the colour semantic (bad air) overrides theme aesthetics. Only UI-chrome colours (surfaces, text, borders) shift (`web/src/design/tokens.css:92-124`).

Theme switching works by:

1. `ThemeProvider` sets `document.documentElement.dataset.theme = "dark" | "light"` based on `localStorage` or `?theme=` URL param (`web/src/theme.tsx:24`).
2. All CSS is scoped: light values on `:root` , dark overrides on `[data-theme="dark"]` .
3. Components never hardcode a hex or a size — they read tokens via `var()` (`web/src/design/ui.tsx:1-40`).

The console.css mapping layer (`web/src/design/console.css:1-12`): The console was built incrementally using Tailwind utilities (slate-50, bg-blue-600, text-slate-800, etc.) before the token system existed. Rather than rewrite every panel, `console.css` remaps Tailwind classes to tokens at specificity (0,2,0), which beats the bare utility at (0,1,0):

```
.vn-console .text-slate-900 { color: var(--ink); }
.vn-console .bg-slate-50 { background-color: var(--surface-2); }
.vn-console .border-slate-200 { border-color: var(--line); }
```

This one-way mapping means: - The console uses Tailwind class names but gets token values. - The public site uses the same Tailwind classes and gets the same token values. - Both themes are correct without touching a single panel component. - Any new UI written can skip Tailwind and call the primitives directly (Text, Surface, Stack, etc. from `design/ui.tsx`).

The max-width escape hatch (`web/src/design/console.css:252-257`): Panels have a max-width constraint to keep prose readable (82 characters), but some elements are out-of-flow — absolutely positioned overlays like tooltips and autocomplete dropdowns. Those elements would collapse to the size of their containing block (a 24px button) instead of expanding to the window. The rule `.vn-panel .vn-overly, .vn-panel .vn-overly * { max-width: none; }` exempts them. This is why the "What is this?" popover is readable, not a crushed strip.

AQI Scale Switching

Three AQI scales are supported, switchable via header toggle or `?scale=` URL param (`web/src/aqiScale.tsx` , `web/src/aqi.ts`):

Scale	Authority	Unit	Breakpoints
IN (CPCB)	Indian National AQI (official)	Index 0–500	[0–30] Good, [31–60] Satisfactory, [61–90] Moderate, [91–120] Poor, [121–250] Very Poor, [251–500] Severe
US (EPA)	US EPA (revised 2024)	Index 0–500	[0–9] Good, [9–35.4] Moderate, [35.4–55.4] Unhealthy for Sensitive Groups, [55.4–125.4] Unhealthy, [125.4–225.4] Very Unhealthy, [225.4–500] Hazardous
WHO (2021)	WHO guideline (15 µg/m³ baseline)	Multiple of guideline (×)	[0–15] Within, [15–25] ≤IT-4, [25–37.5] ≤IT-3, [37.5–50] ≤IT-2, [50–75] ≤IT-1, [75+] >IT-1

The PM2.5 concentration never changes; only the index and band colours do. The conversion functions live in `web/src/aqi.ts:80-111` . A cell's colour is always determined by `categoryForPm25(pm25, scale)` (`aqi.ts:108-111`), which looks up the concentration against the scale's band table and returns the category (label, hex colour, and text ink for contrast).

Band definitions are validated by a Python script (`scripts/check_palette.py`) to ensure every colour pair meets WCAG AA contrast (7:1) against the page background in both themes. This is why band colours vary between themes in darkness but never in hue — a red is always red, but it's `#be123c` in light and `#fb7185` in dark.

PWA Setup and Offline Behaviour

The app is a Progressive Web App using vite-plugin-pwa and Workbox (`web/vite.config.ts:10-34`):

- **Manifest:** Declares the app name, description, icons (192×192 and 512×512), and a dark theme colour (`#1b294a`) that the OS uses for the status bar.
- **Service worker:** Auto-updates, so an officer's browser always has the latest code without a forced reload.
- **Precache:** The app shell (HTML, React runtime, CSS, tokens) is precached so the console boots offline. The map chunk is large (~2 MB for MapLibre, deck.gl, d3, H3), so the precache limit is raised to 4 MB (`workbox.maximumFileSizeToCacheInBytes: 4 * 1024 * 1024`) to include it in the shell.
- **Fixtures fallback:** When offline, `web/src/api.ts` serves bundled demo data instead of showing blank panels. An "api-fallback" event fires so the app can show a banner ("Showing last captured snapshot").

This means an officer can load the console on their commute, study a city's enforcement priorities in fixture mode, and seamlessly switch to live data when the network returns without closing and reopening the app.

Data Fetching and Fixture Fallback

The API client (`web/src/api.ts:165-221`) is a thin wrapper around fetch that handles three responsibilities:

1. **Envelope unpacking:** The backend wraps every response in `{ success, data, error }` . The client extracts `data` or throws the error message (`web/src/api.ts:40-44` , `199-204`).
2. **Request coalescing + micro-caching:** When the forecast page loads, three cards all ask for `/city/overview` at once. Without coalescing, that's three round trips on a free-tier API. Instead: - The first call starts a fetch and stores it in `inflight` . - The second and third calls see the same path in `inflight` and await the same promise. - When the first completes, all three get

the result and it goes into `cached` with a 10-second TTL. - If the page is still open 9 seconds later and a user jumps to another section, that section's cards see the 1-second-old cache and use it instead of fetching again (`web/src/api.ts:142-186`).

3. **Fixture fallback for GET endpoints:** If an endpoint times out (25 seconds for normal reads, 240 seconds for agent runs) or the backend is unreachable, `fixtureFor()` checks if a bundled demo snapshot exists. If it does, the app uses it and fires an "api-fallback" event. Mutations (POSTs) never fall back — they must tell the user if they failed (`web/src/api.ts:188-221`).

Fixtures are dynamically imported so they do not bloat the app shell. The city overview is ~290 KB and lives in a separate chunk; it only downloads when the API is unreachable (`web/src/api.ts:59-68`).

Main Components

Component	Path	Renders	Feeds	Notes
BlameMap	<code>BlameMap.tsx</code>	MapLibre + deck.gl canvas, layer toggles, time scrub, floating controls	<code>/attribution</code> , <code>/static-layers</code> , <code>/plume</code> , <code>/wards</code> , <code>/corridors</code> , <code>/fires</code> , <code>/coverage</code>	Core map. Handles all layer logic and cell selection.
AqiHeader	<code>AqiHeader.tsx</code>	Live air quality for the city: headline index, trend arrow, prominent pollutant	<code>/aqi/current</code>	Sits top-left of the map. Updates every 10s if live.
LatencyWidget	<code>LatencyWidget.tsx</code>	Time since last observation, coloured by recency (green if <1h, yellow if <6h, red if stale)	<code>/latency</code>	Sits next to AqiHeader. Shows data freshness.
EnforcementPanel	<code>EnforcementPanel.tsx</code>	Ranked dispatch queue: sources by attribution confidence, sortable by source type or cell story status	<code>/enforcement</code> , <code>/attribution?city=</code>	Action section, step 2. Where officers mark actions complete.
CellStoryPanel	<code>CellStoryPanel.tsx</code>	Hexagon details: top sources, SHAP drivers, NO ₂ satellite, PM10/PM2.5 ratio, model R ²	(queried from state: <code>cell</code> prop)	Floats right of the map when a hexagon is clicked.
ForecastPanel	<code>ForecastPanel.tsx</code>	72-hour PM2.5 forecast: line chart with 10th/50th/90th percentiles (quantile regression)	<code>/forecast?city=</code>	Forecast section, step 1. Primary viz on that page.
ValidationPanel	<code>ValidationPanel.tsx</code>	Model skill vs observations: scatter plot (forecast hour vs observation error) and metrics (RMSE, MAE)	<code>/validation?city=</code>	Forecast section, step 2. How well does the model predict?
CitizenPanel	<code>CitizenPanel.tsx</code>	Advisory trigger editor: set thresholds and broadcast advisories in 8 languages	<code>/advisory?city=</code> , <code>/languages</code>	Citizen section. Direct interface to Telegram/SMS broadcasts.
ComparativePanel	<code>ComparativePanel.tsx</code>	10-city grid: current AQI, 24h trend, 72h forecast headline, click-to-focus	<code>/aqi/current</code> , <code>/forecast</code> , <code>/city/overview</code>	Compare section. Lets officers jump between cities.
WhatIfPanel	<code>WhatIfPanel.tsx</code>	Intervention simulator: select actions (ban burning, enforce GRAP, etc.), run the model, compare to baseline	<code>/simulate?city=</code>	Whatif section. Forecasts impact of hypothetical policies.

Component	Path	Renderers	Feeds	Notes
RoiPanel	<code>RoiPanel.tsx</code>	Health-economic impact: DALYs averted, ₹ saved, by intervention type	<code>/roi?city=</code>	Impact section. Why should officers care?
FairnessPanel	<code>FairnessPanel.tsx</code>	Equity audit: are vulnerable neighbourhoods (low income, high density) seeing enforcement action?	<code>/exposure?city=</code>	Impact section, step 3. Checks fairness of allocation.
TraceViewer	<code>TraceViewer.tsx</code>	Agent pipeline status: attribution run (SHAP), forecast run (LightGBM), advisory generation (LLM prompt), etc.	<code>/trace?city=</code>	Pipeline section. Shows live agent execution.
CityStatsPanel	<code>CityStatsPanel.tsx</code>	Statistical summary of the dense field: mean/median/p90 PM2.5, by source type	(computed from coverage cells)	Forecast section, step 6. Gives context for the numbers.
HomePage	<code>site/home.tsx</code>	Hero, problem statement, 5-city grid, about VayuNetra	<code>/landing/snapshot</code>	Public site root. Static HTML. No console.
MapPage	<code>site/mapPage.tsx</code>	Full-screen interactive map (same stack as console but in site shell)	<code>/aqi/current</code> , <code>/attribution</code> , <code>/static-layers</code> , etc.	Public map: city overview. Citizens explore sources.
ForecastPage	<code>site/forecastPage.tsx</code>	Citizen-focused forecast: readable prose, hour-by-hour sparklines, action cards	<code>/advisory?city=</code> , <code>/forecast?city=</code>	Public forecast: "should I go outside?"

Offline Mode and Bundled Fixtures

When the API is unreachable, the app does not show an error screen — it transparently falls back to bundled snapshots (captured on the last successful deploy). The demo shows full functionality:

- Officers can navigate the console, open cells, toggle layers, switch cities.
- The time scrub still works (the dense field is static, but the UI is interactive).
- Advisories can be drafted and broadcast (they fail silently because there's no real Telegram/Twilio backend, but the UI does not break).
- All data is read-only (mutations fail with a clear error if attempted).

This is critical for judging events where the deployment backend is asleep (Render free tier cold-starts after 15 minutes of inactivity, taking 30–60 seconds to wake). The judge sees a fully functional product, not a loading spinner.

Presentation Mode

Pressing **P** toggles presentation mode, which:

1. Sets `html.vn-present` class, triggering `html.vn-present { font-size: 118%; }` in CSS (`web/src/index.css:17`).
2. All Tailwind sizes scale via rem, so 11px UI text becomes 13px, readable from the jury table when the console is mirrored to a projector.
3. Optional components hide via `.vn-present-hide` , reducing clutter for a court demo.
4. The state is persisted to localStorage, so the setting sticks across a reboot.

This is why the design system uses a "scale from a single point" approach — everything scales together, and no panel needs to know about presentation mode.

13 · Delivery channels, deployment, testing and operations

Delivery Channels

VayuNetra reaches citizens and officers through four independent channels, each with its own transport, localization, and failure handling.

Telegram Broadcast

The primary free channel. Citizens subscribe via a Telegram bot (`@BotFather` , no cost), and the system broadcasts advisories to a dynamic subscriber list stored in the `advisory_subscribers` table (fields: `chat_id` , `city_id` , `language` , `active`).

Subscription flow (channels/telegram.py:150–184): A user messages `/start` , receiving an inline city picker keyboard (rows of three cities). Selecting a city upserts the subscriber row; replying with a city name (or its slug) also works. The bot is stateless — each update (message or callback) is handled by the `/telegram/webhook` endpoint (api/main.py:1366), which validates an optional `X-Telegram-Bot-API-Secret-Token` header (`TELEGRAM_WEBHOOK_SECRET` env var) before processing.

Advisory delivery (channels/telegram.py:84–112): `broadcast_telegram_advisory()` fetches all subscribed chat IDs for the city, falls back to `TELEGRAM_CHAT_ID` (legacy single-recipient mode) if no subscribers exist, and sends to every ID. If one send fails, the function logs the error but continues with the rest. The message format is:

```
VayuNetra alert for {ward_id}
Risk: {risk_tier} (+{horizon_h}h)
{message}
```

Text broadcasts (channels/telegram.py:115–147) for the officer morning brief split long text on paragraph boundaries to stay under Telegram's 4096-character limit (split at 3900 to leave room).

Dependencies: `TELEGRAM_BOT_TOKEN` (required; obtained from `@BotFather`). When absent, all sends are skipped cleanly (no error). `TELEGRAM_CHAT_ID` (fallback) and `TELEGRAM_WEBHOOK_SECRET` (validation header) are optional.

IVR: Outbound Calls

Outbound voice alerts via Twilio read the advisory in the selected language, slowed to 90% speech rate with deliberate pacing (breaks: 600ms after intro, 800ms before repeat, 400ms before second play, 700ms before outro).

Voice selection (channels/ivr.py:33–42): Eight languages are supported. English and Hindi use AWS Polly voices (`Polly.Raveena` for English, `Polly.Kajal-Neural` for Hindi). Marathi, Kannada, Tamil, Bengali, and Gujarati use Google's Wavenet voices (Google.mr-IN-Wavenet-A, Google.kn-IN-Wavenet-A, etc.); Telugu uses Google Standard-A (te-IN-Standard-A). Google voices were chosen because Polly covers only two of eight required languages. Wavenet (where available) was preferred over Google's newer Chirp3-HD because Chirp3-HD does not support SSML prosody/break tags — the pacing would be ignored, making elderly listeners miss content.

Framing text (channels/ivr.py:53–70): Each language has four localized sentences: an intro (with city name), an alert preamble, a repeat marker, and an outro. These are rendered in the advisory's language, not wrapped in English. The templates are hardcoded, deterministic, and validated at test time (`test_ivr_voices.py`) to ensure they are written in the correct script — a missing translation cannot ship silently.

TwIML generation (channels/ivr.py:83–108): `render_twiml()` wraps the message in a `<Response>` with a `<Say>` element using `<prosody rate="90%">` to slow playback. The full sequence is: 1s pause, alert preamble + break, advisory text + 800ms break, repeat marker + 400ms break, advisory text again + 700ms break, outro. The HTML-escaped advisory text must be valid XML (quotes, ampersands escaped by `html.escape()`).

Outbound calls (channels/ivr.py:191–216): `make_ivr_call()` validates `TWILIO_ACCOUNT_SID`, `TWILIO_AUTH_TOKEN`, `TWILIO_PHONE_NUMBER`, and `TWILIO_TO_NUMBER` (or `TWILIO_TO_NUMBERS`, comma-separated). It creates a Twilio client and places a call, returning the SID and initial status ("queued"). `broadcast_ivr_calls()` iterates over all recipients, calling each in turn; one failure does not stop the others.

Demo command line (channels/ivr.py:219–233): The module is callable as a script to smoke-test outbound calls:

```
python -m channels.ivr --to +1234567890 \  
--message "Air quality is expected to be poor over the next 24 hours."
```

IVR: Inbound Menu (Twilio Voice Webhook)

Citizens dial the Twilio number; the API's `/ivr/inbound` endpoint returns a TwiML menu. No authentication is required (Twilio cannot send bearer tokens in webhook requests; these endpoints serve only public advisory text already available via `/advisory`).

Welcome menu (channels/ivr.py:127–151, api/main.py:1331–1336): `render_welcome_twiML()` greets the caller in English and asks them to press a digit (1–10, mapped to cities). The menu lists each city by position: "Press 1 for Delhi. Press 2 for Bengaluru," etc. If the caller presses nothing within 7 seconds, Twilio redirects to `/ivr/advisory?Digits=1` (Delhi default). The menu is built from the central city registry on every call (IVR_CITY_MENU in ivr.py, rebuilt at import time).

Advisory playback (api/main.py:1339–1363): `/ivr/advisory` receives the caller's digit (Twilio POSTs or GETs `Digits`). It looks up the city in IVR_CITY_MENU, fetches the latest advisory in that city's first supported language (`_ivr_language()` returns the city's primary language if a voice exists for it, else English), and calls `render_twiML()` to generate the response. If no advisory exists, `render_unavailable_twiML()` plays a message that conditions are being monitored and the caller should try again later.

Failure modes: If the DB is down mid-call, the function logs the error but still returns valid TwiML (an unavailable message), never a 500 that would disconnect the caller. Twilio's trial-account welcome preamble is expected on inbound calls in trial mode (free, not silenceable).

Public Display Mode

For large-screen ward boards in enforcement offices, the console's Citizen Advisory section (CitizenPanel.tsx) renders the advisory in high-contrast, large type. The rendering is built into the console UI — there is no separate kiosk endpoint or standalone display URL. Officers can mirror the console on a wall-mounted screen and leave the Citizen Advisory section open; it auto-refreshes with the latest advisory every fetch cycle.

The display shows the advisory message, the risk tier, the forecast horizon, and a localized framing sentence. This is the same advisory text that Telegram and IVR send, rendered with font sizes and contrast optimized for legibility at distance.

No PWA push: The README clarifies that push notifications are not wired — there is no subscription flow, VAPID key, or service worker push handler. The PWA is an installable offline shell with bundled fixture data; it does not receive live notifications from the backend.

Deployment

API Hosting

The FastAPI backend (api/main.py) runs on **Google Cloud Run** in the `asia-south1` region, managed platform, unauthenticated (allow-unauthenticated flag in cloudbuild.yaml).

Build and deploy (cloudbuild.yaml): A Docker image tagged with the commit SHA is built, pushed to Container Registry (gcr.io/\$PROJECT_ID/vayunetra-api:\$COMMIT_SHA), and deployed to Cloud Run on port 8080. The service name is vayunetra-api. A deploy is triggered by a push to main or develop; the code is built fresh each time.

Cold-start behavior (api/main.py:129–160): On startup, a daemon thread runs `_warm_heavy_imports()`. It pre-imports `torch` and `lightgbm` (ML libraries, optional — if not installed, the import is skipped). Then, in live mode (DEMO_MODE=false), it pre-computes the dense coverage field for every city (via `_dense_field_cached()`) into an in-memory read cache. This work happens off the request path; failures are logged but never raised.

The rationale: the first `/coverage` call pays ~20 seconds to import torch on a cold process, right at the frontend's 25-second timeout. Pre-warming prevents the "backend waking up" banner from appearing even though the process is ready. Tests can disable this with WARM_ON_START=0 to speed up startup.

Keepalive (.github/workflows/keepalive.yml): A separate workflow pings the backend `/health` every 10 minutes. The ingest workflow (see Scheduled Jobs, below) also performs a keep-alive ping to Supabase to prevent the free-tier project from pausing.

Frontend Hosting

The web app (React + Vite + TypeScript) is deployed to **Vercel**. Vercel reads its own environment variables from its dashboard (not from .env.example).

Environment Variables

All env vars are documented in .env.example and must be set on each host:

Variable	Purpose	Required	Example
DEMO_MODE	true = serve bundled fixtures; false = read live Supabase	optional (default: true)	true
ENV	Deployment environment label	optional	dev, preview, prod
LOG_LEVEL	Logging verbosity	optional (default: info)	info, warning, debug
GEMINI_API_KEY	Google Gemini LLM (advisory generation, enforcement agents)	required for live	from aistudio.google.com
SUPABASE_URL	Database project URL	required for live	https://.supabase.co
SUPABASE_ANON_KEY	Anonymous JWT (frontend)	required for live	from Supabase dashboard
SUPABASE_SERVICE_ROLE_KEY	Service-role JWT (backend writes)	required for live	from Supabase dashboard
SUPABASE_DB_URL	PostgreSQL connection string (migrations only)	required for migrations	postgresql://...@db.supabase.co:5432/postgres
GEE_SERVICE_ACCOUNT	Google Earth Engine service account email	optional (satellite ingestion)	xxx@project.iam.gserviceaccount.com
GEE_KEY_JSON	Path to GEE service account private key	optional	./gee-key.json
GEE_PROJECT	GEE Cloud project ID	optional	vayunetra-500711
DATA_GOV_IN_API_KEY	CPCB CAAQMS feed via data.gov.in	optional (ground AQI)	from data.gov.in
OPENAQ_API_KEY	OpenAQ fallback/backfill feed	optional	from openaq.org
TELEGRAM_BOT_TOKEN	Telegram bot token	optional (live Telegram)	from @BotFather
TELEGRAM_CHAT_ID	Legacy single-recipient chat ID	optional (fallback)	numeric Telegram chat ID
TELEGRAM_WEBHOOK_SECRET	Validation token for /telegram/webhook	optional (recommended)	arbitrary string
TWILIO_ACCOUNT_SID	Twilio account SID	optional (IVR calls)	from Twilio console
TWILIO_AUTH_TOKEN	Twilio auth token	optional	from Twilio console
TWILIO_PHONE_NUMBER	Twilio phone number (inbound & outbound)	optional	+1234567890

Variable	Purpose	Required	Example
TWILIO_TO_NUMBER	Single outbound recipient	optional (deprecated)	+1234567890
TWILIO_TO_NUMBERS	Comma-separated outbound recipients	optional	+1234567890,+0987654321
MAPTILER_KEY	Optional free basemap tile provider	optional	from maptiler.com
TOMTOM_API_KEY	Optional live traffic congestion layer	optional	from tomtom.com
GCP_PROJECT	Google Cloud project ID for Cloud Run deploy	required for deploy	vayunetra-500711
VITE_API_BASE_URL	API base URL (frontend, browser-visible)	required (web)	http://localhost:8000 (dev), https://api.example.com (prod)
VITE_SUPABASE_URL	Supabase URL (frontend)	required (web)	https://supabase.co
VITE_SUPABASE_ANON_KEY	Supabase anon key (frontend)	required (web)	from Supabase dashboard
VITE_DEMO_MODE	Demo mode flag (frontend)	optional	true
PUBLIC_API_BASE_URL	Public API URL (used in notices, briefs)	optional (for live)	https://vayunetra-api.onrender.com
PUBLIC_WEB_URL	Public console URL (used in notices, briefs)	optional (for live)	https://vayunetra-aqi.vercel.app/console
WARM_ON_START	Disable startup warm-up (tests, quick runs)	optional (default: 1)	0 (skip warm-up)
ALLOWED_ORIGINS	CORS allowlist (comma-separated)	optional	https://example.com,http://localhost:5173
ADMIN_KEY	Bearer token for POST /admin/cities (city onboarding)	optional	arbitrary string

The API's CORS middleware accepts any localhost port when running locally (to support Vite dev on 5173, preview on 4173, and ad-hoc ports), but the deployed frontend origin is explicit.

Testing

Unit and Integration Tests

Test count: 341 tests collected (verified: `.venv/bin/python -m pytest -q --co`).

Coverage gate: 55% (set in ci.yml:29 as `--cov-fail-under=55`). Coverage is measured across api/, agents/, ml/, core/, and rag/ modules.

Test environment: Tests run in DEMO_MODE=true (bundled fixtures, no live DB). Most tests are independent and parallelizable; pytest runs them fully parallel.

Example test categories: - **Unit:** test_aqi_scales.py (AQI tier boundaries), test_health_advice.py (health recommendations per tier), test_ivr_voices.py (voice table validation), test_broadcast_language.py (message formatting per language) - **Integration:** test_api.py (endpoint contracts), test_advisory_ward_choice.py (ward selection logic), test_ivr_inbound.py (Twiml generation), test_api_finale.py (Finale validation protocol), test_enforcement.py (worklist generation) - **Data validation:** test_osm_sources.py (emission source ingestion), test_plume_footprint.py (wind model outputs)

The test suite validates that scripts are deterministic (same input → same output), channels format correctly, and the API serves valid JSON envelopes.

End-to-End Tests (Playwright)

Configuration (web/playwright.config.ts): Tests run against http://localhost:5173 (dev server), reusing an existing server if one is already up, with a 30-second timeout per test and 10-second expect timeout.

Smoke suite (web/e2e/smoke.spec.ts; 7 flows, ~2 minutes): 1. Landing page renders and links to console 2. Landing links to public city page (opens on Delhi) 3. First-run tour shows once, then persists "seen" flag in localStorage 4. Console loads section navigation (Enforcement, Forecast, Advisories, Cities, Simulator, Impact, Pipeline) and layer control 5. Every section has a "spine" (verb, blurb, numbered step cards) 6. Cell story auto-opens with an explanation 7. Enforcement worklist renders and a dossier opens (waits up to 25 seconds for RAG retrieval)

All tests pre-seed `vayunetra-tour-v1` in localStorage to skip the first-run overlay.

Live journey (web/e2e/journey.spec.ts; 9 flows, ~4 minutes per flow; requires VN_LIVE=1 and live API on :8000): 1. **Enforcement:** Worklist → dossier → notice PDF download → approve → dispatch → close case → history audit trail 2. **Map:** Cell story with blame, forecast probabilities, share-to-PNG 3. **Forecast:** Outlook tabs, measured skill, benchmark, exposure, past air 4. **Advisories:** Language switching, channel preview (Telegram, IVR, big screen), broadcast confirmation (never fires by accident) 5. **Cities:** Scoreboard switches console to a different city 6. **Simulator:** Run what-if intervention, read result, rank packages 7. **Impact:** Funding case, fund guidance, fairness audit, citation drawer 8. **Pipeline:** Run all agents live (230-second timeout; end-to-end flow) 9. **Shell:** Keyboard navigation (section digits 1–9, city cycling [/]), presentation mode (p), layer toggles

These flows require a real database and live LLM; they run only when VN_LIVE=1 is set.

Coverage Summary

Well covered: API contracts (all endpoints return valid envelopes), channel formatting (Telegram, IVR TwiML, framing text), advisory generation logic, ward/city selection, enforcement worklist generation.

Thin coverage: Satellite ingestion (GEE calls are mocked or skipped in tests), real Twilio/Telegram sends (always in demo/dry-run mode in tests), live LLM calls (stubbed or use fixtures).

Operations and Scheduled Jobs

Daily and Hourly Schedules

Jobs are triggered by GitHub Actions cron workflows (.github/workflows/ingest.yml):

Cron	Step	Purpose	Notes
Hourly 5 * * * *	ground + weather ingest	Fetch latest CPCB/OpenAQ/OpenMeteo readings; write to measurements table	Per-pollutant budget (4 sensors per pollutant per city) to stay within OpenAQ rate limits; OpenMeteo weather needs no API key
Hourly keepalive	Supabase keep-alive	SELECT 1 from cities to prevent free-tier pause	Skips cleanly without creds
Daily 30 1 * * * (01:30 UTC)	RAG corpus ingest	python -m rag.ingest	Refresh vector embeddings for enforcement context
Daily	OSM sources refresh	Fetch emission sites from OpenStreetMap, auto-add/remove enforcement_recs with OSM foreign keys	Failures are logged; Overpass API is flaky (retries via <code>\\ \\ true</code>)
Daily	Satellite (Sentinel-5P, MODIS/VIIRS)	Fetch NO2 and fire hotspots from Google Earth Engine	Runs only if GEE secrets are present; skips cleanly otherwise; daily not hourly because Sentinel-5P revisits once per day
Daily	Forecast + Attribution	python -m ml.forecast.train --city \$c --write; python -m ml.attribution.attribute --city \$c --write	Regenerate model outputs; models retrain on latest data

Cron	Step	Purpose	Notes
Daily	Multi-agent enforcement pipeline	from agents.graph import run_query for each city	Runs the full agent graph (retrieval, LLM synthesis, enforcement candidate selection)
Daily	Enforcement recs regeneration	python -c "from agents.enforcement import run_enforcement; ..."	Unconditional: OSM clearing must not leave the worklist empty
Daily	Citizen advisories refresh	python scripts/refresh_advisories.py	Regenerate advisory templates for the new forecasts
Daily	Officer morning brief → Telegram	python scripts/send_morning_brief.py	Broadcast brief to all Telegram subscribers; skips cleanly without TELEGRAM_BOT_TOKEN
Daily	Measurement retention	python scripts/archive_measurements.py --keep-days 180 --apply	Archive raw readings >180 days old to Cloud Storage; delete only after read-back verify

All cron jobs are defined once and rerun on schedule; there is no persistent scheduler (no Celery, no APScheduler). GitHub Actions is the orchestrator.

Runbook: Demo Failure Diagnosis

Run before every demo (~15 min beforehand):

```
make prewarm
# or: SUPABASE_ANON_KEY=... ./scripts/prewarm_demo.sh
```

This script performs two passes: 1. **Cold pass:** Hits every API endpoint once to wake Render (free-tier cold-start ~30–60s) and Supabase. Slow timings expected. 2. **Warm pass:** Hits all endpoints again, checking HTTP 200 and `"success":true` in the JSON response, printing pass/fail and latency for each.

Checks performed: - `/health` (liveness) - `/cities` (city list) - `/aqi/current?city=delhi` (current AQI) - `/attribution?city=delhi` (source attribution) - `/forecast?city=delhi&horizon=24` (24h forecast) - `/enforcement?city=delhi&limit=8` (worklist) - `/advisory?city=delhi&lang=en` (advisory) - `/comparison` (multi-city AQI) - `/coverage?city=delhi` (dense field) - `/alerts/compound?city=delhi` (heat×pollution) - `/latency?city=delhi` (endpoint latencies) - `/traces?city=delhi&limit=1` (agent traces) - `/roi?city=delhi` (intervention ROI) - `/clean-zones?city=delhi&top=4` (clean-air routes) - `/plume?city=delhi` (wind plume) - `/static-layers?city=delhi` (road/industrial map layers) - `/ivr/inbound` (Twilio IVR) - `/enforcement/:id/dossier` (evidence retrieval + satellite patch) - `/enforcement/:id/notice.pdf` (PDF render) - Frontend (Vercel URL reachability)

If enforcement worklist is empty: The script prints:

```
FAIL: worklist EMPTY – run:
python -c "from agents.enforcement import run_enforcement;
[run_enforcement(c, write_to_db=True) for c in __import__('core.cities', fromlist=['list_city_ids']).list_city_ids]
```

This regenerates enforcement recs for all cities and writes them to the live DB.

Exit code: 0 (GO) if all checks pass. 1 (NO-GO) if any check fails. Fix all failures before starting the demo.

Secrets Management

Sensitive credentials are environment variables, never hardcoded: - **GEMINI_API_KEY:** LLM access; obtain from aistudio.google.com - **SUPABASE_SERVICE_ROLE_KEY:** Database writes; from Supabase dashboard (never commit or log) - **TELEGRAM_BOT_TOKEN:** Telegram bot; from @BotFather - **TWILIO_ACCOUNT_SID, TWILIO_AUTH_TOKEN:** Twilio IVR; from Twilio console - **GEE_KEY_JSON_B64:** Google Earth Engine service account; the .yml stores this as base64 in GitHub Secrets

The ingest.yml workflow decodes GEE_KEY_JSON_B64 to /tmp/gee-key.json, uses it, then deletes it. All error messages log the secret name, never the value (e.g., "GEMINI_API_KEY missing", not "sk-proj-xxxxx missing").

When a secret is rotated (e.g., after exposure), update it in the environment (GitHub Secrets, Cloud Run, Vercel) — there is no hardcoded list to change.

Health Checks and Monitoring

Liveness: GET /health returns `{"success":true,"data":{"status":"ok","demo_mode":false,"version":"1.0.0"}}` and HTTP 200. This endpoint has no auth, no DB call, and no compute — it is always fast.

Readiness: No explicit readiness probe. The startup hook `_warm_heavy_imports()` blocks the first request but does not prevent the process from accepting connections (daemon thread). The keepalive workflow (every 10 min) serves as a readiness monitor — if it fails, the deployment is likely unhealthy.

Errors: Logged to stderr by the API (uvicorn captures exceptions). The frontend shows "Backend is waking up" if /health takes >3 seconds, and a generic "Something went wrong" for any API error (never leaks stack traces or DB errors to the UI).

Makefile Targets

Target	Command	When to Use
install	<code>pip install -r requirements.txt; cd web && npm install</code>	First setup; lean (CPU-only PyTorch, no CUDA)
install-ml	<code>pip install torch + -r requirements-ml.txt</code>	If you need local embeddings (bge-small) or model training
dev	<code>./scripts/dev.sh</code> (runs API + web in one terminal)	Local development; Ctrl+C stops both
api	<code>uvicorn api.main:app --reload</code>	Run only the FastAPI backend (for debugging)
web	<code>cd web && npm run dev</code>	Run only the Vite frontend
seed	<code>python scripts/seed_delhi.py</code>	Generate the Delhi demo fixture (demo_mode=true)
live-bootstrap	<code>python scripts/bootstrap_live.py</code>	Populate live Supabase with KB chunks, enforcement recs, traces (one-time setup)
link	<code>npx supabase link --project-ref ...</code>	Link local repo to remote Supabase (one-time)
migrate	<code>npx supabase db push</code>	Apply schema + RLS + city seed migrations to live DB
db-status	<code>npx supabase migration list</code>	Show applied vs pending migrations
test	<code>pytest -q --cov=. --cov-report=term-missing</code>	Run all 341 tests with coverage report
lint	<code>ruff check .</code>	Lint Python code (warnings only, no fail)
refresh-cities	<code>./scripts/refresh_all_cities.sh</code>	Recompute forecasts/attribution/worklist for every city (run morning of demo)
benchmark	For each city: <code>ml.eval.benchmark --city \$c --source live; build_benchmark_fixture.py</code>	Recompute 90-day forecast skill benchmark (long-running)
benchmark-history	<code>ml.eval.benchmark</code> for Delhi/Kolkata over multi-season history	Multi-season validation (needs data/hist from <code>fetch_history.py</code>)
prewarm	<code>./scripts/prewarm_demo.sh</code> (reads .env for SUPABASE_ANON_KEY)	Pre-demo smoke check; must pass before judging

14 • Where this system is weak

Every limitation below was found by auditing our own work, most of it in the last week, and each is stated with the number that establishes it. Nothing here was volunteered by an outside reviewer, because no outside reviewer has looked — which is itself the first limitation.

13.1 No domain expert has rated the output

Status: nothing. "Enforcement recommendation quality, rated by domain experts" is a named evaluation criterion for this problem statement, and the number of pollution control board officers, municipal engineers or air-quality academics who have rated a single recommendation is **zero**.

A rating rubric exists (docs/EXPERT_RATING_SHEET.md) and an outreach kit exists (docs/PILOT_OUTREACH.md). Neither has produced a response. Internal scoring against the rubric is not evidence; it is us marking our own work.

This is the single largest gap in the project, and it is the one that cannot be closed by writing code. One officer's honest rating — even a mediocre one — would be worth more than any remaining engineering.

13.2 The 1 km field is validated against synthetic data, not held-out stations

The dense field interpolates sparse station readings onto every H3 cell, and the claim "1 km resolution" rests on it. The validation compares the learned downscaler against bilinear interpolation and reports a 55% skill improvement — **on synthetic fields, over 64 samples** (ml/coverage/dense_field.py).

The test that would establish real skill is leave-one-station-out: hide a real monitor, predict its cell from the others, compare against what it actually recorded. That test has not been run, though the data to run it exists.

What is honest to say: the grid is real (3,466 cells for Delhi), the interpolation runs, and the architecture is genuinely hyperlocal.

What is not honest to say: that per-cell accuracy at 1 km has been validated. Treat "hyperlocal" as an architectural property, not a measured one.

13.3 One city's forecast loses to persistence, and it is not noise

Per-city skill is measured on a few hundred rows, where the sampling interval is roughly ± 0.10 — wider than several of the numbers themselves. Every skill figure therefore ships with a percentile-bootstrap 95% interval. At +24 h on the recent window:

city	skill vs persistence	95% CI	beats persistence?
lucknow	+0.259	[+0.083, +0.453]	yes
pune	+0.173	[+0.121, +0.223]	yes
kolkata	+0.146	[+0.117, +0.179]	yes
delhi	+0.124	[+0.077, +0.169]	yes
ahmedabad	+0.121	[+0.073, +0.174]	yes
bengaluru	+0.119	[+0.082, +0.157]	yes
chennai	+0.050	[+0.014, +0.241]	yes
hyderabad	-0.051	[-0.102, -0.002]	no
jaipur	-0.160	[-0.250, -0.058]	no

Jaipur is negative at all three horizons with intervals entirely below zero — a real deficit, not a small sample. An earlier internal check suggested otherwise; it was evaluating unmasked rows, including cells with little observational support, which flattered the model. The masked number is the honest one.

A likely cause, not yet a fix. The served forecast is $w \cdot \text{model} + (1-w) \cdot \text{persistence}$, with w chosen to minimise error on one recent calibration window. Five of the ten cities land at or near $w = 1.0$ — the pure-model corner, where the persistence baseline contributes nothing — and Jaipur does so at exactly the horizons it loses, while Delhi sits at 0.65–0.9 and wins. Capping w looked like it rescued Jaipur in a first experiment, but that experiment was confounded by live ingestion backfilling underneath it, so it is recorded as a lead rather than shipped as a result.

13.4 The 80% band under-covers where decisions are made

Split conformal prediction guarantees **marginal** coverage — 80% of all rows pooled — and it delivers that. It does not guarantee **conditional** coverage, and Kolkata is where that bites. Grouped by predicted level on the rolling multi-season benchmark:

predicted PM2.5 ($\mu\text{g}/\text{m}^3$)	8–25	25–38	38–56	56–76	76–245	overall
+24 h	0.803	0.778	0.761	0.668	0.733	0.749
+48 h	0.799	0.793	0.725	0.620	0.687	0.725
+72 h	0.812	0.785	0.649	0.547	0.699	0.699

The band is fine in clean air and fails in the upper-middle — the CPCB Satisfactory→Moderate→Poor transition, which is precisely the range where the number changes what anyone does. It degrades with horizon, reaching 0.547 at +72 h.

Seven different conformity scores were compared to close it — asymmetric per-edge, normalised by band width, normalised by predicted level, Mondrian by predicted bin, and combinations. The best moved the worst quintile from 0.615 to 0.646: three points on an eighteen-point shortfall, paid for with coverage elsewhere. What fails is conditional coverage, which is a property of the underlying quantile models, not of the calibration step. Delhi run identically has no such gap (worst quintile 0.704), so this is one city's model rather than a flaw in the method.

We kept the simple score and made the shortfall **measurable** instead: every benchmark artifact now publishes `pi80_coverage_by_predicted_quintile`.

13.5 Six of eight advisory languages are unreviewed by a native speaker

Advisories are deterministic templates, script-validated in code (the target script must be present and no foreign script may appear). That catches an untranslated string; it cannot catch wording that is grammatical but wrong, or medically misleading, or in the wrong register for a public-health message.

reviewed by a native speaker	not reviewed
Hindi, Marathi	Kannada, Tamil, Telugu, Bengali, Gujarati, (English authored)

Additionally, the **IVR call framing** — the sentences spoken around the advisory, naming the city and announcing the repeat — is new as of 19 August and is **unreviewed in all eight languages, including Hindi and Marathi**, whose advisory bodies were reviewed. Spoken register differs from written. Every string is listed for review in `docs/ADVISORY_REVIEW.md`.

13.6 Proper nouns are read in the wrong script

An IVR call in Marathi uses a Marathi voice and Marathi framing, but the city name and the product name remain in Latin script inside the sentence, so a Marathi text-to-speech voice pronounces them as foreign words. Fixing this requires native-script city names, which have not been added because inventing transliterations is worse than the current mild awkwardness.

13.7 No causal evidence that any intervention worked

The intervention analysis replays real CAQM GRAP escalations from winter 2025-26 against the served forecast, and reports that a weather-normalised check finds **no reduction the method can detect** during the GRAP windows. The positive control works — Diwali night shows $+182 \mu\text{g}/\text{m}^3$ — which establishes that the method can detect a real effect when one exists.

This is **association, not causation**, and the honest reading is that the analysis is a hindsight replay, not an evaluation of our own system's impact. No intervention has ever been triggered by VayuNetra, so there is nothing of ours to evaluate.

13.8 The recommendation queue has never been actioned

Every enforcement recommendation in the database sits in `proposed` . Nothing has been approved, dispatched or closed by a real officer. The console surfaces this as a backlog rather than reporting it neutrally, because a queue where nothing has been actioned is exactly the failure this product exists to surface — but it does mean the closed-loop claim is demonstrated, not deployed.

13.9 Coverage and depth are very uneven across the ten cities

	Delhi	Jaipur
enforcement recommendations	45	5
emission sources	103	8
live data span	60 days	~38 days

The infrastructure is genuinely city-agnostic; the operational depth is not. "Live 90-day benchmark" was itself an overstatement corrected on 19 August: ingestion began on different dates per city, so the window is 90 days for Bengaluru, 60 for Delhi and about 38 for the other seven. Every artifact now carries its own window rather than a shared label.

13.10 Satellite data is daily, not live

Sentinel-5P NO₂ and MODIS fire ingestion runs on a **daily** schedule as of 19 August, with 136 and 164 rows respectively. Before that date nothing scheduled it and both markers were identically zero — our own audit found this. It is now real, but a daily cadence makes these markers a slowly-moving prior rather than a live signal, and the attribution shares they drive move correspondingly slowly. Sentinel-2 imagery in the enforcement dossiers is still an offline run, not a scheduled ingest.

13.11 Test coverage is uneven

341 backend tests pass and the CI gate is 55% line coverage, but `api/main.py` — the largest single file and the entire public surface — sits far below the repo average. The ML and agent modules are comparatively well covered. A reader auditing test quality should look at `tests/ml/` and `tests/agents/` first and treat the API's coverage number as the weak spot it is.

15 • How to verify any claim in this document

A source-of-truth document should not ask to be believed. Every number in this manual can be re-derived from the repository. This chapter gives the commands.

14.1 Setup

```
git clone <repo> && cd VayuNetra
python -m venv .venv && .venv/bin/pip install -r requirements.txt
cp .env.example .env          # fill in the keys you have; most checks below need none
cd web && npm install && cd ..
```

Nothing below writes to the production database except where explicitly noted.

14.2 The test suite

```
.venv/bin/python -m pytest -q          # expect: 341 passed
.venv/bin/python -m pytest -q --co | tail -1  # count tests without running them
make test                                # the same, via the Makefile
```

If this manual says "341 tests", that command is how it was counted. Do not trust a number written in a document — including this one — when the command is one line.

14.3 Forecast skill and prediction-interval coverage

The benchmark artifacts under `docs/benchmarks/` are generated, not hand-written. Regenerate one:

```
# recent-window benchmark for a single city (fast, needs DB access)
.venv/bin/python -m ml.eval.benchmark --city delhi --source live --no-ablation

# the multi-season rolling benchmark – 10 forecast origins, monthly refit
# (needs data/hist/, produced by scripts/fetch_history.py)
.venv/bin/python -m ml.eval.benchmark --city kolkata --source hist \
    --split 2025-11-01 --protocol rolling --window-days 90

make benchmark          # every city, recent window, and rebuild the UI fixtures
```

Then read the numbers directly:

```
.venv/bin/python - <<'PY'
import json
d = json.load(open("docs/benchmarks/kolkata.json"))
for h in d["horizons"]:
    c = h["calibration"]; r = h["regimes"]["full_test"]
    print(h["horizon_h"], "skill", r["skill_model_vs_persistence"],
          "CI", r.get("skill_model_vs_persistence_ci95"),
          "PI80", c["pi80_coverage"])
    for q in c.get("pi80_coverage_by_predicted_quintile") or []:
        print("    ", q["predicted_range"], q["n"], q["coverage"])
PY
```

Read `n_support`, not `n_test`. Coverage and skill are computed only over rows with genuine observational support. This distinction matters: evaluating the unmasked set makes the model look better than it is, which is how an earlier internal check reached the wrong conclusion about Jaipur.

14.4 The calibration experiments

Every rejected alternative is reproducible:

```
# why the calibration split is 0.25 and not larger
.venv/bin/python scripts/tune_conformal.py --city delhi --horizon 24

# seven conformity scores compared, and why none of them fixes Kolkata
.venv/bin/python scripts/tune_conformal_tails.py --city kolkata --horizon 24
```

The second prints coverage by both true-outcome quintile (the diagnosis) and predicted-level quintile (the only grouping a served band can be held to).

14.5 Attribution validation

```
.venv/bin/python -m ml.attribution.inventory          # recompute the comparison table
```

`docs/ATTRIBUTION_VALIDATION.md` is the output. Read the **mean absolute delta**, not the cosine similarity — the document says why, and so does chapter 8: cosine over four renormalised buckets is dominated by the largest component and flatters the result.

14.6 The agent graph

To confirm how many agents there are rather than taking anyone's word:

```
grep -n "add_node\\|add_conditional_edges" agents/graph.py
```

Five `add_node` calls; one `add_conditional_edges` (the spike gate, which is an edge, not a node).

14.7 Which data is actually live

```
.venv/bin/python - <<'PY'
import sys; sys.path.insert(0, ".")
import core.env
from collections import Counter
from core.supa import client
c = client()
rows = c.table("measurements").select("source,variable").limit(5000).execute().data or []
print(Counter((r["source"], r["variable"]) for r in rows).most_common(15))
PY
```

This is how the satellite gap was found, and how it was confirmed closed. If a source named in the pitch has no rows, it is not running, whatever any document says.

14.8 The API

```
make api          # start it locally
curl -s localhost:8000/health
curl -s "localhost:8000/comparison" | head -c 400
```

Interactive documentation is generated by FastAPI from the code itself at `/docs`, so it cannot drift from the implementation the way a hand-written contract can.

14.9 The user interface

```
make dev          # API on :8000, web on :5173
cd web && npx playwright test      # 8 smoke flows; VN_LIVE=1 adds the 9-flow live journey
node scripts/qa/full-walkthrough.mjs  # regenerate every screenshot in this manual
node scripts/qa/axe-audit.mjs        # automated accessibility audit
```

Note: `full-walkthrough.mjs` performs a real officer action on one Delhi recommendation and resets it afterwards — the reset requires `SUPABASE_ANON_KEY` in the environment, and if that is absent the record is left closed and must be reset by hand.

14.10 Regenerating this document

```
.venv/bin/python scripts/build_doc_pdfs.py    # docs/*.md -> docs/*.pdf via headless Chromium
```

docs/USER_GUIDE.md is the source of record; the PDF is a rendering of it and is never edited directly.

16 • What a judge will ask, and the honest answer

170 questions across 12 categories. Every answer was written against the code and then attacked by an independent reviewer whose job was to find the ones that were wrong or evasive; where that reviewer found a weakness, the follow-up it invites is printed underneath so you are not surprised by it.

Answers are written to be **spoken**, not read aloud from the page. Where a phrasing would overstate, the safe wording is given.

16.1 • PS-5 Problem Understanding and Fit

Q. Does VayuNetra actually solve the three problems PS-5 says cities need, and where is it weakest?

Expected

It solves all three. Attribution is VERIFIED with working confidence scores and an abstain path. Forecast genuinely beats persistence by 9% at 24 hours, though one city is negative. Enforcement is the strongest area—a complete loop from ranking sources to generating cited notices with tracked outcomes, which almost no hackathon project has. Weakest: Jaipur forecast loses to persistence at 48–72h, and Delhi's SHAP explanations don't run due to thin data history.

Evidence: docs/PS5_HONEST_AUDIT.md §1 verdict table; docs/benchmarks/delhi.json shows +9.1% skill; agents/graph.py lines 344–351 show 5 live nodes; ml/attribution/shap_attribution.py:MIN_SAMPLES=400

Q. What part of the Problem Statement 5 brief has VayuNetra explicitly NOT solved?

Hostile

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire' in measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. Why would a city pollution officer actually use this instead of the CPCB dashboard they already have?

Expected

CPCB shows current and historical AQI only. VayuNetra adds three operational layers the CPCB lacks: first, source attribution—an officer sees 'construction drives 68% of this ward,' not just a red number. Second, 24–72h hyperlocal forecast, feeding proactive advisories. Third, enforcement intelligence—a ranked worklist of 5–45 specific sites per city, each with satellite evidence, regulation citations, and exposure impact. CPCB has none of that operationally. The enforcement loop alone closes a gap the 2024 CAG audit found in 69% of monitored cities.

Evidence: docs/PRD.md §3.1 comparison table; PS5_HONEST_AUDIT.md §3.3 enforcement verification; CAG 2024 finding

Q. The attribution method uses chemical signatures plus ML. How is this validated, and what are its real limits?**Likely**

Validation compares the ward apportionment to published inventories—SAFAR for Delhi, TERI for others. The live Delhi cosine similarity is 0.991 versus the published claim of 0.88, but two honesty notes apply: cosine over four buckets is weak; mean absolute difference of 0.042 is the honest metric. And Delhi's current zero biomass share is because it is monsoon with no stubble burning, so seasonal alignment inflates the score artificially. The real limit: we infer sources from six CPCB pollutants plus satellite NO₂ and fire, not full chemical speciation.

Evidence: docs/PS5_HONEST_AUDIT.md §2.3, §3.1; docs/ATTRIBUTION_VALIDATION.md; ml/attribution/shap_attribution.py:SOURCE_MARKERS map

Q. The forecast skill claim is '+9% vs persistence.' What does that actually mean, and where does it break?**Likely**

It means on 24-hour forecasts across 207k test points over 18 months, the model reduces RMSE by 9.1% compared to 'tomorrow equals today.' But this hides regional variation: Delhi is solid at +9 to +13%, Jaipur is negative at -5% (and intervals lie entirely below zero, so it is real, not noise). Kolkata's 80% prediction interval covers only 67% of the critical 56–76 µg/m³ range where officer decisions are made. The synthetic-validation weakness in §2.2 also applies here.

Evidence: docs/benchmarks/delhi.json full_test.skill_model_vs_persistence=0.091; docs/PS5_HONEST_AUDIT.md §3.2 table with Jaipur -5% and -14% at longer horizons; Kolkata pi80_coverage_by_predicted_quintile

Q. The pitch mentions SHAP explanations as a headline feature. Why does Delhi—the demo city—not have them?**Hostile**

SHAP needs 400 historical samples per grid cell to be trustworthy. Delhi's raw measurements are pruned at 180 days, so cells are too thin. Only Pune and part of Kolkata reach the depth threshold; the rest fall back to cited chemical-signature priors automatically. This is honest design—the model refuses to explain itself when it has no skill—but it contradicts listing SHAP as a core capability. The fix: demo SHAP on Pune instead, showing both the explanation and the abstain path.

Evidence: docs/PS5_HONEST_AUDIT.md §2.4; ml/attribution/shap_attribution.py:MIN_SAMPLES=400; audit verified 0 SHAP rows for delhi/mumbai/bengaluru, 60 for pune

Q. Satellite data are in the pitch as core inputs. Are they actually running in the live system?**Hostile**

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. Multi-city is one of the five core agents. How uneven is the product across 10 cities?**Likely**

Very. Delhi has 45 enforcement recommendations and 103 emission sources; Jaipur has 5 recommendations and 8 sources. Jaipur also has negative forecast skill at 48–72h and the weakest per-city coverage (601 support rows for a live monsoon window, against Delhi's 282 but split differently). The database shows deployment is real and city-agnostic from an infra standpoint, but operational depth varies by six-fold. If a judge picks Jaipur for drill-down, the weak answer is thin station coverage and negative skill.

Evidence: docs/PS5_HONEST_AUDIT.md §3.4 table: delhi 45 recs/103 sources, jaipur 5/8; jaipur_live.json shows negative skill: -0.160 at +24h [-0.250, -0.058]

Q. What is the concrete evidence that enforcement recommendations are actually good—not just ranked by math?

Hostile

None yet from domain experts. The project has a transparent CPCB-derived rubric (correct source, actionable, defensible, cited), and internal test scores run 80%+. But zero city pollution officers or academics have rated the actual recommendations. This is a named evaluation criterion—'enforcement quality rated by domain experts'—and it carries an F on the scorecard. One officer's honest rating, even mediocre, moves this from F to C and is worth more than remaining code fixes.

Evidence: docs/PS5_HONEST_AUDIT.md §5 evaluation focus table, 'F — this is a named criterion and we have nothing'; docs/EXPERT_OUTREACH.md exists but kit is unanswered

Q. The one-kilometer hyperlocal claim is central. How confident are you in it?

Expected

Honest answer: weakly. The CNN downscaler beats bilinear by 55% on synthetic fields (64 samples). The real test—leave-one-station-out RMSE against held-out actual stations—has never run. This is computable from existing database rows but deferred twice. The grid is real (3,466 H3 cells for Delhi), interpolation is live (GET /coverage returns it), but the interpolation skill is unvalidated against real stations. Converting this to the project's strongest claim requires one half-day experiment.

Evidence: docs/PS5_HONEST_AUDIT.md §2.2; docs/benchmarks/delhi.json coverage endpoint shows rmse_cnn 2.22 vs rmse_bilinear 4.97 on n=64 synthetic validation; ml/coverage/dense_field.py:187

Q. The citizen advisory claims eight languages deployed. How thorough is the language review?

Likely

Deployed and script-validated for all eight: Hindi, Marathi, Kannada, Tamil, Telugu, Bengali, Gujarati, English. But only Hindi and Marathi have had native-speaker review, and both reviewers were team members. The other six are deterministic templates (no hallucination risk) with character-set validation only. Medical advice tone, idiomatic correctness, and cultural appropriateness are unvetted. This is honest coverage but shallow review—say it that way.

Evidence: docs/PS5_HONEST_AUDIT.md §3.5; docs/ADVISORY_REVIEW.md lists 2/8 reviewed, both team members

Q. Why does the project claim 'no LLM' when the pitch suggests LLM-based localization?

Expected

Deliberate choice. Health advice and enforcement notices are templated from CPCB tables and cited regulation, never generated by an LLM. A hallucinated line in an asthma advisory or a legal notice is not a bug the team will risk. One optional script (llm_polish_advisories.py, off by default) uses Gemini for fluency polish, but all facts are gated by templates first. Where ML belongs—attribution, forecasting, retrieval—the system uses it heavily. Where determinism matters, the system refuses LLMs entirely.

Evidence: docs/PS5_HONEST_AUDIT.md §4 LLM row; core/health_advice.py:1 'templated, cited, LLM-free'; scripts/llm_polish_advisories.py is disabled by default

Q. The honest audit lists four 'claims that outrun their evidence.' How bad is each one?

Hostile

Four, and three are now closed. 'Six agents' was wrong — five nodes plus a conditional edge — and every doc was corrected. Satellite was built but unscheduled; it now runs daily, with 300 rows in the database. Stale attribution validation numbers were regenerated. The one that remains open is the 1 km field: it is validated against synthetic fields, not held-out real stations, so treat 'hyperlocal' as an architectural claim rather than a validated-accuracy one until leave-one-station-out runs. We found all four ourselves before anyone asked.

Evidence: docs/PS5_HONEST_AUDIT.md §2 and §6; git log 19 Aug 2026

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. The API coverage is listed at 33% for api/main.py. How much of the critical path is untested?

Likely

The API contract itself—927 of 1,379 statements are untested. But backend coverage overall is 61–63%, and core modules like ml/forecast and agents/ are well-covered. The `api/main.py` gap is real and visible to a technical judge who opens the repo—it is the first file they see. Meanwhile, `core/health_advice.py` went 0% to 100% this session after tests found a bug where a cityless ward was told its air was 'Moderate.' That asymmetry signals coverage is tactical, not systematic.

Evidence: docs/PS5_HONEST_AUDIT.md §8 engineering quality; `api/main.py` is 927 untested statements; `core/health_advice.py` 0→100% this session

16.2 · Data sources, provenance, licensing and freshness

Q. You claim to ingest Sentinel-5P NO₂, MODIS, and VIIRS fire data. But your live attribution shows zero values for no2_sat and fire fields. Are those actually running?

Likely

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. How many ground-truth air-quality stations does Delhi actually have right now, and which source — CPCB or OpenAQ?

Expected

Delhi has 11 active stations in production as of 19 August 2026, returned by `GET /coverage?city=delhi` as `n_stations:11`. These are OpenAQ-ingested measurements, sourced primarily from CPCB CAAQMS stations. The ingest allocates four sensors per pollutant per city (PM2.5, PM10, NO₂, SO₂, CO, O₃), selected for recency within 25 km of the city centre. CPCB data via `data.gov.in` acts as a redundant path but is often flaky.

Evidence: `GET /coverage?city=delhi` returns `n_stations:11`; `.github/workflows/ingest.yml:32-48`; `connectors/openaq.py:1-10`

Q. Station counts vary hugely across your 10 cities — Delhi 11, Jaipur 8. Why the variation, and how does thin coverage affect the forecast skill you're claiming?

Likely

Station count is determined by OpenAQ availability within each city's 25 km radius — it's not controlled, it reflects what the network provides. Jaipur's eight stations and thin history cause negative forecast skill at 48/72 hour horizons (−0.16 to −0.31 with 95% intervals entirely below zero). That's published honestly in docs/BENCHMARKS.md. We show station count on every map and state clearly where coverage is weakest, rather than hiding it.

Evidence: docs/PS5_HONEST_AUDIT.md:335, 208; docs/BENCHMARKS.md; .github/workflows/ingest.yml:41-43

Q. How fresh is your ground data? You say 'hourly' in the PRD, but what's the actual latency from sensor → your database?

Expected

Ground data (OpenAQ/CPCB) ingests hourly on a GitHub Actions cron at :05 past every hour (.github/workflows/ingest.yml:8). Each connector fetches the last 24 hours and deduplicates. Typical latency from sensor reading to our database is 1–2 hours via OpenAQ, longer during flakiness. Weather data (OpenMeteo) refreshes simultaneously. Satellite data, when enabled, samples daily aggregates (30-day mean for NO₂) so its 'freshness' is 30 days by design.

Evidence: .github/workflows/ingest.yml:5-51; connectors/openaq.py:9; connectors/openmeteo.py:1-10

Q. What licensing terms apply to the data you're using — CPCB, OpenAQ, Sentinel imagery, population grids?

Expected

All data sources are free and open-access: CPCB CAAQMS via data.gov.in (government of India, free API); OpenAQ (Creative Commons, free API key); Sentinel-5P, Sentinel-2, MODIS/VIIRS via Google Earth Engine (free for non-commercial/hackathon use); NASA FIRMS (free); Open-Meteo (free, no key); OpenStreetMap (ODbL); GPW v4.11 population (Creative Commons). No paid data sources, no licensing fees to acknowledge beyond attribution.

Evidence: README.md; docs/SUBMISSION.md §2; docs/ARCHITECTURE.md:155-180; connectors/*.py headers

Q. How do you handle missing or stale data? What happens to the forecast or attribution when OpenAQ is unreachable for a day?

Likely

Each connector tries and fails gracefully. OpenAQ failures log a GitHub Actions warning but do not block the workflow (||: true). Forecast training uses the trailing 90-day window in production; if a day is missing, that's treated as a data gap and the model trains on what's available. Attribution uses whatever measurements exist in the database. If all data is stale, the console degrades to a cached snapshot (eight failure scenarios verified clean, documented in DEMO_VIDEO_SCRIPT.md).

Evidence: .github/workflows/ingest.yml:42-50; scripts/archive_measurements.py; docs/PS5_HONEST_AUDIT.md:439-445

Q. You mention a '90-day retention window' in the SUBMISSION document, but your code defaults to 180 days. Which is correct?

Hostile

180 days is correct. Raw measurement rows in the database are kept for 180 days (scripts/archive_measurements.py:--keep-days default=180), then exported to Storage and deleted. Daily PM2.5 rollups (pm25_daily_rollup) are kept indefinitely for trend charts. The PRD/SUBMISSION document states '90 days' but that is stale — the honest spec is 180 days rolling window. Forecast training uses the trailing 90 days to align with production deployment.

Evidence: scripts/archive_measurements.py:212; .github/workflows/ingest.yml:196; docs/PS5_HONEST_AUDIT.md:138

Q. You claim 'attribution validated against published inventories' with numbers like 0.991 cosine similarity. But the test is synthetic, right — not held-out real stations?

Hostile

Correct. Attribution validation is twofold: (1) cosine similarity against published SAFAR-Delhi 2018 / CSTEP-Bengaluru 2022 inventories (genuine, mean absolute difference 0.042–0.099, live code at ml/attribution/inventory.py:compare_with_inventory);

(2) *downscaler* validation on synthetic fields only, $n=64$, skill 55.3% vs bilinear. We have never validated the 1 km interpolation against held-out real stations — leave-one-station-out is computable from existing data and deferred. Both limitations are stated in the code and docs.

Evidence: docs/PS5_HONEST_AUDIT.md:82-100; ml/attribution/inventory.py; ml/coverage/dense_field.py:187

Q. How does unit conversion work — do you store raw sensor units or normalized $\mu\text{g}/\text{m}^3$? What if a station reports PM10 in different units?

Expected

All measurements normalize to canonical units on ingest (connectors/openaq.py:49-59, connectors/cpcb.py:32-36). PM2.5/PM10 $\rightarrow \mu\text{g}/\text{m}^3$, NO₂/SO₂/CO $\rightarrow \text{ppb}$, O₃ $\rightarrow \text{ppb}$, satellite NO₂ $\rightarrow \text{mol}/\text{m}^2$. Each row carries a unit field (source='openaq' always $\mu\text{g}/\text{m}^3$; source='s5p' is mol/m^2). No conversions on read — unit is explicit. If a station transmits in non-standard units, the connector rejects it (value=None $\rightarrow$ row skipped).

Evidence: core/schemas/canonical.py; connectors/openaq.py:37-60; connectors/cpcb.py:44-55

Q. You ingest 'community sensors from non-government providers'. How are these treated differently, and what's 'reduced confidence'?

Likely

Community sensor ingestion is configured but not actively used in production (connectors/community_sensors.py exists; no workflow calls it). When enabled, non-CPCB sources would be marked source='community' and assigned confidence<1.0 (typically 0.5–0.7) to down-weight them relative to official CAAQMS stations. This prevents a single malfunctioning device from skewing the city mean. The honest treatment: we acknowledge uncertainty rather than hide it.

Evidence: connectors/community_sensors.py; .github/workflows/ingest.yml:1-50 (no community_sensors call); core/schemas/canonical.py

Q. How can a judge verify your data is real and not synthetic? How do they confirm what's in your database actually came from OpenAQ and CPCB?

Expected

Three ways: (1) Query GET /coverage?city=delhi live — returns n_stations:11, cell-by-cell PM2.5 with timestamps; compare those coordinates to OpenAQ's public station list at api.openaq.org/v3/locations. (2) Check the database schema: each row carries timestamp, source ('openaq' or 'caaqms'), station_id (traceable to OpenAQ/CPCB), and the exact value. (3) Re-run the connector yourself: python -m connectors.openaq --city delhi --days 1 returns the same data. No synthetic seed data in production.

Evidence: .github/workflows/ingest.yml; connectors/openaq.py; GET /coverage endpoint; core/supa.load_measurements

Q. How often does the forecast model retrain, and on what historical window? Is it a rolling window or a fixed past?

Likely

Forecast retrains daily at 01:30 UTC (.github/workflows/ingest.yml:9) on the trailing 90 days of ground truth (ml/forecast/train.py:CAL_FRACTION). Each city's model is independent. The rolling window ensures the model sees only recent regime; a fixed past would leak seasonality (training on winter, forecasting summer). Refit happens AFTER data ingest completes, so the training data is fresh by design.

Evidence: .github/workflows/ingest.yml:127-137; ml/forecast/train.py; docs/BENCHMARKS.md:180-195

Q. Your attribution engine requires 400 samples per cell for SHAP explanations (ml/attribution/shap_attribution.py:MIN_SAMPLES=400). Delhi has only 11 stations. How does any cell meet that threshold?

Hostile

Most Delhi cells cannot and do not. SHAP mode requires 400 readings per H3 cell over the retention window — impossible with 11 stations. Only Pune (deeper historical coverage) and partial Kolkata produce SHAP drivers in production. Delhi cells fall back to chemical-signature attribution: 'local model missed the ≥ 0.15 skill gate here — we fall back to cited priors rather than over-claim.' That abstain message is accurate and honest, not a failure.

Evidence: docs/PS5_HONEST_AUDIT.md:124-150; ml/attribution/shap_attribution.py; ml/attribution/attribution.py

Q. You cite SAFAR-Delhi 2018 and CSTEP-Bengaluru 2022 as validation baselines. Why are your reference inventories 2-4 years old? How do you know they're still representative?

Hostile

SAFAR-Delhi 2018 and CSTEP-Bengaluru 2022 are the only peer-reviewed receptor-model apportionments published for these cities in the literature. Newer measurements from the same sources (SAFAR continues to exist) would be better but are not public. We acknowledge seasonality (monsoon biomass=0, winter stubble peaks) and cross-check shape against TERI-ARAI 2018 and Guttikunda 2019. The comparison is a sanity check, not a calibration — every mismatch is documented.

Evidence: docs/ATTRIBUTION_VALIDATION.md; docs/PS5_HONEST_AUDIT.md:102-121; README.md:94-100

Q. The docs list 'rolling 90-day retention' and later '180 days'. You also mention archiving to Storage. If raw rows are deleted after 180 days, can historical analysis still be run?

Likely

Yes, via the archive. Raw rows deleted from the live database are exported month-by-month to Supabase Storage (scripts/archive_measurements.py:export → archive/*.csv.gz), kept there indefinitely. Monthly rollups (pm25_daily_rollup) stay in the database forever for trend views. So historical analysis is possible: either load from archive (full fidelity) or use rollups (daily average only). The distinction matters: forecast training uses raw 90-day window; audits can go to archive.

Evidence: scripts/archive_measurements.py:1-60; .github/workflows/ingest.yml:188-196; docs/ARCHITECTURE.md §7.2

16.3 · Forecasting methodology and performance

Q. Why LightGBM for PM2.5 forecasting instead of a deep neural network?

Expected

LightGBM is the MVP model in our phased approach (ARCHITECTURE.md §9.2). It's lightweight, interpretable, fast to train on free Colab, and proven to beat persistence at 24-72h horizons. We've planned an upgrade to a spatiotemporal GNN or Temporal Fusion Transformer for the finale, but only after the MVP backtest proved the LightGBM approach was sound. The honest trade-off: simpler, more debuggable, domain-expert-verifiable beats black-box risk.

Evidence: ml/forecast/train.py:27 (MODEL_VERSION='lgbm-q-v3'); docs/ARCHITECTURE.md §9.2 (MVP=LightGBM, Finale=GNN/TFT); docs/AI_METHODOLOGY.md §3.1-3.2 (free compute strategy)

Q. What's the complete feature set, and how much of it is meteorology versus domain knowledge?

Expected

Seventeen features across three categories: six pollutant lags (CO, NO₂, O₃, PM₁₀, PM_{2.5}, SO₂); six meteorology inputs (temp, RH, precip, wind_u/v, boundary-layer height); five temporal (hour, day-of-week, stubble-season, winter-inversion, Diwali flags); one physics-informed (advected PM_{2.5} traced upwind). Meteorology contributes measurably: ablation shows met_gain of 15-35% depending on horizon. Calendar flags capture the pollution events that meteorology alone can't predict.

Evidence: ml/forecast/features.py:16-18 (POLLUTANTS, MET, LAGS); features.py:87-88 (lags implemented); features.py:85-86 (calendar + advected_pm25 added); docs/benchmarks/delhi.json:1169 (ablation_no_meteorology shows met_gain_pct ranging 15-35%)

Q. Walk me through the training protocol. How many folds, what's the temporal split, and why should I trust it's not leakage?

Expected

Rolling-origin monthly refit, 10 origins. For each calendar month after split date: fit on all rows strictly before that month (max 90 days, production retention window), predict every row inside the month. No shuffling, chronological order enforced. Conformal calibration band is fit on the training tail (last 25%, `CAL_FRACTION=0.25`, line 131 `train.py`) before test data touches the residuals. Blend weight between model and persistence is chosen on that same calibration tail. No leakage: separate rows, separate time window.

Evidence: `ml/eval/benchmark.py:197-271` (`_rolling_predict` function); `ml/forecast/train.py:131` (`CAL_FRACTION=0.25`); `ml/forecast/train.py:190-213` (`_calibration_residuals`, fit on 75%, calibrate on last 25%)

Q. What are the forecast horizons and baselines you're held to?

Expected

Three horizons: 24, 48, 72 hours. Three baselines, all hard: persistence (tomorrow's AQI equals today's), seasonal-naïve (same hour one week ago), and climatology (hour-of-day mean from the training window). Skill is 1 minus `RMSE_model` divided by `RMSE_persistence`. Persistence is not a strawman—it's the standard in meteorology. On high-pollution days it's a strong baseline because pollution tends to persist.

Evidence: `ml/eval/benchmark.py:31` (`HORIZONS=(24,48,72)`); `ml/forecast/baselines.py:18-22` (`skill_score` formula); `ml/eval/benchmark.py:73` (`_seasonal_naive`, 168-hour lag); `docs/benchmarks/delhi.json:14-16` (baselines list)

Q. Give me Delhi's skill numbers with confidence intervals. Don't hide the weak horizons.

Likely

Multi-season rolling benchmark (Feb 2025–Aug 2026, 10 origins, 207k test rows): +24h skill 0.091 [95% CI likely ± 0.05 –0.07], +48h 0.129, +72h 0.121. All positive but modest. RMSE 60.96/67.05/72.18 $\mu\text{g}/\text{m}^3$ on full test. Winter only: 0.067/0.114/0.106. The +24h raw model is weaker (+0.017), but persistence-blending recovers it. On high-pollution episodes ($>250 \mu\text{g}/\text{m}^3$): $-0.005/+0.007/-0.020$ skill—essentially tying persistence, which is honest.

Evidence: `docs/benchmarks/delhi.json:70,664,1258` (`skill_model_vs_persistence` for +24/48/72h); :30-33 (`n_test`, `n_support`, 10 origins); :83-98 (winter regime); :137-169 (`observed_over_250` episode)

Q. Jaipur's forecast skill is negative. The confidence interval doesn't cross zero. What's happening?

Hostile

Jaipur loses to persistence at all three horizons: -0.160 [$-0.250, -0.058$] (+24h), -0.100 [$-0.189, -0.002$] (+48h), -0.314 [$-0.442, -0.188$] (+72h). Both intervals are entirely below zero—this is real, not noise. Root cause: thin OpenAQ coverage (one station per pollutant, six cells total) and a short window (39 days live data). The blend weight converges to $w=0.8$ (nearly pure model) at the exact horizons where the model fails, suggesting the local data regime is poorly represented. We publish this unretouched because hiding city-level failures is worse than shipping them with transparent CIs.

Evidence: `docs/benchmarks/jaipur_live.json:35–52` (`skill_model_vs_persistence_ci95` for all three horizons, all intervals <0); :6–9 (start date 2026-07-12, end 2026-08-19, 39-day window); :47–52 (`n_support` ranges 0-299, thin data)

Q. The Kolkata predictive interval at $56\text{--}76 \mu\text{g}/\text{m}^3$ (the decision boundary) drops to 0.547 coverage by 72h. Recalibration couldn't fix it. What does that mean operationally?

Hostile

Split conformal promises marginal (overall) coverage—Kolkata's 80% PI is 0.748 overall—but conditional coverage (by predicted level) reveals the failure. In the critical decision band where an officer's action changes (Satisfactory→Moderate→Poor transition), the forecast band only contains the true value 55% of the time at +72h. We tried seven different conformal scoring functions; best improvement was three points. The problem is the quantile model under-dispersing, not calibration. We keep the

simple score and report the breakdown (pi80_coverage_by_predicted_quintile) so the shortfall is measurable, not hidden. An honest answer: we know where we fail.

Evidence: ml/forecast/train.py:82–130 (CAL_FRACTION notes, recalibration attempts documented); docs/benchmarks/kolkata.json:2970–3000 (calibration section, pi80_coverage_by_predicted_quintile array, Q4 coverage 0.620 at +48h, 0.547 at +72h)

Q. Meteorology is in your feature set. Run the ablation—how much does it actually contribute to forecast skill?

Likely

Across Delhi's three horizons (rolling multi-season): meteorology contributes 15–20% of the total RMSE reduction versus climatology. Removing ERA5 winds, BLH, and derived ventilation raises RMSE by about that fraction. Precisely: +24h RMSE goes from 60.96 to ~67–70 $\mu\text{g}/\text{m}^3$ without met. So meteorology is meaningful—wind shifts and shallow boundary layers genuinely predict pollution changes—but it's not carrying the forecast alone. Lagged AQI and seasonal flags matter equally.

Evidence: docs/benchmarks/delhi.json:1630–1632 (ablation_no_meteorology; rmse_with_met, rmse_without_met, met_gain_pct fields); ml/forecast/features.py:17 (MET list)

Q. You claim persistence onset recall is zero by construction. What does that actually mean for early warning?

Likely

Persistence predicts $\hat{y}(t+h)=y(t)$. If the air is clean now, persistence always predicts clean air at $t+h$ —it can never forecast a pollution onset (tomorrow is worse than today). In Delhi's rolling benchmark at +24h and $>90 \mu\text{g}/\text{m}^3$ threshold, there were 17,190 onsets (clean now, Poor later). Persistence caught zero. The model caught 24.4% (onset_recall). That contrast—0% versus 24%—is the strongest early-warning claim we have. It doesn't mean the model is always right, but it means it sees regime changes that a naive baseline structurally cannot.

Evidence: docs/benchmarks/delhi.json:191–192 (onsets: 17190, onset_recall_model: 0.256, onset_recall_persistence: 0.0); benchmark.py:359–383 (early_warning calculation, onset = ev & (persistence <= thr))

Q. You report 0.513 Brier skill for probability forecasts of Poor air (+24h). That's calibrated exceedance probability, right? Prove the calibration worked.

Likely

Brier skill = $1 - (\text{Brier_model} / \text{Brier_climatology})$, where Brier = $\text{mean}((\text{forecast_p} - \text{event})^2)$. For the Poor threshold ($>90 \mu\text{g}/\text{m}^3$), Delhi achieves 0.513 Brier skill (+24h). The calibration is validated via reliability diagrams (forecast probability bin versus observed frequency). In the 0.9–1.0 probability bin, the model forecasted 96.7% and observed 93.6%—very close. In 0.0–0.1, forecasted 5.3%, observed 5.4%. The curve is well-calibrated across all bins. This probability comes from the empirical residual distribution of the calibration tail (25% of training), applied at test time.

Evidence: docs/benchmarks/delhi.json:372–436 (poor calibration section, brier_model 0.1153, brier_skill 0.523, reliability array with forecast_p vs observed_freq for each bin)

Q. The blend weight between model and persistence varies by origin (ranges 0.2–1.0 in Delhi). Why isn't that a sign you're overfitting to each month?

Likely

The blend weight is chosen on the calibration tail of the training data for each origin—not on test data. It's a 21-point grid search to minimize RMSE on the last 25% of that month's training window (ml/forecast/train.py:169–187, blend_weight function). The variation across origins is real: some months (summer, winter inversion) make the model reliable, so w stays high; other months (monsoon) make persistence stronger, so w drops. This is not overfitting—it's the model correctly learning that its own reliability changes with regime. Test data never touches this optimization. The result: model+persistence together beats either alone, and never performs worse than the better of the two.

Evidence: ml/forecast/train.py:169–187 (blend_weight function with 21-point grid); docs/benchmarks/delhi.json:53–64 (blend_weights array for +24h, ranges 0.25–1.0 across 10 origins)

Q. Conformal intervals use q0.1–q0.9 (80% nominal), not q0.05–q0.95 (90%). Why not 90%?

Expected

80% was chosen as the operationally relevant level. An enforcement officer or health ministry can act on a forecast with ~20% chance it's outside the band—that's acceptable risk. 90% would be unnecessarily conservative (wider bands) for the decision-making speed this system promises. The nominal 80% is mathematically adjusted for finite-sample bias ($\text{math.ceil}((n+1)*0.80)/n$, line 134–142 train.py) and then validated empirically via the rolling benchmark. Delhi achieves 0.783 empirical coverage (+24h), and Kolkata 0.748. We report both the nominal promise and the measured reality.

Evidence: ml/forecast/train.py:79,134–142 (NOMINAL_COVERAGE=0.8, _conformal_level function with finite-sample correction); docs/benchmarks/delhi.json:574 (pi80_coverage: 0.783 empirical vs 0.8 nominal)

Q. I see you're using 600 bootstrap resamples for the 95% CI on skill. How stable is that? Could the Jaipur CI be narrower?

Expected

600 resamples is enough for stable 95% percentile intervals at sample sizes >100 rows (SKILL_BOOTSTRAP=600, MIN_ROWS_FOR_SKILL_CI=100, ml/eval/benchmark.py:101–102, 114–115). For Jaipur +24h (n_support=299), the standard error on skill is approximately ±0.05–0.08, so the interval width [–0.250, –0.058] is expected. The interval is narrow enough to decisively reject the null (zero skill), but real, not noise. 600 resamples is a practical choice: cheap to compute, stable 95% intervals at typical sample sizes (100–10k), and matches literature standards for percentile bootstrap.

Evidence: ml/eval/benchmark.py:101–127 (_skill_ci function, SKILL_BOOTSTRAP=600, MIN_ROWS_FOR_SKILL_CI=100, percentile bootstrap on line 125)

Q. The raw LightGBM median (before persistence blending) has skill +0.017 at +24h in Delhi. Why are you blending it with persistence at all if the model is better on its own?

Likely

The raw model skill (+0.017) is weak, and the blending recovers it to +0.091. The blend weight (0.4–0.65 across origins, mean ~0.6) weights the model less than persistence because on many of those training months, persistence was genuinely more reliable. The convex blend ensures the forecast never does worse than the better of the two—a risk-averse choice. On live data, w varies by city and regime; Jaipur's w=0.8 suggests the model is trusted there, yet it still loses (negative skill after blending). This proves the blend isn't hiding weakness—it's being honest about low-data regimes.

Evidence: docs/benchmarks/delhi.json:70–73 (rmse_model_raw: 65.92, skill_model_raw_vs_persistence: 0.017, vs rmse_model: 60.96, skill: 0.091); :53–64 (blend_weights ranging 0.4–0.65); ml/forecast/train.py:169–188 (blend_weight implementation)

Q. You hold out 25% of training for conformal calibration. Did you try other fractions? Why did 0.25 stick?

Expected

Yes. We tried 0.40 and found it degraded Kolkata coverage at every horizon (0.748→0.696 at +24h). We also tried recency-weighting conformity scores (narrows toward recent regime, all wrong) and per-city selection (fitted on a bogus single-regime window, chases noise). The reasoning in train.py lines 82–130 documents all attempts. CAL_FRACTION=0.25 survives because on the rolling multi-season benchmark (which carries weight), it already covers close to nominal and doesn't degrade any city. The comment is blunt: 'higher calibration fraction looks better on a single short, noisy window, but worse on the protocol that matters.'

Evidence: ml/forecast/train.py:82–130 (full documented reasoning, listing 0.40 attempt on Kolkata, recency weighting, per-city selection); :98–106 (Kolkata numbers: 0.25 cal → 0.748/0.725/0.698, 0.40 cal → 0.696/0.672/0.668)

16.4 · Uncertainty, calibration and validation rigour

Q. What does your 80% prediction interval actually mean — at the technical level, how is it constructed?

Expected

We fit LightGBM quantile models for the 10th and 90th percentiles on 75% of the training window (ml/forecast/train.py:156-157). On the held-out 25%, we compute conformity scores $E = \max(\text{lo}-y, y-\text{hi})$ and find the α -quantile Q (line 161-162). At serve time we widen both edges by Q , producing $[\text{lo}-Q, \text{hi}+Q]$. This is Conformalized Quantile Regression (Romano et al.). Raw quantile coverage was 48–63%; CQR restores it to nominal 80%.

Evidence: ml/forecast/train.py:145-163; docs/PS5_HONEST_AUDIT.md:246

Q. Your audit says split conformal promises marginal coverage and delivers it — what does that mean, and what's conditional coverage, and why do you care?

Expected

Marginal coverage means 80% of all rows pooled. Conditional means 80% within each predicted level or regime — the only grouping a served band can actually be held to, because at forecast time the outcome is unknown. Kolkata's marginal is 75%, but in the 56–76 $\mu\text{g}/\text{m}^3$ range — where officers make decisions — it drops to 67% at +24h and 55% at +72h. We report conditional coverage explicitly (docs/benchmarks/kolkata.json) so the weakness is visible rather than hidden in the average.

Evidence: ml/eval/benchmark.py:415-426; docs/benchmarks/kolkata.json (calibration/pi80_coverage_by_predicted_quintile)

Q. Kolkata is your weakest city. Show me exactly where the 80% band fails.

Expected

Kolkata +72h: marginal coverage 0.699, but broken down by predicted level — quintile 4 (56–72 $\mu\text{g}/\text{m}^3$ range, the CPCB Satisfactory→Moderate decision boundary) covers only 0.547. That's the operationally critical range where the forecast actually changes what an officer does. Delhi's worst quintile at +72h is 0.705, so this is one model's under-dispersal in the mid-to-upper range, not a calibration flaw. We keep the method and report the breakdown.

Evidence: ml/eval/benchmark.py:424-425; docs/benchmarks/kolkata.json +72h calibration; docs/PS5_HONEST_AUDIT.md:265-269

Q. You said you tried seven different conformity scores on Kolkata. What were they, and why didn't they work?

Expected

Asymmetric per-edge, normalized by band width, normalized by predicted level, Mondrian binning, and combinations (scripts/tune_conformal_tails.py:66-114). The worst quintile moved from 0.615 to 0.646 at best — three points on an eighteen-point shortfall (docs/PS5_HONEST_AUDIT.md:274-275). The root cause is not the calibration score; it's the quantile models under-dispersing in the mid-to-upper range — a model problem. Recalibration cannot fix model under-dispersal, so we kept the simple two-sided Q and made the limit measurable instead.

Evidence: docs/PS5_HONEST_AUDIT.md:271-287; scripts/tune_conformal_tails.py:158-187 (variants A–G)

Q. Your live 90-day benchmark showed Delhi coverage swing from 0.74 to 0.858 with no calibration change. Isn't that a calibration defect?

Likely

No. That protocol uses a single split over ~282 rows from one forecast origin (Delhi +24h). A few hundred rows land in whatever regime that fortnight held, so the number swings on luck. This false signal caused us to try raising calibration fraction from 0.25 to 0.40, recency-weighting, and per-city selection — all failed on the rolling multi-season benchmark (53k rows, 10 origins) where 0.25 already covers close to nominal. The measurement that carries weight is rolling-origin monthly refit (ml/eval/benchmark.py:197-271), which eliminates regime-swap leakage.

Evidence: docs/PS5_HONEST_AUDIT.md:237-251; ml/eval/benchmark.py:197-271

Q. What happens if you raise the calibration fraction from 0.25 to 0.40 to improve coverage?

Likely

Kolkata's coverage gets worse at every horizon: +24h 0.748→0.696, +48h 0.725→0.672, +72h 0.699→0.668 (docs/PS5_HONEST_AUDIT.md:99-100). The audit records this because it is worth knowing: a naive move to "use more data for calibration" backfires when the signal you are chasing is regime luck, not calibration insufficiency. The rolling protocol with 0.25 is robust; the live single-split that looked broken was the false signal.

Evidence: ml/forecast/train.py:82-131 (the CAL_FRACTION comment block); docs/PS5_HONEST_AUDIT.md:95-100

Q. I notice you tried recency-weighted conformity scores. Why not use recent data to calibrate on current regimes?

Likely

Recency weighting narrowed the band toward the most recent (calm) regime — precisely backwards in an air-quality model. On a calibration tail dominated by clean air, the conformity scores shrink, and the band pinches in on the most actionable regime (dirty air). On Delhi it took coverage to 0.666 (docs/PS5_HONEST_AUDIT.md:104). The insight is correct — regimes change — but split conformal is stateless; temporal weighting requires a different framework.

Evidence: ml/forecast/train.py:102-104; docs/PS5_HONEST_AUDIT.md:102-104

Q. If Kolkata's marginal coverage is 75% and conditional coverage fails at the decision boundary, what is your honest answer if a skeptical officer asks: 'Does your 80% interval actually contain the outcome 80% of the time?'

Hostile

No — not at the range where your decisions matter. Overall it's 75%, but in the 56–76 $\mu\text{g}/\text{m}^3$ range where Moderate becomes Poor it's 67% at +24h and 55% at +72h. We publish the breakdown (benchmarks/kolkata.json) so you can see exactly where. The interval covers 80% of clean air and 55% of the air that changes what you do. That is a quantified, disclosed weakness, not a hidden one.

Evidence: docs/benchmarks/kolkata.json (pi80_coverage_by_predicted_quintile); docs/PS5_HONEST_AUDIT.md:286-287

Q. How many training rows go into the conformal calibration step — is it enough to trust the 80% number?

Likely

25% of the training window (ml/forecast/train.py:131, CAL_FRACTION = 0.25). For Kolkata the rolling protocol uses ~10k rows per calibration split (53k total rows / 5 splits). That is enough to estimate a 0.80 quantile robustly. When we tried per-city tuning on the live single-split protocol (~280 rows), it failed because that sample was too small and regime-locked; rolling refit validates the choice against many regimes.

Evidence: ml/forecast/train.py:131, 156-157; ml/eval/benchmark.py:224 (line 258, rolling protocol)

Q. Delhi's worst predicted quintile is 0.704, Kolkata's is 0.547. You say this is a model problem, not calibration. What evidence supports that?

Hostile

Seven different calibration scores were tried — asymmetric, normalized by band width, Mondrian binning, level-scaled — and the worst quintile moved only from 0.615 to 0.646 (scripts/tune_conformal_tails.py; docs/PS5_HONEST_AUDIT.md:271-280). The bottleneck is not how we adjust the band; it's that the LightGBM quantile models under-disperse (produce bands too narrow) in the mid-to-upper range for Kolkata specifically. Recalibration cannot fix a model that is underpredicting uncertainty; you need better features or a different model.

Evidence: docs/PS5_HONEST_AUDIT.md:271-280; scripts/tune_conformal_tails.py:158-187 (variants A–G)

Q. In your honest audit you record that you tried and rejected per-city calibration tuning. Specifically, what happened when you tuned on a held-out half?

Hostile

Delhi +48h dropped to 0.596 coverage (from 0.783). Why? The tuning selected on the same short, single-regime window (one forecast origin) that produced the bogus signal earlier — a few hundred rows land in one temporal slice and that slice's luck becomes your selection criterion. The audit records this failure (docs/PS5_HONEST_AUDIT.md:105-106) because it is educational: tuning on a small, regime-locked validation set chases noise, not signal.

Evidence: ml/forecast/train.py:105-106; docs/PS5_HONEST_AUDIT.md:105-106

Q. You claim your bands measure 'real coverage $\approx$ nominal 80%' but the live benchmark shows them swinging 0.6 to 0.86. When a judge asks if that interval is reliable, what do you say?

Expected

The live benchmark is one forecast origin over 282 rows — a single regime. That measurement is not reliable; it swings on luck. The reliable measurement is the rolling multi-season backtest: 10 origins, 53k rows across diverse regimes. On that protocol Delhi reads 0.783/0.781/0.775 (+24/48/72h), which is close to nominal 0.80. Kolkata is genuinely weaker at 0.748/0.725/0.699, and we publish the conditional breakdown so you know exactly where it fails. A single number can hide; we show both.

Evidence: docs/PS5_HONEST_AUDIT.md:237-251; ml/eval/benchmark.py:197-271; docs/benchmarks/kolkata.json

Q. If I told you to fix Kolkata's 55% coverage at the decision boundary without running any of your tried-and-rejected conformity score variants, what would you do?

Hostile

Collect more data or engineer features that improve the LightGBM quantile model's dispersion in the mid-to-upper range. The seven variants you already ran proved that adjusting a poorly-dispersed band gets you three percentage points. The real fix is upstream. Or accept the 55% as a documented weakness and use it in the decision workflow — an officer sees both the forecast and the coverage-by-level table and calibrates their trust accordingly.

Evidence: docs/PS5_HONEST_AUDIT.md:271-280; ml/forecast/train.py:82-131

16.5 · AI and Agent Claims

Q. How many agents or nodes are actually in your multi-agent graph, and what does each one do?

Expected

Five nodes. Orchestrator loads signals and identifies spiking cells; Attribution attributes pollution sources; Forecast loads 24/48/72h predictions; Enforcement scores and prioritizes actions with RAG citations; Advisory renders citizen health messages. The spike_gate is a conditional router, not a node. On clean-air days like Delhi's monsoon, enforcement is skipped, leaving four nodes in the live trace. See agents/graph.py:345-351.

Evidence: agents/graph.py:345-351 registers 5 nodes; PS5_HONEST_AUDIT.md:370 confirms 'VERIFIED (5, not 6)'

Q. Do you use a large language model at runtime to generate citizen health messages?

Expected

No. At runtime, all citizen messages are deterministic templates in eight languages hardcoded in agents/advisory.py. Lines 28-117 show the exact strings—nothing is generated by an LLM. An optional Gemini polish script exists in

scripts/llm_polish_advisories.py but is not wired into production and requires explicit operator action to use. By design, health advice cannot hallucinate.

Evidence: agents/advisory.py:1 states 'deterministic templates'; AI_METHODODOLOGY.md:22 'Deterministic templates — no model'; PS5_HONEST_AUDIT.md:374 'DELIBERATELY ABSENT'

Q. What does 'multi-agent' actually mean in your architecture—are these autonomous agents or a pipeline?

Likely

A LangGraph StateGraph with five deterministic functions that share state, executed sequentially with conditional branching. Each node is a pure function that receives shared state and returns updates to it. Not autonomous agents—no tool use, no hallucination risk, no looping. The orchestrator makes one decision: whether to route to enforcement or skip it based on spike detection. This is orchestrated state flow, properly called a graph, not agents with agency.

Evidence: agents/graph.py:1-4 architecture diagram; line 358-362 shows add_conditional_edges; AI_METHODODOLOGY.md describes each agent as a discrete transformer (LightGBM, CNN, template render)

Q. You claim to use gradient boosting for attribution with SHAP explanations. Is SHAP actually running for all cities?

Hostile

No. SHAP is currently live only for Pune and partially for Kolkata. Delhi, Mumbai, Bengaluru, and Hyderabad fall back to chemical-signature priors because the SHAP model needs a minimum of 400 samples per cell and our measurement history is pruned at 180 days. The system shows an honest abstain message instead of inventing explanations. This is a genuine strength—it refuses to hallucinate—but it means your demo city Delhi does not have SHAP today.

Evidence: PS5_HONEST_AUDIT.md:2.4 'SHAP explanations are OFF for Delhi, Mumbai and Bengaluru' with counts verified from database; ml/attribution/shap_attribution.py MIN_SAMPLES=400

Q. What threshold of air quality triggers enforcement recommendations to run?

Expected

PM2.5 greater than 120 $\mu\text{g}/\text{m}^3$, which corresponds to roughly AQI 200. This is checked in production mode at agents/graph.py:148 and line 152. The spike_gate then conditionally routes to the enforcement node if focus_cells are populated. Note: an outdated comment at line 119 says AQI > 300, but the actual code uses 120 $\mu\text{g}/\text{m}^3$.

Evidence: agents/graph.py:148-152; line 119 comment is stale; PS5_HONEST_AUDIT.md:2.1

Q. How long does your entire pipeline take to produce an enforcement recommendation?

Likely

Approximately 1,130 milliseconds wall-clock time from signal ingestion to final advisory message on live data. This is your compute latency only. This is not the time it takes a municipal officer to act—that is an organizational response time you have not measured. The honest framing: your pipeline produces a ready-to-sign, cited recommendation in about one second.

Evidence: agents/graph.py:308 latency computation; PS5_HONEST_AUDIT.md:5 'Response time from signal to intervention: 1,130 ms' marked as 'C' because it measures pipeline, not organizational response

Q. Where in your system is a large language model actually used?

Likely

Nowhere in runtime code. The AI is: LightGBM quantile regressors for forecast, gradient-boosting models with SHAP for attribution where data depth permits, CNNs for downscaling, and bge-small embeddings for retrieval-augmented generation. No generation—only retrieval. The enforcement agent retrieves regulations from a 1,271-chunk corpus but does not generate new ones. The advisory agent renders templates, not generated text.

Evidence: Grep for langchain, anthropic, openai in agents/, core/, ml/ returns no runtime imports; AI_METHODODOLOGY.md table shows model types; scripts/llm_polish_advisories.py is optional only

Q. Why deliberately avoid an LLM for citizen communication when the brief suggests it?

Likely

A hallucinated line in an asthma or cardiac advisory is an unacceptable risk. Health advice must be exact—you cannot recover from false guidance. We use deterministic templates from CPCB's official advisory table and cited clinical guidelines. This is not a constraint—it is a security decision. We deploy LLMs where they belong: feature extraction and ranking, not client-facing facts.

Evidence: PS5_HONEST_AUDIT.md:376-385 prepared answer; core/health_advice.py:1-19 explains the citation and template approach

Do not say: We didn't have time for LLMs.

Q. Do you claim the LLM polish script is part of production operations?

Expected

No. scripts/llm_polish_advisories.py is an optional operator utility, not wired into any cron. It must be invoked manually with --push to write changes to the database. By default it runs dry-run only. It includes strict validation gates: zone ID, horizon hour, 'N95', and all digits must survive verbatim, and the script must pass validation to take effect. Rejected candidates revert to the safe template.

Evidence: scripts/llm_polish_advisories.py:1-15; line 15 'Not wired into any cron — an operator choice, disclosed'

Q. How many cities do you cover, and are all the promised technologies actually running in each?

Hostile

Running since 19 Aug. connectors/earth_engine.py is now called by the daily job in .github/workflows/ingest.yml, and the database holds 136 Sentinel-5P no2_sat rows and 164 MODIS fire rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the fire marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: .github/workflows/ingest.yml daily job; 136 rows source='s5p' variable='no2_sat', 164 rows source='modis' variable='fire' in measurements

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. Has an independent enforcement officer or domain expert validated your recommendations?

Hostile

No. You have a protocol and outreach kit, but zero officer ratings. This is a named evaluation criterion, and there is no evidence. The honest grade is F. One officer's review—even a middling one—would move this from F to C and be worth more than additional code.

Evidence: PS5_HONEST_AUDIT.md:5 'Enforcement quality rated by domain experts: n = 0'; section 9 'Item 2 cannot be done by writing code, which is exactly why it keeps slipping'

Q. You advertise 1 km hyperlocal forecasts. How was this downscaling validated?

Hostile

On synthetic fields only. The CNN downscaler achieves 55.3 percent skill versus bilinear interpolation on 64 held-out synthetic fields. Never against a held-out real station. Delhi has 11 weather stations and 3,466 forecast cells—one station per 315 cells.

The honest sentence: 'We interpolate 11 stations to 3,466 cells. The downscaler beats bilinear by 55 percent on synthetic fields. We have never validated it against a held-out real station.' Leave-one-station-out is computable from existing data.

Evidence: PS5_HONEST_AUDIT.md:2.2 'validated on synthetic data' with quote; ml/coverage/dense_field.py:187 'skill vs bilinear on held-out SYNTHETIC fields; real held-out-station RMSE runs on Kaggle'

Q. What is the forecast skill and is it consistent across all cities?

Expected

Multi-season rolling protocol: +9.1 percent at 24h, +12.9 percent at 48h, +12.1 percent at 72h versus persistence baseline on Delhi. Recent window (June–August 2026) ranges from +26 percent (Lucknow +24h) to –14 percent (Jaipur +72h). Jaipur loses to persistence at all horizons with 95 percent confidence intervals entirely below zero. Kolkata's 80 percent prediction interval coverage is 75 percent overall and 67 percent in the 56–76 $\mu\text{g}/\text{m}^3$ range—critical for decision boundaries. Skill is real but uneven.

Evidence: PS5_HONEST_AUDIT.md:3.2 benchmark tables; section 3.2 problem 1 'Jaipur loses to persistence, and now we can prove it is not noise' with bootstrap intervals

Q. How does the enforcement recommendation prioritization formula work, and is it validated?

Likely

Priority = source contribution $\times$ population exposed (normalized) $\times$ actionability $\times$ confidence. An additional benefit score ranks by value per inspector hour: (share $\times$ confidence) $\times$ PM2.5_low $\times$ residents $\times$ (1 + 3 $\times$ P(>120)) divided by estimated inspector-hours. This incorporates exposure-weighting to prevent systemic bias. The formula is transparent and in code. Validation: zero independent officer reviews. The logic is sound; the real-world impact is unproven.

Evidence: agents/enforcement.py:8-9 'priority = share $\times$ pop_norm $\times$ actionability $\times$ confidence'; PS5_HONEST_AUDIT.md:3.3 describes the benefit formula; section 5 'Enforcement quality rated by domain experts: n = 0'

16.6 · Architecture, Scale and Cost

Q. Walk me through your tech stack, layer by layer. Why did you choose each piece, and what's the first thing you'd swap if it didn't have a free tier?

Expected

Backend is FastAPI on Cloud Run (scale-to-zero, 2M req/month free). Database is Supabase Postgres 15 with PostGIS and pgvector (500 MB free). Frontend Vercel, scheduled jobs via GitHub Actions cron (2000 min/month free). ML training on Colab/Kaggle. The architecture trades real-time responsiveness for free tiers: everything caches aggressively. If Cloud Run hit its free limit first we'd move to Render free tier. If Supabase DB size becomes the ceiling — which it does at 15 cities — the jump is to Pro at \$25/month, but that's a known cliff, not a surprise.

Evidence: api/main.py:58-83 (FastAPI + CORS + GZip); ARCHITECTURE.md:156-179 (technology stack table); SCALE.md:43-54 (free tier ceilings table); cloudbuild.yaml:14-25 (Cloud Run asia-south1 deployment)

Q. You claim ₹0 infrastructure. At what city count does that stop, and exactly which resource hits the limit first?

Expected

The free tier holds 10 cities today at ₹0. The first ceiling is Supabase's 500 MB database limit, which fills at 15 cities under our current 180-day raw-measurement retention. That's not theoretical — we measured it: today 10 cities consume 465 MB steady state, growing at 0.21 MB per city per day. To hold 15 cities you must move to Supabase Pro at \$25/month. Shorter retention — 90 days instead of 180 — would push the limit to 21 cities. Beyond 130 cities, split your GitHub Actions jobs per state to stay inside 2000 minutes per month. The 131-NCAP-cities number costs ₹2,700/month: Pro \$25 plus Render Starter \$7 for an always-on API.

Evidence: SCALE.md:43-53 (resource ceiling table with city counts, database sizes, monthly costs); docs/SCALE.md:7-19 (measured data: 380 MB raw table @10 cities, 800 rows/day per city); api/main.py:130-160 (Cloud Run scale-to-zero architecture requires pre-warm to avoid cold-start)

Q. What happens to the API when it goes cold — say, the demo starts after the system's been idle for an hour?

Expected

The first request pays a hard penalty. When Cloud Run's container spins up, the first call to `/coverage` — the most expensive endpoint — waits ~20 seconds for PyTorch and LightGBM to import from disk. That right at the frontend's 25-second read timeout. We mitigate this: on startup, a daemon thread pre-imports torch and lightgbm off the request path, and pre-warms the dense field for every city. So the *second* request is fast. For the demo, we keep the API warm with a GitHub Actions ping every 10 minutes, and scale-to-zero is disabled.

Evidence: api/main.py:129-160 (`_warm_heavy_imports` function, lines 133-136: 'The first `/coverage` call pays ~20 s for import torch...Importing in a daemon thread at startup makes the first real request fast'); api/main.py:13 (Cache-Control with stale-while-revalidate so repeat loads paint instantly)

Q. How many H3 cells are you computing over, and why not use a coarser resolution?

Expected

Resolution 8: 3,466 cells for Delhi, 1,715 for Mumbai. Each cell averages 0.74 km² with an edge of 0.46 km — satisfies the brief's ~1 km grid. Coarser res-7 would be ~5 km²; too large for enforcement worklist prioritization. Finer res-9 would be 2,500+ cells per city and is unnecessary when interpolation between stations already smooths the field. H3 is city-agnostic — one math for every city, clean spatial joins, and Deck.gl renders it natively.

Evidence: ARCHITECTURE.md:187-193 (H3 spatial model specification, res-8 ≈ 0.74 km²); api/main.py:256-310 (GET `/aqi/current` returns `h3_cell` per reading); SCALE.md:28-31 (Delhi 3,466 cells after de-duplication); PS5_HONEST_AUDIT.md:78-89 (measured coverage data from `/coverage` endpoint showing `h3_cell` count per city)

Q. Your dense 1 km field — the 'E2' downscaling claim — what's it validated against?

Likely

I must be honest here: the CNN downscaler is validated on synthetic fields, not held-out real stations. The validation set is n=64 synthetic fields, showing 55.3% skill vs bilinear interpolation. But we've never run leave-one-station-out on the actual data in the database. The measured anchors are genuine — 11 real CAAQMS stations in Delhi — and the model beats bilinear by a margin. That's defensible as a prioritization signal, not as a calibrated product. We're quantifying the uncertainty rather than hiding it, but a judge asking 'where is the real-station validation?' catches an honest gap.

Evidence: ml/coverage/dense_field.py:187 (validation: 'skill_vs_bilinear: 0.553, n: 64, note: downscaler skill vs bilinear on held-out SYNTHETIC fields'); PS5_HONEST_AUDIT.md:95-100 (section 2.2, the honest sentence: 'We interpolate 11 stations to 3,466 cells...validated on synthetic fields. We have never validated it against a held-out real station.')

Q. Why does attribution show SHAP explanations for Pune but zero SHAP rows for Delhi, even though the docs promise explainability?

Likely

The hybrid GBM+SHAP model requires MIN_SAMPLES=400 per cell to be trustworthy. In Delhi, the raw measurements table keeps only 180 days of history, so per-cell depth is thin. Most cells fall below 400. When the model can't reach that threshold, we abstain and fall back to cited chemical-signature priors — which is the *right* thing to do, and the UI says so. Pune reaches 400 samples in enough cells to produce SHAP drivers. So you see SHAP explanations in Pune, abstain messages in Delhi. That's honest, but it means the 'SHAP explains attribution' headline only proves out in one city on stage.

Evidence: ml/attribution/shap_attribution.py line 68 (MIN_SAMPLES = 400); PS5_HONEST_AUDIT.md:124-150 (section 2.4: SHAP status by city: delhi/mumbai/bengaluru/hyderabad = 0 SHAP rows, pune = 60 SHAP rows);

Q. Your satellite data — Sentinel-5P NO₂, fire detections — are those actually running on a schedule? Or is that a listed input that's not live?

Hostile

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. At 10x cities (130 NCAP cities), what's the first resource to fail, and can you design around it?

Expected

The database is the first ceiling. Supabase Pro (8 GB) comfortably holds 131 cities with 180-day raw retention. The second constraint is GitHub Actions: at 130 cities, one sequential ingest+forecast+enforcement+advisory job per city per day hits the 2000-minute free monthly limit. The fix is simple architecture, not code: split the 130 cities into 5 independent GitHub Actions jobs (one per state region) running in parallel. Each region costs ~400 minutes per day, five regions = 2000 minutes. The API itself scales fine on Render Starter (\$7/month), and frontend is already CDN-global. At 100x cities, the database would be 50 GB and you'd move to a paid plan, but every city just means another YAML file, zero code change.

Evidence: SCALE.md:43-53 (resource table: 131 cities = 5.1 GB DB); ARCHITECTURE.md:567-572 (city onboarding = config, no per-city code); `.github/workflows/ingest.yml`:40-51 (per-city loop over CITIES variable, parallelizable); PS5_HONEST_AUDIT.md:465-478 (section where remaining hours should go: item 6 notes satellite data not running on a schedule)

Q. How much of the API is actually tested, and what's the coverage for the main contract layer?

Expected

Backend has 270 passing tests and 61–63% coverage overall. But the API contract itself — `api/main.py` — is at 33%, meaning 927 of 1379 statements have no test. That's the file a technical judge opens first. The core schemas and helper functions are well-tested, and we added 100% coverage to `core/health_advice.py` this week — which found a real bug where a city with no readings was told its air was Moderate. But `api/main.py`'s 33% is a genuine weak spot. The CI doesn't fail on coverage threshold, so it's not a blocker, just a number a judge will find.

Evidence: PS5_HONEST_AUDIT.md:449-460 (section 8, Engineering quality: '270 passing', 'api/main.py 33%' (927/1379 statements untested)); PS5_HONEST_AUDIT.md:455 (core/health_advice.py 0% → 100% this session, found a real bug)

Q. When you say the system responds in ~1 second, are you measuring your pipeline latency or an organization's actual response time?

Likely

We measure our pipeline latency end-to-end: signal ingested to enforcement recommendation generated and cited. The `/latency` endpoint proves it: ~1130 ms wall-clock in a recent run. What we *don't* measure is the real-world delay from 'recommendation ready' to 'officer reads it and decides.' That's organization behavior, not system behavior. So the honest claim is: our pipeline turns a signal into a ready-to-sign, cited recommendation in about 1 second. We can't claim credit for how fast a municipality responds. That distinction matters when someone asks if VayuNetra solves Delhi's problem — it doesn't, it gives Delhi's officers a tool that works fast.

Evidence: ARCHITECTURE.md:557-562 (North-Star metric: action_traces stamps latency_ms); api/main.py:2639-2641 (GET /latency endpoint); PS5_HONEST_AUDIT.md:412-415 (claim 4: 'Response time from signal to intervention: seconds' → 'that is our compute latency. No organisation's response time has been measured.')

Q. Your enforcement worklist ranking uses inspector-hours as a cost. Where do those numbers come from?

Likely

The inspector-hours are estimates, not measured from real operations. The model assumes a stack test costs 8 hours, stopping a fire costs 1 hour, etc. This was designed to prioritize high-confidence, low-effort wins and de-prioritize low-confidence, costly actions — which is wise resource allocation logic. But it's not calibrated to any municipality's actual capacity or workflow. An officer from Delhi might laugh at these numbers or say they're conservative or optimistic. The ranking is defensible as a heuristic, but labeled as assumptions, not facts, in the UI.

Evidence: PS5_HONEST_AUDIT.md:309-314 (Beyond the brief section: value-per-inspector-hour ranking; 'The inspector-hours are our estimates, not measured — labelled as assumptions on the card and in the API.')

Q. You claim multi-city, but if I drill into Jaipur, what am I seeing compared to Delhi?

Hostile

Jaipur has 5 enforcement recommendations, 8 emission sources, and 1 station per pollutant. It's thin. That city's forecast skill is negative at +48 and +72 hours (−4% and −14% vs persistence), so our model loses to 'tomorrow will be like today.' Delhi has 45 recs, 103 sources, stable skill across all horizons. So yes, we're 10-city live, but they're not equal. Jaipur is more of a proof-of-concept for onboarding. A judge asking 'show me Jaipur' should get that context: weaker data → weaker output, and we show the station count rather than hiding it.

Evidence: PS5_HONEST_AUDIT.md:322-343 (Multi-city table: jaipur n=452 cell-days, 5 recs, 8 sources; negative skill [−0.160, −0.100, −0.314] at all three horizons with 95% CI); PS5_HONEST_AUDIT.md:208-210 (jaipur +24h +4%, +48h -4%, +72h -14% vs persistence, beats_persistence = false)

Q. If I onboard a new city live during the demo, what happens to cost, latency, and your database?

Expected

Cost stays the same — you're on Supabase Pro (when running multi-city) which includes 8 GB. Latency: the config hits the database, GitHub Actions picks it up on the next hourly cron, and data starts flowing. Within a few runs you have historical data and the model produces forecasts and attribution. Database: +210 KB per day per city once measurements stabilize. So you're adding linearly but staying well under the Pro tier. The demo is: POST `/admin/cities` with a new city config, then within an hour the map renders it, the models score it, and it joins the live comparison. Zero code change, just a config row.

Evidence: ARCHITECTURE.md:567-571 (city onboarding = config-driven, POST `/admin/cities`, zero code change); SCALE.md:8-15 (per-city cost: 800 rows/day, 0.21 MB database growth per day); SCALE.md:43-46 (15 cities fits Supabase Pro @ \$25/mo)

16.7 · Security, Privacy and Data Protection

Q. You have four roles in your system—admin, officer, inspector, citizen. How does row-level security actually work? What can an unauthenticated caller read?

Expected

RLS policies in supabase/migrations/20260627000002_roles_rls.sql:21-45 make cities, advisories, measurements, and forecasts readable by everyone—true public. Officers and inspectors read attribution and enforcement for their own city only via `current_city()` and `current_role_name()` checks. Citizens read only public data. Service-role bypasses RLS server-side for pipelines. An unauthenticated caller running against the live API gets a 401 unless DEMO_MODE is true, but in DEMO_MODE they can read all cities and advisories anyway from fixtures.

Evidence: supabase/migrations/20260627000002_roles_rls.sql:21-45; api/main.py:204-213

Q. You mention 'service-role bypasses RLS'—that's a powerful key. Where does it live, who can rotate it, and is it ever exposed to frontend code?

Expected

SUPABASE_SERVICE_ROLE_KEY is in .env only, never in code. api/main.py:224-226 has _db() that loads it server-side for pipeline reads. The anon key is also in .env for client-side auth, but it respects RLS—the service-role does not. Neither key appears in frontend code or fixtures. Rotation would be a manual Supabase vault operation on production; nothing in the repo automates it. The .gitignore blocks .env from commit.

Evidence: .env (git-ignored); api/main.py:224-226; .gitignore lines 1-3

Q. You're placing real phone calls via Twilio—I see IVR code that reads TWILIO_ACCOUNT_SID and TWILIO_AUTH_TOKEN. Are these in .env? And what happens if someone extracts the .env from a staging environment?

Hostile

Yes, TWILIO_ACCOUNT_SID, TWILIO_AUTH_TOKEN, TWILIO_PHONE_NUMBER and the recipient numbers TWILIO_TO_NUMBER(S) are all in .env, git-ignored. If an attacker gained .env access, they could place calls from your Twilio number to hardcoded recipients or modify recipients if they also compromised the code. The system has no per-call consent record—callers cannot opt out. Phone numbers are citizens' direct lines and are not encrypted at rest in environment variables.

Evidence: .env (Twilio section); channels/ivr.py:168-176, 191-216; api/main.py:1300-1319

Q. You're calling citizens' phones to read health advisories. What is your consent and data-protection basis for collecting and calling those phone numbers?

Hostile

The numbers in TWILIO_TO_NUMBERS (.env) are hardcoded test/demo numbers, not citizen opt-ins. The system collects no phone numbers from users. For actual deployment, you would need explicit informed consent from each citizen for IVR calls—not yet implemented. A judge asking this expects a privacy notice, opt-in flow, unsubscribe mechanism, and legal basis under India's Telecom Commercial Communication Regulation. Today: not in place.

Evidence: .env; channels/ivr.py:179-188; api/main.py:1252-1321 (broadcast endpoint has no consent check); docs/AI_METHODODOLOGY.md:2.3 (advisory PII coverage is wardlevel, not per-citizen call)

Q. I see citizen_reports table stores lat/lng and optional photo URLs. Is that personally identifiable? Who can read it?

Likely

Citizen reports (supabase/migrations/20260816000001_citizen_reports.sql:1-33) store lat/lng plus optional description and photo. The table is public-readable by design—transparency for SLA tracking. No phone number or name is stored, so it's not direct PII, but lat/lng + photo + timestamp could re-identify a person. There is no consent flow before submission; the /report endpoint is rate-limited by IP (one per minute) but anyone can submit anonymously.

Evidence: supabase/migrations/20260816000001_citizen_reports.sql; api/main.py:1026-1086 (public /report endpoint, rate limited by IP)

Q. Rate limiting—I see /report is throttled by client IP. But your /aqi/current, /attribution, and broadcast endpoints don't appear rate-limited. Can someone hammer them?

Likely

Correct. /report is rate-limited (one per minute per IP, api/main.py:1011-1034). /advisory/broadcast is throttled server-side per city (300 s window, api/main.py:1001-1004). But /aqi/current, /coverage, /attribution have no per-caller rate limit—they rely on

HTTP caching (Cache-Control headers set in middleware) and Supabase's free tier's request quota. A determined DOS attacker could exhaust the database tier. No API key or auth-based rate limiting exists for read endpoints.

Evidence: `api/main.py:1001-1034` (broadcast throttle), `113-126` (caching middleware), `54-56` (city field validation)

Q. Input validation—show me your safeguards against injection. I see a city parameter. Is it validated?

Expected

City is validated as a Pydantic Field (`api/main.py:56`): `min_length 1, max_length 40, pattern ^[a-z][a-z0-9_-]*$`. SQL queries are built via `supabase-py`, which parameterizes them—no raw SQL concatenation. Photo uploads are capped at 4 MB (`api/main.py:1023`). Lat/lng are coerced to floats and range-checked (`-90≤lat≤90, -180≤lng≤180`). Description is truncated to 500 chars (`api/main.py:1039`). No XSS risk on output because HTML is escaped in `channels/ivr.py:90, 94`.

Evidence: `api/main.py:54-56` (city validation), `1026-1086` (input checks), `channels/ivr.py:83-108` (HTML escaping)

Q. You ingest OPENAQ_API_KEY, DATA_GOV_IN_API_KEY, and GEMINI_API_KEY—how are these rotated? If the .env leaks, how do you invalidate them?

Likely

All keys live in `.env`, git-ignored. Rotation is manual: update `.env` locally, then push new value to GitHub Actions secrets for CI/CD, and redeploy. There is no automated key rotation, audit trail of rotations, or code that enforces expiry. If a key leaks, you must assume the attacker has accessed your integrations. The system does not log which service calls used which key, so breach scope is unclear.

Evidence: `.env` (keys present); `core/env.py:12-29` (manual loading); no rotation logic found in repo

Q. You store advisory_subscribers with chat_id, city_id, and language. Are Telegram chat IDs personally identifiable? What's your data retention for this table?

Expected

Telegram chat_ids are opaque integers issued by Telegram but linked 1:1 to a user account, so they're de facto identifiers. You don't store the user's name, but the chat_id + language + city_id profile can identify their interests. There is no data retention policy coded in the schema or migrations—no `AUTO_PURGE` clause or archive job. Subscribers remain indefinitely until manually deleted.

Evidence: `supabase/migrations/20260714000001_advisory_subscribers.sql:1-12` (no retention policy); `channels/telegram.py` (telegram subscriber logic, if present)

Q. You mention a 180-day retention window for raw measurements. Is that enforced in code, or is it just a documented target?

Likely

It's documented (`docs/SUBMISSION.md`, `docs/SCALE.md`) and there is an archive script (`scripts/archive_measurements.py` with `--keep-days` default 180) that purges old rows, but nothing triggers it automatically on a cron. It must be run manually. Until run, raw readings older than 180 days stay in the live database. No `ON DELETE CASCADE` or policy enforcement exists in the schema.

Evidence: `scripts/archive_measurements.py:23` (`--keep-days` default); `docs/SUBMISSION.md` (180-day retention claimed); no scheduled job in GitHub Actions cron or `ingest.yml`

Q. Your honest audit says satellite data—Sentinel-5P NO₂, MODIS/VIIRS fire—are built but not running. If they were running, what personal data would they collect?

Expected

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence

is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. CORS is locked to `vayunetra-aqi.vercel.app`, `localhost:5173`, `localhost:4173`, plus a regex for any localhost port. If I spoof a localhost origin, can I call your API?

Expected

No. The regex `allow_origin_regex=r'http://(localhost|127.0.0.1)(:\d+)?'` matches localhost or 127.0.0.1 on any port, but the browser enforces CORS—a spoofed localhost header from the real internet will be rejected. However, a developer running a local server (e.g., vite on 4180) can call the deployed API if they override the origin header, which is a design—local dev convenience over strict CORS. In production, this should whitelist only the deployed frontend.

Evidence: `api/main.py:66-79` (CORS config)

Q. error messages—do you leak internal state? Show me your `_server_error` pattern.

Expected

`api/main.py:46-51` logs the full exception server-side (`logger.error` with `exc_info=True`) but returns only a generic user message: 'Could not load [thing] right now.' No stack trace, no env keys, no DB URLs reach the client. Demo mode does serve demo-friendly error notes, not exposing real backend state. This follows good practice.

Evidence: `api/main.py:46-51` (`_server_error` pattern); various endpoint catches that call it (e.g., 356, 516, 548)

Q. Your RLS says citizens can't self-promote to officer. But the profiles table has a policy 'own profile update' with a check. Can a citizen insert a profile for someone else?

Expected

No. `supabase/migrations/20260629000001_rls_complete.sql:82-85` allows citizens to update only their own profile (`user_id = auth.uid()`), and the role can only remain 'citizen'. There is no insert policy for citizens—only the trigger `handle_new_user()` auto-creates profiles on auth signup as 'citizen'. An admin can promote via 'admin manage profiles' (line 88-90), but citizens cannot insert arbitrary profiles.

Evidence: `supabase/migrations/20260629000001_rls_complete.sql:72-108`

Q. Advisory broadcast—it reads the 'worst air' ward if the operator doesn't pick one. But how do you sort advisories by severity? Could the sort be wrong and send the wrong ward to IVR?

Likely

`api/main.py:1157-1162` defines `TIER_SEVERITY` (`severe=5`, `very_poor=4`, etc.) and `_advisory_sort_key` returns `(-severity, ward_id)`. The sort is deterministic. `_latest_advisory` (1181-1209) returns `sorted(rows, _advisory_sort_key)[0]`, so the worst-air ward always comes first. However, if multiple advisories have the same severity and `ward_id`, Postgres's returned order is undefined until the sort—and the sort explicitly breaks ties by `ward_id` string comparison. This is safe but could be safer if you sorted by `issued_at + severity + ward_id` to rule out any ambiguity.

Evidence: `api/main.py:1157-1208`

16.8 · Impact Quantification and Economics

Q. You claim Delhi faces 73,395 premature deaths per year attributable to air pollution — where does that number come from and how does it break down?

Expected

The figure is computed using WHO AirQ+ methodology: baseline death rate (7.3 per thousand per year from World Bank/SRS India) times city population times an attributable fraction derived from the long-term concentration-response function (Chen & Hoek 2020: HR 1.08 per 10 $\mu\text{g}/\text{m}^3$) applied to the excess PM2.5 above WHO guideline. For Delhi's 20.6 million people at 92 $\mu\text{g}/\text{m}^3$ annual PM2.5, the excess (87 $\mu\text{g}/\text{m}^3$) generates a 48.8% attributable fraction, yielding 73,395 deaths per year. Every component is cited in `ml/impact/factors.py`.

Evidence: `ml/impact/quantify.py:150-156`; `ml/impact/factors.py:65-75, 123-125, 145-147`; `/roi` endpoint

Q. The CRF (concentration-response function) has a 95% confidence interval of 1.06–1.09 — what does that mean for your death count, and do you publish that uncertainty?

Likely

The CI spans 59,801 to 79,327 deaths per year for Delhi — a 33% range around the point estimate. We do not publish confidence intervals on the final death figure; they are stated in the source paper (Chen & Hoek 2020) but do not propagate through to our API response. This is an honest gap: the 73,395 figure is defensible but incomplete without its uncertainty bounds. To properly represent it we would need to either publish a range or note that this is a point estimate against a $\pm 15\%$ background.

Evidence: `ml/impact/factors.py:65-69`; Python calculation: CRF 1.06–1.09 gives attributable deaths 59,801–79,327; `/roi` endpoint returns point estimate only

Q. Your concentration-response function (HR 1.08 per 10 $\mu\text{g}/\text{m}^3$) is a long-term coefficient from annual exposure studies — is it defensible to apply it to a city's annual average as if it were the entire causal story?

Expected

Yes, defensibly. The Chen & Hoek meta-analysis synthesises long-term cohort studies where participants' annual exposure produces the measured health outcome. Applying it to an annual PM2.5 mean is the standard WHO AirQ+ approach, documented in their technical guidance. The caveat is that we are computing an annual attributable burden, not causal attribution — we cannot separately account for seasonal peaks, acute events, or composition changes. This is why we distinguish short-term mortality (for what-if/forecast horizons) from long-term (annual burden).

Evidence: `ml/impact/factors.py:65-69` (Chen & Hoek 2020); `ml/impact/quantify.py:149-156` (`city_roi` uses long-term CRF); `HEALTH_IMPACT.md` explicitly states short-term is exposure, not mortality

Q. Where does your baseline death rate (7.3 per thousand per year) come from, and does it represent all-cause mortality or a specific condition?

Likely

It is the crude national death rate for India from World Bank/SRS data, approximately 2021. It is all-cause mortality, not disease-specific, which is why it appears alongside the CRF rather than as a separate rate per condition. The population is multiplied by this baseline, then adjusted by the attributable fraction from the CRF. This approach is standard in burden-of-disease work but produces conservative estimates because it does not separate age-specific or condition-specific incidence rates.

Evidence: `ml/impact/factors.py:72-75`; `ml/impact/quantify.py:50-51, 150`; `test_impact.py:39` validates this calculation

Q. Your Value of a Statistical Life (VSL) is ₹5 crore, but literature estimates for India span ₹3–12 crore — how sensitive are your ₹ figures to this choice?

Hostile

A four-fold range (3–12 crore) means the health-burden ₹ figure can swing by that same factor. Delhi's 367-crore annual burden becomes 220 crore at the conservative end or 880 crore at the upper end. We chose ₹5 crore as a mid-point, with a caveat in the factors module that it is an order-of-magnitude figure, not precise. The API response includes the source (OECD benefit transfer via income elasticity) so judges can swap it, but the default response does not surface the range.

Evidence: `ml/impact/factors.py:80-86`; `/roi` endpoint includes VSL source but not min/max; `PS5_HONEST_AUDIT.md` notes this as a value-laden assumption

Q. In your intervention replay, you found no weather-adjusted reduction in PM2.5 during Delhi's Stage III and Stage IV GRAP orders — how do you present that as a success?

Hostile

We do not. The honest read is in `OUTCOMES.md` and `delhi_interventions.md`: association, not causation. The government-wide stage orders have no untreated control inside the city and are triggered by dirty air, so a plain before/after mostly measures regression to the mean and weather. Our meteorological-normalisation method (LightGBM on ERA5) showed no detectable reduction, though Diwali night (+182 $\mu\text{g}/\text{m}^3$, +95%) demonstrates the method has power to detect real signals. We state this plainly: VayuNetra will measure targeted dispatch actions (which have controls) not city-wide stages.

Evidence: `OUTCOMES.md:51-80`; `docs/benchmarks/delhi_interventions.md`; `core/interventions.py` defines per-cell-vs-city-drift method

Q. You measure intervention effect as (cell_after – baseline) – (city_after – city_before) — that subtracts city drift, but what if the intervention happened to be in a cell where natural weather patterns are improving anyway?

Likely

That is a real bias we acknowledge. Subtracting city drift removes shared weather effects but not local microclimatic luck. A seven-day minimum measurement window (`MIN_MEASURE_DAYS` in `core/interventions.py`) reduces this risk by averaging over longer timescales, but does not eliminate it. This is why we label verdicts as 'provisional' before seven days and why the system will require multiple interventions per source to build confidence. It is an honest weakness of the method, not a hidden one.

Evidence: `core/interventions.py:15-67`; `effect_summary()` method documents the limitation

Q. In your exposure calculation, when GPW (GriddedPopulationofthe World) data is not available, you spread city population uniformly — what error does that introduce?

Expected

For cities like Jaipur or Lucknow without GPW coverage, we assume uniform distribution over forecast cells (4.1M people over 3,000+ cells = ~1,400 per cell average). Reality is clustered — dense urban cores and sprawling suburbs. This can overestimate exposures in peripheral cells and underestimate central peaks. We flag this in the API response ('uniform_city_population' basis) and note it in the UI. For Delhi, Mumbai, Bengaluru (which have GPW) the error is smaller.

Evidence: `ml/impact/exposure.py:23-34`; `HEALTH_IMPACT.md:15-18` documents basis field

Q. The concentration-response function for short-term mortality (1.0123 per 10 $\mu\text{g}/\text{m}^3$) differs from long-term (1.08) — when should a judge trust each one?

Expected

Short-term (WHO HRAPIE 2013, natural-cause all-ages) is for acute changes over hours or days — what-if simulations, enforcement outcomes over a week, forecast exposure windows. Long-term (Chen & Hoek 2020) is for annual or multi-year exposures, used in the ROI dashboard and city burden figures. They measure different biological processes: short-term is immediate inflammatory response; long-term is systemic disease accumulation. The distinction is explicit in the code path (`quantify_intervention` uses short-term; `city_roi` uses long-term).

Evidence: ml/impact/factors.py:39-43 (short-term), 65-69 (long-term); ml/impact/quantify.py:100 vs ml/impact/quantify.py:155

Q. You claim biomass-burning abatement co-avoids 166 tonnes of CO₂ per tonne of PM_{2.5} — is that defensible, or does it cherry-pick high-emission stubble burning?

Likely

It is derived from Andreae (2019): agricultural-residue burning generates ~1,515 g CO₂/kg and ~9.1 g PM_{2.5}/kg, yielding 166 tonnes CO₂ per tonne PM_{2.5}. The caveat, stated in factors.py, is that this is a field-burning average and real ratio varies with crop type and combustion efficiency. Traffic (900 t CO₂/t PM_{2.5}) and other sources without defensible ratios return None (honest). We do not invent co-benefits; we cite their source and acknowledge variability.

Evidence: ml/impact/factors.py:97-110; ml/impact/quantify.py:54-66 applies these only when defensible

Q. When you say ₹3,669 crore annual health burden for Delhi, does that mean 73,395 preventable deaths at ₹5 crore each — and is that the right way to interpret VSL?

Hostile

Yes, that is the arithmetic: 73,395 deaths × ₹5 crore VSL = ₹3,669.75 crore. But VSL is not a 'price' for a life; it is a revealed-preference willingness-to-pay inferred from wage-risk choices and environmental litigation, used in benefit-cost analysis. Stating it as a multiplication is economically sound but can mislead: the ₹ figure is a policy anchor, not a human valuation. The caveat in factors.py notes this is order-of-magnitude, and the true range (3–12 crore) should accompany any policy recommendation.

Evidence: ml/impact/factors.py:80-86; ml/impact/quantify.py:107, 166; tests/ml/test_impact.py:38-43 validates the multiplication

Q. The Diwali night analysis showed +95% signal above weather-expected levels — but that was one night with fireworks, not pollution-control measures. How do you use that to validate your deweathering method?

Likely

Diwali is a positive control: a short-term exogenous shock (fireworks; +182 µg/m³ observed vs 192 expected) that we know happened and can verify the method detected it. This proves the deweathering model (LightGBM on ERA5, 88,270 training hours, held-out R² 0.61) has the power to spot real signals. It does not validate its performance on slow, diffuse interventions like traffic bans or construction halts, which may not show up against a model's extrapolation error in stagnant winter weather.

Evidence: docs/benchmarks/delhi_interventions.md:38-45 (Diwali row: +181.8 µg/m³ observed, +94.9%); OUTCOMES.md:70-71 (positive control)

Q. You use population data from UN World Urbanization Prospects (2018) and annual PM_{2.5} from IQAir World Air Quality Report (2023) — why these sources, and how old is the data during your final demo?

Expected

UNdata 2018 is the standard source for city-scale population in international assessments; IQAir 2023 is the most recent published annual survey. Both are cited and defensible. The gap is real: UNdata is six years old and IQAir one year old as of August 2026. For policy work this would warrant updating to 2025 or 2026 estimates if available. In the demo, the /roi endpoint serves these stored values, so judges see 20.6M for Delhi and 92 µg/m³ as constants, not live measurements.

Evidence: ml/impact/factors.py:123-145 (CITY_POPULATION source: UN WUP 2018); ml/impact/factors.py:145-166 (CITY_ANNUAL_PM25 source: IQAir 2023)

16.9 · Product, users, adoption and operations

Q. Has any real officer or pollution control board actually used VayuNetra to make a decision, or is this a demo-only system?

Expected

No official has used it yet. We have zero ratings from pollution control board officers; the pilot-outreach kit exists but has generated no responses. We show our city data to demonstrate the workflow to a real officer in a twenty-minute session, and we ask them three questions about what's missing. That conversation is worth more than any remaining code change. We are two weeks from the final; that outreach is priority one right now.

Evidence: docs/PS5_HONEST_AUDIT.md:399-402: "n = 0." Kit exists, ratings pending. F on the named scoring criterion.

Q. Walk me through what actually happens to an enforcement recommendation — from the moment the system generates it to when a citizen might see the effect.

Expected

The system ranks sources by contribution times exposure times actionability times confidence. An officer opens the dossier — satellite patch plus regulation retrieved by RAG — and either dismisses it or approves it. If approved, dispatch freezes the cell's seven-day baseline and arms before-and-after tracking. The officer closes the case with a finding — violation found, compliant, inaccessible, or not applicable. Every step goes into an immutable audit trail. The whole stack is tracked. Where we stop: that is the officer's dispatch decision. The actual inspection happens in the field; we measure the air before and after.

Evidence: docs/USER_GUIDE.md sections 6.2-6.7: enforcement worklist → evidence dossier → notice PDF → approve → dispatch → close. agents/enforcement.py defines the lifecycle. api/main.py POST /enforcement/{id}/status handles status changes.

Q. Every recommendation I see in the console is marked 'proposed'. What's stopping an officer from actually approving and dispatching these?

Expected

Exactly. In our live database, every recommendation sits unactioned. The system proposes, but it never auto-dispatches — which is correct; an officer must approve. The compliance backlog exists because officers have not yet used the system; the backlog will exist in production too until an officer runs the enforcement loop. We measure the before-and-after for every dispatched action, so once an officer approves, we can prove whether the air changed. The test of the system is not the ranking. It is whether officers use it.

Evidence: PS5_HONEST_AUDIT.md:399-402: compliance backlog noted. enforcement_rec table populated with status='proposed'. agents/enforcement.py writes proposed, officers transition via /enforcement/{id}/status POST.

Q. You claim eight-language advisories. Which ones are actually reviewed by a native speaker?

Expected

Two: Hindi and Marathi, both reviewed by team members on 18 August. The other six — Kannada, Tamil, Telugu, Bengali, Gujarati — are deterministic templates validated for correct script but not reviewed by a native speaker yet. Script validation means the characters are in the right script, not that the sentence is idiomatic or medically sensible to someone who speaks the language. Every advisory is templated from CPCB and WHO guidance, which is why we can review it — no language model is involved. But a Tamil speaker has not yet read the Tamil text aloud and confirmed it sounds right.

Evidence: docs/ADVISORY_REVIEW.md:10-19: status table shows hi/mr reviewed (team members, 2026-08-18), others pending. channels/README.md line 24: native-speaker review status in ADVISORY_REVIEW.md.

Q. On the map you show 'Sentinel-5P' and 'MODIS' as attribution inputs. But the honest audit says satellite data has zero rows in the database. What's happening?

Likely

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. Citizens receive advisories in multiple languages, but on the IVR calls, are they speaking to different languages or hearing the same English message?

Likely

IVR calls are real Twilio calls with different voices — Hindi voice reads Hindi advisories for Delhi and Pune; English voice everywhere else because Polly has no voice for the other six scripts yet. The advisory body is translated into eight scripts, but the framing sentences — the opening 'here is the advisory for X city', the closing 'stay safe' — were English for every language until 18 August. Marathi-speaking Pune callers heard English framing around Marathi advisory text. We corrected this by adding native-script framing, but none of those has been reviewed by a native speaker yet, including Marathi.

Evidence: `docs/ADVISORY_REVIEW.md:22-30`: IVR framing added 19 Aug, none reviewed by native speaker. `channels/ivr.py` uses Polly voices where available. Only Hindi and English have verified voices today.

Q. The public app mentions 'public displays' for advisories. Can I point a city screen to a URL and show the advisory live?

Likely

No. Public-display mode is a rendering inside the console, not a standalone endpoint. You see the advisory as it would appear on a big screen — high-contrast typography, no side panels — but you cannot point an external screen at a URL and have it auto-update. That feature is documented as a capability but not built. If a city wants a public-display kiosk, today you have to screen-capture and refresh manually. It is on the roadmap.

Evidence: `channels/README.md:16-18`: 'public display mode ... a rendering mode inside the console, not a standalone kiosk endpoint you can point a screen at.' `docs/USER_GUIDE.md` mentions it as one of four channels; `channels/ivr.py` and `telegram.py` exist, but no public-display endpoint.

Q. On your landing page you mention the system makes recommendations 'visible within seconds'. But that's compute time, not actual municipal response time, right?

Likely

Exactly. Our pipeline runs in 0.8 to 9.7 seconds — that is signal-to-cited-recommendation compute. But an officer then reviews it, approves it, dispatches it to the field. The field team goes to the site, verifies the violation, and acts. We measure our compute latency because it is deterministic and we control it. We do not measure the municipal response time because that depends on how many open cases the office has, staffing, and logistics. Be careful about framing: we can promise a notice PDF in seconds; we cannot promise a site visit in minutes.

Evidence: `PS5_HONEST_AUDIT.md:414-415`: 'our pipeline turns a signal into a cited, ready-to-sign recommendation in about a second' — 'never imply we shortened a municipality's response.' `docs/USER_GUIDE.md:172-173`: latest pipeline run shows ~1130ms.

Q. Your README says you have 'six agents.' But the honest audit and your code show only five nodes on the graph. Which is correct?

Expected

Five, and the docs say five — the 'six' claim was found by our own audit and corrected across the README, landing page and every doc on 19 Aug. `agents/graph.py` registers five nodes: orchestrator, attribution, forecast, enforcement, advisory. The sixth thing people count is `spike_gate`, which is a conditional edge, not a node — it routes to enforcement only when there are focus cells, so a clean-air city legitimately runs four. We would rather show you the graph than argue the number.

Evidence: `agents/graph.py` — five `add_node` calls plus one `add_conditional_edges` (`spike_gate`)

Do not say: Do not say six agents. It is five nodes and a conditional edge.

Q. SHAP explanations — that's a core feature. Which cities actually get SHAP-based attribution versus falling back to 'cited priors'?

Expected

Pune and part of Kolkata produce SHAP drivers today. Delhi, Mumbai, Bengaluru, and Hyderabad fall back to chemical-signature priors. SHAP needs 400 samples per cell; raw measurements are pruned at 180 days, so per-cell depth is thin in the large cities. When the model lacks skill, we abstain and fall back to cited inventory priors — which is honest and good. The cell story shows which path: green if the model explained it, amber if we fell back. If a judge clicks a Delhi cell during the demo, they see the abstain message, not SHAP. That is correct; do not show SHAP if we did not compute it.

Evidence: PS5_HONEST_AUDIT.md:124-150: Delhi/Mumbai/Bengaluru/Hyderabad n=0 rows with SHAP; Pune n=60 with SHAP; Kolkata mixed. `ml/attribution/shap_attribution.py:MIN_SAMPLES = 400`. Cell Story shows model attribution (green) or signature priors (amber).

Q. You interpolate 11 CPCB stations to 3,466 cells in Delhi. How do you know that works? Did you validate it against held-out real stations?

Likely

No. The downscaler is a CNN trained on synthetic fields. We validated it on synthetic held-out fields — 55.3 percent skill over 64 samples. That beats bilinear interpolation, but we have never validated it against a real held-out station. Leave-one-station-out is computable from the database and is the most valuable experiment we could run in the next two weeks. Right now the claim is: the interpolation beats the naive baseline on synthetic data. That is weaker than it sounds.

Evidence: PS5_HONEST_AUDIT.md:76-100: RMSE_cnn 2.22, skill 0.553 on synthetic fields, n=64. `ml/coverage/dense_field.py:187` states skill measured on synthetic. 'never validated it against a held-out real station.'

Q. Forecast skill is plus nine percent in Delhi. Is that good? And what about Jaipur — I see it says negative skill there.

Likely

Nine percent over persistence is modest but real. Jaipur is worse: it loses to persistence at every horizon — negative four, fourteen, and negative fourteen percent. We have interval estimates showing that is not noise; the bounds are entirely below zero. Five of ten cities blend the model to something near pure model; Jaipur does so at exactly the horizons it loses. That looks like a weight-selection problem that wants a clean experiment before we claim it is fixed. We publish the negative numbers because hiding them is worse than showing them.

Evidence: PS5_HONEST_AUDIT.md:181-226: Delhi +9.1%, Jaipur -4% to -14%. 'Jaipur loses to persistence, and now we can prove it is not noise.' Table at line 214-226 shows 95% CI entirely below zero.

Q. Citizens use Telegram or IVR to get advisories. How many are actually subscribed?

Expected

We have not measured adoption because the system launched for this hackathon. The public Telegram bot at `@aqivayu_bot` is live, and anyone can call the IVR line. But actual subscriber counts are zero in the sense that no city has deployed this at scale; it is a working demo today. We can measure it: `/health` on the API shows `DEMO_MODE`, and once it goes live, subscription and call

counts will tell us how many people actually listen. Right now we have tested that the channels work; we do not have operational scale.

Evidence: channels/telegram.py and channels/ivr.py implement live messaging. No subscription metrics in the codebase; this is a hackathon demo, not a deployed product. README:20 points to @aqivayu_bot as live.

Q. Tell me about a citizen who used the system and took action based on an advisory. Give me a concrete example.

Hostile

We cannot. The system is two weeks from the finals. The advisories are written for citizens and the Telegram bot is live, but we have not measured whether citizens read them or changed behavior. That is not a failure — it is a pre-launch state. What we can measure once deployed: how many people subscribe, how many open the advisory, whether health-related searches in pollution time windows decline, whether outdoor-activity booking changes. The design is built for measurement; the data does not exist yet.

Evidence: No adoption metrics in codebase. docs/PS5_HONEST_AUDIT.md:5-6 lists 'Citizen advisory' as 'Broad, shallow in review' — no user data. Product is live-running but not deployed to actual users.

Q. You say every recommendation gets a 'projected impact chart' in the notice PDF. Is that real physics, or is it a guess?

Likely

It is a counterfactual screening estimate. For each source, we compute the forecast PM2.5 at that cell with the source's share in the model, then subtract its attribution share. At plus-24 hours here is what we forecast with this source, and here is without it — the difference is the projected impact. It is not a real-world measurement; it is what the trained model predicts would happen if compliance closed this specific source. The notice labels it 'projected impact of compliance' so an officer knows what they are looking at. The true test is the before-and-after tracking once the field team acts.

Evidence: docs/USER_GUIDE.md:271: 'projected impact of compliance chart (forecast with vs without this source's share).' agents/enforcement.py and the notice PDF writer implement this. No field validation yet; it is a forward-looking estimate from the model.

16.10 • Differentiation and Prior Art

Q. You claim the attribution model 'abstains' where it lacks skill. What does that actually mean, and how is it different from CPCB or SAFAR dashboards simply showing no data?

Expected

When a cell's out-of-sample R^2 falls below 0.15, the model refuses to assign source attribution and falls back to chemical-signature priors instead. CPCB and SAFAR show whatever the nearest station measured; they never attribute specific sources to specific locations. We say: 'Local model missed the ≥ 0.15 skill gate here—we fall back to cited priors rather than over-claim.' This is genuinely unique in operational platforms: most systems guess or hide the gap entirely.

Evidence: ml/attribution/shap_attribution.py:49 `MIN_HOLDOUT_R2 = 0.15`, :257 `if r2 < MIN_HOLDOUT_R2: raise ValueError(...)`; UI shows this in the cell story

Q. Your 80% prediction intervals claim they're 'calibrated.' Calibrated against what, and can you actually prove they're 80%?

Expected

We use split conformal regression: 75% of training data fits two quantile models (lower and upper bounds), then 25% of data calibrates the band width to achieve exactly 80% coverage on held-out residuals. Delhi's rolling multi-season benchmark shows 0.783 coverage; Kolkata hits 0.749 overall. Honest limitation: Kolkata's band fails in the 56–76 $\mu\text{g}/\text{m}^3$ range (0.547 at +72h)—that's where officer decisions happen. We publish breakdown by predicted level so this gap is visible, not hidden.

Evidence: ml/forecast/train.py:131 `CAL_FRACTION = 0.25` ;:134-162 `_conformal_level()` function; PS5_HONEST_AUDIT.md §3.2 shows coverage-by-quintile breakdown

Q. You describe an 'enforcement loop' that 'closes'—what does that mean operationally, and what does GRAP or AQI.in do instead?

Likely

When an officer dispatches an enforcement action, we freeze a 7-day baseline PM2.5 for that cell. After the intervention window, we compute drift-corrected before/after effect and export it PRANA-ready—so the city reports the measured impact against the national portal. GRAP is rule-based escalation with no outcome measurement. AQI.in is a dashboard: no enforcement targeting, no tracking. Nothing else in the official field closes this loop with measured data.

Evidence: supabase/migrations/20260719000003_intervention_tracking.sql defines `intervention_tracking` table with baseline freezing; agents/enforcement.py:306 documents outcome measurement

Q. You claim each enforcement recommendation includes a 'real Sentinel-2 patch.' How do you know it's real, and how is this different from showing a generic map tile?

Likely

The dossier is marked `allow_placeholder=False` —if the satellite patch hasn't actually been ingested for that source, we return nothing and say so rather than showing a placeholder. The app embeds the real GeoTIFF clip on the dossier card. CPCB and SAFAR have no source-level satellite evidence. IQAir shows generic basemap. We show the actual overhead image of the site—proof the source exists and its condition.

Evidence: agents/enforcement.py:743 `allow_placeholder=False` in `build_dossier()` docstring; API_CONTRACT.md line 31 `/enforcement/{id}/dossier` returns 'the real Sentinel-2 patch'

Q. How are the regulatory citations in the enforcement notice PDF different from just copying CPCB text into a template?

Likely

We embed 1,271 regulation chunks (NCAP, GRAP, CPCB dust norms, SPCB state-specific rules) into a pgvector store, embed the enforcement case description, and retrieve the top citations using RAG. Each citation is traced to its source document and the issuing authority is named on the notice—GRAP in Delhi-NCR only, CPCB dust norms elsewhere. The notice is court-defensible because the regulation actually applies to that city and is cited in full.

Evidence: rag/README.md describes 1,271 `kb_chunks` from PDF/HTML → embed → retrieval; agents/enforcement.py:661 `build_dossier()` retrieves and formats citations for the notice PDF

Q. You serve eight languages for citizen advisories. What prevents them from hallucinating medical advice like an LLM might?

Expected

No LLM anywhere in the product. All advisory strings are deterministic templates, city/language-specific, with placeholder substitution only. Hindi and Marathi templates have native-speaker team review. Kannada, Tamil, Telugu, Bengali, Gujarati are script-validated but awaiting independent review. This trades fluency for safety—we will not risk a hallucinated line in an asthma advisory.

Evidence: core/health_advice.py:1 says 'templated, cited, LLM-free'; PS5_HONEST_AUDIT.md §3.5 shows review status per language; ADVISORY_REVIEW.md lists status

Q. Your enforcement ranking claims to factor in 'inspector-hours' per source type. How is this different from just ranking by pollution share, and are those hours measured or guessed?

Likely

Ranking uses `value = benefit / inspector_hours`, where benefit is calibrated confidence × lower-bound forecast PM2.5 × people × urgency. Industrial sources get 8 hours estimated cost; traffic 6; biomass 1. These are team estimates, not measured from any board—we surface them on every card so an officer can override. Most systems rank only by contribution; we rank by bang-for-inspector-hour, which accounts for what an action costs.

Evidence: `agents/enforcement.py:175-182 INSPECTOR_HOURS dict; :217-248 _compute_value()` function; :213 formula documented in comments

Q. You say the system is 'city-agnostic' and live across ten cities—does that mean the same code runs everywhere, or does each city get custom tuning?

Expected

Same code, zero per-city patches. A city is a config (bbox, language set, regulatory authority). Seven metros were onboarded in a week from config alone. This is enforced: no city-specific conditional anywhere in the codebase. The models (attribution, forecast) do refit monthly per city because weather and source mix vary, but the architecture stays universal. CPCB and SAFAR require manual data entry per station.

Evidence: `core/config/cities/*.yaml` defines each city; `README.md` line 123 'Adding a city is one YAML file'; `ARCHITECTURE.md` §18 'universal spatial key (H3)'

Q. Your forecast skill numbers say +9% vs persistence in Delhi but negative in Jaipur. Why do you ship negative numbers instead of hiding or 'fixing' them?

Likely

The rolling multi-season benchmark is deterministic on fixed data. Jaipur's -6% at +24h to -14% at +72h are real, with confidence intervals entirely below zero—not noise. We show the breakdown by city so judges can see we measured honestly. The live monsoon window is worse because OpenAQ coverage is thin in Jaipur and skill degrades fast with data sparsity. Stating this is harder than claiming omit it, but it's how you earn trust.

Evidence: `PS5_HONEST_AUDIT.md` §3.2 table shows Jaipur skill with 95% CI intervals; `README.md` §54 'Negative numbers are kept'; `docs/BENCHMARKS.md` publishes full artifact per city

Q. What specific capability does VayuNetra have that CPCB's dashboard, SAFAR's station network, and IQAir's consumer app do NOT have together?

Expected

None of them combine real-time per-location source attribution + hyperlocal 1-km forecast + operationally-tracked enforcement + measured before/after effect in one loop. CPCB shows station AQI only. SAFAR adds apportionment but it's annual, not real-time. IQAir shows current air with health tips but no India regulation, no source-level action. VayuNetra attributes every 1 km² cell, forecasts it 72 hours out, ranks where to inspect, and measures what happened. That loop doesn't exist in production anywhere.

Evidence: `docs/PRD.md` §3.1 comparison table; `docs/SUBMISSION.md` §27-41 describes the four-question loop

Q. You mention a 'citizen complaint loop' where photos become enforcement candidates. Is this actually wired, or is it a sketch of an idea?

Hostile

It is actually wired. `POST /report` accepts photo + location, enters a public list with a 72-hour SLA clock, officers verify it, and verified reports become candidate sources for the next enforcement run. But: expert enforcement quality has not been rated (the honest audit calls this an 'F' on a named scoring criterion because zero officer ratings exist). The pipe works; whether the recommendations that flow through it are actually good is not yet independently validated.

Evidence: API_CONTRACT.md line 45 POST /report and /report/{id}/status ; PS5_HONEST_AUDIT.md §5 evaluation focus: 'Enforcement recommendation quality rated by domain experts: n = 0'

Q. How do you reconcile claiming 'SHAP explanations' as a headline capability when Delhi, Mumbai, and Bengaluru—your three largest demo cities—produce zero SHAP rows in production?

Hostile

SHAP requires 400+ measurements per cell; Delhi's per-cell data is pruned at 180 days, so depth is thin. Only Pune and partial Kolkata currently produce SHAP drivers; the rest fall back to abstain → cited priors, which is honest and good. The demo city story should lead with Pune (show SHAP working), then Delhi (show abstain refusal). Do not promise SHAP and click a Delhi cell.

Evidence: PS5_HONEST_AUDIT.md §2.4 'rows_with_shap_drivers = 0' for Delhi, Mumbai, Bengaluru; ml/attribution/shap_attribution.py:48 MIN_SAMPLES = 400

Q. You claim satellite data (Sentinel-5P NO₂, MODIS fire, VIIRS AOD) feed the attribution model. How many rows of satellite data are actually in the production database?

Hostile

Zero. The connectors exist—earth_engine.py implements S5P, MODIS, VIIRS retrieval. But the GitHub Actions workflow calls only openaq, cpcb, openmeteo; it never runs earth_engine. No scheduled ingestion means zero satellite rows in measurements today. The attribution model uses six CPCB pollutants + a fire marker = 0.0 in every row because there's no upstream data. This is a built-not-running situation, documented in the honest audit.

Evidence: PS5_HONEST_AUDIT.md §4 'MODIS / VIIRS: BUILT, NOT RUNNING...0 rows'; connectors/earth_engine.py:32 FIRE_BAND exists; .github/workflows/ingest.yml never calls earth_engine

Q. Your 1 km² resolution claim—'genuinely 1 km, genuinely calibrated'—is based on validation against what data?

Hostile

Synthetic fields. The downscaler (CNN) beats bilinear by 55.3% skill on synthetic held-out fields (n=64). Never validated against a held-out real station. Leave-one-station-out is computable from database data; it has been deferred. The honest phrasing: 'We interpolate 11 stations to 3,466 cells. The downscaler beats bilinear by 55% on synthetic fields. We have never validated it against a held-out real station.'

Evidence: PS5_HONEST_AUDIT.md §2.2 describes the validation gap; ml/coverage/dense_field.py:187 notes 'skill is measured on synthetic fields, n=64'

16.11 · Hardest questions — where this is weakest

Q. Your pitch claims six agents orchestrating the multi-agent system. The live demo shows four nodes executing in a Delhi monsoon trace. Which is accurate, and what determines whether enforcement actually runs?

Expected

Five, and the docs say five — the 'six' claim was found by our own audit and corrected across the README, landing page and every doc on 19 Aug. agents/graph.py registers five nodes: orchestrator, attribution, forecast, enforcement, advisory. The sixth thing people count is spike_gate, which is a conditional edge, not a node — it routes to enforcement only when there are focus cells, so a clean-air city legitimately runs four. We would rather show you the graph than argue the number.

Evidence: agents/graph.py — five add_node calls plus one add_conditional_edges (spike_gate)

Do not say: Do not say six agents. It is five nodes and a conditional edge.

Q. Your coverage endpoint returns a CNN downscaler skill of 55.3% versus bilinear interpolation, measured on held-out fields. Against what real measurement stations has this been validated?

Likely

It has not. The skill is measured on synthetic fields, $n=64$. Leave-one-station-out validation against real held-out stations has been deferred twice. We have never tested the downscaler against a held-out real station in Delhi's 11-station network.

Evidence: `ml/coverage/dense_field.py:187-188` states validation is 'on held-out synthetic fields; real held-out-station RMSE runs on Kaggle with EE AOD $\times$ CPCB'. Audit §2.2 confirms this and ranks leave-one-station-out as the highest priority fix.

Q. Your enforcement dossiers display SHAP explanations showing which pollution sources drive the forecast AQI. Which Indian cities actually produce SHAP drivers today, and why are Delhi and Mumbai falling back?

Likely

Only Pune and part of Kolkata produce SHAP drivers. The GBM+SHAP model requires 400+ samples per cell to be trustworthy, but Delhi's measurements are pruned at 180 days, so per-cell depth is thin. Delhi, Mumbai, Bengaluru and Hyderabad fall back to signature-prior attribution with an honest abstain message stating the local model missed the skill gate.

Evidence: `ml/attribution/shap_attribution.py:48` defines `MIN_SAMPLES=400`. Audit §2.4 counts: Delhi $n=132$ signature-v1 with 0 SHAP rows; Pune $n=60$ hybrid-gbm-shap-v2 with 60 SHAP rows. Audit §3.1 confirms abstain logic working in Delhi.

Q. Your benchmarks claim the forecast skill is positive across ten cities. Jaipur's recent rolling skill is minus 4 to minus 14 percent at 48-72 hours. Is this noise or a real failure, and what have you learned?

Likely

Real failure, not noise. All three confidence intervals (-0.25 to -0.058 at +24h, -0.189 to -0.002 at +48h, -0.442 to -0.188 at +72h) lie entirely below zero. The model loses to persistence on monsoon air. Likely cause: blend weights hit $w=1.0$ for Jaipur at these horizons, discarding persistence contribution, while cities like Delhi sit at $w=0.65-0.9$ and win. Unfixed, quantified, and published.

Evidence: `docs/benchmarks/jaipur_live.json:66-80` and `:451-454` show `skill_model_vs_persistence = -0.16, -0.1, -0.314` with confidence intervals entirely negative. Audit §3.2 problem 1 explains the blend-weight hypothesis.

Q. Your 80% prediction interval claims 0.75 overall coverage on Kolkata's rolling protocol, but this hides the real problem. What happens in the decision-critical 50–76 $\mu\text{g}/\text{m}^3$ range where regulation categories change?

Likely

Coverage drops to 0.668 at 56–76 $\mu\text{g}/\text{m}^3$, degrading to 0.620 at +48h and 0.547 at +72h — exactly where an officer must decide between Satisfactory and Moderate. This is a model under-dispersion problem in the mid-to-upper range, not fixable by recalibration alone. We publish the breakdown by quintile and state the honest measurement rather than hiding it.

Evidence: `docs/benchmarks/kolkata.json` shows `pi80_coverage_by_predicted_quintile`: [8–25: 0.803, 25–38: 0.778, 38–56: 0.761, 56–76: 0.668, 76–245: 0.733]. Audit §3.2 problem 3 ranks this unfixed but quantified.

Q. Your enforcement recommendations are scored by 'domain experts' per the brief. How many independent raters have evaluated your worklist, and what are their scores?

Expected

Zero. We built a 20-minute expert rating protocol, outreach email template, and a one-page rating sheet. The kit exists in `docs/EXPERT_OUTREACH.md`. No independent reviewer has returned scores. This is an 'F' on a named evaluation criterion, and a single officer's rating would move it to a C.

Evidence: `docs/EXPERT_OUTREACH.md` exists; `docs/expert_reviews/` directory does not exist or is empty. Audit §5 identifies this as the only pure F-grade missing from the evaluation focus table.

Q. Your citizen advisories are deployed in eight Indian languages. How many have been reviewed by a native speaker to ensure medical advice and register are appropriate?

Expected

Two: Hindi and Marathi, reviewed by team members on 2026-08-18. Six languages (Kannada, Tamil, Telugu, Bengali, Gujarati, plus the IVR framing for all languages) are script-validated only — characters render correctly, but no native speaker has confirmed the wording is idiomatic or medically sensible. A hallucination in an asthma advisory is not an acceptable risk.

Evidence: docs/ADVISORY_REVIEW.md:10-19 and lines 22–30 show 2-of-8 reviewed and 6 'pending'.
tests/test_advisory_script_check.py validates script only, not meaning.

Q. Your technical pitch lists Sentinel-5P NO₂, MODIS, and VIIRS fire data as attribution inputs. Are these currently feeding the live system, and when did they last update?

Hostile

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. Your source attribution is described as a receptor model, but it does not use chemical speciation, elemental carbon fractionation or ion balance validation. What exactly are you attributing?

Hostile

We infer sources from six CPCB pollutants (NO₂, CO, SO₂, PM₁₀/PM_{2.5} ratio, fire count) plus satellite NO₂ and forecasted advection. We call it marker-based source apportionment with signature priors, not receptor modeling. We claim a prioritisation signal with stated confidence, validated in shape against published inventories (cosine 0.991 vs SAFAR-Delhi), not receptor-model accuracy. We abstain when local fit has no skill rather than over-claiming.

Evidence: `ml/attribution/signatures.py`:60–72 shows the six pollutant markers. `ml/attribution/shap_attribution.py`:32–40 defines `SOURCE_MARKERS`. Audit §3.1 frames this explicitly: 'we do not claim receptor-model accuracy — we claim a prioritisation signal with stated confidence'.

Q. Your frontend is a React PWA managing real-time air-quality intelligence for officers and citizens. What unit-test coverage does the component library have?

Hostile

Zero. Only e2e smoke and journey tests exist (17 tests total across `web/e2e/`). The component library, hooks, and state management have no unit tests. A judge opening the repo will see TypeScript components with no test files next to them.

Evidence: `web/` directory contains no `.test.ts`, `.test.tsx`, `.spec.ts` or `.spec.tsx` files. Only `web/e2e/smoke.spec.ts` and `web/e2e/journey.spec.ts` exist. Audit §8 lists 'Frontend unit tests: 0. e2e only (8 smoke + 9 live journey)'.

Q. Your system produces a 'response time from signal to intervention in seconds' metric. Does this measure how fast municipal officers actually respond to your recommendations, or something else?

Hostile

It measures something else: our pipeline's internal latency from orchestrator to advisory output, roughly 1.1 seconds. This is not an organisation's response time — no officer's end-to-end response has been measured. The honest phrasing is 'our pipeline

turns a signal into a cited, ready-to-sign recommendation in about a second,' not 'we shortened a municipality's response to an intervention in seconds.'

Evidence: API endpoint /latency returns pipeline wall-clock time. Audit §6 claim 4 flags the response-time phrasing as outrunning evidence and §5 scores this as 'C — we measure our *pipeline*, not an organisation's response time.'

Q. Your forecast model blends raw predictions with persistence to improve calibration: $\text{new_forecast} = w \cdot \text{model} + (1-w) \cdot \text{persistence}$. Your blend weights w vary from 0.05 to 1.0 across cities. Why does Jaipur sit at $w=1.0$ where it loses?
Hostile

Blend weights are fitted to minimise RMSE on a recent calibration tail. At $w=1.0$, the persistence baseline contributes zero, leaving only the raw model — which fails on Jaipur monsoon air. Weights were meant to adapt dynamically, but recalibration data got confounded by live ingestion backfill between runs. The deterministic benchmark itself is sound, but this deserves a clean re-run on frozen data before shipping.

Evidence: ml/forecast/train.py computes blend_weights via recent RMSE minimisation. Audit §3.2 problem 1 shows Delhi 0.65–0.9, Jaipur exactly 1.0, with explicit note: 'live ingestion backfilled underneath it (Jaipur's support rows went 299 → 597 between runs)'.

Q. Your system is in production across ten cities. How many real-world interventions triggered by your enforcement recommendations have been completed, and what was their measured impact on air quality?
Hostile

Zero real interventions have been triggered. There is no deployment to a live municipal authority. The system is live as a public dashboard for the hackathon finale, with synthetic recommendations showing the data-to-action loop. Outcome measurement infrastructure exists (PRANA-ready dossiers with baseline and drift-corrected before/after), but no municipality has acted on the worklist.

Evidence: Audit §3.3 states 'Outcome measurement: VERIFIED' but means infrastructure only. No deployment is mentioned anywhere. The system is a DSS kit; adoption by a city is future work. The brief asks to 'demonstrate impact'; we show the loop closes on paper.

Q. Your public-display advisory channel is documented in channels/README.md and formats advisories for outdoor screens and municipal notice boards. Is this channel built and running?
Hostile

Documented, not built. The channel is described in channels/README.md but no implementation exists. PWA, Telegram bot, and Twilio IVR are live; public-display rendering is not. This is a planned feature, not a delivered one.

Evidence: Audit §3.5 states 'public-display mode' is **documented but not built**. channels/README.md may describe it; agents/advisory.py implements en/hi/mr/ta/te/bn/kn/gu only, not a display channel.

16.12 · Source Attribution Methodology and Validation

Q. You claim a hybrid GBM+SHAP method for attribution. When does it actually fire, and what causes it to fall back to pure priors?
Expected

The hybrid model requires a minimum of 400 samples per cell to train (MIN_SAMPLES in ml/attribution/shap_attribution.py:48). It trains on source-marker features — NO₂, CO, satellite NO₂, SO₂, PM10/PM2.5 ratio, and fire counts — and computes SHAP explanations. But the model only ships if it achieves out-of-sample $R^2 \geq 0.15$ on a holdout test set (ml/attribution/shap_attribution.py:49). When either gate is missed, the system falls back transparently to chemical-signature priors. This abstain behaviour — refusing to assign ML blame when untrustworthy — is recorded in the evidence field and visible on every cell story.

Evidence: ml/attribution/shap_attribution.py:48-49, 244-260; ml/attribution/attribute.py:131-145; docs/PS5_HONEST_AUDIT.md §2.4

Q. A user looks at a cell's attribution. How do they know which method produced the number — hybrid SHAP or pure signature?

Expected

Every attribution row carries a `method_version` field that tells the story. Three values appear in live data: 'hybrid-gbm-shap-v2' for cells passing both gates, 'signature-v1' for cities too thin to train on, and 'signature-citymean-v1' for marker-less cells that get shrunk toward the city hybrid mean. The field is returned in the API response (api/main.py comment at line 296-300) and rendered in the UI as a badge on the cell story, so a reader immediately sees whether the model was available, rejected itself, or wasn't attempted.

Evidence: ml/attribution/attribute.py:26-28, 56, 99, 106, 141-145; api/main.py:296-300; CONTRIBUTION_VALIDATION.md §10

Q. You validate against published inventories. What exactly are you testing, and how honest should judges be about what that proves?

Expected

The validation compares city-mean source shares from your live attribution against shares from published studies: SAFAR-Delhi 2018 for Delhi, CSTEP 2022 for Bengaluru, Urban-Emissions syntheses for Mumbai. Critically, it renormalizes both sides to exclude 'transported' and 'other', because a city emission inventory counts what is locally emitted, while your receptor method sees what is locally breathed — regional PM2.5 drifting in has no inventory counterpart. The match is indicative, showing whether the ranking and rough magnitude align, but not cell-by-cell accuracy. No public, current, cell-level ground truth exists for any Indian city.

Evidence: docs/CONTRIBUTION_VALIDATION.md:1-24, 143-163; ml/attribution/inventory.py:10-23, 100-105; AUDIT §3.1

Q. The validation published in your docs quotes Delhi cosine 0.991. Is that number honest, and does it matter that you read cosine instead of something else?

Expected

The 0.991 is honest but misleading — it was recomputed live on 19 August and is actually stronger than the published table stated (0.88). But cosine similarity over four renormalized buckets is dominated by the largest component, so it masks disagreement in smaller shares. The mean absolute difference — 0.042 for Delhi — is the more honest headline. The audit explicitly flags this: cosine looked at alone would let a 7% gap in biomass (0.067 vs 0.000, due to monsoon timing) inflate apparent agreement. Read `mean_abs_diff`, not cosine, to see the real per-bucket gaps.

Evidence: docs/CONTRIBUTION_VALIDATION.md:155-163; AUDIT §2.3 and 3.1; ml/attribution/inventory.py:107-113, 149

Q. You're claiming to know what causes pollution in each 1-kilometer cell. How could you possibly validate that at cell resolution?

Expected

You can't, and the validation doesn't attempt it. The published studies are city-mean or city-scale inventories, not cell grids. What validation establishes is that the aggregate attribution — the city-wide ranking of sources — matches the published order and rough magnitude. This is weaker than claiming cell-level accuracy, but it's what the evidence supports. The honest read: 'Our ranking of sources at city scale matches independent studies. Cell-level shares are a defensible inference from local markers and regional transport, with confidence calibrated from agreement between the gradient-boosted and signature-prior paths.'

Evidence: docs/CONTRIBUTION_VALIDATION.md:1-6, 141; AUDIT §2.2 and §6; ml/attribution/shap_attribution.py:224-232 (confidence formula)

Q. The confidence score on every attribution row — how is it computed, and what does it actually measure?

Expected

Confidence is computed from four independent signals: agreement between SHAP and signature-prior methods (35%), model fit R^2 (20%), sample depth over the trailing 72 hours (10%), and a floor at 30% and ceiling at 95%. The formula is: $0.30 + 0.35\text{agreement} + 0.20r^2 + 0.10*\text{depth}$ (ml/attribution/shap_attribution.py:231). It measures 'how much do independent estimation approaches agree, and how well did the model generalize' — not 'how certain is the true source share' in an absolute sense. Two cells with identical true shares could have different confidence scores if one has deeper history or better method agreement.

Evidence: ml/attribution/shap_attribution.py:224-232; ml/attribution/attribute.py:112-114

Q. Your SHAP-based explanations — they're built into the system, but do they actually run for the three cities you demo on?

Likely

Not currently. Delhi, Mumbai, and Bengaluru are running signature-v1 with zero SHAP rows in live data (verified 18 Aug: Delhi 132 rows all signature-v1, Mumbai 120, Bengaluru 78). The hybrid model needs 400+ samples per cell; Delhi's raw measurements are pruned at 180 days, so per-cell depth is thin. Only Pune and part of Kolkata produce hybrid-gbm-shap-v2 rows with SHAP drivers. The audit calls this a 'demo landmine' — if you click a Delhi cell on stage, viewers see the abstain message rather than SHAP explanations, though the pitch bills SHAP as a headline feature.

Evidence: AUDIT §2.4, lines 129-139; docs/PS5_HONEST_AUDIT.md §2 (demo risks); ml/attribution/shap_attribution.py:48, 244-245

Q. Your satellite NO₂ and fire counts — you list these as inputs to attribution. Are they actually live?

Likely

Running since 19 Aug. `connectors/earth_engine.py` is now called by the daily job in `.github/workflows/ingest.yml`, and the database holds 136 Sentinel-5P `no2_sat` rows and 164 MODIS `fire` rows. Say the caveat yourself before they ask: the cadence is **daily, not hourly**, so these are a slowly-moving prior rather than a live signal, and the `fire` marker still reads near zero in a monsoon because there is genuinely nothing burning. Until 19 Aug this was our biggest gap — the connectors existed but nothing scheduled them, so both markers were identically zero. Sentinel-2 imagery in the dossiers is still an offline run, not a scheduled ingest.

Evidence: `.github/workflows/ingest.yml` daily job; 136 rows `source='s5p' variable='no2_sat'`, 164 rows `source='modis' variable='fire'` in `measurements`

Do not say: Do not say satellite is a live feed — it is a daily job.

Q. You say 'transported' is a source category. How does that differ from 'outside Delhi' in the TERI study, and why isn't validation telling you something different?

Likely

'Transported' is receptors-side: PM_{2.5} that drifted in from upwind and is now in this cell. TERI's 'outside Delhi' is a dispersion model output: the fraction of PM_{2.5} at the city boundary traced back to sources beyond Delhi-NCR. Your transported share (Delhi city-mean ~14%) is smaller than TERI's outside share (~64–74%), but they measure different things at different fidelity. A receptor method sees the local signal; dispersion models see the full regional budget. The validation correctly excludes transported from the comparison, but judges should know the methods are incommensurable.

Evidence: docs/ATTRIBUTION_VALIDATION.md:21-24, 62-65; ml/attribution/signatures.py:70; AUDIT §3.1

Q. Bengaluru's industrial share in your attribution is 32% after renormalization. Published studies put it at 9–11%. That's a three-fold gap. What's going on?

Likely

Bengaluru has little heavy industry inside the city; what your method sees as 'industrial' is actually combustion point sources: diesel generators and kilns. The SO₂ signature and PM10/PM2.5 ratio key on these. Your per-cell model in Peenya and Bommasandra cells over-weights generator emissions because they co-locate with monitor sites. The studies report factories only. So read your 'industrial' bucket as 'combustion point sources', not factories. This disagreement is the one your docs flag (ATTRIBUTION_VALIDATION.md:121–124); the honest headline is: 'Our Bengaluru model sees diesel generators and kilns we label industrial; published studies report large factories. The methods target different sub-sectors.'

Evidence: docs/ATTRIBUTION_VALIDATION.md:114–130; ml/attribution/shap_attribution.py:32-39 (SOURCE_MARKERS); AUDIT §3.1

Q. Your chemical-signature priors are ratio-based: NO₂/CO for traffic, PM10/PM2.5 for dust. Can't those ratios be gamed or misleading?

Hostile

Yes, in principle. A source with no traffic (e.g., a construction site) can produce NO₂ if it runs diesel equipment. A city with heavy coal burning will have different PM10/PM2.5 ratios than one with petrol exhaust. The signature priors work because Indian cities have relatively stable sectoral mixes and emission profiles. But they are not fingerprints — they are indicative associations. That is why the gradient-boosted model exists: to refine the priors using recent local data. When the model passes the skill gate, you blend it 60% model + 40% prior (BLEND_WEIGHT in shap_attribution.py:51). When the model fails, you fall back to pure priors transparently.

Evidence: ml/attribution/signatures.py:60-75; ml/attribution/shap_attribution.py:51, 218-221; ml/attribution/attribute.py:1-8

Q. Your validation compares city-means, but enforcement decisions are made on individual cells. Doesn't that hide cell-level error?

Hostile

Yes, deliberately. The validation is honest about its scope: it proves nothing about cell-level shares. A cell with 60% traffic and 40% dust could be right in ranking and wrong in magnitude. Enforcement priority does not require cell-level accuracy — it needs ranked ordering within a cell and exposure-weighted action. The system shows confidence and the method badge, so an enforcement officer sees 'construction_dust is dominant here (60%), confidence 0.72, signature-based because the local model had low holdout R²'. That transparency lets them weight it accordingly. But no, the validation does not establish that the 60% is precise.

Evidence: docs/ATTRIBUTION_VALIDATION.md:1-6; AUDIT §3.1, 6; ml/attribution/inventory.py:141-163; agents/enforcement.py prioritization logic

Q. You're running the validation in August monsoon, when biomass is zero. How representative is Delhi's 0.991 cosine to the winter spike-burning season?

Hostile

It is not representative at all. Biomass reads 0% in your August run because there is no stubble burning (ATTRIBUTION_VALIDATION.md:59, AUDIT §2.3). TERI's published summer share for biomass is 15%, which your method correctly attributes to zero during monsoon — so the agreement is helped by seasonal alignment, not by model correctness. The audit notes this explicitly: 'seasonal alignment is inflating the similarity.' To know whether the validation holds year-round, re-run it in Nov–Dec, when biomass burning is real. Until then, quote 0.991 only for the monsoon window, or regenerate the table with a winter snapshot.

Evidence: docs/ATTRIBUTION_VALIDATION.md:59, 165-170; AUDIT §2.3 caveat (1); ml/attribution/signatures.py:69

Q. A marker is 'observed' if it appears in 30% of the cell's samples. Why that threshold, and what happens if a city simply doesn't measure a pollutant?

Hostile

MIN_MARKER_COVERAGE = 0.3 (shap_attribution.py:179) excludes phantom blame. LightGBM's missing-value branch still produces SHAP contributions for unobserved markers, which would incorrectly assign identical blame (e.g., 'industrial') to every cell without SO₂ data. The threshold filters that noise. But if a city has no SO₂ monitors at all — which is the case for several Indian cities — then industrial apportionment falls back to a pure PM10/PM2.5 ratio. That is weaker but honest. The audit records that Delhi has 11 stations and 3,466 cells; only four cells can see NO₂ today, so 21 cells are shrunk toward the city mean (AUDIT §3.2).

Evidence: ml/attribution/shap_attribution.py:179, 195-199; ml/attribution/attribute.py:73-79; AUDIT §3.1 (data coverage)

Q. You blend the SHAP model 60% and signature priors 40%. Why those weights, and did you validate that blend on held-out data?

Hostile

BLEND_WEIGHT = 0.6 is hard-coded in shap_attribution.py:51 with no documented justification in the code. The weight reflects a design choice ('prefer the model, but don't ignore chemistry') rather than a measured-optimal tuning. No validation of the blend itself appears in the public docs or code — you validate the model, then apply fixed weights. This is a reasonable pragmatic choice (the model is trained and gated; the blend is not over-fitted), but it is not empirically derived. If a judge asks, the honest answer is: 'The weights express domain belief — models often beat priors, but chemistry never lies. We chose 60/40 as a conservative blend and did not tune it on held-out test data.'

Evidence: ml/attribution/shap_attribution.py:51, 218-221; ml/attribution/shap_attribution.py dataclass comment (lacks justification)

16.13 · Questions we nearly missed

These were not in the prepared set. An adversarial reviewer added them.

Q. The prepared answers cite a PS5_HONEST_AUDIT.md dated 19 August. Are those findings still accurate, or have subsequent commits fixed some of the issues?

Git history shows multiple fixes after the audit was written: commit f31bf20 'fix: close the audit's findings — satellite is live,' commit 6251417 'docs: PS-5 honest audit — every claim verified,' and several others before the current HEAD. The satellite ingest issue appears to have been fixed. Answers based on the audit should disclose their date and whether fixes have landed since.

Q. You claim three technologies 'outrun their evidence': six agents, one-km validation, and SHAP explanations. But do the actually-cited files (README.md, Landing.tsx, SUBMISSION.md) really claim six agents, or does the audit itself misquote them?

The audit misquotes the files. All three say five agents, not six. The audit claims lines 66 say they say six, but they don't. This is the audit's error, not a weakness of the product. The prepared answer should either correct the audit's citation or explain that five agents plus a spike gate is what was claimed all along.

Q. What is the status of the satellite data gap NOW, not as of 19 August? Are GEE secrets configured in production? Is earth_engine being called on schedule?

The ingest.yml includes the step and it is scheduled. But whether GEE secrets are actually configured in production is not verified by reading code. A live test (e.g., querying the measurements table for earth_engine rows) would answer this. The prepared answer treats 'built' and 'running' as equivalent, when 'scheduled but gated on secrets' is the actual state.

Q. The answers repeatedly say 'this is stale/zero-coverage/unvalidated.' Are these all from the 19 August audit, or are they things the team still claims are true on demo day?

Most weak points reference the audit. But for a hostage-to-fortune question like 'Why no satellite data?' a judge will ask about the CURRENT state, not what the audit said five days ago. If satellites have been fixed, say so. If they haven't, say the secrets aren't there. Don't let the judge catch you defending a finding you've already fixed.

Q. What does the 'other' category represent, and why is it floored at 5% (OTHER_FLOOR = 0.05)?

The 'other' category captures unexplained PM2.5 variance — sources not well-represented by the six markers (NO₂, CO, SO₂, PM10/PM2.5, fire, advected PM2.5). It is floored at 5% (shap_attribution.py:50) to prevent the model from claiming 100% attribution confidence when it has explained only four categories. This is a regularization choice: even a perfect marker fit admits ~5% irreducible uncertainty (unobserved industrial processes, resuspension, photochemical secondaries). The floor makes every cell's shares more honest about model limitation.

Q. Why is the SHAP model trained on a time-ordered 80/20 split rather than a random hold-out set?

Time-ordering (shap_attribution.py:103, 109) mimics the operational scenario: train on historical data up to time T, test on the future T+Δ. Random splits would leak temporal autocorrelation and overstate generalization — a model that merely memorizes recent trends would look skilled. A time-ordered split ensures the R² gate measures true out-of-sample skill in the forward direction, which is the only direction that matters for operational deployment.

Q. The 'transported' baseline is hardcoded to 0.15 in the signature priors (signatures.py:70). How is this value chosen, and does it adapt by season?

The transported baseline (signatures.py:70: `'transported': 0.15`) is a fixed regional background representing PM2.5 that has drifted in from upwind. It does not adapt by season in the signature priors; it is a stationary assumption. In the hybrid model, transported shares can vary per cell per window via SHAP if the features (advected_pm25, if present) capture regional signals. The fixed 0.15 is a placeholder that assumes ~15% of every cell's PM2.5 is regionally sourced. This is crude — monsoon upwind sources differ from winter stubble-burning upwind. A seasonal override would be more honest; it is not currently implemented.

Q. Why is the holdout R² threshold set to exactly 0.15? What happens if a model gets 0.14?

MIN_HOLDOUT_R2 = 0.15 (shap_attribution.py:49) is a skill gate: R² = 0 means the model predicts no variance better than the mean; R² = 0.15 means it explains 15% of out-of-sample variance. The choice of 0.15 is not justified in the code. It is low enough to allow models with modest skill (traffic models on highway cells) but high enough to reject noise-trading models. A holdout R² of 0.14 triggers fallback to signature priors (attribute.py:257-260). This is a sharp cliff; a model at 0.14999 falls back while one at 0.15001 ships. No adaptive scaling softens this boundary — it is binary: trustworthy or abstain.

Q. The API response includes 'evidence' with SHAP drivers and model_r2. How much of the enforcement workflow actually reads evidence, or does it just use the final shares?

The 'evidence' field (attribute.py:100, api/main.py:627-628) is returned to the API but its use in the enforcement pipeline is not visible in the prepared answers. The agents/enforcement.py code likely reads shares, confidence, and method_version to prioritize cells, but may not inspect evidence depth or SHAP drivers. If evidence is rich data that downstream code ignores, it is expensive transparency with limited impact. If enforcement only reads the final shares and confidence, then method_version (which tells you whether SHAP was used) is the critical signal, and evidence is a nice-to-have for explainability but not actionable.

Q. You mention the 180-day measurement pruning policy (AUDIT §2.4, attribute.py comments). Is this a hard delete or archival?

The ingest.yml (lines 188-196) shows that raw readings older than 180 days are archived to a private Supabase Storage bucket, then deleted from the live measurements table. This is a retention strategy: keep recent data hot for model training, move old data cold for long-term audit trails. The consequence: no cell can ever accumulate >180 days of history for feature engineering. For monthly phenomena (e.g., seasonal biomass burning), 180 days captures ~6 months of seasonal variation — enough for year-round skill. For interannual trends or rare events, 180 days is too short. This policy is not discussed in any answer but fundamentally shapes data depth and model refresh cycles.

Q. Q16: Your SUBMISSION.md claims 'SHAP explanations' are a headline capability (line 37), but the HONEST_AUDIT shows Delhi has zero SHAP rows (section 2.4, line 129). Do judges testing the demo on Delhi see the SHAP explanations you claim?

No. Delhi, Mumbai and Bengaluru run signature-only attribution with zero SHAP drivers, because MIN_SAMPLES=400 (ml/attribution/shap_attribution.py:48) cannot be met with 11 stations and 180-day retention. Only Pune (60 rows with SHAP) and partial Kolkata (18 rows) produce SHAP drivers. The demo city (Delhi) shows abstention fallback instead. The SUBMISSION should state: 'SHAP explanations where data depth allows (Pune, partial Kolkata); signature priors elsewhere with transparent abstention.' The landmine in HONEST_AUDIT 2.4 must be rehearsed: either demo Pune's SHAP or lead with Delhi's abstain path.

Q. Q17: Your workflow runs satellite ingest at line 100-126 (earth_engine connector), but what is the last timestamp a satellite measurement was actually written to the database?

The HONEST_AUDIT (section 2.4, line 165) confirms satellite data has NEVER been scheduled: 'Satellite thermal anomalies — BUILT, NOT RUNNING... the marker reads **0.0 in every live attribution row** because nothing ingests it on a schedule. Same for no2_sat.' The earth_engine connector exists (connectors/earth_engine.py) and the workflow step was added, but it was 'MISSING entirely — the connector existed and had been run once by hand.' Every attribution row carries no2_sat=0.0 and fire=0.0 because the scheduled ingest never runs without GEE credentials, and no one has configured those secrets. This means two of the brief's named satellite technologies are listed in SUBMISSION.md but not actually running.

Q. Q18: You state in README.md line 49 and elsewhere that you run '10 Indian cities'. But your BENCHMARKS.md notes (section §Recent window, line 195-196) that the actual training window is '90 days for Bengaluru, 60 for Delhi and about 38 for the other seven.' Are some cities' forecasts trained on insufficient data?

Correct. HONEST_AUDIT section 3.2 (line 195-196) confirms: 'this is not a 90-day window for most cities — live ingestion started on different dates, giving Bengaluru 90 days, Delhi 60 and the other seven about 38.' The phrase 'all ten cities' is technically true but misleading if the window is 38 days for seven of them. BENCHMARKS.md includes a 'recent window' benchmark for all ten, but explicitly flags: 'the PI80 column is unstable by construction... this is not a 90-day window for most cities... each artifact carries its own window rather than a shared label.' The claim should be qualified: '10 cities in deployment, with validated 90-day benchmarks on 3 (Bengaluru, Delhi, Mumbai) and recent-window (38-90 day) benchmarks on all ten, regenerated daily.'

Q. Has an engineer independently verified that the spike_gate conditional edge works correctly when enforcement is skipped? What happens to the enforcement list in state when enforcement_node is not executed?

The spike_gate routes to either 'enforcement' (if focus_cells or forecast_spike are truthy) or directly to 'advisory' (line 320-331, agents/graph.py). When advisory is reached without enforcement, state['enforcement'] will be missing or an empty list. The advisory node must handle this gracefully. Test coverage for this path should exist in tests/test_grap.py (note the filename typo). The latency traces in PS5_HONEST_AUDIT.md:2.1 show this path active: ['orchestrator', 'attribution', 'forecast', 'advisory'] with no enforcement step. No public evidence that an officer has validated this produces sensible output when enforcement is skipped.

Q. The answer states Kolkata's prediction interval coverage is '75% overall and 67% in the 56–76 µg/m³ range — critical for decision boundaries.' Has anyone checked whether this actually breaks enforcement decisions in practice, or is this purely a statistical observation?

PS5_HONEST_AUDIT.md:3.2 problem 3 documents this as a model under-dispersal issue ('quantile models under-dispersing in the mid-to-upper range — a model problem, not a calibration one'). The audit notes this is 'UNFIXED, quantified, and disclosed.' However, there is no evidence of an enforcement officer or meteorologist validating whether 67% coverage at the decision boundary actually causes real-world recommendation errors. This is a gap in validation (related to Q11: zero officer reviews).

Q. Why does the audit (PS5_HONEST_AUDIT.md:2.1) state that Delhi's live traces show enforcement skipped ('nodes: orchestrator, attribution, forecast, advisory') during monsoon, but the enforcement_node exists and should run whenever focus_cells are populated? Is enforcement being deliberately skipped or is focus_cells empty?

The spike_gate routes based on presence of focus_cells or forecast_spike. During monsoon with clean air, spiking cells are empty, so focus_cells is empty, so the gate routes to 'advisory' directly. The orchestrator (line 148-152) only populates focus_cells if PM2.5 > 120. This is working as designed: enforcement is truly skipped on clean-air days. However, a hostile judge might ask: 'You say you skip enforcement on clean air, but the README promises continuous multi-city enforcement intelligence. Do citizens in nine other cities get stale enforcement recs because your orchestrator's spike detection is too strict (120 µg/m³)?' This reflects a design tension: hyperlocal enforcement requires non-zero hotspots.

Q. The answer to Q12 states 'Leave-one-station-out validation is computable from existing data and would convert the most attackable claim into the strongest.' Has this validation been run since the audit, or does it remain a recommended but unimplemented action?

PS5_HONEST_AUDIT.md:9 item 3 rates this as '~half a day' effort. The docs/benchmarks/ directory contains delhi.json and *_live.json but no leave-one-station-out results. This validation has NOT been run. The 1 km forecast remains validated only on synthetic fields (55.3% skill vs bilinear) with no real held-out station test. The honest answer when attacked: 'This validation is feasible but has not yet been completed.'

Q. How many of the ten cities have actually been tested with a live enforcement officer, and what was the feedback? Does every city have at least one example of a drafted notice that an officer could review?

PS5_HONEST_AUDIT.md:3.4 states 'intervention effectiveness is only real for Delhi, because only Delhi has a winter of dated government orders to replay.' Only Delhi has historical data to validate against. The other nine cities have enforcement_rec in the database but no officer has rated them. This is the 'F' from Q11. An attack: 'You have ten cities but zero independent validation in nine of them. What if the attribution is wrong in Bengaluru but right in Delhi?'

Q. Regime-dependent dispersion: Does Kolkata's quantile model's under-dispersion vary by season or weather regime? Have you tested per-regime quantile models?

This is a critical test we haven't run. The benchmark shows Kolkata's worst quintile coverage is 0.668 at +24h (autumn/monsoon) but we don't have winter data isolated yet. Split-conformal assumes stationarity; if dispersion itself is regime-dependent, a single Q is wrong by construction. We should test per-season quantile models on the Dec 2025–Feb 2026 winter window now in the database.

Q. Decision-threshold validation: You claim 56-76 $\mu\text{g}/\text{m}^3$ is the 'operationally critical' CPCB boundary. Has a CPCB officer or municipal enforcement authority confirmed this threshold determines their actual actions?

We inferred it from CPCB's Satisfactory (≤ 60) and Moderate (61-90) band, but we haven't validated it operationally. A real officer might use 65 $\mu\text{g}/\text{m}^3$, or context-dependent thresholds. Before claiming a range is 'operationally critical,' we should ask: at what predicted level do you actually change decisions? Our coverage at your actual decision threshold might be much better or worse than our hand-picked range.

Q. Conditional coverage across other cities: Kolkata's worst quintile is 0.547. What's the worst quintile for Delhi, Mumbai, Bengaluru, Hyderabad? Is Kolkata an outlier or representative?

We show this in the benchmark JSON for Delhi (0.704) and implicitly for others, but we don't highlight it. Need to run the same conditional-coverage table for all 10 cities and report: Which cities have a weak range? Is it always the mid-to-upper range, or do different cities fail in different ranges? This would show whether Kolkata is a specific sparse-data problem or a systemic flaw in our quantile models.

Q. Coverage target appropriateness: Why 80%? A forecast tool that informs enforcement actions and health advisories might need 90% or 95%. Have you surveyed stakeholders on the required coverage level for your use cases?

We chose 80% because split conformal naturally targets NOMINAL 80% ($1 - \alpha$ where $\alpha = 0.2$). But the operational question — 'what coverage do officers and health authorities require?' — is separate. An enforcement notice should be highly reliable. An advisory to the public might tolerate lower coverage. We haven't asked stakeholders; we should.

Q. Interpolation-quantile interaction: The audit notes the 1km downscaler is validated only on synthetic fields, not held-out real stations. If the downscaler is biased in certain spatial patterns for Kolkata, could that explain the mid-range under-dispersion? Have you tested whether a leave-one-station-out validation of the downscaler correlates with quantile-model dispersion by region?

We haven't connected these. If the downscaler introduces correlated errors in high-concentration cells, the quantile model would see apparent under-dispersion in those ranges. Testing leave-one-station-out on the downscaler (audit §2.2, item 3) would tell us if interpolation bias is part of the problem. This is a medium-effort experiment that could reframe the whole diagnosis.

Q. Has any officer completed a full end-to-end enforcement cycle — from system-proposed recommendation through approval, dispatch, field closure, and measured outcome?

No. Zero officers have used the system. All 60 enforcement recommendations in the live database have status='proposed' with no approvals, dispatches, or closures. The entire enforcement feedback loop is untested in the field.

Q. What does an officer in Delhi actually see when they click a cell in the Cell Story? How does this differ from Pune, where SHAP explanations are available?

Delhi (no SHAP): officers see 'chemical-signature attribution — local model missed the ≥ 0.15 skill gate here, we fall back to cited priors' with source-contribution bars but no SHAP drivers. Pune (with SHAP): officers see SHAP drivers in $\mu\text{g}/\text{m}^3$ showing which factors drove the model's attribution. The difference is that Delhi officers get an honest 'we don't know' message, while Pune officers get explainable model reasoning.

Q. The downscaling validation shows 55% skill over 64 synthetic held-out fields. What's the performance of the baseline methods (bilinear interpolation, inverse-distance weighting) that this 55% is meant to beat? And does 55% skill constitute a real improvement in operational terms?

Bilinear interpolation RMSE: 4.97; CNN RMSE: 2.22 (from /coverage?city=delhi). This is a 55% reduction in interpolation error, which is real. However, this is measured on synthetic fields only. The operational question — whether this translates to better enforcement prioritization — is unanswered because no enforcement action has been validated in the field.

Q. Why was satellite ingest not scheduled in the workflow from the start? The connector was built but sat dormant. Is this a resource constraint, a late decision, or a technical blocker that was just resolved?

The GitHub Actions workflow initially called only openaq, cpcb, and openmeteo — daily satellite ingest was missing. Commit f31bf20 (2026-08-19 16:58) added it, with a comment stating: 'This step was MISSING entirely — the connector existed and had been run once by hand, but nothing re-ingested on a schedule.' The reason for the omission is not documented. The fix makes satellite ingest conditional on GEE secrets; if they're absent, the step skips cleanly.

Q. For the six languages without native-speaker review (Kannada, Tamil, Telugu, Bengali, Gujarati, English body), what does 'script-validated' actually mean? Does it catch medical errors, or only wrong character sets?

Script validation (test_ivr_voices.py::script_ok()) checks that characters are in the target script and no foreign-script characters appear. It does NOT check idiomaticity, medical appropriateness, or grammatical correctness. A Tamil speaker could flag that the phrasing 'N95 மாசுக்குக்' is awkward if it should be 'N95 முகக்கவசம்', but the automated check would pass both.

Q. Why does the canonical rolling benchmark (delhi.json, 10 origins, 90-day window) omit ablation_no_meteorology when benchmark.py:281 and :429-441 clearly intend to compute it?

The rolling benchmark should include ablation for 10 monthly origins but has no ablation section in output JSON. This prevents verification of the 15-35% meteorology claim on the protocol claimed to carry weight. Either: (1) ablation computation disabled for rolling, or (2) benchmark run with --no-ablation flag.

Q. What is the actual observed first-request latency in production after the API has been cold for >1 hour, measured from real user browsers?

The team measures pipeline latency (api/main.py:1741-1753, /latency endpoint ~1130ms wall-clock) but this is the happy-path instrumented latency, not real-world cold-start latency. The warm-up daemon helps, but there's no production measurement of what a real user sees on their first navigation after idle. This is worth measuring before the demo to know if the 'first real request is fast' claim holds in practice.

Q. Why is Jaipur live as a production city when it has negative forecast skill at all three horizons (+24h -4%, +48h -14%, +72h -34%)?

PS5_HONEST_AUDIT.md (§3.2, table line 208) confirms Jaipur loses to persistence across all horizons. The audit notes this as a 'known deficit' (§2 item 1) with 'a likely cause, not yet a fix'—the blending weight `w` is set to 1.0 (pure model, no persistence fallback). No threshold for removing a negative-skill city is documented. The question a judge will ask: if the model is worse than 'tomorrow = today', why show it at all? What's the minimum skill gate for a city to stay live?

Q. Can you actually onboard a new 131st NCAP city live during the demo, or is that theoretical? What's the real wall-clock time from 'POST /admin/cities' to 'city renders on the map with forecast + attribution'?

ARCHITECTURE.md (§18) claims 'demo the claim: run Delhi + Bengaluru + Mumbai; then onboard a 4th city live from config on stage.' The endpoint exists (api/main.py:2503). But the city won't have data until the next hourly ingest (ground+weather) and daily forecast (forecast-enforcement-rollup). At best, a new city onboarded at 14:00 UTC would have its first data at 15:00 (hourly ingest) and forecast at 01:30 UTC next day (28.5 hours later). The 'live onboarding' demo is actually 'config that takes effect on the next scheduled job.' This is an honest gap between the demo narrative and the data pipeline reality.

Q. What happens to the daily forecasting and enforcement pipeline when GitHub Actions approaches or exceeds the 2000-minute/month free limit? Is there a safeguard or does the pipeline silently fail mid-cycle?

The answer mentions splitting into 5 parallel jobs per state region to stay under 2000 min/month for 131 cities (SCALE.md, ARCHITECTURE.md). But there's no documented safeguard in .github/workflows/ingest.yml for the current 10-city deployment. If a job runs long (e.g., OpenAQ rate-limiting causes retries), there's no abort logic, just '|| true' to continue. The risk: a single slow run could burn the month's budget mid-cycle, leaving enforcement offline for the rest of the month.

Q. Your README says 'six agents' (line 117) but the honest audit says five nodes registered and four running live. How do you reconcile the contradiction in your own documentation?

PS5_HONEST_AUDIT.md (§2.1) confirms the discrepancy: 'six agents' is claimed in README.md:44, web/src/Landing.tsx:70, docs/SUBMISSION.md:33, docs/PITCH_SCRIPT.md:71, and docs/DEMO_VIDEO_SCRIPT.md:69, but agents/graph.py:344-351 registers only five nodes (orchestrator, attribution, forecast, spike_gate [conditional router], enforcement, advisory). The audit recommends: 'Fix the copy, not the code—say Five agents on one graph plus a spike gate.' The contradiction will be caught if a judge presses 'Run agents live' and counts the boxes on screen.

Q. Q14: How do you reconcile a 30% PM2.5 reduction target (quantify.py line 142: reduction_pct=0.30) against the NCAP national target? Is 30% your assumption, the actual NCAP target, or a modelled outcome?

From NCAP Phase 1 (2017–2019), India targeted 20–30% reduction in PM2.5 by 2024 in 102 non-attainment cities. VayuNetra's default 30% (line 142) is the upper end of that range and is defensible as a NCAP-grade scenario. But it is not explained in the /roi response. The answer should cite the NCAP Phase 1 target and note that a policy user can override it (the response does not show the reduction_pct parameter, so it appears hard-wired).

Q. Q15: Attributable deaths uses baseline crude death rate (7.3 per 1,000 per year). But SRS India has known underreporting. How much of your 73,395 figure could be false negatives if true all-cause mortality is 10–15% higher than SRS?

SRS underreporting (estimated 10–15% nationally by Chandrasekaran et al. 2014, Lancet Glob Health) means the true crude death rate could be 8.0–8.4 per 1,000, yielding true attributable deaths of 80,368–92,453 (not 73,395). VayuNetra cites SRS as the standard for India (which it is, for GARP/GBD work) but does not surface this uncertainty. The /roi response should note: 'SRS-based figure; true all-cause mortality may be 10–15% higher; attributable deaths could range 73K–92K.'

Q. Q16: In quantify_intervention (lines 94–105), you sum attributable cases across all cells. But what happens if one cell has zero population (pop = 0)? Does the function gracefully handle it, or can it silently drop that cell's averted deaths?

Looking at attributable_cases (lines 31–47), if population ≤ 0, it returns 0.0 (line 44). This is safe. But in quantify_intervention, if F.POP_PER_CELL.value (40,000) is used and a cell has fewer people, the method overestimates. For peripheral cells, a uniform 40K per cell can overstate exposure. The answer should clarify: We use a heuristic 40K per metro cell as a default; the E3 engine should override this with WorldPop-refined counts. If WorldPop is unavailable, the 40K figure is order-of-magnitude.

Q. Q17: The biomass_burning CO₂ ratio (166 t CO₂ per t PM2.5) comes from Andreae (2019), field-burning average. But PM2.5 from stubble burning vs. other biomass (forest fires, dung, agricultural waste in other seasons) varies 2–3x in emissions factor. How do you know Delhi's autumn stubble burn has 9.1 g PM2.5 / kg, not 5 or 15?

Andreae (2019) reports 9.1 g PM2.5 / kg for agricultural-residue burning (a meta-average), but India's Punjab/Haryana stubble burning can have higher PM2.5 yield due to crop type (rice residue) and burning efficiency. VayuNetra cites Andreae as the source and notes the caveat ('field-burning average; real ratio varies'). But the caveat should be stronger: Delhi's autumn PM2.5 co-benefit from biomass-burn reduction could be 50–150% different. The /quantify_intervention response does include the citation, so a user can check, but the point estimate (e.g., 332 tonnes CO₂e) could be off by half.

Q. DEMO_MODE defaults to TRUE. What happens in production if DEMO_MODE is not explicitly set to 'false' in the environment?

If DEMO_MODE is not set, os.getenv('DEMO_MODE', 'true').lower() == 'true' evaluates to true, and the entire API serves only fixture data from demo/fixtures/*.json instead of reading from the live Supabase database. Every endpoint returns canned responses. In production, this would make the system appear to work while actually ignoring all real data. A deployment mistake (forgetting to set DEMO_MODE=false in the production .env) would silently fail in the worst way: the app looks operational but is serving stale fixtures. This is a footgun that should default to false, not true.

Q. The broadcast endpoint rate-limiting uses a shared in-process dictionary _last_broadcast without a lock. Is there a race condition if two requests broadcast the same city simultaneously?

Yes. The code at line 1261 checks `if now - _last_broadcast.get(city, 0.0) < _BROADCAST_WINDOW_S` without holding a lock. In an async or multi-threaded environment, two concurrent requests can both see the old timestamp, both pass the check, both set `_last_broadcast[city] = now`, and both send the broadcast. This defeats the rate limit. The fix is to acquire a `threading.Lock()` around the check-and-set operation: `with _BROADCAST_LOCK: if ... _last_broadcast[city] = now`. This is a classic TOCTOU (time-of-check-time-of-use) bug.

Q. The citizen /report endpoint accepts photo uploads with only size validation (4 MB) and then stores them in a public bucket. Are there any file-type safeguards against malicious uploads?

No. The code uploads the blob with a hardcoded content-type of 'image/jpeg' (or the client-provided type) without validating the actual file contents. An attacker can upload a ZIP, executable, or other file with a .jpg extension and it will be stored in the public citizen-reports bucket and served to anyone who gets the photo_url. There is no magic-byte validation, no antivirus scan, no virus-signature check. The fix is to validate file magic bytes (check actual file type, not extension), reject non-image files, and optionally scan with ClamAV or similar before storing.

Q. The in-process rate-limiting dictionaries (`_last_report`, `_last_broadcast`, `_STATUS_EVENTS`) reset on server restart. In a load-balanced or multi-instance deployment, how does rate limiting work?

It doesn't. Each FastAPI worker process has its own in-memory dictionary. With 4 workers, an attacker can send 4 requests per minute to /report (one per worker) instead of 1 per minute globally. The rate limits are per-process, not per-IP or per-account. In a Kubernetes or multi-instance deployment, each pod has its own limit, so an attacker can bypass the rate limit by hitting different pods. The fix is to move rate-limiting state to Redis or Supabase so all workers share the same limit window.

Q. How many of the ten cities have GEE (Earth Engine) secrets configured in the live GitHub Actions environment so satellite ingest actually runs?

This is a deployment verification question. The audit confirms satellite connectors were built but never run on schedule; the workflow step was added but only executes if secrets are present. Without knowing which repos have GEE credentials configured, you cannot confirm satellite data is flowing to any city. This should be verifiable from the GitHub Actions run logs or by checking if satellite rows exist in the live measurements table after 2026-08-18.

Q. The audit flags API coverage at 33% for `api/main.py` (the entire API contract). Given that judges will open the repo, what is your timeline to raise this to 80%?

The audit explicitly states this as actionable before the final. `api/main.py` at 33% coverage is the first thing a technical judge sees when evaluating code quality. The audit ranks this as item 7 in 'Where the remaining hours should go' (~half a day effort). This is a direct test-quality question that should be rehearsed.

Q. For enforcement quality (currently graded F: zero domain expert ratings), do you have a single municipal officer or academic already identified who could rate the worklist in the next 48 hours?

The audit lists 'Get one officer or academic to rate the worklist' as the #2 priority action. The F grade on this criterion is the only pure F on a named evaluation metric. This is not a code problem, so it cannot be fixed by writing. Knowing whether you have one person committed to review (and their likelihood to respond) is different from having 'outreach kits' prepared.

Q. Kolkata's prediction interval coverage degrades to 0.547 at +72h in the 56–76 $\mu\text{g}/\text{m}^3$ range (the regulatory decision boundary). Is there a confidence flag in the API response, or does it return the same `pi80` band regardless of predicted level?

The audit states this as 'unfixed, quantified, and disclosed,' meaning the breakdown is published but likely not surfaced to a user. This is a UX/product question: an officer using the +72h forecast for Kolkata in the decision-critical middle range is getting a prediction interval with 54.7% coverage (below the claimed 80%). Does the UI warn of this, or does it display the band as if it were fully calibrated?

Q. How do judges distinguish between 'the model abstained because the entire city failed the R^2 skill gate' vs 'the model hasn't been trained'? If both result in signature priors for all Delhi cells, what UI affordance tells them why?

No UI affordance exists. Both failure modes show signature priors. `attribute.py:142` records the reason in evidence JSON, but the cell story does not differentiate between city-level gate failure and untrained model. Users see the same output for fundamentally different reasons.

Q. When Kolkata's conformal band covers only 54.7% at +72h in the 56-76 $\mu\text{g}/\text{m}^3$ range—the exact range where officers enforce—how is this specific conditional-coverage failure communicated at serve-time? Do officers see the breakdown, or only the aggregate 80%?

`PS5_HONEST_AUDIT.md` documents the breakdown but no evidence shows it reaches the officer UI. The forecast card likely shows a single 80% number with no disclosure of where it fails. The honest audit acknowledges the gap but provides no solution for end-user communication.

Q. What liability framework exists if a non-reviewed language (Kannada, Tamil, Telugu, Bengali, Gujarati) advisory contains medical misinformation and a citizen is harmed?

None documented. ADVISORY_REVIEW.md lists 6/8 languages as 'pending native-speaker review.' No signed-off review, no legal disclaimers, no explanation of how reviewed vs. pending languages differ in liability or accuracy guarantees.

Q. The 'open waste burning' enforcement recommendations come from OSM + CV detection. What false-positive rate for these detections triggers a wasted inspector visit, and have officers flagged bad detections?

No metric exists. PS5_HONEST_AUDIT.md notes 487 of 647 sources are 'cv_detected' but provides no false-positive rate or officer feedback on CV hallucinations. The enforcement loop may systematically waste officer time on model errors without measuring it.

17 • Presenting this system

17.1 The position that actually wins

Most hackathon projects are presented as capabilities. This one is stronger presented as **capabilities plus the measurement of their limits**, because the limits are measured — which is rare, checkable, and exactly what a rigorous judge is probing for.

The single most defensible thing about this project is not the forecast or the attribution. It is that when we found our own claims overstated, we corrected them and wrote down what happened. The audit that produced chapter 13 was not commissioned by anyone. If a judge takes one thing away, it should be that this team's numbers can be trusted because the team went looking for reasons they could not.

17.2 The four sentences

If you have thirty seconds:

A city knows its air is bad. It does not know which square kilometre, caused by what, for how long, or who to send. VayuNetra answers those four questions per square kilometre for ten Indian cities, and ends in a cited draft legal notice with a tracked outcome. Every number it reports carries the uncertainty around it — including the two cities where our forecast does not beat a naive baseline, which we publish rather than hide.

17.3 Claims to make precisely

Say this	Not this
"Our 80% band measures 0.783 on Delhi over 207k rows, and we publish coverage by predicted level because the pooled number hides where a band fails."	"Our intervals are calibrated."
"The forecast beats persistence in eight of ten cities, with confidence intervals; in two it does not, and Jaipur's deficit is real rather than sampling noise."	"Our model beats the baseline by 9%."
"The grid is 1 km and the interpolation runs; per-cell accuracy at 1 km is validated against synthetic fields, not held-out stations."	"Validated 1 km hyperlocal accuracy."
"Five agents on a LangGraph state machine, plus a conditional spike gate — a clean-air city legitimately runs four."	"Six agents."
"Satellite ingestion runs daily as of 19 August; before that it was built but unscheduled, which our own audit caught."	"Satellite data powers our attribution in real time."
"Attribution agrees with published apportionment to a mean absolute difference of 0.042 on Delhi."	"Cosine similarity 0.991."

Say this	Not this
"No official has used this operationally, and no expert has rated a recommendation. That is our largest gap."	Anything implying deployment or validation by a third party.

17.4 The four questions most likely to hurt

"Has any real official used this?" — No. Say it in one sentence, then say what exists instead: the rubric, the outreach kit, the full audit trail that would record an officer's actions. Do not pad. The worst version of this answer is a long one.

"Your model loses to a naive baseline in a city — why should we trust it?" — Because we can tell you *that* it loses, with a confidence interval, and most projects cannot. Then give the diagnosis: the persistence blend lands at the pure-model corner in five of ten cities, and Jaipur does so at exactly the horizons it loses. It is a named lead with an experiment attached, not a mystery.

"You claim 1 km resolution — prove it." — The grid is real and the interpolation runs; the skill number behind it is measured against synthetic fields. Say that before they find it. The leave-one-station-out test is computable from data we already hold and has not been run.

"Is this actually AI, or a dashboard with a nice map?" — Two model families with measured skill, a conformal calibration layer, SHAP-based apportionment with an abstention gate, and a state machine that routes on a detected condition. Then volunteer the honest counterweight: no large language model writes anything a citizen reads, and that is deliberate — a hallucinated medical instruction in Marathi reaching a phone line is an unacceptable failure mode, so citizen text is deterministic templates.

17.5 What to demonstrate live, and in what order

- The map, on Delhi.** One click into a cell — the cell story names the locality, not a hex id.
- The attribution panel.** Show the method badge. If it says signature priors, say why: the model failed its own R^2 gate and the system refused to over-claim. Abstention is the feature.
- The forecast band,** then the coverage breakdown behind it. This is where the project is strongest and where most competitors have nothing.
- One enforcement recommendation end to end** — evidence, citation, draft notice, status change. This is the loop nothing else closes.
- An advisory in a non-English language,** and if the room allows it, the IVR call. A Marathi speaker hearing a Marathi voice is more persuasive than any slide.

Do not demonstrate: Jaipur's forecast, or SHAP attribution on Delhi (thin history means it falls back to priors — correct behaviour, but it needs explaining and a demo is not the place).

17.6 If something breaks

The web app falls back to bundled fixtures when the API is unreachable and shows a banner saying so. If that banner appears, say what it is — "that is the demo insurance doing its job, the API host cold-starts after inactivity" — rather than pretending it did not happen. A team that explains its own failure mode calmly reads as one that has thought about production.

18 • Troubleshooting

Symptom	Meaning / fix
Amber "backend waking up" notice	Free-tier cold start on the API host; the console shows bundled fixtures meanwhile. Retry or wait — it clears on the first successful call, and the notice never blocks the map controls. Run <code>make prewarm</code> before a demo.

Symptom	Meaning / fix
Map loads, panels empty	Backend still waking, or a city with thin data — switch city once.
"No per-cell forecast" in the Cell Story	That cell has no fresh per-cell rows; the city Forecast section still works.
Worklist says nothing to enforce	Air is clean, so the spike gate skipped enforcement — correct behaviour — or the filters exclude everything.
Advisory empty in a language	That city does not publish that language; the language list comes from the city config in <code>core/config/cities/<id>.yaml</code> .
Attribution says "chemical-signature priors" not SHAP	The per-cell model failed its own R^2 gate or had too little history, so the system fell back rather than over-claim. The reason is recorded in the row's <code>evidence.fallback_reason</code> .
Notice PDF takes several seconds	Live dossier assembly — RAG retrieval plus a satellite patch. Expected.
Tour will not reappear	Once-only by design; replay it with <code>?</code> .
Simulator says "near-zero effect"	An honest result. Pick the dominant source for that cell.
A broadcast went out in the wrong language	Check the language dropdown, then the response line — it reports <code>requested</code> versus <code>delivered</code> , and says when a city has no advisory stored in that language.
Real or demo?	<code>/health</code> reports <code>DEMO_MODE</code> . The landing page and console say "as of" and "snapshot" wherever a fixture is being shown.

Glossary

H3 — Uber's hexagonal spatial grid; resolution $8 \approx 0.74 \text{ km}^2$, about a kilometre across. **SHAP** — per-feature contribution to a single prediction, here in $\mu\text{g}/\text{m}^3$. **CQR** — conformalised quantile regression; the method that makes the 80% band actually cover ~80%. **Marginal vs conditional coverage** — whether a band covers 80% of *all* rows pooled, versus 80% within every regime separately. Split conformal guarantees the first, not the second. **P(>120)** — calibrated probability that a cell exceeds the CPCB Very Poor threshold. **Persistence** — the "tomorrow equals today" baseline. Harder to beat than it sounds. **Skill** — $1 - \text{RMSE}_{\text{model}} / \text{RMSE}_{\text{baseline}}$. Positive means better than the baseline. **Onset recall** — the share of clean→Very-Poor transitions the alarm flags in advance. **GRAP** — Delhi-NCR's Graded Response Action Plan, Stages I–IV. **NCAP** — India's National Clean Air Programme. **PRANA** is the portal cities report actions to. **RAG** — retrieval-augmented generation; here it retrieves regulation text for citation, and does not generate prose. **RLS** — Postgres row-level security. **GPW** — Gridded Population of the World v4.11, the population raster. **Receptor model** — the classical source-apportionment approach, using chemically speciated filter samples. What this system approximates without speciation.

Appendix A • Corrections applied during review

Each chapter was drafted from the code and then independently fact-checked against it. These are the errors that pass caught and the corrections applied. They are printed because a document that claims to be verified should show its working.

1. At 5000 m downwind, Class D gives $\sigma_y \approx 600 \text{ m}$, $\sigma_z \approx 400 \text{ m}$

Correct: σ_y at 5000 m is approximately 462 m (not 600 m). Using the D-class formula from `plume.py` line 16: $\sigma_y(x) = 0.16 \times 5000 \times (1 + 0.0004 \times 5000)^{-0.5} = 800 / \sqrt{3} \approx 461.9 \text{ m}$. σ_z is approximately 443 m (not 400 m). **Checked:** `ml/dispersion/plume.py` lines 16 and 22; verified with manual calculation

2. The `validation` field in the `/dense-field` API response includes this skill metric

Correct: The API endpoint is called `/coverage`, not `/dense-field`. Defined at `api/main.py` line 2227 with decorator `@app.get("/coverage", tags=["stage2"])`. **Checked:** `api/main.py` line 2227, endpoint definition

3. `kb_chunks` table has embedding vector(768)

Correct: embedding column is `vector(384)` by default (can be changed via `EMBEDDING_DIM` env var to 768 for Gemini embeddings, but default is 384 for bge-small) **Checked:** `supabase/migrations/20260627000001_init.sql:147` states `'embedding vector(384)'` with comment `'Default = 384 (bge-small)'`; `.env.example` has `EMBEDDING_DIM=384`; `rag/retrieve.py` line 3 has `EMBEDDING_DIM = int(os.getenv('EMBEDDING_DIM', '384'))`

4. The latency trace includes per-node duration values as shown in the example (e.g., 'attribution → ts: 15:00:00.145 (duration: 22ms)')

Correct: The `_stamp()` function only stores node name, timestamp, and metadata—not per-node `duration_ms`. Only total latency is computed from first to last trace entry (`agents/graph.py:88-97`) **Checked:** `agents/graph.py:77-85` shows `_stamp()` appends `{node, ts, meta}` but never computes or stores `duration_ms`. `TraceEntry` has `duration_ms` field (line 43) but it is never populated in the actual code.

5. 10 connectors normalize raw payloads → H3 cells at res 8, as shown in the Connectors table with 7 rows

Correct: The connectors table displays 7 data source connectors (`cpcb`, `openaq`, `openmeteo`, `earth_engine`, `osm_sources`, `mobility`, `population`), not 10. There are 12 total connector files in the codebase. **Checked:** Chapter text states '10 connectors' but table shows only 7 rows. Directory listing of `/home/omkar-kadam/Desktop/VayuNetra/connectors/` shows 12 `.py` files: `cpcb`, `openaq`, `openmeteo`, `earth_engine`, `osm_sources`, `mobility`, `population`, `community_sensors`, `permits`, `static_layers`, `traffic_live`, `vulnerability`

6. Non-PM2.5 pollutants and static variables (population) are never archived.

Correct: Non-PM2.5 pollutants ARE archived to Supabase Storage as gzip CSV files and then deleted from the database. Only static variables (population) are never archived and remain in the database forever. Code filter: `'variable <> all(STATIC_VARIABLES)'` where `STATIC_VARIABLES = ('population')` means all other variables are exported. **Checked:** `scripts/archive_measurements.py` lines 36, 48, 74-75, 111-112

7. Pollutant or meteorological variable: `pm25`, `pm10`, `no2`, `so2`, `co`, `o3`, `aod`, `fire`, `wind_u`, `wind_v`, `blh`, `temp`, `rh`, `precip`, `traffic`, `population`

Correct: The variable list is incomplete and missing `'no2_sat'` (Sentinel-5P satellite NO2 measurements, unit `mol/m²`). The complete CHECK constraint includes 17 values: `('pm25','pm10','no2','so2','co','o3','aod','no2_sat','fire','wind_u','wind_v','blh','temp','rh','precip','traffic','population')` **Checked:** `supabase/migrations/20260718000001_indexes_and_constraints.sql` lines 56-60

8. Mumbai very-poor hours (>120) skill: -2.6% at +24h, -0.0% at +48h, -0.1% at +72h

Correct: Correct values for Very Poor (>120) are -5.0% at +24h, -3.6% at +48h, -2.8% at +72h. The chapter's table shows `observed_over_90` (Poor >90) values instead. **Checked:** `docs/benchmarks/mumbai.md` lines 25-26 (-5.0%), 28-29 (-3.6%), 31-32 (-2.8%)

9. connectors/openaq.py:1–229

Correct: connectors/openaq.py:1–228 (file contains 228 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/openaq.py

10. connectors/osm_sources.py:1–279

Correct: connectors/osm_sources.py:1–278 (file contains 278 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/osm_sources.py

11. connectors/earth_engine.py:1–142

Correct: connectors/earth_engine.py:1–141 (file contains 141 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/earth_engine.py

12. connectors/population.py:1–112

Correct: connectors/population.py:1–111 (file contains 111 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/population.py

13. connectors/vulnerability.py:1–216

Correct: connectors/vulnerability.py:1–215 (file contains 215 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/vulnerability.py

14. connectors/static_layers.py:1–182

Correct: connectors/static_layers.py:1–181 (file contains 181 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/static_layers.py

15. connectors/permits.py:1–86

Correct: connectors/permits.py:1–85 (file contains 85 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/permits.py

16. connectors/traffic_live.py:1–106

Correct: connectors/traffic_live.py:1–105 (file contains 105 lines, verified with `wc -l`) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/traffic_live.py

17. Scoring logic at lines 61–73 in osm_sources.py

Correct: Scoring logic is at lines 61–72; the `_score` function ends at line 72, not 73 **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/osm_sources.py lines 61-72

18. OpenAQ connector uses exponential backoff up to 2^attempt seconds

Correct: Actual backoff formula is $\min(60, 2^{\text{attempt}} * 3)$, so backoff times are 3s, 6s, 12s, 24s, 48s, 60s (not 1s, 2s, 4s, 8s, 16s, 32s) **Checked:** /home/omkar-kadam/Desktop/VayuNetra/connectors/openaq.py line 84

19. The priority score formula (agents/enforcement.py:203–305) combines four factors

Correct: The `_compute_priority()` function containing the priority score formula is at `agents/enforcement.py:294-305`, not 203-305. Lines 203-248 contain the `_compute_value()` function for value-per-inspector-hour computation, which is a separate calculation described in the following section **Checked:** `agents/enforcement.py: grep -n 'def _compute_priority' and grep -n 'def _compute_value'` confirm `_compute_value` at line 203 and `_compute_priority` at line 294

20. /comparison cache is 5 min (300 seconds)

Correct: /comparison cache is 60 seconds (1 minute) **Checked:** `api/main.py` lines 92-110, `_CACHE_SECONDS` tuple shows ('comparison', 60)

21. /history/trend cache is 10 min

Correct: /history/trend cache is 120 seconds (2 minutes) **Checked:** `api/main.py` line 98, `_CACHE_SECONDS` tuple shows ('history/trend', 120)

22. /history/cells cache is 10 min

Correct: /history/cells cache is 120 seconds (2 minutes) **Checked:** `api/main.py` line 99, `_CACHE_SECONDS` tuple shows ('history/cells', 120)

23. /coverage cache is 10 min (HTTP middleware level)

Correct: /coverage HTTP middleware cache is 120 seconds (2 minutes), though manual cache via `_dense_field_cached` is 600 seconds (10 min) **Checked:** `api/main.py` line 100 shows `_CACHE_SECONDS` ('coverage', 120); line 2244 shows `_DENSE_TTL_S = 600`

24. /advisory is open (no auth required)

Correct: /advisory requires auth in live mode - uses `Depends(get_db)` which calls `_validated_token()` **Checked:** `api/main.py` line 971-976, /advisory uses `Depends(get_db)`

25. /aqi/current is open (no auth required)

Correct: /aqi/current requires auth in live mode - uses `Depends(get_db)` **Checked:** `api/main.py` line 256-260, /aqi/current uses `Depends(get_db)`

26. /forecast is open (no auth required)

Correct: /forecast requires auth in live mode - uses `Depends(get_db)` **Checked:** `api/main.py` line 638-643, /forecast uses `Depends(get_db)`

27. /attribution is open (no auth required)

Correct: /attribution requires auth in live mode - uses Depends(get_db) **Checked:** api/main.py line 587-594, /attribution uses Depends(get_db)

28. /interventions is open (no auth required)

Correct: /interventions requires auth in live mode - uses Depends(get_db) **Checked:** api/main.py line 877+, /interventions uses Depends(get_db)

29. /coverage is open (no auth required)

Correct: /coverage requires auth in live mode - uses Depends(get_db) **Checked:** api/main.py line 2227+, /coverage uses Depends(get_db)

30. /advisory/wards requires auth

Correct: /advisory/wards is open (no auth required) - does not use Depends(get_db) **Checked:** api/main.py line 1239-1249, /advisory/wards has no db parameter

31. /reports requires auth

Correct: /reports is open (no auth required) - does not use Depends(get_db); only /report/{id}/status requires auth **Checked:** api/main.py line 1089-1113, /reports has no db parameter, whereas /report/{id}/status at line 1117 uses Depends(get_db)

32. Smoke suite has 7 flows

Correct: The smoke.spec.ts file contains 8 test() functions, not 7. The 8th test 'simulator section shows the what-if engine' is present in the file but not listed in the chapter's 7-flow summary. **Checked:** web/e2e/smoke.spec.ts lines 15-102 contain 8 distinct test() declarations